

# **Foundations of Computer Architecture and Cyber-Physical Systems 2.0**

**A Deterministic Approach Using VirtMCU**

**Refract Systems**

*A comprehensive graduate-level text bridging the gap between discrete digital logic  
and continuous-time physical dynamics.*

---

**Version: 1.0.0 Compiled: 2026**

# Preface: How to Use This Book

The exponential increase in computing power and the pervasive integration of digital systems into the physical world have fundamentally transformed engineering. We are no longer simply designing isolated embedded systems; we are engineering **Cyber-Physical Systems (CPS 2.0)** that require a deep, integrated understanding of both hardware microarchitecture and continuous-time physical dynamics.

This textbook was written to bridge the gap between traditional computer architecture and modern cyber-physical system design. It is designed to take students on a journey from the fundamental digital logic that makes up a processor, through the operating systems that schedule its tasks, and finally to the distributed networks, deterministic co-simulations, and domain-specific architectures that define the modern computing landscape.

## Target Audience and Prerequisites

This book is intended for senior undergraduate and graduate students in computer engineering, electrical engineering, and computer science. Because modern CPS design sits at the intersection of hardware and software, the text assumes readers are familiar with the basics of computer science and have completed a first programming course in C or C++. Furthermore, since we explore how hardware controls physical actuators, a basic understanding of continuous-time mathematics and introductory electronics will be highly beneficial.

## Organization of the Book

The material is systematically divided into distinct parts to manage the complexity of the subject matter:

- **Part I: Machine Organization and Microarchitecture (Chapters 1–6)** builds the foundation. It covers digital logic, the ARM AArch64 Instruction

Set Architecture (ISA), pipelined microarchitecture, the memory hierarchy, and System-on-Chip (SoC) interconnects.

- **Part II: The Hardware-Software Interface (Chapters 7–8)** crosses the boundary into software, focusing on bare-metal execution, exceptions, interrupts, and the critical scheduling algorithms of Real-Time Operating Systems (RTOS).
- **Part III: Cyber-Physical Systems and Co-Simulation (Chapters 9–11)** introduces the CPS 2.0 paradigm. It explores the 3Cs (Computation, Communication, and Control), the necessity of deterministic virtual time, and the Sensor/Actuator Abstraction Layers (SAL/AAL) required to couple discrete-event cyber simulators with continuous-time physics engines.
- **Part IV: Advanced Architectures and Security (Chapters 12–14)** looks to the future, examining the open-source RISC-V architecture, cyber-physical threat intelligence and resilience, and Domain-Specific Architectures (DSAs) that accelerate AI at the edge and in the cloud.

## Using This Book in a Course

This book contains enough material for a rigorous, rapid-paced, single-semester introduction to computer architecture and CPS, or it can be expanded into a two-semester sequence to allow students more time to experiment in the lab.

### For a Single-Semester Course (14-15 Weeks):

- **Weeks 1–5:** Focus heavily on **Chapters 1 through 5** to establish a strong hardware and microarchitecture foundation.
- **Weeks 6–8:** Transition to software with **Chapters 6 through 8**, emphasizing interrupts, memory-mapped I/O, and RTOS determinism.
- **Weeks 9–12:** Introduce the CPS paradigm using **Chapters 9 through 11**. This is a great time to introduce mid-term laboratory projects involving co-simulation.
- **Weeks 13–15:** Conclude with advanced topics based on instructor interest. An instructor focused on hardware design might prioritize **Chapter 12 (RISC-V)** and **Chapter 14 (DSAs)**, while an instructor focused

on networks and infrastructure might focus on **Chapter 13 (Security and Resilience)**.

### **For a Two-Semester Sequence:**

- **Semester 1 (Computer Architecture):** Dedicate the entire first term to **Chapters 1 through 8**, allowing students to deeply digest the instruction set principles, pipelining, and memory hierarchies. The labs can focus on translating C code to ARM assembly, observing cache hits/misses, and writing bare-metal device drivers.
- **Semester 2 (Cyber-Physical Systems):** Dedicate the second term to **Chapters 9 through 14**. This allows for a much deeper exploration of control algorithms, distributed optimal power flow, and modifying the RISC-V instruction set.

## **Laboratory and Software Tools**

The best way to learn digital design and system architecture is to do it. However, verifying hard real-time CPS constraints on standard operating systems is often frustrated by wall-clock scheduling jitter. To resolve this, this book includes companion homework assignments designed around the **VirtMCU FirmwareStudio**.

Detailed instructions on how to use the software tools in a course are provided alongside the text. The VirtMCU framework allows students to execute ARM and RISC-V binaries in a deterministic, `slaved-icount` mode. By pairing VirtMCU with external physics engines (such as MuJoCo) via zero-copy shared memory, students can gain hands-on experience balancing an inverted pendulum or stabilizing a drone using the precise cycle-accurate timings required in the real world.

# Chapter 1: Fundamentals of Computer Architecture and Digital Logic

## 1.1 Introduction: The Art of Managing Complexity

The design of modern computer systems is a remarkable feat of engineering. A contemporary microprocessor contains billions of microscopic transistors, each switching billions of times per second. No human being could possibly understand or design such a system by writing continuous mathematical equations describing the movement of electrons through each individual transistor. To build these vastly complex systems, computer scientists and engineers rely on a systematic approach to managing complexity: the use of strict **hierarchical abstraction**.

Abstraction involves hiding lower-level details when they are not important to the current problem, allowing engineers to construct complex systems layer by layer. At the very bottom of the computing hierarchy lies solid-state physics, which dictates how electrons move through semiconductor materials. Above physics is the device level, consisting of individual transistors. Transistors are wired together to form analog and digital circuits, which are in turn organized into logic gates (such as AND, OR, and NOT). These logic gates are combined to build microarchitectural components like adders, multiplexers, and memories, which ultimately implement the Instruction Set Architecture (ISA)—the hardware-software interface. Operating systems and high-level application software sit at the very top of this abstraction stack.

## 1.2 The Digital Abstraction

In the physical world, variables such as voltage, frequency, and mass are continuous. However, designing a computer to perform precise mathematical calculations using continuous analog voltages is highly prone to noise and signal degradation. To overcome this, architects use the **digital abstraction**. Digital systems intentionally restrict their signals to a discrete set of values.

In a binary digital system, information is represented using just two states: 0 and 1, which correspond to the Boolean logic values FALSE and TRUE, or electrical states LOW and HIGH. By disciplining ourselves to use only these two discrete states, we can combine simple components into highly reliable and sophisticated systems.

Beneath this digital abstraction, the physical work is performed by Complementary Metal-Oxide-Semiconductor (CMOS) transistors. A transistor behaves fundamentally as an electrically controlled switch.

- **nMOS Transistors:** Turn ON (conduct current) when a high voltage (logic 1) is applied to the gate, and turn OFF when a low voltage (logic 0) is applied.
- **pMOS Transistors:** Turn ON when a low voltage (logic 0) is applied to the gate, and turn OFF when a high voltage (logic 1) is applied.

By arranging these nMOS and pMOS switches into complementary pull-up and pull-down networks, engineers create robust logic gates. Because digital logic circuits interpret continuous voltages falling within specific ranges as either a strict 0 or 1, they inherently reject minor analog noise, ensuring perfect signal reconstruction across billions of operations.

## 1.3 Historical Milestones in Computing

To appreciate modern architectures, it is helpful to look back at the technological milestones that brought us here.

The conceptual foundation of the digital computer predates the discovery of the electron. In the 1830s, Charles Babbage designed the **Analytical Engine**, widely considered the first attempt to build a general-purpose digital computer. The machine was designed to be entirely mechanical, constructed of brass and powered by a steam engine. It featured a mechanism for executing algorithms and outputting results, foreshadowing modern programmable architecture, though it was ultimately defeated by the inadequate manufacturing technology of its time.

Modern computer history truly began in 1946 with the completion of the **ENIAC** (Electronic Numerical Integrator and Computer) by J. Presper Eckert and John Mauchly. Composed of 18,000 vacuum tubes, ENIAC was the first general-purpose electronic computer. However, programming ENIAC required physically rewiring huge patch panels, a tedious and inflexible process.

The solution to this inflexibility came from John von Neumann, who helped develop the **Stored-Program Concept** with the EDVAC and IAS machines. Von Neumann realized that a program's instructions could be represented in digital form and stored in the computer's memory alongside the data it manipulated. This architecture—where memory holds both instructions and data, and a central processing unit fetches and executes those instructions sequentially—remains the foundation of almost every modern processor.

Over subsequent decades, the industry evolved rapidly from room-sized mainframes to minicomputers, and eventually to personal computers (PCs) and mobile devices. Today, computing spans an enormous continuum: from miniature embedded microcontrollers operating in the Internet of Things (IoT), to Personal Mobile Devices (PMDs) like smartphones, all the way up to massive Warehouse-Scale Computers (WSCs) powering global cloud infrastructure.

## 1.4 Moore's Law and the Power Wall

The astonishing proliferation of computing power over the past half-century was driven by two empirical observations: Moore's Law and Dennard Scaling.

In 1965, Intel co-founder Gordon Moore predicted that the number of transistors that could be economically integrated onto a single silicon chip would double approximately every 18 to 24 months. For decades, **Moore's Law** held true, granting architects a seemingly infinite budget of transistors to implement advanced features like superscalar execution and deep caches.

Simultaneously, the industry benefited from **Dennard Scaling**. In 1974, Robert Dennard observed that as the linear dimensions of a transistor shrunk, the operating voltage and current could scale proportionally. Since the dynamic power consumption of a CMOS circuit is proportional to  $C \times V^2 \times f$  (where  $C$  is capacitance,  $V$  is voltage, and  $f$  is frequency), Dennard scaling implied that computer architects could continuously shrink transistors, pack more of them onto a die, and aggressively increase the clock frequency *without* increasing overall power density.

However, exponential laws eventually hit physical limits. Dennard scaling officially ended in the mid-2000s. As transistors became microscopic, threshold voltages could not be lowered further without causing massive leakage currents, establishing a non-scaling baseline for power consumption. The result was the **Power Wall**: chips became thermally limited, meaning architects could no longer simply increase clock frequencies to gain performance.

Moore's Law has also slowed significantly; the cadence between new semiconductor process nodes has stretched from 2 years up to 3 to 5 years. To continue improving performance within strict power limits, modern architectures rely heavily on multicore designs (placing multiple processors on a single chip) and Domain-Specific Architectures (DSAs) that aggressively optimize efficiency for specific workloads.

## 1.5 Quantitative Principles of Computer Design

Because transistor budgets are constrained by power limits and cost, modern computer architecture is an empirical science. Design decisions are made by analyzing the quantitative trade-offs between performance, price, and energy consumption.



When analyzing and optimizing architectures, computer scientists rely on several fundamental, quantitative principles:

### 1.5.1 The Principle of Locality

Programs do not access their code and data uniformly. Empirical observations reveal a widely held rule of thumb: a program typically spends 90% of its execution time in only 10% of its code. This is known as the **Principle of Locality**, and it manifests in two distinct ways:

1. **Temporal Locality:** If a memory location is accessed, it is highly likely to be accessed again in the near future (e.g., instructions inside a loop, or variables used as counters).
2. **Spatial Locality:** If a memory location is accessed, locations with adjacent or nearby addresses are likely to be accessed soon (e.g., sequentially executing code, or traversing an array of data).

Computer architects aggressively exploit locality by implementing memory hierarchies. Fast, small, and expensive SRAM caches are placed close to the processor to hold the active “working set” of data, while slower, cheaper DRAM holds the remainder.

### 1.5.2 Amdahl's Law

When considering an architectural enhancement (such as adding a specialized hardware accelerator), engineers must calculate the actual speedup it will provide to the overall system. This is governed by **Amdahl's Law**, which states that the overall performance improvement gained by optimizing a specific part of a system is strictly limited by the fraction of the execution time that the optimized part is actually used.

If an enhancement makes a specific operation  $N$  times faster, but that operation only accounts for 20% of the program's execution time, the remaining 80% of the execution time remains entirely unaffected. Ignoring Amdahl's Law is a common pitfall; engineers must always measure where a program spends its time *before* expending the resources to optimize a specific

hardware feature, lest they fall prey to what is often called “Amdahl’s heartbreaking law”.

### 1.5.3 Dependability via Redundancy

Computers must not only be fast and efficient; they must also be dependable. Any physical device will eventually fail. Therefore, architecture dictates that dependability is achieved through **redundancy**. If a single component (such as a power supply or a memory chip) is prone to failure, designers must include redundant components that can automatically take over. In modern systems, this ranges from Error-Correcting Codes (ECC) in memory arrays to fully redundant servers in a warehouse-scale datacenter.

With these foundational rules established, we are ready to venture into the details of digital logic. In Chapter 2, we will explore the mathematical language of computers—Boolean algebra and binary numbering—and demonstrate how physical transistors are combined into the logic gates that perform computation.

# Chapter 2: Data Representation and Machine Organization

## 2.1 Introduction: The Language of the Machine

In digital electronics, a circuit is a network that processes discrete-valued variables. While at the lowest level of physics these variables represent continuous electrical voltages, computer architects abstract these continuous signals into just two discrete binary states: 0 and 1. By disciplining ourselves to use only these two states, we can combine simple components into highly reliable and sophisticated systems,.

Before we can architect a processor capable of executing complex instructions, we must first establish how data is represented mathematically using these binary digits (bits), and how physical transistors are arranged to perform logical computations on those bits. This chapter explores the fundamentals of number systems, combinational logic, and sequential logic, serving as the bridge between raw electricity and organized computation.

## 2.2 Data Representation: Number Systems

Because digital systems consist entirely of 1s and 0s, understanding binary arithmetic is essential. In the binary (base-2) number system, each bit represents a power of 2, just as each digit in the decimal system represents a power of 10.

Because reading and writing long strings of binary digits is tedious and error-prone, assembly language programmers heavily utilize the **hexadecimal (base-16)** number system. Hexadecimal simplifies binary representation by grouping every four bits (a nibble) into a single alphanumeric character from 0 to 9 and A to F. By padding a binary number with leading zeros to make its length a

multiple of four, one can effortlessly translate long bit strings into compact hexadecimal equivalents.

### 2.2.1 Signed vs. Unsigned Integers

The interpretation of a binary sequence depends entirely on the hardware acting upon it. For signed integers, modern computers universally employ the **two's complement** representation.

To find the two's complement (the negative) of a binary number, one simply inverts all the bits and adds 1. This representation has a compelling architectural advantage: **addition works properly and identically for both positive and negative numbers**. A microprocessor does not need separate hardware adders for signed and unsigned math; the exact same logic gates compute the correct sum in both cases, though the processor flag used to detect an arithmetic overflow will differ,.

## 2.3 Boolean Algebra and Logic Gates

Computers process binary signals using logic gates. The fundamental logic gates—**AND**, **OR**, **NOT**, and **XOR**—are rigorously defined by their truth tables,,. A truth table is a tabular representation of a logical expression's output as a function of all possible combinations of its inputs.

### 2.3.1 CMOS Transistor Implementation

Beneath the digital abstraction, logic gates are constructed from **Complementary Metal-Oxide-Semiconductor (CMOS)** transistors. CMOS technology utilizes two types of transistors that act as voltage-controlled switches:

- **nMOS Transistors:** Turn ON to conduct current when a high voltage (logic 1) is applied to the gate.

- **pMOS Transistors:** Turn ON to conduct current when a low voltage (logic 0) is applied to the gate.

By arranging these nMOS and pMOS switches into complementary pull-up and pull-down networks, engineers create robust logic gates,. Because pMOS transistors are effective at passing high voltages and nMOS transistors are effective at passing low voltages, **CMOS circuits naturally favor the implementation of inverting gates like NAND and NOR** over non-inverting gates like AND and OR,,.

### 2.3.2 Boolean Algebra and Minimization

To manage the complexity of millions of interconnected gates, engineers use **Boolean algebra** to mathematically analyze and simplify logic circuits. Any digital logic function can be expressed in a sum-of-products canonical form, where each row of a truth table that results in a TRUE output corresponds to a specific “minterm”.

To reduce the physical footprint, cost, and power consumption of a circuit, designers employ Boolean theorems to minimize these equations before translating them into physical hardware,. A particularly powerful tool is **De Morgan’s Theorem**, which dictates that the complement of the product of all terms is equal to the sum of their individual complements. De Morgan’s Theorem physically proves that a NAND gate is logically equivalent to an OR gate with inverted inputs.

## 2.4 Combinational Logic Circuits

A digital circuit is classified as **combinational** if its outputs depend exclusively on the current values of its inputs,. Combinational circuits are memoryless; a change in the inputs immediately flows through the logic gates to alter the outputs,.

## 2.4.1 Multiplexers and Decoders

Engineers group combinational logic into higher-level building blocks. Two of the most ubiquitous are multiplexers and decoders:

- **Multiplexers (Mux):** Choose an output from among several possible inputs based on the binary value of a select (or control) signal.
- **Decoders:** Perform a translation from a binary input code to a one-hot output. An  $N$ -input decoder asserts exactly one of its  $2^N$  output lines depending on the specific input combination.

## 2.4.2 Arithmetic Circuits

Arithmetic circuits are the central building blocks of the processor. The simplest arithmetic block is the **half adder**, which adds two 1-bit inputs to produce a sum bit and a carry-out bit, typically constructed from an XOR gate and an AND gate,.

Because the half adder lacks an input to accept a carry from a previous column, it is insufficient for multi-bit addition. The **full adder** solves this by incorporating an incoming carry bit ( $C_{in}$ ) into the calculation,. By chaining multiple full adders together—wiring the  $C_{out}$  of one stage into the  $C_{in}$  of the next—architects create a **ripple-carry adder** capable of summing multi-bit data words.

Because electrical signals take finite time to travel through logic gates—a phenomenon known as **propagation delay**—the speed of a ripple-carry adder is limited by how long it takes for the carry signal to cascade from the least significant bit to the most significant bit,.

## 2.5 Sequential Logic: Introducing Memory

While combinational circuits are essential for computing values, a computer must be able to store those values. This requires **sequential logic**, where the

outputs depend on both current and prior input values,. Sequential logic inherently possesses memory.

### 2.5.1 Latches and Flip-Flops

The fundamental building block of memory is the bistable element, a circuit with two stable states that can retain a value indefinitely. By cross-coupling two NOR gates or two NAND gates, engineers create an **SR (Set-Reset) Latch**,. Pulsing the Set input forces the output to 1, and pulsing the Reset input forces the output to 0.

By adding an enable signal and ensuring the Set and Reset inputs are always opposites, the circuit becomes a **gated D latch**, which passes the data input (D) to the output (Q) when the enable is high, but locks and holds the previous value when the enable goes low,.

To construct highly reliable synchronous systems, architects rely on **edge-triggered D flip-flops**,. Unlike a latch, which is level-sensitive, a flip-flop only samples its input and changes its output at the exact microscopic instant that the clock signal transitions (e.g., on the rising edge). This isolates the inputs from the outputs for the remainder of the clock cycle, allowing combinational logic sitting between flip-flops to safely compute the next state without causing chaotic feedback loops.

### 2.5.2 Finite State Machines (FSM)

Complex sequential circuits are formalized as **Finite State Machines (FSMs)**. An FSM consists of a set of states, a state transition function, and an output function.

- In a **Moore machine**, the outputs depend solely on the current state.
- In a **Mealy machine**, the outputs depend on both the current state and the current inputs.

FSMs form the core control logic for processors, stepping the hardware through the sequential phases of instruction fetching, decoding, and execution.

## 2.6 Hardware Description Languages (HDLs)

Modern microprocessors contain billions of logic gates, rendering traditional schematic diagrams completely obsolete for large-scale design,. Instead, computer architects specify the behavioral and structural properties of digital systems using **Hardware Description Languages (HDLs)**,.

The two globally prevalent HDLs are **SystemVerilog** and **VHDL**. These languages strongly resemble traditional programming languages like C, but they are uniquely designed to model the parallel nature of hardware. Using HDLs, engineers can simulate the timing and functionality of a circuit on a workstation before it is manufactured. Once the code is verified, logic synthesis tools automatically translate the HDL text into a highly optimized netlist of physical logic gates and flip-flops,.

For example, in VHDL, a hardware component is cleanly partitioned: the `entity` block defines the external inputs and outputs (the interface), while the `architecture` block defines the internal combinatorial and sequential behavior of the circuit.

With the digital logic foundations established, we now have the building blocks required to design an actual computer processor. In Chapter 3, we will assemble these multiplexers, adders, and registers to form the core execution units of an Instruction Set Architecture.



# Chapter 3: The Instruction Set Architecture (ISA) - ARM AArch64

## 3.1 Introduction: The Hardware-Software Interface

In the previous chapters, we examined the digital logic that physically processes binary signals. We now move up the hierarchy to the **Instruction Set Architecture (ISA) level**. The ISA is the formal boundary between the hardware and the software. It encompasses the complete set of machine instructions, the programmer-visible register set, the memory addressing modes, and the data types understood by a specific microprocessor. For software engineers and compiler writers, the ISA level is what is typically thought of as “machine language”.

An architecture is completely defined by its instruction set and architectural state. While different microarchitectural implementations can vary wildly in cost, performance, and complexity (e.g., adding deeper pipelines or out-of-order execution), they must all successfully execute the same binary programs if they share the same ISA. In this chapter, we will explore the 64-bit ARM architecture (AArch64) as our primary educational ISA and introduce the fundamentals of assembly language programming.

## 3.2 Instruction Set Design: RISC vs. CISC

When designing an instruction set, computer architects must decide how to encode the operations (opcodes) and their operands. Architectures generally fall into two broad philosophies: Complex Instruction Set Computers (CISC) and Reduced Instruction Set Computers (RISC).

CISC processors, such as the Intel x86 family, provide highly complex instructions that can span varying lengths (up to 17 bytes) and can often perform arithmetic operations directly on memory addresses. RISC processors, such as ARM and MIPS, prioritize simplicity and regularity. To simplify hardware decoding, RISC architectures typically encode all instructions in a fixed 32-bit width format.

A defining feature of RISC architectures like the ARM is their use of a **load/store architecture**. All computational activity strictly takes place within the processor's general-purpose registers; the only instructions permitted to interact directly with main memory are those that explicitly load a value from memory or store a value into memory. Furthermore, RISC processors typically employ a three-address machine format for data processing, where an arithmetic instruction explicitly specifies two source registers and one destination register.

### 3.3 The AArch64 Programmer's Model

To write assembly language for the ARM processor, a programmer must understand its architectural state. The 64-bit ARM architecture (AArch64) provides many general-purpose registers for data manipulation. These registers are 64 bits wide and are named **X0** through **X30**. The lower 32 bits of these registers can be accessed independently using the names **W0** through **W30**.

While the CPU hardware treats most of these registers as functionally interchangeable, the software Application Binary Interface (ABI) imposes strict rules on how they are used. Certain registers are dedicated to special purposes:

- **The Link Register (LR / X30):** The ARM processor reserves this register to hold the return address during subroutine and function calls. When a program executes a branch and link ( `bl` ) instruction, the CPU automatically saves the address of the next sequential instruction into the LR.
- **The Stack Pointer (SP):** Used to maintain the call stack in memory, the SP points to the current top of the stack and is heavily utilized for pushing

and popping variables, preserving volatile registers, and saving return addresses during nested function calls,.

- **The Program Counter (PC):** The PC maintains the memory address of the instruction currently being executed by the processor.

## 3.4 Memory Addressing and Endianness

Modern processors utilize byte-addressable memory, meaning each byte in memory is assigned a unique address. When the CPU needs to access a multibyte data type (such as a 32-bit word or a 64-bit doubleword), it must be aware of the system's **endianness**.

In a **little-endian** memory organization, the lowest-order byte of a multibyte value is stored at the lowest memory address,. Conversely, a **big-endian** architecture stores the highest-order byte at the lowest memory address. The ARM architecture actually supports selecting between big- or little-endian modes under software control—a feature known as bi-endianness—though most modern operating systems execute in little-endian mode by default.

### 3.4.1 Addressing Modes

To load data from memory into a register (or store data from a register to memory), the processor must compute the target memory address. ARM supports several addressing modes:

- **PC-Relative:** Computes the address as an offset relative to the current Program Counter. The `adr` and `adrp` instructions use this mode to securely obtain the address of labels and constants located nearby in the program's code section.
- **Register Indirect:** The memory address is provided directly inside a base register (e.g., `ldr r0, [r1]` ).
- **Register Indirect with Offset:** An offset is added to the base register to compute the address. This offset can be an immediate numerical constant (e.g., `ldr r0, [r1, #32]` ) or the value of another register.

- **Pre-increment and Post-increment:** These modes automatically update the base register after computing the address, making them exceptionally useful for iterating through arrays or implementing push and pop stack operations,. For example, `ldr x0, [sp], #16` pops data off the stack and subsequently adds 16 to the stack pointer. Conversely, `str lr, [sp, #-16]!` pushes data by first subtracting 16 from the stack pointer and then storing the Link Register at that new address.

## 3.5 ARM Assembly Language Basics

Assembly language programs use **directives** (or pseudo-instructions) to instruct the assembler on how to organize the program in memory. A typical program is divided into multiple distinct sections:

- **The `.text` section:** This section contains the actual executable machine instructions. For security, operating systems typically map the `.text` section into memory pages that are marked as read-only and executable, meaning any attempt to store data here will result in a segmentation fault,.
- **The `.data` section:** Used to declare initialized static variables and embed lists of data.
- **The `.rodata` section:** Used specifically for read-only constant data, such as fixed string sequences.
- **The `.bss` section:** Reserved for uninitialized global variables. Variables placed in the `.bss` section are zeroed out at startup and consume very little disk space within the compiled executable file.

Because the Memory Management Unit (MMU) controls access permissions (like read/write vs. read-only) with page-level granularity, the linker ensures that each of these sections begins on a new MMU page boundary.

Data processing in AArch64 centers around a core group of instructions. Data is transferred between registers and memory using `ldr` (load register) and `str` (store register), or between registers using `mov` ,. Basic arithmetic is handled by instructions such as `add` , `sub` , and `mul` ,. Bitwise logical instructions, including

`and`, `orr` (OR), and `eor` (Exclusive-OR), are provided to manipulate specific bits within a word.

## 3.6 Low-Level Control Flow

To emulate the control flow structures found in High-Level Languages (HLLs)—such as `if...then...else` blocks, `switch` statements, and `while` loops—assembly language relies heavily on conditional branches and indirect jumps.

This process begins with the **compare instruction ( `cmp` )**. The `cmp` instruction evaluates two operands by effectively subtracting the second operand from the first. However, rather than saving the mathematical result into a destination register, it uses the result solely to update the processor's status flags. These flags record whether the result was Zero, whether it generated a Carry, whether it was Negative, or whether it caused an Overflow.

Once the flags are updated, a **conditional branch instruction** is used to transfer control depending on the specific states of those flags.

- `beq` (Branch if Equal) and `bne` (Branch if Not Equal) evaluate the Zero flag to jump if the operands were identical or different.
- `blo` (Branch if Lower) and `bhi` (Branch if Higher) are used to branch based on unsigned magnitude comparisons.

To emulate an `if...then...else...endif` statement in assembly language, the programmer uses a `cmp` instruction followed by a conditional branch that evaluates the *opposite* of the high-level condition. If the condition evaluates to false, the CPU branches *over* the `then` block instructions directly into the `else` block. An unconditional branch ( `b` or `bal` ) is placed at the end of the `then` block to jump over the `else` statements so that both blocks are not executed sequentially,.

For `switch...case` statements with many potential paths, assembly programmers often use **jump tables**. Instead of writing a long chain of `cmp` and `beq` instructions, the program calculates an index, loads an execution

address from an array of addresses in memory, and uses an indirect branch (like `br`) to jump instantly to the appropriate case handler.

Finally, for transferring control to a subroutine or function, the CPU uses the **bl (Branch and Link)** instruction, which jumps to the target label while simultaneously saving the return address into the Link Register (LR). When the subroutine finishes its execution, it issues a **ret (Return)** instruction, which unconditionally copies the address held in the LR back into the Program Counter, safely returning control to the caller,.

# Chapter 4: Processor Microarchitecture

## 4.1 Introduction: Architecture vs. Microarchitecture

To a software engineer, the Instruction Set Architecture (ISA) defines the entirety of the processor. The ISA specifies the registers, the memory model, and the instructions available to the programmer. However, the ISA does not dictate how those instructions are physically executed.

The specific physical arrangement of registers, memories, Arithmetic Logic Units (ALUs), and multiplexers used to implement an ISA is known as the **microarchitecture**. A single architecture, such as ARMv8, can be implemented by many different microarchitectures. One microarchitecture might be optimized for ultra-low power consumption in a battery-operated cyber-physical system, while another might be optimized for massive throughput in a cloud server. Though they share the same ISA and run the exact same compiled C code, their internal hardware designs differ vastly.

At the microarchitecture level, a processor is fundamentally a highly complex synchronous sequential digital circuit—essentially a giant finite state machine. The processor's job is to read instructions and data, compute results, and update the architectural state on the rising edge of the system clock.

## 4.2 The Physics of Performance

Before designing our processor, we must define how we measure its performance. The execution time of a program is determined by the following fundamental performance equation:

$$\text{Execution Time} = \text{Instructions} \times \text{CPI} \times T_c$$

1. **Instructions:** The total number of instructions executed by the program. This is determined by the compiler and the ISA.
2. **CPI (Cycles Per Instruction):** The average number of clock cycles required to execute each instruction.
3.  **$T_c$  (Clock Period):** The duration of a single clock cycle (in seconds), which is the inverse of the clock frequency.

The challenge of the microarchitect is to choose a design that minimizes the overall execution time. As we will see, this is a delicate balancing act. Making the hardware simpler might decrease the Clock Period ( $T_c$ ) but increase the CPI. Conversely, adding complex hardware to execute multiple instructions simultaneously will decrease the CPI, but the added logic delay might require a longer Clock Period.

## 4.3 The Single-Cycle Datapath and Control Unit

The simplest way to build a processor is to design a microarchitecture that executes one entire instruction in a single clock cycle. This is known as a single-cycle processor. To build it, we separate the design into two interacting parts: the **datapath**, which contains the structural components that hold and manipulate data, and the **control unit**, which decodes the instruction and commands the datapath.

### 4.3.1 The Fetch-Decode-Execute Cycle

The execution of any instruction on a von Neumann architecture universally follows a sequence of discrete steps known as the fetch-decode-execute cycle:

1. **Fetch:** The processor uses the Program Counter (PC) to fetch the 32-bit instruction from the instruction memory,. Simultaneously, the processor calculates the address of the next instruction (typically  $PC + 4$ ).
2. **Decode:** The processor parses the instruction bits to determine the operation type. Simultaneously, it reads the required source operands



from the register file.

3. **Execute:** The Arithmetic Logic Unit (ALU) performs the requested operation, such as addition, subtraction, or logical AND, on the operands. For load/store instructions, the ALU adds the base register value to a sign-extended immediate offset to calculate the effective memory address.
4. **Memory:** If the instruction is a load or a store, the processor reads from or writes to the data memory at the address computed by the ALU.
5. **Write-back:** The result of the ALU computation or the data read from memory is written back to the destination register in the register file.

### 4.3.2 The Single-Cycle Control Unit

The datapath contains multiple interconnected hardware blocks, and multiplexers are used to route data between them. The control unit acts as the “brain,” driving the select lines of these multiplexers.

The control unit is purely combinational logic. It receives the opcode and function fields of the fetched instruction and instantly generates the necessary control signals. For example, if the instruction is a memory store ( `sw` ), the control unit will assert the `MemWrite` signal to enable writing to the RAM, but it will de-assert `RegWrite` because a store instruction does not update the processor’s registers. If the instruction is a conditional branch ( `beq` ), the control unit commands the ALU to subtract the two source registers. If the ALU outputs a `Zero` flag (indicating the registers are equal), the control unit switches a multiplexer to load the PC with the computed branch target address instead of the standard  $PC + 4$ .

### 4.3.3 The Single-Cycle Timing Bottleneck

In a single-cycle microarchitecture, the architectural state (the PC, the register file, and memory) is updated strictly on the rising edge of the clock. This means that the clock period ( $T_{c\$}$ ) must be long enough to allow electrical signals to propagate through the entire datapath.

The critical path—the longest and slowest sequence of logic—typically occurs during a memory load instruction. The signal must travel through the PC, into the instruction memory, through the register file, through the ALU to calculate the address, into the data memory to fetch the value, through a multiplexer, and finally back to the setup pins of the register file.

Because the clock cycle must be stretched to accommodate the slowest possible instruction, all other, faster instructions (like a simple addition) are artificially delayed. In modern deep sub-micron silicon, this makes the single-cycle processor incredibly inefficient, limiting the maximum clock speed to a few tens of megahertz,.

## 4.4 Pipelining: Increasing Throughput

To overcome the bottleneck of the single-cycle design, computer architects employ **pipelining**. Pipelining divides the instruction execution process into a sequence of smaller, faster stages.

Consider the analogy of an automobile assembly line. If a team of mechanics builds one entire car from start to finish before beginning the next, the factory produces one car per day. However, if the factory is divided into stations—chassis, engine, doors, wheels, paint—multiple cars can be assembled simultaneously. A finished car rolls off the line every hour, even though any individual car still takes a full day to build.

### 4.4.1 The Classic Five-Stage RISC Pipeline

Modern processors slice the datapath into multiple stages separated by hardware registers known as pipeline registers. The classic RISC pipeline consists of five stages,:

1. **Fetch (IF):** Read the instruction from the L1 Instruction Cache.
2. **Decode (ID):** Decode the instruction, resolve register addresses, and read operands from the Register File.
3. **Execute (EX):** The ALU performs the computation or address calculation.

4. **Memory (MEM):** Access the L1 Data Cache for load/store operations.
5. **Write-back (WB):** Write the final result back to the Register File,.

Because the pipeline registers isolate the stages, five different instructions can be in flight simultaneously.

#### 4.4.2 Latency vs. Throughput

It is critical to understand that pipelining *does not* reduce the execution time of an individual instruction; in fact, the addition of pipeline registers introduces a slight propagation delay that actually increases the latency of a single instruction.

However, microprocessors process billions of instructions per second, so we prioritize **throughput** over single-instruction latency. In a perfectly balanced 5-stage pipeline, the clock frequency can be increased to nearly five times that of a single-cycle processor, because the clock period is now determined only by the delay of the single slowest pipeline stage, rather than the sum of all stages. Ideally, one instruction completes and retires every single clock cycle, achieving an ideal CPI of 1.

### 4.5 Pipeline Hazards and Mitigation

In reality, the CPI of a pipelined processor is always greater than 1. This is because instructions in a program are rarely completely independent. When one instruction depends on the result of a previous instruction that is still in the pipeline, a **hazard** occurs.

#### 4.5.1 Data Hazards and Forwarding

A **data hazard** occurs when an instruction attempts to read a register that has not yet been written back to the register file by a preceding instruction. The most common data hazard is the Read-After-Write (RAW) dependency,.

Consider this sequence:

```
add x1, x2, x3    // x1 = x2 + x3
and x4, x1, x5     // x4 = x1 AND x5
```

The `add` instruction calculates the result in the EX stage, but it does not write the value into `x1` until the WB stage. The subsequent `and` instruction attempts to read `x1` during its ID stage. If we allow the pipeline to proceed normally, the `and` instruction will read the stale, outdated value of `x1` from the register file, causing catastrophic computational failure.

To resolve this, modern processors use **forwarding** (or bypassing),. The processor adds a Hazard Detection Unit that monitors the pipeline. If it detects that an instruction in the EX stage needs a value currently held in the MEM or WB pipeline registers, it activates multiplexers that intercept the data and route it directly into the ALU input, bypassing the register file entirely,.

### 4.5.2 Load-Use Stalls (Bubbles)

Forwarding solves most data hazards, but it cannot solve all of them. Consider a **Load-Use hazard**:

```
ldr x1, [x2]      // Load memory at x2 into x1
add x4, x1, x5     // x4 = x1 + x5
```

The `ldr` instruction does not actually retrieve the data from memory until the end of the MEM stage. However, the subsequent `add` instruction requires that data at the beginning of the EX stage. Because the data physically does not exist inside the processor yet, it is impossible to forward it back in time.

In this scenario, the hardware must physically **stall** the pipeline. The Hazard Detection Unit forces the `add` instruction to wait in the Decode stage for an extra clock cycle. To prevent the Execute stage from executing garbage, the control unit zeroes out the control signals, inserting a “bubble” (a hardware No-Operation) into the pipeline. Stalls degrade performance, increasing the CPI above 1,.

### 4.5.3 Control Hazards and Branch Prediction

The most severe threat to pipeline performance is the **control hazard**. When the processor fetches a conditional branch instruction (like `beq`), it does not actually know if the branch will be taken until the ALU evaluates the condition in the Execute stage. But by that time, the processor has already fetched two subsequent instructions into the pipeline.

If the pipeline simply pauses to wait for the branch to resolve, it would suffer massive delays. If it guesses the branch is not taken and fetches sequentially, it will execute the wrong code if the branch is actually taken. When this happens, the processor must **flush** the pipeline, throwing away the mistakenly fetched instructions and converting them into bubbles,. The cycles wasted fetching the wrong instructions are called the branch misprediction penalty.

To mitigate this, processors utilize complex **Branch Predictors**. A dynamic branch predictor utilizes a Branch Target Buffer (BTB) to cache the destination addresses of recent branches, alongside a state machine that tracks the history of the branch. A common design uses a 2-bit finite state machine for each branch. The FSM transitions between “Strongly Not Taken”, “Weakly Not Taken”, “Weakly Taken”, and “Strongly Taken”. This hysteresis ensures that a loop that executes 100 times will only mispredict on the very last iteration when the loop finally exits. Modern branch predictors achieve accuracy rates exceeding 90%.

## 4.6 Advanced Microarchitecture: Superscalar and Out-of-Order Execution

The 5-stage pipeline represents a scalar processor, meaning it issues exactly one instruction per clock cycle. The absolute maximum throughput of a scalar processor is an IPC (Instructions Per Cycle) of 1. To achieve greater performance, engineers exploit **Instruction-Level Parallelism (ILP)**, executing entirely independent instructions simultaneously,.

### 4.6.1 Superscalar Issue

A **superscalar** microarchitecture duplicates the internal datapath hardware to fetch, decode, and issue multiple instructions in a single clock cycle,. A modern high-performance processor (such as the ARM Cortex-A series or Intel Core architectures) might fetch up to four or six instructions per cycle.

To support this, the execution unit is populated with multiple ALUs, Floating-Point Units (FPUs), and load/store units,. Superscalar processors boast an ideal IPC greater than 1 (or a CPI of a fraction of a cycle),.

### 4.6.2 Out-of-Order (OoO) Execution

In a strictly in-order pipeline, if an instruction stalls due to a cache miss, every single instruction behind it in the program is blocked, even if those subsequent instructions have their operands ready and are completely independent,.

**Out-of-Order (OoO)** execution allows the processor to look ahead across a vast “instruction window”. When a micro-operation is decoded, it is dispatched to a queue known as a Reservation Station. The scheduler monitors the reservation stations and dispatches instructions to the ALUs the exact moment their data operands are ready, regardless of their original sequence in the compiled program code,.

Because instructions finish in a chaotic order, the processor utilizes a **Reorder Buffer (ROB)**. Instructions are retired from the ROB strictly in their original program order. This ensures that if a hardware exception or interrupt occurs, the architectural state of the machine remains perfectly precise and easily recoverable,,.

### 4.6.3 Register Renaming

Out-of-order execution introduces new hazards known as Write-After-Write (WAW) and Write-After-Read (WAR) dependencies. These are “false” dependencies caused by the compiler reusing a limited number of architectural registers (like `x0` through `x30` in ARM).

To eliminate these false dependencies, OoO processors implement **Register Renaming**. The processor contains a massive pool of hidden, internal physical registers. When an instruction decodes, a hardware mapping table dynamically assigns the architectural destination register (e.g., `x1`) to an unused physical register. This allows multiple instructions to write to `x1` simultaneously in the pipeline without corrupting each other's data, allowing true ILP to flourish,.

## 4.7 Multithreading and Multi-core Architectures

There is a hard physical limit to how much ILP exists in a single program thread. Even with massively wide superscalar execution and deep OoO windows, processors eventually run out of independent instructions to schedule, leaving expensive functional units completely idle,.

### 4.7.1 Hardware Multithreading

To keep the execution units saturated, microarchitects implemented **hardware multithreading**, allowing a single physical processor core to present the illusion of multiple logical processors to the operating system. The processor duplicates the Program Counter (PC) and the Register File for each thread, but shares the heavy execution hardware (the ALUs and memory buses),.

- **Coarse-grained Multithreading:** The processor runs one thread until it suffers a major stall (like an L2 cache miss), at which point it instantly switches context to a second thread to hide the memory latency,.
- **Fine-grained Multithreading:** The processor switches between threads every single clock cycle in a round-robin fashion, inherently masking data and control hazards.
- **Simultaneous Multithreading (SMT):** Found in almost all modern high-performance cores, SMT allows the superscalar scheduler to issue instructions from multiple independent threads into the execution units during the *exact same clock cycle*,.

As thermal dissipation limits (the “power wall”) prevented further increases in single-core clock frequencies, the industry transitioned to **multi-core**

architectures, placing multiple independent CPU cores onto a single silicon die. Modern SoCs utilize a combination of multicore chips, superscalar pipelines, and SMT to achieve vast computational throughput.

## 4.8 VirtMCU Homework: Analyzing Pipeline Stalls and Deterministic Execution

The advanced microarchitectural optimizations we just discussed—deep pipelining, out-of-order execution, branch prediction, and multi-level caching—are brilliant for increasing average throughput. However, for Cyber-Physical Systems (CPS), these features create a nightmare scenario: **timing unpredictability**.

When controlling an aircraft actuator or a robotic joint, missing a control loop deadline by a few microseconds can cause catastrophic physical failure. In a standard OS running on a modern OoO core, it is virtually impossible to guarantee exact microsecond execution bounds due to the dynamic nature of branch mispredictions and cache misses.

In this homework, you will use the **VirtMCU FirmwareStudio** to analyze instruction latency at the bare-metal microarchitectural level.

1. **Code Profiling:** You will be provided with two versions of a C array-processing loop. The first is an unoptimized sequence containing severe Read-After-Write (RAW) Load-Use data hazards. The second version utilizes loop unrolling and instruction scheduling to place independent math instructions between the `ldr` and data consumption, eliminating the pipeline stall.
2. **Simulation under Jitter:** First, compile the unoptimized binary and run it in a standard, free-running QEMU emulator. Observe how the host operating system's scheduler creates wall-clock jitter, making it difficult to measure the exact nanosecond penalty of the data hazards.
3. **Deterministic `slaved-icount` Execution:** You will then load the binaries into VirtMCU configured in `slaved-icount` mode. Because VirtMCU strictly lock-steps the virtual clock to the exact number of retired



instructions (bypassing host OS jitter), you will use the emulated hardware timer registers to measure the exact cycle count of the execution loops.

4. **Pipeline Analysis:** Calculate the empirical CPI of both the unoptimized and optimized assembly loops. You must submit a report identifying the exact assembly instructions where the processor was forced to insert a pipeline bubble due to a Load-Use data hazard, verifying your analysis against the cycle counts retrieved from the VirtMCU deterministic timer.

# Chapter 5: The Memory Hierarchy

## 5.1 Introduction: The Memory Wall

In the preceding chapter, we examined the internal microarchitecture of the processor, demonstrating how pipelining, superscalar execution, and out-of-order scheduling can dramatically increase the instruction execution rate. However, a processor can only execute instructions as fast as it can be supplied with code and data. Historically, microprocessor speeds have doubled every two years, whereas the access latency of main memory has decreased at only a fraction of that rate. This ever-expanding divergence between CPU execution speed and memory latency is known as the “memory wall”.

If a modern processor running at 3 GHz (where one clock cycle takes 0.33 nanoseconds) were forced to access main memory directly for every instruction, it would have to wait hundreds of clock cycles for the data to arrive. The execution pipeline would starve, and the massive throughput gains of the microarchitecture would be entirely wasted.

To overcome this bottleneck, computer architects do not rely on a single, monolithic pool of memory. Instead, they construct a **memory hierarchy**. The hierarchy is organized as a pyramid: at the very top, residing inside the CPU core, are the processor registers, offering instant access. Below the registers is the **cache memory**, which is small but extremely fast, operating at speeds comparable to the processor. Below the cache is the main memory, which is much larger but significantly slower. At the base of the pyramid lies secondary storage (magnetic hard disk drives or solid-state drives), offering vast capacity but agonizingly slow access times measured in milliseconds. By intelligently migrating data up and down this hierarchy, the memory system provides the software engineer with an illusion: a memory pool that possesses the speed of the fastest cache and the immense capacity of the largest disk.

## **5.2 Primary Storage Technologies: SRAM and DRAM**

To understand the trade-offs in the memory hierarchy, we must first look at the physical semiconductor technologies used to implement it. There are two primary types of volatile memory used in computer architectures: Static Random Access Memory (SRAM) and Dynamic Random Access Memory (DRAM).

### **5.2.1 Static RAM (SRAM)**

SRAM is the fastest memory technology available and is utilized primarily for processor caches. A standard SRAM bit cell is constructed from six MOSFET transistors arranged as two cross-coupled inverters (forming a flip-flop) plus two access transistors,. Because the data is actively driven by a continuous feedback loop between the inverters, SRAM retains its state as long as power is applied without any need for external intervention,.

The flip-flop design provides incredibly low latency, allowing read and write operations to keep pace with high-frequency CPU clocks. Furthermore, reading an SRAM cell is non-destructive. The major downside to SRAM is its size and cost. Because it requires six transistors per bit, an SRAM array consumes a massive amount of silicon die area,. Additionally, the transistors in the cross-coupled inverters continuously leak a small amount of current, resulting in static power dissipation that grows severely as transistor sizes shrink.

### **5.2.2 Dynamic RAM (DRAM)**

Main memory is almost universally implemented using DRAM. A DRAM bit cell is astonishingly simple: it consists of exactly one transistor and one microscopic capacitor,. A logical 1 or 0 is represented by the presence or absence of an electrical charge on the capacitor. Because it only uses one transistor per bit, DRAM achieves enormous data density, making it vastly cheaper per megabyte than SRAM.

However, this density comes with severe performance penalties. The capacitor naturally leaks charge over time, meaning the memory will “forget” its data within milliseconds,. To prevent data loss, the memory controller must continually read and rewrite the contents of the entire memory, a process known as **refreshing**,. Furthermore, reading a DRAM cell requires the bit-line to share charge with the microscopic capacitor, a delicate analog process that takes significant time and actually destroys the charge in the capacitor,. As a result, every DRAM read must be immediately followed by a hardware rewrite to restore the data. To manage this, DRAM chips use a multiplexed addressing scheme where a row (page) is activated first, followed by column accesses. This multi-step process makes DRAM latency hundreds of times slower than SRAM,.

## 5.3 Cache Memory Fundamentals

Because SRAM is too large and expensive to serve as main memory, architects insert small SRAM caches between the CPU and the DRAM,. The effectiveness of a cache relies entirely on the mathematical probability that a required piece of data is already present in the SRAM. This probability is driven by two empirical principles of software behavior known as **locality of reference**:

1. **Temporal Locality**: If a memory location is accessed, it is highly likely to be accessed again in the near future,,. A software engineer iterating a counter variable inside a `for` loop is demonstrating temporal locality. The cache exploits this by keeping a copy of recently accessed data.
2. **Spatial Locality**: If a memory location is accessed, memory locations with nearby addresses are likely to be accessed soon,,. A software engineer traversing an array of integers demonstrates spatial locality. The cache exploits this by never fetching just a single byte from main memory; instead, it fetches an entire contiguous **cache line** (or cache block), typically 32 or 64 bytes wide,.

When the processor requests data, it searches the cache. If the data is found, a **cache hit** occurs, and the data is served at SRAM speeds,. If the data is absent, a **cache miss** occurs, and the processor pipeline must stall while the hardware fetches the entire cache line from the slower main memory,.

The overall efficiency of the memory hierarchy is quantified by the **Average Memory Access Time (AMAT)**:  $AMAT = Hit\_Time + (Miss\_Rate \times Miss\_Penalty)$  To reduce AMAT, modern SoCs utilize multiple levels of cache. The **Level 1 (L1)** cache is tightly coupled to the processor pipeline, incredibly fast (often 1-3 cycles), but very small (e.g., 32 KB). It is usually split into an L1 Instruction Cache (I-cache) and L1 Data Cache (D-cache) to prevent structural hazards and enable simultaneous instruction fetches and data loads—a hallmark of the Harvard architecture. Below the L1 is the **Level 2 (L2)** cache, which is larger (e.g., 256 KB to 1 MB), somewhat slower, and usually unified (holding both instructions and data). Many high-performance systems add a massive **Level 3 (L3)** cache shared among multiple processor cores.

## 5.4 Cache Organization and Placement Policies

When a cache fetches a 64-byte block from main memory, it must decide exactly where in its internal SRAM array to place it. Because the main memory is gigabytes in size and the cache is only kilobytes, countless memory blocks must map to the same cache slots. The hardware rules governing this mapping dictate the cache's placement policy.

### 5.4.1 Direct-Mapped Caches

A direct-mapped cache is the simplest organization. The cache is divided into a one-dimensional array of "sets," where each set can hold exactly one cache line. A physical memory address is mapped to a specific set using a hardwired modulo operation.

To look up data, the cache controller splits the memory address into three fields:

- **Byte Offset:** The lowest bits specify which exact byte within the 64-byte cache line is requested.
- **Index (Set):** The middle bits determine which specific set in the cache array to check.

- **Tag:** The uppermost bits identify the exact main memory block currently residing in that set,.

Because each memory address maps to exactly one set, finding data is fast; the cache checks the set's Tag against the address's Tag. If they match, and the "Valid" bit is set, it is a hit,.

However, direct-mapped caches suffer from severe **conflict misses**. If a program actively bounces between two different memory addresses that happen to map to the exact same cache set (e.g., array elements spaced exactly by a multiple of the cache size), the cache will continually evict and reload the same set,. This pathological condition, known as thrashing, destroys performance.

### 5.4.2 Fully Associative Caches

To eliminate conflict misses entirely, a cache can be made **fully associative**. In this design, a memory block can be placed in *any* available slot in the cache. The address is only split into a Tag and a Byte Offset; there is no Index.

While this prevents arbitrary collisions, looking up data becomes a hardware nightmare. The cache controller must compare the requested Tag against the Tags of *every single line in the entire cache simultaneously*. This requires massive amounts of comparator logic (AND-gates and OR-gates), which consumes excessive die area and power, making it entirely impractical for large L1 or L2 caches,,.

### 5.4.3 N-Way Set-Associative Caches

Modern processor caches strike a compromise using **set-associative** mapping. The cache is divided into sets, but each set contains  $N$  separate slots (or "ways"),. For example, in an 8-way set-associative cache, a memory address still hashes to one specific set, but the cache line can be placed into any of the 8 available ways within that set.

During a lookup, the hardware reads all  $N$  tags in the indexed set in parallel and compares them to the requested address,. This effectively mitigates the “birthday paradox” of cache collisions without requiring the massive hardware overhead of a fully associative search. Most L1 caches today are 4-way or 8-way set-associative, offering an optimal balance between hit rate, hardware complexity, and access latency.

## 5.5 Cache Replacement and Write Policies

When a set-associative cache experiences a miss and all  $N$  ways in the target set are currently full, the cache controller must choose one existing line to evict.

The gold standard for eviction is the **Least Recently Used (LRU)** policy,. LRU relies on the principle of temporal locality, assuming that the cache line which has gone the longest without being accessed is the safest to discard,. True LRU is easy to implement for 2-way caches but becomes exponentially complex for 8-way or 16-way caches. Therefore, modern processors often use **Pseudo-LRU** (using a tree of binary flags to approximate age) or a simple Not-Recently-Used (NRU) clock algorithm,.

Handling memory stores (writes) introduces further complexity. If the CPU updates a variable, it modifies the copy in the cache. The cache and main memory are now out of sync. There are two primary write policies to resolve this:

- **Write-Through:** Every write to the cache is simultaneously propagated out to the main memory,. This ensures memory is always perfectly up-to-date, making recovery easy. However, it severely wastes memory bus bandwidth. If a program increments a counter 1,000 times, write-through initiates 1,000 slow DRAM transactions.
- **Write-Back:** Writes only update the cache,. A special metadata flag called the **dirty bit** is flipped to indicate the cache line contains modified data,. The line is only written back to DRAM when it is eventually evicted. This technique, called write coalescing, drastically improves performance by absorbing thousands of intermediate writes,.

Furthermore, if a CPU attempts to write to an address that is *not* currently in the cache, it must decide whether to fetch the block first. Modern processors usually employ a **write-allocate** policy, dragging the missing cache line from DRAM into the cache before modifying it, assuming the software will likely write to adjacent variables in the same block shortly.

## 5.6 Cache Coherence in Multicore Systems

As discussed in Chapter 4, to increase throughput, modern architectures place multiple processor cores on a single silicon die (Symmetric Multiprocessing or SMP). Each core possesses its own private L1 cache. If we utilize a high-performance write-back policy, we inevitably trigger the **cache coherence problem**.

Suppose Core A and Core B both load a shared memory variable  $X$  (currently 0) into their respective L1 caches. Core A executes  $x = 5$ . Because Core A uses a write-back cache, the new value 5 is stored only in Core A's L1 cache; main memory remains 0, and Core B's cache remains 0. If Core B subsequently reads  $X$ , it will see stale, invalid data. This violates the sequential consistency expected by software engineers writing multithreaded applications.

To prevent this, the hardware implements a **cache coherence protocol**. In smaller SMP systems, this is achieved via **snooping**. All cache controllers constantly monitor (snoop) the shared memory bus. When a core wishes to write to a shared variable, it broadcasts a signal.

The baseline snooping protocol is **MSI**, which assigns one of three states to every cache line:

1. **Modified (M)**: The line is dirty. This core owns the only valid copy in the system.
2. **Shared (S)**: The line is clean. Multiple cores may have copies. The core can read it, but cannot write to it without asking permission.
3. **Invalid (I)**: The line contains garbage.



If a line is Shared, and Core A wants to write to it, Core A must broadcast an “Invalidate” message on the bus (BusRdX). Core B hears this snoop message and downgrades its copy of  $XX$  to Invalid. Core A then transitions to Modified and performs the write. If Core B later wants to read  $XX$ , it broadcasts a read request. Core A snoops the request, pauses, flushes its Modified line back to DRAM, and then both cores transition back to Shared.

Modern processors (like the Core i7) use the extended **MESI** protocol, which adds an **Exclusive (E)** state. Exclusive indicates the line is clean, but this core is the *only* one holding it. If the core decides to write, it can transition from E to M silently, without broadcasting an expensive invalidate message on the bus.

Advanced ARM processors use the **MOESI** protocol, which adds an **Owned (O)** state,. The Owned state solves a major inefficiency: if Core A has a Modified line and Core B wants to read it, standard MESI forces Core A to flush the line to slow DRAM before Core B can read it. Under MOESI, Core A simply transmits the dirty data directly to Core B over the fast internal bus. Core A becomes the “Owner” (responsible for eventually writing it back to DRAM), and Core B enters the Shared state. This peer-to-peer cache transfer entirely bypasses the main memory bottleneck.

For enormous systems with hundreds of cores where broadcast snooping would saturate the bus, architects use **directory-based coherence**,. A centralized directory tracks exactly which cores hold which cache lines, allowing point-to-point invalidation messages rather than global broadcasts,.

## 5.7 Virtual Memory

In the early days of computing, software ran directly on physical memory. If an application needed 64 MB of RAM but the machine only had 32 MB, the programmer had to manually chop the program into “overlays,” shuttling segments of code back and forth from the disk—a tremendously tedious process.

To solve this, architects introduced **Virtual Memory**. The operating system and hardware conspire to give each running process the illusion of an immense,

private, contiguous address space (e.g., 4 GB on a 32-bit system or 256 TB on a 64-bit system), regardless of how much physical RAM is actually installed,.

This illusion is sustained via **paging**. The vast virtual address space is sliced into fixed-size blocks called **virtual pages** (typically 4 KB),. The physical RAM is similarly sliced into 4 KB chunks called **page frames**. At any given time, only the actively used virtual pages (the “working set”) reside in physical page frames. The rest are parked on the incredibly slow secondary storage (the hard disk or SSD), in an area known as the swap file.

If the CPU attempts to read a virtual address whose corresponding page is currently on disk, the hardware triggers an exception called a **page fault**,. The operating system traps the fault, pauses the program, evicts an old physical page frame (often using the Clock or LRU algorithm), fetches the requested 4 KB page from the hard drive, and then restarts the instruction,. From the software engineer’s perspective, the memory read simply took a very long time; the underlying mechanics are entirely transparent.

Beyond giving the illusion of infinite memory, paging provides strict **memory protection**. Because each process has its own isolated mapping, Process A cannot physically formulate an address that maps to Process B’s page frames. The MMU simply forbids it, throwing an access violation exception if attempted,.

## 5.8 Address Translation and the MMU

The translation of a programmer’s virtual address into a hardware physical address is performed by the **Memory Management Unit (MMU)**.

For a 32-bit system using 4 KB ( $2^{12}$  bytes) pages, a virtual address is mathematically divided into two parts:

- **Page Offset:** The lowest 12 bits. This defines the exact byte within the 4 KB page. Because the page size and frame size are identical, the offset requires absolutely no translation. It passes directly to the physical address,.

- **Virtual Page Number (VPN):** The upper 20 bits. The MMU uses the VPN as an index into a **Page Table** to find the corresponding Physical Page Number (PPN),.

The Page Table is a data structure maintained by the Operating System residing in physical memory. Each Page Table Entry (PTE) contains the mapped Physical Page Number, alongside metadata bits: a Valid bit (is it in RAM or on disk?), a Dirty bit (has it been modified?), and protection bits (Read/Write/Execute permissions),.

A flat page table mapping 4 GB of space requires  $2^{20}$  (over 1 million) entries per process. To avoid wasting megabytes of contiguous RAM just for mapping tables, modern architectures use **multilevel page tables**,. The 20-bit VPN is split again (e.g., into a 10-bit Page Directory Index and a 10-bit Page Table Index),. The CPU first looks up the Page Directory, which points to a secondary Page Table, which finally points to the physical frame. This allows sparse allocation, where page tables for unused memory regions do not need to be instantiated,.

## 5.9 The Translation Lookaside Buffer (TLB)

Hierarchical page tables solve the memory footprint issue but introduce a catastrophic performance problem. Because the page tables are stored in main memory, translating a single virtual address requires two separate physical memory reads (one for the page directory, one for the page table) *before* the CPU can even attempt to read the actual data,. This would triple the execution time of every instruction.

To bypass this penalty, the MMU relies on a specialized cache called the **Translation Lookaside Buffer (TLB)**,,. The TLB is an ultra-fast, fully associative cache placed directly inside the MMU that stores the most recently used Virtual-to-Physical translations,.

When the CPU generates a virtual address, the MMU first blasts the Virtual Page Number into the TLB,. If a TLB Hit occurs (which happens >99% of the time due to spatial and temporal locality), the MMU instantly retrieves the Physical Page

Number and synthesizes the physical address in a single clock cycle, completely skipping the page table walk,.

If a **TLB Miss** occurs, the hardware (or a microcode software trap, depending on the architecture) must perform the expensive “page walk” through the hierarchical tables in DRAM, locate the translation, and load it into the TLB, evicting an older entry,. Because a TLB is fully associative to maximize hit rates and minimize conflict misses, it is hardware-intensive and usually quite small (holding only 64 to 512 entries),.

For a software engineer, understanding the TLB is critical for high-performance code. If a program randomly accesses memory across thousands of distinct 4 KB pages (e.g., traversing a poorly constructed linked list spread across a large heap), it will cause “TLB thrashing.” The TLB will constantly miss, forcing the hardware to perform page table walks, reducing the CPU’s throughput to a crawl even if the data itself resides comfortably in the L2 cache.

## 5.10 VirtMCU Homework: Memory Access Profiling

In standard software development on a host operating system, observing the exact microarchitectural penalties of cache misses and TLB thrashing is incredibly difficult. The host OS scheduler interrupts your process randomly, and background tasks continually pollute the L1/L2 caches and TLB, injecting massive amounts of wall-clock jitter into any performance measurement.

For this chapter’s homework, you will bypass these abstractions using **VirtMCU FirmwareStudio**. You will analyze memory hierarchy latency utilizing VirtMCU’s deterministic `slaved-icount` mode, which guarantees exact, jitter-free cycle counts.

1. **Cache Thrashing:** You will be provided with a C program that performs matrix multiplication. You will write two versions of the inner loop: one that iterates through the matrices in row-major order (respecting spatial locality) and one that iterates in column-major order (violating spatial locality).

2. **TLB Miss Measurement:** You will write a second program that initializes a large array of structs. One version of your code will traverse the structs sequentially. The second version will traverse the structs using a randomized stride that deliberately jumps across 4 KB page boundaries on every access, intentionally exceeding the TLB's 64-entry capacity.
3. **Deterministic Cycle Profiling:** You will compile these binaries for bare-metal ARM execution and load them into VirtMCU. By reading the virtual hardware timers in `slaved-icount` mode before and after your execution loops, you will measure the exact nanosecond penalty introduced by cache line fetches and page-table walks. You will document the performance divergence, proving empirically how software data structures interact with physical cache associativity and MMU translation buffers.

# Chapter 6: SoC Interconnects and Peripherals

## 6.1 Introduction: Breaking Out of the Core

Welcome back. In our previous chapters, we constructed the processor pipeline and designed a deep memory hierarchy to keep the execution units fed with instructions and data. However, a processor that only talks to its own memory is effectively a brain in a jar; it can compute, but it cannot interact with the universe. For a Cyber-Physical System (CPS), interaction is the entire point. The system must sense physical properties, process them, and actuate mechanical responses.

To bridge this gap, modern microprocessors are integrated with numerous communication interfaces and hardware controllers onto a single silicon die, creating a System-on-Chip (SoC). A SoC is not merely a collection of isolated parts; it is a complex, distributed network of functional blocks. These blocks, ranging from low-speed sensors to high-speed graphics accelerators, must all communicate flawlessly.

In this chapter, we explore how data physically moves across the silicon. We will abandon the simplified abstraction of a monolithic “system bus” and dive into the mechanics of modern SoC interconnects. We will examine the Advanced Microcontroller Bus Architecture (AMBA), detailing the APB, AHB-Lite, and AXI protocols. Next, we will transition to the hardware-software interface, explaining how your C code interacts with these hardware peripherals through Memory-Mapped I/O (MMIO), interrupts, and Direct Memory Access (DMA). Finally, we will look at the Universal Asynchronous Receiver-Transmitter (UART) as a fundamental case study for external serial communication.

## 6.2 The Principles of SoC Interconnects

Historically, computer components were distributed across multiple chips on a printed circuit board and connected via shared, parallel copper wires—a traditional “bus”. In a classic shared bus, multiple devices physically share the same wires. To prevent electrical short circuits (a “bus fight”), devices use tri-state logic gates to disconnect themselves from the wires when not transmitting. However, shared buses suffer from severe physical and performance limitations. Only one device can transmit at a time, meaning bandwidth is inherently limited. Furthermore, as clock frequencies increase, the capacitance and physical length of shared wires cause signal degradation and propagation delays that cap the maximum speed.

In a modern SoC, the “bus” is no longer a single set of shared wires; it is an active **interconnect fabric**. The interconnect fabric multiplexes and routes transactions between components using dedicated, unidirectional point-to-point wiring.

### 6.2.1 Initiators and Targets

Communication across an interconnect always involves two roles:

- **Initiator (or Manager / Requester):** The device that starts a transaction by supplying an address and command. The CPU core and the DMA controller are the most common initiators.
- **Target (or Subordinate / Completer):** The device that responds to the transaction. Memory controllers, UART peripherals, and timers are common targets.

The interconnect fabric sits between the initiators and the targets. It intercepts all transactions issued by the initiators and routes them to the appropriate target based on the address.

## 6.2.2 Address Decoding and Routing

Every target connected to the fabric is assigned a specific range of addresses, defining the system's **memory map**. When an initiator issues a transaction, it drives the target address onto the interconnect.

The fabric contains an internal **address decoder**. The decoder compares the high-order bits of the requested address against a routing table to determine which target the initiator is trying to reach. For example, if a SoC places a UART controller at address `0x4000_0000`, any read or write starting with `0x40` will be routed specifically to the UART's hardware port.

Once the target is identified, the interconnect fabric uses multiplexers to physically connect the initiator's data and control lines to the target's receiving lines. To avoid complex hardware subtractors in the decoding logic, SoC memory maps generally place peripherals at boundaries that are multiples of powers of two. This allows the interconnect to route traffic simply by looking at a specific prefix of the address bits, functioning much like an IP subnet mask in computer networking.

## 6.2.3 Backpressure and Flow Control

In a heterogeneous SoC, an initiator running at 1 GHz may attempt to send data to a peripheral running at 10 MHz. If the initiator sends data faster than the target can process it, data will be lost. To prevent this, interconnect protocols implement **backpressure**.

Backpressure allows a slower component to stall a faster component. This is universally handled using a "ready-valid" handshake mechanism. The initiator asserts a `VALID` signal when its data and address are stable. The target asserts a `READY` signal when it is capable of accepting the data. The transaction only completes on the exact clock edge where both `VALID` and `READY` are asserted simultaneously. If the target is busy, it deasserts `READY`, forcing the initiator to hold its data on the lines and wait, effectively stalling the initiator's pipeline.



## 6.3 The AMBA Interconnect Protocols

To allow intellectual property (IP) blocks from different vendors (e.g., an ARM CPU, a Synopsys memory controller, and a custom accelerator) to communicate, the semiconductor industry relies on standardized interconnect protocols. The most ubiquitous standard is ARM's **Advanced Microcontroller Bus Architecture (AMBA)**. AMBA is not a single bus; it is a family of protocols designed for different performance, power, and area trade-offs. We will examine three critical AMBA protocols: APB, AHB-Lite, and AXI.

### 6.3.1 APB: The Advanced Peripheral Bus

Not all peripherals require high-bandwidth communication. Devices like timers, GPIO controllers, and UARTs have very low data throughput requirements. Providing these slow devices with complex, high-speed interconnect logic would waste precious silicon area and increase static power consumption.

The **Advanced Peripheral Bus (APB)** is designed specifically for these low-bandwidth control interfaces. APB is a non-pipelined, state-machine-based protocol. Every transaction takes a minimum of two clock cycles to complete.

The APB protocol utilizes a strict three-state machine: **IDLE**, **SETUP**, and **ACCESS**:

1. **IDLE**: The bus is inactive.
2. **SETUP**: When a transaction begins, the requester asserts a select signal ( **PSEL** ) for the target peripheral, along with the address ( **PADDR** ) and a read/write control signal ( **PWRITE** ). The bus remains in this state for exactly one clock cycle, giving the target peripheral time to wake up and prepare its internal logic.
3. **ACCESS**: On the next clock edge, the requester asserts the enable signal ( **PENABLE** ). If it is a write, the data is driven onto **PWDATA** . The target evaluates the request. If the target needs more time, it holds its **PREADY** signal low, extending the **ACCESS** state and applying backpressure. Once the target is finished, it drives **PREADY** high (and supplies read data on

`PRDATA` if applicable). The transaction completes, and the bus returns to `IDLE` or begins a new `SETUP` phase.

By guaranteeing that signals remain completely stable during the `SETUP` to `ACCESS` transition, APB eliminates unwanted circuit switching (glitches), significantly reducing dynamic power consumption.

### 6.3.2 AHB-Lite: The Advanced High-performance Bus

While APB is excellent for slow peripherals, it is far too slow for main memory or high-speed DMA transfers. For high-bandwidth components, designers use the **Advanced High-performance Bus (AHB-Lite)**.

AHB-Lite achieves high throughput by abandoning the slow state-machine approach of APB in favor of a strictly **pipelined** architecture. In AHB-Lite, every transaction is split into two distinct phases: an **Address Phase** and a **Data Phase**.

- **Address Phase:** Lasts exactly one clock cycle. The manager drives the address ( `HADDR` ), the transfer direction ( `HWRITE` ), and the size of the transfer ( `HSIZE` ).
- **Data Phase:** Begins immediately after the Address Phase. The manager drives write data ( `HWDATA` ), or the subordinate provides read data ( `HRDATA` ).

Because of the pipelining, the *Data Phase of the current transaction overlaps exactly with the Address Phase of the next transaction*. This means that while the memory is fetching data for address `$A$`, the manager is already transmitting address `$A+4$` on the bus. This overlapping allows AHB-Lite to achieve 100% bus utilization during sequential reads or writes, effectively transferring one word per clock cycle.

To apply backpressure, an AHB-Lite subordinate pulls the `HREADY` signal low during the Data Phase. Because the pipeline is locked together, stalling the Data Phase of the current transaction inherently forces the Address Phase of the subsequent transaction to stall as well.

AHB-Lite also supports **Bursts** ( `HBURST` ). A burst allows a manager to inform the subordinate that it plans to perform a sequence of 4, 8, or 16 consecutive transfers. This allows advanced memory controllers, like SDRAM, to optimize their internal row-buffer accesses before the subsequent addresses are even placed on the bus.

### 6.3.3 AXI: The Advanced eXtensible Interface

As processor speeds pushed into the gigahertz range and multi-core SoCs emerged, even the pipelined AHB-Lite protocol became a bottleneck. In AHB-Lite, the Address and Data phases are rigidly locked together; if a memory read stalls, the entire bus pipeline is blocked, preventing any other transactions from occurring. Furthermore, AHB-Lite uses a single set of shared control signals for both reading and writing, meaning a system cannot read and write simultaneously (half-duplex).

To solve this, ARM introduced the **Advanced eXtensible Interface (AXI)**. AXI is a radical departure from traditional bus architectures. Instead of a single monolithic bus, AXI defines **five completely independent, unidirectional channels** between a manager and a subordinate:

1. **Read Address Channel (AR):** Carries the address and control info for a read.
2. **Read Data Channel (R):** Carries the requested data from the subordinate back to the manager.
3. **Write Address Channel (AW):** Carries the address and control info for a write.
4. **Write Data Channel (W):** Carries the payload data from the manager to the subordinate.
5. **Write Response Channel (B):** Carries a success/failure acknowledgment from the subordinate back to the manager.

Each of these five channels operates completely independently, using its own `VALID` and `READY` handshake signals. This physical decoupling yields immense performance benefits.

First, it enables **Full-Duplex** operation; a DMA controller can write data on the W channel at the exact same time it is receiving completely unrelated read data on the R channel, doubling peak throughput.

Second, AXI supports **Outstanding Transactions** and **Out-of-Order Completion**. A CPU can blast five different read addresses down the AR channel without waiting for the data to return. The memory controller receives all five addresses and can reorder the reads based on which physical DRAM banks are currently open. When the memory controller finds the data, it sends it back on the R channel. Because the data might return out of order, every transaction is tagged with a unique Transaction ID ( `ARID` / `RID` ). The CPU matches the incoming data to the original request using this ID, much like a restaurant patron matching their receipt number to an order.

Because AXI is complex, ARM also provides **AXI-Lite**, a stripped-down version that removes burst capabilities and IDs. AXI-Lite provides AXI's channel decoupling but is simple enough to interface with standard registers.

## 6.4 The Hardware-Software Interface

Having established how data traverses the silicon via interconnects like AXI and APB, we must examine how the software engineer actually commands this hardware. In high-level languages like C, we manipulate variables in memory. But how do we tell a physical UART peripheral to transmit a byte?

### 6.4.1 Memory-Mapped I/O (MMIO) vs. Port-Mapped I/O

There are two architectural paradigms for interacting with peripherals: Port-Mapped I/O and Memory-Mapped I/O.

In older architectures like the Intel x86, peripheral devices exist in a completely separate, isolated address space. To access them, the CPU provides special, dedicated assembly instructions like `IN` and `OUT`. When the CPU executes an `OUT` instruction, it asserts a specific physical pin on the processor (the `M/IO`

pin) to tell the hardware, “I am talking to a peripheral, not main memory”. This is known as **Port-Mapped I/O**.

The ARM architecture (and almost all modern RISC designs) abandons this approach in favor of **Memory-Mapped I/O (MMIO)**. In MMIO, there are no special I/O instructions. The peripheral’s control and data registers are mapped directly into the standard, unified physical memory address space. From the CPU’s perspective, a peripheral register is indistinguishable from a standard RAM location. To send data to a peripheral, the software simply executes a standard assembly `STR` (Store Register) instruction targeting the specific memory address assigned to that device.

This allows the C programmer to interact with hardware simply by casting an integer address to a volatile pointer:

```
#define UART_TX_REG (*(volatile uint32_t*) 0x40001000)
UART_TX_REG = 'A'; // Hardware automatically routes this to the UART
```

The `volatile` keyword is critical here. It explicitly commands the C compiler *never* to optimize away the memory access or cache the value in a CPU register, guaranteeing that every read or write genuinely reaches the hardware. Furthermore, MMIO regions are explicitly marked as “Device Memory” in the MMU page tables, bypassing the L1/L2 caches to ensure immediate, strongly-ordered electrical delivery to the peripheral.

### 6.4.2 Programmed I/O (Polling)

The simplest method to manage a peripheral is **Programmed I/O** (often called polling). In this approach, the CPU executes a continuous loop, reading the peripheral’s status register until a specific bit indicates the device is ready.

For example, to receive a byte from a UART, the CPU repeatedly reads the UART’s line status register checking the `Data Ready` flag. Once the flag becomes `1`, the CPU reads the data register to retrieve the byte. While trivial to program, polling is highly inefficient. The CPU wastes thousands or millions of clock cycles spinning in a tight loop, unable to perform useful computation or

enter low-power sleep states. In modern, multitasking operating systems, polling is strictly avoided for routine operations.

### 6.4.3 Interrupt-Driven I/O

To eliminate the massive waste of CPU cycles caused by polling, systems rely on **Interrupt-Driven I/O**. Instead of the CPU constantly asking the peripheral if it is ready, the peripheral notifies the CPU when an event occurs.

Peripherals are wired directly to the CPU's **Nested Vector Interrupt Controller (NVIC)**. When a UART receives a byte, its hardware asserts an interrupt request (IRQ) line. The NVIC detects this, immediately pauses the CPU's current execution thread, pushes the CPU state to the stack, and forces the Program Counter to jump to an **Interrupt Service Routine (ISR)**. The ISR reads the byte from the UART's data register, places it into an operating system queue, clears the hardware interrupt flag, and returns execution to the original thread.

Interrupt-driven I/O allows the CPU to execute completely unrelated applications, or drop into deep sleep modes, waking only for a few microseconds when a physical event requires attention.

### 6.4.4 Direct Memory Access (DMA)

While interrupts solve the problem of polling, they are still too slow for massive data transfers. If a gigabit Ethernet controller receives a 1,500-byte packet, generating 1,500 individual CPU interrupts to copy the data byte-by-byte into RAM would overwhelm the processor, a problem known as "interrupt storming".

To handle bulk data, SoCs utilize a **Direct Memory Access (DMA) Controller**. A DMA controller is an independent hardware engine that sits on the interconnect (typically as an AXI Manager) capable of generating its own read and write transactions.

When the CPU needs to transfer data, it configures the DMA controller's MMIO registers, providing a source address, a destination address, and a block size.

The CPU then goes back to its normal tasks. The DMA controller autonomously arbitrates for interconnect bandwidth and transfers the data directly between the peripheral and main memory, completely bypassing the CPU execution pipeline. Because the DMA controller shares the memory interconnect, it performs **cycle stealing**, slipping its transactions onto the bus during cycles when the CPU is idle or accessing L1 cache. Once the entire multi-kilobyte block is successfully transferred, the DMA controller fires a single interrupt to notify the CPU that the memory buffer is ready for software processing.

## 6.5 Core SoC Peripherals

With an understanding of how they connect to the system, we can briefly examine the internal architecture of the most ubiquitous CPS peripherals.

### 6.5.1 General-Purpose Input/Output (GPIO)

The most fundamental interface is the **GPIO** controller. GPIOs allow software to directly manipulate the physical voltage levels on the external pins of the SoC chip. A GPIO block contains several MMIO registers:

- **Direction Register:** Sets each pin as an input (read sensor) or output (drive actuator).
- **Data Register:** If configured as an output, writing a 1 or 0 physically drives the pin to  $V_{DD}$  or GND. If configured as an input, reading this register samples the current digital voltage on the pin.
- **Pull-up/Pull-down Configuration:** Activates internal resistors to prevent “floating” inputs that could cause CMOS shoot-through current.

### 6.5.2 The Universal Asynchronous Receiver-Transmitter (UART)

While GPIO is useful for simple logic, exchanging complex data with external microcontrollers or diagnostic terminals requires serial communication. The

**UART** is a ubiquitous, legacy protocol that transmits data sequentially over a single wire without a shared clock signal.

Because there is no clock wire to synchronize the sender and receiver, a UART relies on strict timing agreements. Both ends must be configured to the exact same **baud rate** (symbols per second).

A standard UART transmission frame operates as follows:

1. **Idle State:** The line is held at logic HIGH (1).
2. **Start Bit:** The transmitter pulls the line LOW (0) for exactly one baud period. This alerts the receiver that a frame is beginning, allowing it to synchronize its internal clock timers.
3. **Data Bits:** The transmitter sends 5 to 8 bits of payload data, usually Least Significant Bit (LSB) first.
4. **Parity Bit (Optional):** A hardware-generated bit used for basic error detection (e.g., ensuring the total number of 1s is even).
5. **Stop Bit:** The line is driven HIGH (1) to return to the idle state and guarantee a clear transition for the next start bit.

Because modern CPUs run billions of times faster than standard UART baud rates (e.g., 115,200 bps), UART IP blocks incorporate hardware **FIFO (First-In, First-Out) buffers**. When the CPU wishes to print a string, it blasts the entire string into the UART's TX FIFO using MMIO in just a few nanoseconds. The hardware state machine then sequentially shifts the bits out over the physical wire at the slow baud rate, eventually raising an interrupt when the FIFO is empty and ready for more data.

### 6.5.3 Synchronous Serial: SPI and I2C

While UART is asynchronous, faster peripheral communication is achieved using synchronous protocols that include a dedicated clock wire.

- **SPI (Serial Peripheral Interface):** A full-duplex protocol utilizing four wires: Master Out Slave In (MOSI), Master In Slave Out (MISO), Serial Clock (SCK), and Chip Select (CS). Because the master explicitly drives the clock,



SPI can easily run at tens of megahertz, making it ideal for interfacing with external flash memory or LCD screens.

- **I2C (Inter-Integrated Circuit):** A half-duplex, multi-master protocol utilizing only two wires: Serial Data (SDA) and Serial Clock (SCL). I2C uses open-drain signaling with pull-up resistors and requires devices to broadcast target addresses over the bus before data transfer. It is slower than SPI but drastically reduces pin count, making it standard for reading low-speed sensor arrays (e.g., accelerometers, temperature probes).

## 6.6 VirtMCU Homework: Building a Deterministic UART Driver

In traditional embedded systems education, writing a bare-metal UART driver introduces students to the harsh realities of physical timing. If your C code relies on polling, and the host operating system scheduling your emulator preempts your thread for a few milliseconds, the simulated UART peripheral will drop characters, leading to corrupted strings and immense frustration. Standard virtual machines fail to replicate the strict, continuous-time relationships required for asynchronous serial communication.

To resolve this, this chapter's homework relies on the **VirtMCU** framework. Because VirtMCU utilizes a dynamic QOM plugin architecture to enforce deterministic virtual time via the `hw/rust/backbone/clock` module, time *stops* for the simulated SoC when the simulator pauses.

1. **Memory Map Investigation:** You will inspect the VirtMCU hardware definition files to locate the base address of the APB-attached UART peripheral.
2. **MMIO Driver Implementation:** You will write a C driver utilizing volatile pointers to interact with the UART's memory-mapped control, status, and data registers. First, you will implement a polled transmission function, writing a string to the TX register and spinning until the `TX_EMPTY` flag is asserted.
3. **Deterministic Evaluation:** You will launch the VirtMCU instance and connect a simulated serial terminal to the `hw/rust/comms/chardev`

deterministic backend. Because VirtMCU strictly lock-steps the UART baud-rate generation to the exact number of retired instructions in the `slaved-icount` mode, you will observe flawless character delivery. You will then intentionally configure a baud rate mismatch between your driver and the terminal to observe framing errors, empirically proving how asynchronous protocols fail without precise timing synchronization.

# Chapter 7: Interrupts, Exceptions, and Bare-Metal Execution

## 7.1 Introduction: The Bare-Metal Environment

In a high-level programming environment, a software engineer rarely concerns themselves with the physical state of the machine at power-on. When a user runs a C program, the OS allocates virtual memory, sets up the execution environment, and smoothly transfers control to the `main()` function.

In bare-metal programming, however, `main()` does not execute by magic. The hardware strictly defines what happens the instant power is applied to the processor. A system must undergo a highly choreographed sequence of hardware and software initializations—collectively known as the boot process—to prepare the environment for high-level language execution.

Furthermore, once the application is running, it must interact with the continuous-time physical world. In Chapter 6, we introduced the concept of programmed I/O, or polling, where the CPU continuously loops to check the status of a peripheral. Polling is incredibly inefficient because it wastes thousands or millions of clock cycles spinning in a tight loop, preventing the CPU from performing useful computation or dropping into low-power sleep states. To create responsive, real-time Cyber-Physical Systems, we must abandon polling in favor of asynchronous, interrupt-driven event handling, where the hardware forcibly alerts the software that an event has occurred.

## 7.2 Booting: What Happens at Power-On

When an embedded processor is first powered up, or when a hardware reset button is pressed, the physical reset signal (often an active-low signal denoted

as `nRST` or `RST#` ) is asserted. This forces the processor's internal state machine into a known, rigid initial state.

But where does the processor find its first instruction? It cannot rely on the main memory (DRAM) because volatile memory contains garbage data at power-on. Instead, the processor is hardwired to look at a highly specific, predefined address in a non-volatile memory region, such as an on-chip Flash memory or a Boot ROM.

### 7.2.1 The Vector Table

Rather than placing the raw executable code directly at the boot address, modern architectures typically place a data structure known as the **Vector Table** at this location. The vector table is an array of addresses (pointers) that dictate where the processor should jump to handle specific hardware events.

For example, in the ARM Cortex-M architecture, the vector table is located at the very beginning of memory. The hardware expects the first two entries of this table to contain the most critical information needed to boot the system:

1. **Initial Stack Pointer (Offset 0x00):** The first 32-bit entry is not an instruction; it is the physical memory address of the top of the stack. When the CPU powers on, it immediately reads this value and loads it directly into the hardware Stack Pointer (SP) register. This is a brilliant architectural optimization, as it guarantees that a valid stack exists before the processor executes a single line of code, allowing the boot code to safely make function calls.
2. **Reset Vector (Offset 0x04):** The second entry contains the memory address of the **Reset Handler**. The hardware automatically loads this address into the Program Counter (PC) and begins fetching instructions.

Because the vector table is typically stored in read-only, non-volatile memory, its contents are established at compile-time. In some advanced implementations, the base address of the vector table can be dynamically relocated by software later in the execution cycle by writing to a specific control register, allowing the operating system to install its own set of handlers.

## 7.2.2 The C Runtime Initialization (CRT0)

The Reset Handler is a small, specialized block of assembly code—often called the startup code or `crt0` (C runtime zero)—responsible for preparing the environment so that C/C++ code can run safely.

Software engineers are used to writing C code containing global variables with initial values (e.g., `int x = 5;`). However, non-volatile Flash memory cannot be modified at runtime, and volatile RAM contains random garbage at power-on. To solve this, the linker organizes the compiled program into distinct segments:

- **.text:** Contains the executable machine instructions.
- **.data:** Contains initialized global and static variables.
- **.bss:** Contains uninitialized global and static variables (which the C standard dictates must default to zero).

During the boot process, the startup code must physically copy the initial values for the `.data` segment from the non-volatile ROM into the volatile RAM. Next, it must loop through the entire `.bss` segment in RAM and write zeros to every memory location. Finally, the startup code configures the start and limit addresses for the heap memory (used for dynamic allocations like `malloc`), and only then does it execute a branch instruction to transfer control to the user's `main()` function.

## 7.3 The Exception Model

In the normal course of execution, the processor fetches instructions sequentially, occasionally altering the Program Counter (PC) due to software branch instructions. However, the processor must possess mechanisms to handle unplanned, asynchronous events that demand immediate attention. This mechanism is known as the **Exception Model**.

An exception is conceptually similar to an unscheduled, forced function call. Exceptions can be broadly divided into two categories based on their origin: hardware and software.

### 7.3.1 Hardware Interrupts

A hardware exception, commonly referred to as an **interrupt**, is triggered by physical events external to the CPU execution pipeline. Examples include a user pressing a key on a keyboard, a timer reaching zero, or a network interface controller receiving a packet.

When a peripheral requires attention, it asserts a physical interrupt request (IRQ) line wired to the processor. Because these events happen independently of the software's execution flow, they are asynchronous. The processor must finish executing the current instruction, halt its normal sequence, and jump to an **Interrupt Service Routine (ISR)** to handle the device.

### 7.3.2 Software Exceptions (Traps and Faults)

Software exceptions are synchronous events generated directly by the execution of a specific instruction within the pipeline.

- **Faults:** Occur when the processor encounters an error condition it cannot resolve natively. Examples include attempting to execute an undefined or illegal instruction opcode, attempting to divide a number by zero, arithmetic overflow, or an MMU page fault (attempting to access virtual memory that is not currently mapped to physical RAM).
- **Traps:** Are intentional, scheduled exceptions triggered by specific assembly instructions (such as `SVC` or `SWI` in ARM, or `SYS / TRAP` in other architectures). Traps are the fundamental mechanism used by user-space applications to request services from the operating system, a concept we will explore further when discussing privilege levels.

## 7.4 Handling Exceptions: The Hardware Perspective

When an exception occurs, the processor hardware must perform several immediate, automated steps before any software handler can run.

### 7.4.1 Recording the Exception State

To handle the exception and eventually return to the interrupted program, the hardware must save the state of the machine. At a minimum, this involves saving the current Program Counter (so the CPU knows where to return) and a code indicating the reason for the exception.

Architectures handle this differently. In the MIPS architecture, for example, the processor utilizes a dedicated set of system management registers known as **Coprocessor 0**. When an exception fires, the hardware automatically copies the current PC into the **Exception Program Counter (EPC)** register and writes an identification code into the **Cause** register. The processor then forces the PC to jump to a hardcoded global exception handler address (e.g., `0x80000180`), where the OS must read the Cause register to determine what actually happened.

In contrast, the ARM architecture utilizes a vectored approach via the Nested Vectored Interrupt Controller (NVIC). Each specific type of exception or peripheral interrupt is assigned a unique index number. The hardware uses this index to directly look up the specific handler address in the Vector Table we discussed earlier, branching immediately to the correct software routine without requiring a massive, centralized `switch/case` statement to decode the cause.

### 7.4.2 Precise Interrupts and Pipeline Flushing

Handling exceptions becomes incredibly complex in modern microarchitectures featuring deep pipelines, superscalar issue, and out-of-order execution.

If a processor is executing five instructions simultaneously in a pipeline and an external interrupt fires, the processor cannot simply stop instantly. A fundamental requirement of modern systems is the **precise interrupt**. For an interrupt to be precise, the hardware must guarantee that:

1. All instructions appearing in the program code *before* the interrupted instruction have fully completed and retired their results to the

architectural state.

2. All instructions appearing *after* the interrupted instruction have completely discarded any partial computations, leaving no trace in the architectural state.

To achieve this in an out-of-order processor like the Intel Core i7, instructions write their intermediate results to hidden physical registers, but they only commit their results to the visible architectural registers (a process called retirement) in strict, sequential program order. If a fault or interrupt occurs, the Retirement Unit halts, preserving the precise state of the machine up to that exact instruction, and the pipeline is flushed of all subsequent speculative instructions. This ensures that when the ISR finishes, the program resumes exactly where it left off, oblivious to the fact that it was ever interrupted.

## 7.5 The Interrupt Service Routine (ISR)

Once the hardware has verified the precise state and branched to the correct address in the vector table, the software **Interrupt Service Routine (ISR)** takes over. Writing an ISR requires extreme discipline, as the routine is hijacking the CPU from an active thread.

As outlined in fundamental operating systems theory, a robust ISR must perform the following distinct software actions:

1. **Context Save:** The very first action the ISR must take is to save the contents of all general-purpose registers it intends to use. If the ISR modifies a register (like `x0` or `R1`) without saving it, it will corrupt the mathematical calculations of the interrupted background program. Registers are typically pushed onto the current stack. *(Note to students: While the specific ARM AArch64 assembly instructions for context switching are standard industry practice and draw on our knowledge from outside the provided foundational texts, you are encouraged to independently verify the exact syntax in the ARM Architecture Reference Manual. Typically, this involves using the `STP` (Store Pair) instruction to push registers to the stack).*
2. **Acknowledge and Clear:** The ISR must communicate with the hardware peripheral (e.g., the UART or Timer) that triggered the interrupt, reading



its status registers to identify the specific sub-event, and explicitly writing a command to clear the hardware interrupt flag. If the flag is not cleared, the CPU will immediately re-trigger the interrupt the moment the ISR returns, resulting in an infinite loop.

3. **Process the Data:** The ISR executes its specific logic—such as copying a received byte from a UART data register into an in-memory ring buffer, or decrementing a counter.
4. **Context Restore:** Once the work is complete, the ISR pops the saved values off the stack, restoring the general-purpose registers to their exact state prior to the interrupt.
5. **Return:** The ISR executes a special machine instruction (often called `RTI`, `IRET`, or a specific branch return in ARM) that atomically restores the processor's status flags from the hardware exception registers and copies the Exception Program Counter back into the active Program Counter, resuming the background task.

### 7.5.1 Interrupt Priorities and Nesting

In a complex system, multiple peripherals may attempt to interrupt the CPU simultaneously. Furthermore, a high-priority event (like an emergency motor shutoff) might occur while the CPU is already in the middle of executing a low-priority ISR (like a keyboard press).

If an ISR disables all other interrupts during its execution to prevent race conditions, the system becomes completely sequential, risking the loss of time-critical data if the first ISR takes too long. To resolve this, modern architectures assign strict mathematical priorities to different interrupts.

When an interrupt fires, the CPU elevates its current execution priority. If a new interrupt arrives with a *lower* priority, it is held in a pending state until the current ISR finishes. However, if an interrupt arrives with a *higher* priority, the CPU immediately preempts the active ISR, pushes its state to the stack, and services the more critical event—a concept known as **Nested Interrupts**. The ARM Nested Vectored Interrupt Controller (NVIC) handles these complex priority evaluations entirely in hardware, significantly reducing software latency.

## 7.6 Privilege Levels and Protection

Interrupts and exceptions are not only used for handling hardware events; they form the absolute foundation of system security and operating system architecture.

Modern processors implement multiple hardware privilege levels, or execution modes. The Intel Core i7, for instance, supports four privilege rings (Ring 0 to Ring 3), though typical operating systems like Linux and Windows only utilize two: Ring 0 (Kernel Mode) and Ring 3 (User Mode).

When the processor is in User Mode, the hardware Memory Management Unit (MMU) strictly restricts what the software can do. User Mode code is absolutely prohibited from executing privileged machine instructions (like disabling hardware interrupts) or accessing memory-mapped I/O peripheral registers. If a User Mode application attempts to directly modify a hardware timer, the MMU intervenes, throwing a synchronous fault exception that terminates the offending program.

But if a user program cannot access hardware, how does it ever print to the screen or save a file? It must ask the highly trusted Operating System to do the work on its behalf.

### 7.6.1 The System Call

Because User Mode code cannot simply execute a `CALL` instruction to branch directly into the Kernel's memory space (which is protected by the MMU), it must use a controlled hardware gateway.

To request a service, the user program loads a specific service number into a general-purpose register and executes a software trap instruction (e.g., `SYS` or `svc`). This instruction intentionally causes a hardware exception. The processor instantly elevates its privilege level to Kernel Mode and forces the PC to jump to a highly secured, predetermined vector address containing the OS System Call Handler.

The OS examines the requested service number, securely validates any passed pointers, performs the privileged hardware I/O operation, and then executes a return-from-exception instruction to demote the processor back to User Mode and resume the application. This strict separation ensures that buggy or malicious user software cannot crash or compromise the underlying Cyber-Physical System.

## 7.7 Context Switching and Parallel Processing

By combining hardware timer interrupts with the privilege and exception model, operating systems can create the illusion of true parallel processing on a single CPU core.

If multiple programs (processes or threads) reside in memory, an RTOS (Real-Time Operating System) configures a hardware timer peripheral to generate a high-priority interrupt at a fixed interval (e.g., every 1 millisecond).

1. The CPU is running Task A.
2. The timer fires, preempting Task A and forcing a jump to the OS Scheduler ISR.
3. The OS saves the exact register state of Task A onto Task A's private stack.
4. The OS changes the CPU's hardware Stack Pointer to point to Task B's private stack.
5. The OS pops Task B's previously saved registers off Task B's stack.
6. The OS executes the return-from-exception instruction, which restores Task B's Program Counter.
7. The CPU resumes executing Task B.

This incredibly rapid process is known as a **Context Switch**. Because it relies entirely on the precise interrupt mechanism preserving and restoring architectural state, neither Task A nor Task B is aware that it was paused and swapped out of the execution pipeline.

## 7.8 VirtMCU Homework: Building a Deterministic Interrupt Handler

In a standard embedded systems curriculum, writing bare-metal interrupt handlers introduces massive frustration. If a student configures a timer interrupt in a standard software emulator (like a vanilla QEMU instance), the emulator relies on the host operating system's wall-clock to trigger the event. Because the host OS (Linux/Windows) introduces unpredictable scheduling delays, context switches, and network latencies, the simulated timer interrupt suffers from immense **wall-clock jitter**,.

This jitter makes it impossible to reliably debug precise hardware timing dependencies, a fatal flaw when designing Cyber-Physical Systems where missing a control deadline by microseconds can crash a drone.

To solve this, your homework relies on the **VirtMCU FirmwareStudio**. VirtMCU completely abandons the host OS wall-clock. Instead, it utilizes a dynamic QOM plugin architecture to enforce the **Big QEMU Lock (BQL)** constraints and implements strict cooperative time slaving.

1. **Architecture Review:** You will begin by reviewing the VirtMCU design rationale in `docs/architecture/01-system-overview.md` to understand how the five implementation pillars guarantee cycle-accurate deterministic execution regardless of host computer speed.
2. **Writing the Handler:** You will be provided with a bare-metal C framework containing an uninitialized vector table. You must write an assembly language Interrupt Service Routine (ISR) that adheres to the strict context save/restore requirements discussed in Section 7.5.
3. **Peripheral Interfacing:** Using the `hw/rust/common/rust-dummy/` module as a template, you will instantiate a simulated push-button peripheral that triggers an IRQ line mapped to the NVIC.
4. **Execution:** You will compile your binary and run it within VirtMCU in `slaved-icount` mode. Because VirtMCU strictly lock-steps the virtual clock to the exact number of retired instructions (bypassing host OS jitter), you will use the `gdb` debugger to step through your ISR cycle-by-cycle. You will empirically verify that your interrupt fires at the exact same

deterministic virtual nanosecond across multiple simulation runs, proving the stability necessary for rigorous CPS 2.0 design.

# Chapter 8: Real-Time Operating Systems (RTOS)

## 8.1 Introduction: Beyond the Super-Loop

In Chapter 7, we breached the hardware-software interface, examining how a bare-metal processor boots and how it utilizes the exception model to asynchronously respond to physical events via hardware interrupts. For highly simplistic embedded devices—such as a microwave oven controller—software engineers often rely on a “super-loop” architecture. In a super-loop, the `main()` function consists of an infinite `while(1)` loop that continuously polls sensors, computes logic, and writes to actuators, occasionally pausing to service a hardware interrupt.

However, as we transition into modern Cyber-Physical Systems (CPS)—such as autonomous drone flight controllers, automotive anti-lock braking systems, and interconnected smart manufacturing robotics—the super-loop rapidly collapses under the weight of its own complexity. When a system must simultaneously sample fifty different sensors at varying frequencies, manage an encrypted TCP/IP network stack, and execute complex motor control algorithms, a single execution thread riddled with interrupt service routines (ISRs) becomes impossible to maintain, debug, or verify.

To manage this immense software complexity, we must introduce an Operating System (OS). For software engineers accustomed to writing applications for Windows or Linux, an OS is viewed as a massive resource manager that abstracts the hardware and provides virtual memory, file systems, and fair multi-tasking. However, general-purpose operating systems (GPOS) are fundamentally designed to maximize **average throughput** and maintain fairness among users. If a desktop PC delays the rendering of a web browser by 200 milliseconds to flush a disk cache, the human user barely notices.

In a Cyber-Physical System, time is a logical correctness criterion. If an automotive airbag controller computes the mathematically perfect deployment force but delivers that signal 50 milliseconds too late, the result is catastrophic. Therefore, CPS architectures rely on a **Real-Time Operating System (RTOS)**. An RTOS strips away the heavy abstractions of a GPOS and focuses obsessively on **determinism**, guaranteeing that critical tasks will execute within strict, mathematically provable time constraints.

## 8.2 The Real-Time Paradigm: Soft, Firm, and Hard

Before examining the internal architecture of an RTOS, we must establish the vocabulary used to classify real-time constraints. A real-time system is one in which the correctness of a computation depends not only upon the logical result but also upon the time at which the result is produced. Real-time systems are categorized by the severity of the consequences when a deadline is missed:

- **Soft Real-Time:** Missing a deadline degrades the performance of the system, but the late result still retains some value, and the system can recover. An example is a user interface update on a smart thermostat or a streaming video decoder; a dropped frame is annoying but not fatal.
- **Firm Real-Time:** Missing a deadline means the computed result is now entirely useless and must be discarded. However, missing the deadline does not cause catastrophic system failure. An example is a high-frequency financial trading algorithm; if the execution engine misses the optimal microsecond window to execute a trade, the opportunity is gone, but the server itself remains operational.
- **Hard Real-Time:** Missing a deadline constitutes an absolute, catastrophic system failure, often resulting in physical destruction or loss of life. In these systems, a late answer is identically equivalent to a wrong answer. Examples include aerospace fly-by-wire controls, pacemaker timing circuits, and industrial robotic arm actuation.

An RTOS is specifically engineered to support hard real-time requirements. It achieves this not necessarily by being “fast” (a 3 GHz desktop processor running Linux will execute raw math much faster than a 100 MHz RTOS-driven

microcontroller), but by being **predictable**. An RTOS minimizes scheduling jitter, bounded interrupt latency, and context-switch overhead to ensure that worst-case execution times (WCET) can be strictly calculated and guaranteed.

## 8.3 Anatomy of an RTOS: Tasks, Threads, and States

In a modern GPOS, the fundamental unit of execution is the process, which possesses its own isolated virtual memory address space mapped by the MMU. To switch between processes, the OS must flush the Translation Lookaside Buffer (TLB) and swap page tables, an incredibly expensive operation.

In a lightweight RTOS, memory protection is often relaxed to achieve deterministic speed. The RTOS typically operates in a single, unified physical address space. Therefore, the fundamental unit of execution is the **thread** or **task**. Multiple threads run in the same address space, meaning they can trivially share global variables, but they maintain their own independent execution flow.

### 8.3.1 The Task Control Block (TCB)

To sustain the illusion that multiple tasks are executing simultaneously on a single CPU core, the RTOS kernel must track the exact machine state of every task. It does this using a fundamental data structure known as the **Task Control Block (TCB)**. The TCB contains the task's unique ID, its scheduling priority, a pointer to its private stack memory, and its current execution state.

At any given moment in time, an RTOS task exists in one of several states:

1. **Running:** The task is currently executing instructions in the CPU pipeline. On a single-core processor, only one task can be in the Running state at any given nanosecond.
2. **Ready:** The task is fully prepared to execute and is simply waiting for the scheduler to grant it access to the CPU.



3. **Blocked (Waiting):** The task cannot execute because it is waiting for a specific event to occur. This could be waiting for a hardware timer to expire, waiting for a packet to arrive on a network interface, or waiting for a shared resource to become available.
4. **Suspended:** The task has been explicitly halted by the programmer or the system and will not be scheduled until it is explicitly resumed.

### 8.3.2 The Context Switch

When the RTOS decides to stop running Task A and start running Task B, it performs a **context switch**. This is an intricately choreographed sequence of assembly language instructions that relies heavily on the processor's exception model.

First, a hardware timer—often called the SysTick timer in ARM architectures—fires a periodic interrupt (e.g., every 1 millisecond). This interrupt preempts the Running task (Task A) and forces the CPU to jump to the RTOS Scheduler ISR. The scheduler must preserve the exact architectural state of Task A so it can be resumed later. It does this by pushing Task A's general-purpose registers (X0-X30 in AArch64), its Program Counter (PC), and its Process State Register (PSTATE) onto Task A's private stack.

The scheduler then updates Task A's TCB with the new stack pointer address. Next, the scheduler selects Task B from the Ready queue, reads Task B's TCB to find its stack pointer, and points the hardware Stack Pointer (SP) register to Task B's stack. Finally, the scheduler pops Task B's saved registers off its stack and executes a return-from-exception instruction. The CPU pipeline seamlessly resumes executing Task B, entirely unaware that a switch occurred.

Because context switching consumes CPU cycles without performing any application-level work, RTOS developers optimize this assembly code ruthlessly, ensuring the context switch overhead is highly deterministic and consumes only a few dozen clock cycles.

## 8.4 RTOS Scheduling Algorithms

The brain of the RTOS is the scheduler, the algorithm that decides which Ready task should transition to the Running state.

### 8.4.1 Round-Robin and Time-Slicing

In a pure **Round-Robin** scheduling algorithm, every task is assigned an equal, fixed time quantum (e.g., 5 milliseconds). The scheduler cycles through a circular queue of Ready tasks, giving each task the CPU for its allotted quantum before preempting it and moving to the next. While this guarantees fairness and prevents any single task from starving, it is disastrous for hard real-time systems. If a critical motor-control task becomes Ready, it cannot execute until all other tasks in the queue have exhausted their time slices, destroying determinism.

### 8.4.2 Priority-Based Preemptive Scheduling

Virtually all modern RTOS kernels utilize **Priority-Based Preemptive Scheduling**. The software engineer explicitly assigns a strict mathematical priority to every task. The scheduling rule is absolute: *The CPU will always execute the highest-priority task that is in the Ready state.*

If a low-priority task is currently Running, and a hardware interrupt wakes up a high-priority task (moving it from Blocked to Ready), the scheduler will immediately preempt the low-priority task and context-switch to the high-priority task. The low-priority task will only resume once the high-priority task voluntarily yields the CPU or enters the Blocked state.

### 8.4.3 Rate Monotonic Scheduling (RMS)

Assigning priorities manually can lead to missed deadlines if the software engineer guesses incorrectly. To mathematically guarantee that a set of

periodic tasks will always meet their hard real-time deadlines, systems architects use analytic algorithms like **Rate Monotonic Scheduling (RMS)**.

RMS assigns static priorities to tasks based on their frequency of execution: the shorter the task's period (the more frequently it must run), the higher its priority. To determine if a set of  $n$  independent, preemptable tasks is schedulable under RMS without missing any deadlines, we calculate the total CPU utilization  $U$ . Let  $C_k$  be the maximum worst-case execution time of task  $k$ , and  $T_k$  be its execution period. The system is provably schedulable if the following condition is met:

$$U = \sum_{k=1}^n \frac{C_k}{T_k} \leq n(2^{1/n} - 1)$$

Consider an example of three critical CPS threads:

- Thread 1: Execution Time  $C_1 = 50$ , Period  $T_1 = 100$
- Thread 2: Execution Time  $C_2 = 100$ , Period  $T_2 = 500$
- Thread 3: Execution Time  $C_3 = 120$ , Period  $T_3 = 1000$

We evaluate the left side of the RMS formula:  $U = \frac{50}{100} + \frac{100}{500} + \frac{120}{1000} = 0.50 + 0.20 + 0.12 = 0.82$  Next, we evaluate the theoretical upper bound for  $n=3$  threads on the right side of the equation:  $3(2^{1/3} - 1) \approx 3(1.2599 - 1) = 0.7798$  Because  $0.82 \not\leq 0.7798$ , this specific set of threads is mathematically *not* guaranteed to be schedulable under pure RMS, meaning under worst-case phasing, a deadline will be missed. By performing this analysis at compile-time, engineers can redesign the system before deploying it to physical hardware.

## 8.5 Inter-Task Communication and Synchronization

Because threads in an RTOS share the same memory space, they can easily pass data back and forth by reading and writing global variables. However, uncontrolled interaction between concurrent processes leads to catastrophic **race conditions**.

Suppose Task A and Task B both increment a global counter tracking physical motor rotations. The C code `counter++` translates to three machine instructions: Load the variable from RAM to a register, increment the register, and Store the register back to RAM. If Task A is preempted by Task B precisely between the Load and Store instructions, both tasks will read the exact same original value, increment it, and write back identical values. One physical rotation is permanently lost in the software state.

The segment of code where a task accesses a shared variable is known as a **critical section** or critical region. To prevent race conditions, the RTOS must provide synchronization primitives that enforce **mutual exclusion**, ensuring that if one task is inside the critical section, no other task can enter it.

### 8.5.1 Semaphores and Mutexes

The most fundamental synchronization primitive is the **semaphore**, invented by E.W. Dijkstra in 1965. A semaphore is an integer variable managed securely by the RTOS kernel. It supports two atomic operations: `down` (historically *P*) and `up` (historically *V*).

When a task wishes to enter a critical section, it executes a `down` operation on the semaphore. If the semaphore's value is greater than zero, the RTOS decrements the value and allows the task to proceed. If the value is 0, the RTOS immediately removes the task from the Running state, places it in the Blocked queue, and schedules a different task. When the task finishes its critical section, it executes an `up` operation, which increments the semaphore. If one or more tasks are sleeping in the Blocked queue waiting for this semaphore, the RTOS wakes up one of them, moving it back to the Ready queue.

A **Mutex** (Mutual Exclusion object) is a specialized binary semaphore initialized to 1. It operates as an electronic "key." A task must acquire the mutex before accessing a shared memory block and must release it when finished. Mutexes often include concept of "ownership," meaning the exact task that acquired the mutex is the only one legally permitted to release it, preventing rogue tasks from unlocking resources they do not own.

Semaphores are also heavily used for the **producer-consumer problem**, where one task generates data (e.g., an ADC sampling voltages) and places it into a shared ring buffer, while another task consumes the data. Semaphores inherently track the number of full slots and empty slots in the buffer, elegantly putting the producer to sleep if the buffer fills up, and waking it the instant the consumer removes an item.

## 8.6 The Hazards of Concurrency: Deadlocks and Priority Inversion

While semaphores solve the race condition problem, they introduce entirely new classes of systemic failure.

### 8.6.1 Deadlock

A **deadlock** occurs when a set of tasks are permanently blocked, each waiting for a resource that is held by another task in the set. This is famously illustrated by the Dining Philosophers problem, where multiple philosophers sit at a circular table requiring two forks to eat, but there is only one fork between each plate. If every philosopher picks up their left fork simultaneously, they will all block forever waiting for the right fork, creating a circular wait dependency.

In CPS architecture, deadlocks are identified by the presence of the four Coffman conditions: Mutual exclusion, hold and wait, no preemption, and circular wait. RTOS designers must break these conditions structurally—for instance, by mathematically mandating a strict global hierarchy for acquiring multiple mutexes, preventing a circular dependency chain from ever forming.

### 8.6.2 Priority Inversion

Priority Inversion is perhaps the most notorious software hazard in Cyber-Physical Systems, famously causing the near-loss of NASA's Mars Pathfinder rover in 1997.

Priority inversion occurs under a specific combination of priority-based preemptive scheduling and mutexes. Consider three tasks: High (H), Medium (M), and Low (L).

1. Task L begins executing and acquires a Mutex protecting a shared communications bus.
2. Task H wakes up, preempting Task L. Task H attempts to acquire the same Mutex. Because Task L holds it, Task H is put into the Blocked state.
3. This is normal so far; however, Task M suddenly wakes up. Because Task M has a higher priority than Task L, and does *not* need the Mutex, the scheduler grants the CPU to Task M.

Task M begins executing a long computation. Task L cannot run because it is preempted by Task M. Task H cannot run because it is blocked waiting for Task L to release the Mutex. Thus, the High-priority task is effectively suspended indefinitely while the Medium-priority task dominates the CPU. The priorities have been inverted!

To solve this, modern RTOS kernels implement the **Priority Inheritance Protocol (PIP)**. When Task H blocks on a Mutex held by Task L, the RTOS temporarily elevates the priority of Task L to match the priority of Task H. Because Task L now has “High” priority, it cannot be preempted by Task M. Task L finishes its critical section, releases the Mutex, and its priority instantly drops back to “Low.” Task H acquires the Mutex and resumes execution, preventing the inversion.

## 8.7 Industry Standards: CMSIS-RTOS

In the embedded ecosystem, learning a new proprietary RTOS API for every microchip vendor creates severe friction. To standardize software portability across the vast array of ARM Cortex-M microcontrollers, ARM developed the Cortex Microcontroller Software Interface Standard (CMSIS).

CMSIS provides a hardware abstraction layer that defines standardized functions for core CPU access (CMSIS-Core), DSP math functions (CMSIS-DSP), and importantly, a standardized RTOS API called **CMSIS-RTOS**.

CMSIS-RTOS is not an RTOS itself, but rather a wrapper API that sits on top of commercial or open-source kernels like FreeRTOS or Keil RTX. It defines a universal set of C macros and functions to manage RTOS task parameters, inter-process communication, synchronization, and scheduling configurations. By adhering strictly to the CMSIS-RTOS API, a software engineer can write a complex multithreaded drone controller on a small Cortex-M3 processor, and effortlessly recompile and port that exact same codebase to a massive multi-core Cortex-M7 SoC without rewriting the thread management or semaphore logic.

## 8.8 VirtMCU Homework: Deterministic Context-Switch Analysis

Understanding the theoretical concepts of priority inversion and RMS scheduling is trivial; identifying and debugging them in an active execution pipeline is enormously difficult.

If you attempt to trace a context switch or a priority inversion bug using a standard QEMU emulator running on Linux or Windows, you will fail. Standard emulators execute instructions using a free-running host clock. The host OS scheduler interrupts the emulator unpredictably to service its own background tasks, injecting massive amounts of wall-clock scheduling jitter into the emulation. Measuring the exact microsecond penalty of an RTOS context-switch becomes impossible because the measurement is polluted by the host machine's background noise.

To overcome this, you will utilize the **VirtMCU FirmwareStudio**. Because VirtMCU uses a dynamic QOM plugin architecture to lock-step the virtual clock precisely to the number of retired instructions ( `slaved-icount` mode), it creates a perfectly deterministic simulation environment devoid of host OS jitter.

1. **RTOS Porting:** You will be provided with a bare-metal C implementation of a lightweight RTOS kernel. Your first task is to examine the context

switch assembly routine, validating how it pushes X0-X30, the FP, and the LR onto the task stack.

2. **Jitter Measurement:** You will configure two threads to toggle a GPIO pin. You will run this in standard QEMU and observe the massive timing variance using a virtual logic analyzer. Then, you will run the identical binary in VirtMCU `slaved-icount` mode, using the deterministic virtual hardware timers to measure the *exact* number of clock cycles consumed by the RTOS scheduler during the context switch.
3. **Provoking Priority Inversion:** You will write an application containing three tasks (High, Medium, Low) and a single shared Mutex. You will intentionally orchestrate a priority inversion scenario. By pausing the VirtMCU simulation at exact cycle counts, you will use GDB to dump the RTOS Ready queues and empirically verify that the High-priority task is incorrectly starving.
4. **Protocol Verification:** Finally, you will enable the Priority Inheritance Protocol feature in your RTOS configuration. Rerunning the deterministic simulation, you will observe the kernel actively promoting the Low-priority task's TCB priority field, proving how software mechanisms resolve the hazards of asynchronous CPS concurrency.



# Chapter 9: Introduction to Cyber-Physical Systems (CPS 2.0)

## 9.1 Introduction: Beyond the Embedded System

Welcome to Part III of our course. Up to this point, we have focused extensively on the internal architecture of the computing node. We progressed from the physics of transistors to the pipeline of the ARM AArch64 processor, navigated the complexities of the memory hierarchy, and finally, in Chapter 8, deployed a Real-Time Operating System (RTOS) to guarantee deterministic thread scheduling.

If our curriculum ended here, you would be well-equipped to design a traditional embedded system. However, the modern engineering landscape demands more. We are no longer simply building isolated microcontrollers that blink LEDs or manage singular engines. We are engineering **Cyber-Physical Systems (CPS)**.

Cyber-Physical System theory emerged as a direct response to the increasingly complex interactions between computer algorithms and physical infrastructure. The goal of a CPS is to seamlessly bundle sensory data and advanced mathematical models to improve the control, automation, and decision-making of physical assets. In the late 2000s, as these systems took the form of next-generation embedded networks, the pioneer Wayne Wolf coined the term “control-computing codesign” to describe this paradigm. Wolf argued that distributed computing and control elements needed a collective, integrated form of examination rather than being isolated into separate hardware and software layers.

In a cyber-physical system, computation is inherently tied to the continuous-time physics of the real world. A desktop computer computes data to produce more data (e.g., calculating a spreadsheet). A CPS computes data to produce *physical action* (e.g., actuating an aircraft’s rudder, deploying a medical

pacemaker, or braking an autonomous vehicle). Consequently, a CPS is evaluated not just on its logical correctness, but on its physical timeliness and safety.

This chapter introduces the evolution of CPS from its foundational state to what is now known as **CPS 2.0**. We will explore the interplay of communication and computational technologies that underpin the fabric of CPS 2.0. We will dissect the anatomy of these systems, examining the feedback loops that bind the cyber and physical domains. Finally, we will outline the “3Cs” that define all modern CPS architectures: Computation, Communication, and Control.

## 9.2 The Evolution to CPS 2.0

To understand where we are, we must understand how we got here. The evolution of cyber-physical systems can be broadly categorized into distinct generational phases.

### 9.2.1 CPS 1.0: Isolated Automation and the Industrial IoT

First-generation industrial CPS—often retrospectively labeled CPS 1.0—focused heavily on the basic automation and control of manufacturing machines and isolated lines. These systems were typically characterized by localized Programmable Logic Controllers (PLCs) governing physical processes via closed-loop feedback.

As networking technologies advanced, these systems transitioned into the Industrial Internet of Things (IIoT). The primary objective during this phase was data telemetry: capturing vast amounts of industrial sensor data and transmitting it to remote cloud services to be visualized, logged, and analyzed by human operators. While this provided unprecedented visibility into physical operations, the systems remained largely reactive. The cloud was used for storage and offline analysis, not for real-time, closed-loop physical control, largely due to the prohibitive latency and jitter of internet-based communication.

### 9.2.2 CPS 2.0: Intelligence, Convergence, and Metasystems

As related technologies matured—specifically in edge computing, deterministic networking, and artificial intelligence—new opportunities opened, leading to the second generation: **CPS 2.0**.

CPS 2.0 requires entirely new software platforms and architectures to be deployed. The defining characteristic of CPS 2.0 is the profound integration of three traditionally isolated domains:

1. **Operational Technology (OT):** The hardware and software that detects or causes a change through the direct monitoring and/or control of physical devices, processes, and events.
2. **Information Technology (IT):** The networking, storage, and processing infrastructure used to manage enterprise data.
3. **Artificial Intelligence (AI):** The application of machine learning, neural networks, and high-frequency data models to enable predictive and autonomous decision-making.

In CPS 2.0, the system transcends traditional boundaries, giving rise to “metasystems” that permanently blur the lines between the physical and virtual worlds. Rather than merely sending data to the cloud for offline analysis, a CPS 2.0 metasystem utilizes high-frequency (HF) data streams to generate AI models on the fly, pushing these models down to the edge (directly onto the local microcontrollers) to influence real-time physical control.

These advanced architectures span multiple domains. We see CPS 2.0 manifesting in smart health-care systems deployed across smart cities, in cloud-enabled smart manufacturing processes (often referred to as Industry 4.0), and in distributed infrastructure like Virtual Power Plants (VPPs). In autonomous vehicles (AVs), CPS 2.0 allows the vehicle to integrate tangible, physical kinematics with advanced cyber decision-making, navigating complex environments autonomously.

## 9.3 Anatomy of a Cyber-Physical System

At the core of every cyber-physical system, regardless of its generational label, is the **cyber-physical feedback loop**. This loop is the fundamental mechanism by which a system observes the environment, reasons about it, and takes action to alter it.

### 9.3.1 Sensing the Physical World

The loop begins in the physical domain. The physical world is governed by continuous-time dynamics—classical mechanics, Newtonian motion, thermodynamics, and electromagnetics. To bring these continuous physical phenomena into the discrete cyber domain, the system relies on sensors (transducers) that convert physical energy into electrical signals.

In a modern CPS, sensing goes far beyond reading a simple analog voltage. Consider the field of Structural Health Monitoring (SHM) for civil infrastructure, such as bridges. In an SHM-oriented CPS, a network of accelerometers continuously monitors the vibration response of the bridge. This raw sensory data is processed using signal feature extraction techniques, such as Fast Fourier Transforms (FFT) or Discrete Wavelet Transforms (DWT). By analyzing high-frequency data, the cyber system can identify decaying modal frequencies—for instance, observing a frequency drop from 2.93 Hz to 1.56 Hz—which physically translates to a reduction in the bridge's structural stiffness due to earthquake damage or aging.

### 9.3.2 Model-Driven Processing and Decision Making

Once the physical data is ingested, the computing core must process it. In CPS 2.0, this relies heavily on **model-driven** processing. The raw data is fed into a mathematical representation (a “digital twin”) of the physical asset.

Returning to our bridge example, the cyber system maintains a baseline finite element model of the bridge. The sensor data is used to continually calibrate this model, updating uncertain parameters such as mass-stiffness features and

boundary conditions to reflect the real-world ground truth. By linking the physical asset to the digital model via this calibration protocol, the CPS can perform structural reliability assessments and risk analyses. If the model predicts an impending failure under a simulated seismic load, the system can autonomously trigger mitigation decisions—such as dynamically optimizing transportation networks to route heavy traffic away from the compromised structure,.

This demonstrates the power of a CPS: it does not just report data; it extrapolates existing sensory knowledge through physics-informed strategies to make intelligent decisions.

### **9.3.3 Actuation and Physical Control**

The final stage of the feedback loop is actuation. Based on the decisions made by the cyber algorithms, the system must interact with the physical world to change its state. This is achieved through actuators—such as electrostatic or electromagnetic transducers, variable reluctance motors, and hydraulic valves.

Actuation is fundamentally tied to control laws. As the cyber system commands an actuator, it alters the physical domain, which in turn alters the readings of the sensors, closing the feedback loop. System functionality and capability are accomplished by means of rigorous optimization and control engineering to ensure physical controllability.

## **9.4 The 3Cs of CPS: Computation, Communication, and Control**

To formalize the architecture of a CPS 2.0 metasystem, engineers analyze the design through the lens of the “3Cs”: Computation, Communication, and Control. Let us examine how each pillar underpins the fabric of modern cyber-physical infrastructure.

### 9.4.1 Computation

Computation in a CPS is the processing engine that executes the control laws and AI models. However, computation in the physical world is vastly different from computation in standard IT.

In a CPS, the state of the system is often defined by non-linear differential equations representing the physical kinematics (e.g., the rigid body dynamics and equations of motion for an aircraft). The processor must solve these equations and execute tracking control laws in real-time. For example, controlling the multi-axis coupling, resonance, and flexible modes of a robotic arm requires matrix mathematics and Proportional-Integral-Derivative (PID) control algorithms to be executed with strict microsecond deadlines,.

In CPS 2.0, computation is heavily decentralized. The architecture often employs **Edge Computing**, where the heavy computational lifting (such as high-frequency signal feature extraction and dimensionality reduction) is performed directly on the local microcontroller rather than being offloaded to the cloud. This guarantees that the control loop maintains its real-time determinism. The cloud is reserved for global orchestration, training heavy AI models, and performing long-term reliability assessments,.

### 9.4.2 Communication

Communication is the lifeblood of any distributed Cyber-Physical System. In a massive system like a Virtual Power Plant (VPP)—which coordinates hundreds of distributed energy resources, servers, and loads—the architecture relies entirely on the timely exchange of telemetry and command data.

Traditional TCP/IP networks are optimized for average throughput, making them highly susceptible to packet loss and latency spikes (jitter). In a hard real-time CPS, a delayed network packet containing a braking command is functionally equivalent to a lost packet, potentially resulting in catastrophic physical failure.

To resolve this, CPS 2.0 embraces **Software-Defined Networking (SDN)**. SDN enables a new era of flexibility by separating the network's control plane from

the data plane. It provides dynamic network management, enhancing system resilience while addressing the imperative requirement to maintain high-performance, deterministic communication in real-time operations. By utilizing coordinated communication control strategies, engineers can systematically analyze and mitigate the adverse effects of packet loss and latency, optimizing performance under real-world physical constraints.

### 9.4.3 Control

Control is the application of cybernetics principles to govern the behavior of the physical asset. The objective of the control system is to force the physical plant to follow a desired reference trajectory, despite external physical disturbances or internal system noise.

In a closed-loop system, the control law minimizes the tracking error—the difference between the desired physical state and the actual measured state. For instance, in the torque control of an electromagnetic induction motor, engineers design optimal tracking control laws that rely on system state matrices and feedback gains. By minimizing quadratic performance functionals and tuning the proportional, integral, and derivative ( $k_p$ ,  $k_i$ ,  $k_d$ ) gains of a PID controller, the cyber system guarantees that the motor's settling time remains within strict physical bounds.

In the CPS 2.0 paradigm, control systems are evolving into **self-adapting systems**. Using holonic reference architectures, the control layers are becoming decentralized and autonomous, capable of dynamically adjusting their own control laws in response to changing physical environments or hardware degradation.

## 9.5 Security and Resilience in CPS 2.0

As we transition to CPS 2.0 and deeply integrate Operational Technology (OT) with internet-facing Information Technology (IT), the attack surface of our physical infrastructure expands exponentially.

In isolated embedded systems (CPS 1.0), security was largely physical; you protected the PLC by putting a lock on the factory door. In CPS 2.0, systems are connected to cloud environments, exposing critical infrastructure—such as power grids, medical cyber-physical systems, and transportation networks—to sophisticated cyber threats,. A cyber-attack on a traditional IT database results in a loss of privacy or data. A cyber-physical attack on a CPS results in a loss of physical control, which can lead to equipment destruction or human casualties.

Therefore, CPS 2.0 architecture mandates “security by design.” This requires the implementation of Cyber-Physical Threat Intelligence, anomaly detection models, and robust access controls. It also requires the system to possess physical resilience—the ability of the control system to detect a cyber-induced hardware fault (such as an attacker feeding spoofed sensor data to the PID controller) and safely degrade operations to prevent physical catastrophe,. We will explore these advanced security topics deeply in Chapter 13.

## 9.6 VirtMCU Homework: CPS 2.0 Architecture Design

In a standard university software environment, designing and testing a true Cyber-Physical System is exceptionally difficult. If you attempt to simulate a drone flight controller relying on a host operating system’s networking stack, the wall-clock scheduling jitter will destroy your PID control loops. Your drone will crash in simulation not because your control mathematics were wrong, but because the host OS delayed a sensor packet by 50 milliseconds to process a background task.

To bridge this gap, this chapter’s homework relies on the **VirtMCU FirmwareStudio**. Because VirtMCU enforces deterministic virtual time and physics engine lock-stepping via the `hw/rust/backbone/clock` module, it completely eliminates wall-clock jitter. Time inside the VirtMCU environment moves strictly in accordance with retired instructions and precise hardware timers.



**Your Assignment:** You are tasked with designing the high-level CPS 2.0 architecture for a multi-rotor autonomous drone.

1. **Map the 3Cs:** You must provide a formal block diagram mapping the Computation, Communication, and Control layers.
2. **Sensing and Actuation (Control):** Identify the physical continuous-time variables you must measure (e.g., gyroscopic pitch, roll, yaw) and the actuators you will drive (e.g., PWM motor controllers). Map these physical properties to the VirtMCU Sensor/Actuator Abstraction Layer (SAL/AAL).
3. **Network Design (Communication):** The drone will utilize a distributed architecture where each motor has its own microcontroller, communicating with a central flight computer. You will design the message format required to transmit control data over the VirtMCU deterministic bus, ensuring you account for the timing strictness required by the physics engine.
4. **Edge AI (Computation):** Describe how a predictive AI model (simulated within the VirtMCU environment) could be deployed at the edge to dynamically adjust the PID feedback gains in response to a simulated physical disturbance, such as a sudden loss of lift on one rotor.

By mapping this architecture into the deterministic confines of VirtMCU, you will transition your understanding from isolated software algorithms to holistic, continuous-time control-computing codesign.

# Chapter 10: Co-Simulation and Virtual Time

## 10.1 Introduction: The Illusion of Time in Computing

In the first nine chapters of this book, we constructed a comprehensive understanding of computing nodes from the digital logic level up to the Real-Time Operating System (RTOS) and the overarching architecture of Cyber-Physical Systems (CPS 2.0). We established that in a true CPS, time is not merely a performance metric—it is a strict criterion for logical correctness. If an autonomous vehicle's braking controller computes the mathematically perfect stopping force but delivers it 50 milliseconds too late, the result is catastrophic physical failure.

However, a fundamental impedance mismatch exists between the cyber and physical domains regarding the concept of time. In the physical world, time is a continuous, monotonically advancing variable governed by differential equations and the laws of physics. Physical components—such as masses, springs, induction motors, and fluids—react instantly and continuously to applied forces.

Conversely, the cyber domain is entirely discrete. Software engineers generally view time as a sequence of logically ordered events. A program executes line one, then line two, then line three. The exact wall-clock time it takes to transition between these lines is abstracted away by the hardware pipeline, the memory hierarchy, and the operating system scheduler. When software is tested in standard development environments, it runs as fast as the host processor allows, completely decoupled from the real-world timeline.

When designing a CPS 2.0 architecture, we cannot test the cyber algorithms and the physical mechanics in isolation. We must evaluate them together to observe their interactions. This requirement introduces the discipline of **co-**

**simulation**, which fuses a discrete-event software simulator with a continuous-time physics engine. To achieve this without introducing massive timing errors, we must strip the concept of “time” away from the host computer’s operating system and replace it with a mathematically rigorous, deterministic **virtual time** clock. This chapter explores the theory of discrete-event simulation, the implementation of Transaction Level Modeling (TLM), and the mechanisms VirtMCU uses to achieve deterministic co-simulation for CPS.

## 10.2 The Problem with Standard Emulation

To test embedded firmware before physical hardware is manufactured, engineers have traditionally relied on software emulators, the most famous being QEMU (Quick Emulator). QEMU utilizes Dynamic Binary Translation (DBT) to convert blocks of guest instruction set architecture (e.g., ARM AArch64) into the host instruction set architecture (e.g., x86\_64) on the fly. This provides incredibly fast execution, allowing developers to boot entire guest operating systems in seconds.

However, standard emulation environments are built for functional software testing, not for Cyber-Physical System validation. They suffer from a critical flaw: **wall-clock scheduling jitter**.

In a standard QEMU setup, the simulated processor’s timer peripherals and interrupt controllers are tied to the host operating system’s real-time clock. When the guest firmware configures a hardware timer to fire an interrupt every 10 milliseconds, QEMU asks the Linux or Windows host OS to set a software timer. The host OS is a general-purpose, multitasking system. It is simultaneously managing disk I/O, network traffic, background services, and graphical user interfaces. If the host OS scheduler decides to preempt the QEMU process to handle an incoming gigabit network packet, the emulator is physically paused.

When QEMU is eventually rescheduled and resumes execution, it checks the host clock, realizes that 15 milliseconds have passed, and immediately fires the guest’s 10ms timer interrupt. From the perspective of the guest firmware, the hardware timer took 15 milliseconds to count to 10 milliseconds. This timing

violation propagates through the guest's RTOS scheduler, delaying the execution of critical control loops.

If this emulator is loosely coupled to a physics engine simulating a drone, the physics engine will continue to advance its continuous-time dynamics while the emulator is paused. When the emulator finally sends its delayed actuator command, the drone's physical state has already degraded, causing the simulated drone to crash. The drone crashed not because the student's Proportional-Integral-Derivative (PID) control math was incorrect, but because the host OS injected scheduling jitter into the simulation. This lack of reproducibility renders standard emulators useless for hard real-time CPS verification.

## 10.3 Principles of Discrete-Event (DE) Simulation

To build a simulation environment devoid of wall-clock jitter, we must formalize how digital hardware simulates time. At the system level, hardware is typically modeled using **Discrete-Event (DE) Simulation** semantics. SystemC, the industry standard C++ library for hardware modeling, utilizes a specialized DE simulation kernel.

Unlike continuous-time physics, where variables change constantly, digital hardware state changes only at discrete points in time, typically triggered by clock edges or signal transitions. A Discrete-Event kernel models this by maintaining a centralized **Event Queue** ordered strictly by virtual timestamps.

The simulation does not execute continuously; it jumps from one event timestamp to the next. The DE simulation semantics rely heavily on a two-phase execution model known as the **Evaluate-Update paradigm**. This paradigm is designed to accurately simulate the parallel nature of hardware on a sequential host CPU.

1. **Evaluate Phase:** The simulation kernel selects all processes (threads or methods) that are sensitive to the current timestamp or triggered by events at this exact virtual time. The kernel executes these processes one by one. Crucially, during this phase, any assignments made to hardware

signals *do not immediately update the signal's value*. Instead, the new values are scheduled to be updated in the future.

2. **Update Phase:** Once all runnable processes have finished evaluating, the kernel applies all the scheduled signal updates simultaneously.

This two-phase approach guarantees deterministic behavior. If Process A and Process B both read Signal X and write to Signal Y at the exact same virtual nanosecond, the Evaluate-Update paradigm ensures that both processes read the old value of X before either can update Y. The order in which the host CPU executes Process A and Process B during the Evaluate phase becomes completely irrelevant, ensuring reproducibility.

### 10.3.1 Delta Cycles

A unique artifact of the DE simulation kernel is the **delta cycle**. A delta cycle represents an infinitesimal step in virtual time. If the Update phase alters a signal, that alteration might trigger a *new* event that is scheduled to occur at the *exact same virtual timestamp* (zero-delay).

When this happens, the simulation kernel does not advance the virtual clock. Instead, it initiates a new Evaluate phase followed by a new Update phase, all within the same virtual nanosecond. This loop of Evaluate-Update stages at a constant virtual time is called a delta cycle. Delta cycles allow the simulator to resolve zero-delay combinational logic paths (such as a cascade of logic gates) in a causally correct sequence without advancing physical time. Only when the event queue contains no more zero-delay events does the kernel advance the virtual clock to the timestamp of the next scheduled event.

## 10.4 Transaction Level Modeling (TLM) and Time

While cycle-accurate DE simulation is highly precise, simulating every individual clock edge, pin transition, and delta cycle for a massive System-on-Chip (SoC) interconnect is incredibly slow. To achieve the simulation speeds necessary to boot an RTOS and run complex CPS software, architects rely on **Transaction Level Modeling (TLM)**.

TLM abstracts away the low-level physical pins and wires of the SoC interconnect (such as AMBA AXI or APB). Instead of simulating the clock-by-clock toggling of `VALID` and `READY` signals, TLM uses object-oriented method calls (like C++ functions) to transfer data between initiators (e.g., CPUs, DMA controllers) and targets (e.g., memories, UARTs).

In the widely adopted OSCI TLM 2.0 standard, these interactions are modeled as transactions containing a payload (the data, address, and command). TLM provides multiple levels of timing abstraction:

- **Untimed (UT):** Used for pure functional software development. Transactions execute instantly in zero virtual time. There is no concept of bus contention or propagation delay.
- **Loosely Timed (LT):** Used for RTOS and driver development. Transactions are annotated with estimated delays, but the simulation relies on **temporal decoupling**. In LT, the CPU is allowed to run ahead of the global virtual clock for a specific “quantum” of time. It accumulates its own local time and only synchronizes with the rest of the system when its quantum expires. This provides immense simulation speedups by avoiding constant context switching with the DE kernel.
- **Approximately Timed (AT):** Used for architectural exploration. This model breaks transactions into multiple phases to accurately simulate bus contention, pipelining, and backpressure. For example, a write transaction is modeled using a four-phase protocol: *Begin Request*, *End Request*, *Begin Response*, and *End Response*. The AT model interacts heavily with the DE kernel's event queue to enforce these timing boundaries, resulting in highly accurate but slower simulations.

For VirtMCU and CPS co-simulation, we must strike a careful balance. We need the speed of Loosely Timed TLM to execute millions of instructions, but we require strict timing boundaries to interact with continuous-time physics.

## 10.5 Co-Simulation: Bridging the Cyber and Physical

Co-simulation in CPS 2.0 is the simultaneous execution of two fundamentally different simulation domains: the discrete-event cyber simulator (e.g., VirtMCU/QEMU executing firmware) and the continuous-time physical simulator (e.g., a MuJoCo physics engine simulating kinematics).

Because these two engines possess entirely different mathematical foundations and concepts of time, they cannot run asynchronously. They must be coupled through an **External Clock Authority**.

In a co-simulation environment, time is quantized into strictly bounded **macro-steps** (e.g., 1 millisecond).

1. The External Clock Authority grants permission for the physics engine to integrate its differential equations forward by 1ms.
2. Simultaneously, it grants the cyber simulator permission to execute exactly 1ms worth of virtual CPU instructions and timer ticks.
3. At the end of the 1ms macro-step, both simulators pause.
4. The physical domain samples its state (e.g., joint angles, velocities) and maps these continuous values to discrete sensor registers (the Sensor Abstraction Layer, SAL).
5. The cyber domain samples its memory-mapped output registers (the Actuator Abstraction Layer, AAL) and maps these discrete commands (e.g., PWM duty cycles) into physical forces.
6. The state variables are exchanged via zero-copy shared memory, and the External Clock Authority issues the command to begin the next 1ms macro-step.

This lock-step synchronization guarantees strict causality. The software cannot run ahead of the physics, and the physics cannot outpace the software. However, enforcing this macro-step boundary on a fast binary-translating emulator like QEMU requires aggressive structural modifications to how the emulator manages its internal execution loop.

## 10.6 VirtMCU Clock Slaving

VirtMCU solves the emulator timing problem by entirely overriding QEMU's internal clock infrastructure. Instead of querying the host OS wall-clock to manage guest timers, VirtMCU utilizes a dynamic QOM (QEMU Object Model) plugin architecture to enforce deterministic virtual time via the `hw/rust/backbone/clock` module.

VirtMCU implements a concept known as **Cooperative Time Slaving**. Under this architecture, the virtual CPU execution loop is subordinated to an external synchronizer. VirtMCU provides two primary modes of operation to achieve this determinism: `slaved-suspend` and `slaved-icount`.

### 10.6.1 Slaved-Suspend Mode

In `slaved-suspend` mode, VirtMCU is designed for high-throughput temporal decoupling. The CPU executes instructions as fast as the host hardware allows. However, the `hw/rust/backbone/clock` module intercepts the QEMU main loop and enforces a strict barrier at the macro-step boundary. If the external physics engine dictates a 1ms step, the CPU is allowed to execute instructions until its accumulated internal virtual time reaches 1ms. Once it hits this exact boundary, the virtual CPU thread is physically suspended (blocked) by the VirtMCU backbone. It will remain suspended indefinitely until the external physics engine completes its own 1ms integration and sends an IPC (Inter-Process Communication) signal to resume for the next step.

### 10.6.2 Slaved-Icount Mode

While `slaved-suspend` forces alignment at the macro-step, the virtual time within that step might still be somewhat abstracted. For cycle-accurate, microsecond-precision validation, VirtMCU utilizes `slaved-icount` mode.

In `slaved-icount` (Instruction Count) mode, virtual time is completely decoupled from any concept of physical time or host execution speed. Instead, time advances purely as a strict mathematical function of the number of retired



instructions. If the simulated embedded processor is defined to run at 100 MHz, the execution of exactly one machine instruction advances the global virtual clock by precisely 10 nanoseconds.

Because the `icount` virtual clock is the absolute source of truth for the entire SoC, all internal hardware timers, watchdogs, and peripheral baud-rate generators are driven directly by this instruction-derived clock. If the firmware configures a 10ms hardware timer, the VirtMCU backbone calculates that exactly 1,000,000 instructions must retire before the interrupt can fire. The execution loop will execute exactly 1,000,000 instructions, stop, inject the hardware IRQ into the NVIC, and then execute the first instruction of the ISR.

This creates a perfectly deterministic, 100% reproducible execution trace. If you run the simulation on a massive server, it might finish in two seconds. If you run it on a slow, 10-year-old laptop, it might take five minutes. However, in both scenarios, the simulated control loop will execute the exact same number of instructions, and the hardware timer will fire on the exact same instruction boundary.

## **10.7 Determinism and Reproducibility in CPS Validation**

Why go to such extreme lengths to eliminate jitter? In the context of Cyber-Physical Systems, non-deterministic timing behaviors can easily mask fatal architectural flaws.

Consider the operational dynamics of a Cyber-Physical Virtual Power Plant (VPP), which coordinates hundreds of Distributed Energy Resources (DERs) via complex communication networks. To maintain grid stability during an emergency frequency support event, a centralized VPP must issue precision demand response commands to isolated loads and energy storage systems. The timing of these command deliveries is critical. If a command packet to a storage inverter is delayed by communication jitter or processing stalls, the response power changes unexpectedly.

As highlighted in distributed system literature, due to the heterogeneity of the information transmitted, changes in the information arrival state result in unequal covariate coefficients in the system's economic and stability models. If a response command to DER A is delayed, it might cause a system cost increment of \$30, whereas delaying a command to DER B might only cost \$10. If the validation environment suffers from wall-clock scheduling jitter, the simulations will produce different economic and physical stability profiles on every single run. It becomes impossible to mathematically verify the control algorithms or optimize the communication scheduling policies because the baseline data is corrupted by simulation noise.

By deploying VirtMCU's deterministic clock slaving, engineers can create a multi-node simulation where the network delays, interrupt latencies, and RTOS context switches are absolutely rigid. If a priority inversion bug exists in the RTOS that causes a 50-microsecond delay in the VPP frequency response task, VirtMCU guarantees that this 50-microsecond delay will appear consistently in every simulation, allowing the engineer to attach a debugger, halt time, and inspect the exact state of the pipeline and memory map.

## 10.8 VirtMCU Homework: Control Loop Simulation

Understanding the theoretical impact of simulation jitter is different from observing it directly in an active execution pipeline. To fully internalize the difference between traditional software emulation and strict co-simulation, this chapter's homework requires you to deploy a real-time control loop within the VirtMCU environment.

1. **Firmware Implementation:** You will be provided with a bare-metal C framework. You must write an Interrupt Service Routine (ISR) that is triggered by a hardware timer configured to fire exactly every 10.00 milliseconds. Inside the ISR, you will increment a global volatile counter and read the current value of a free-running microsecond hardware counter to measure the elapsed time since the last invocation.

2. **Execution under Jitter:** First, compile the firmware and load it into a standard, unmodified QEMU ARM emulator running on your host OS. To simulate a noisy industrial environment, you will concurrently run a heavy, multi-threaded CPU stress test in the background of your host OS. You will print the delta between hardware timer invocations. You will observe significant variance (jitter) as the host OS scheduler randomly preempts the emulator to service the stress test.
3. **Deterministic `slaved-icount` Execution:** Next, load the exact same compiled binary into the VirtMCU framework configured in `slaved-icount` mode. Launch the same background CPU stress test on your host OS.
4. **Analysis:** You will observe that regardless of the immense load on your physical host machine, the delta between ISR invocations in VirtMCU will remain mathematically perfect. Because the virtual clock is slaved exclusively to the instruction count and not the wall clock, the VirtMCU environment completely isolates the cyber-physical state from external computing noise. You must submit a brief report containing the histogram of your timing measurements, proving empirically the necessity of virtual time in CPS 2.0 design.

# Chapter 11: The Sensor/Actuator Abstraction Layer (SAL/AAL)

## 11.1 Introduction: The Cyber-Physical Boundary

In Chapter 10, we established the temporal foundation required for Cyber-Physical System (CPS) co-simulation. We demonstrated that standard software emulators fail to provide the rigorous timing guarantees required for CPS validation due to host operating system scheduling jitter. By implementing Cooperative Time Slaving and deterministic `slaved-icount` execution in VirtMCU, we successfully lock-stepped the discrete-time cyber domain (the processor) with the continuous-time physical domain (the physics engine).

However, synchronizing time is only half of the co-simulation equation. Even if the virtual processor and the physics engine advance their clocks in perfect harmony, they do not speak the same language. A physics engine like MuJoCo calculates the state of the world using continuous floating-point variables representing Newtonian mechanics—positions, velocities, accelerations, and forces. Conversely, the microcontroller firmware operates purely in the discrete digital domain, reading and writing binary integers to memory-mapped I/O (MMIO) registers.

A simulated drone’s flight controller cannot magically “read” a floating-point quaternion representing its physical orientation directly from the physics engine’s memory. In reality, the drone’s processor communicates with a physical Micro-Electro-Mechanical System (MEMS) gyroscope over an I2C or SPI serial bus. The gyroscope senses the continuous physical rotation, converts it to an analog voltage, digitizes it, and stores it in an internal hardware register, which the CPU then reads.

To bridge this semantic gap in our simulation environment, we must construct a rigorous translation interface. In the VirtMCU framework, this is achieved through the **Sensor Abstraction Layer (SAL)** and the **Actuator Abstraction**

**Layer (AAL).** These layers sit between the raw MMIO hardware emulation and the external physics engine, translating continuous-time physical variables into discrete-time cyber registers (and vice versa) while injecting real-world artifacts such as quantization error, sensor noise, and actuator saturation.

## 11.2 Mathematical Modeling of Physical Systems

Before we can translate physical phenomena into digital data, we must formalize how the continuous physical world is modeled. The physical domain is governed by the laws of physics, which yield governing differential equations and constitutive relationships.

For mechanical systems, engineers typically employ lumped-parameter models analyzed via classical mechanics, specifically Newtonian, Lagrangian, or Hamiltonian equations of motion. In a lumped-parameter model, the spatially distributed properties of a physical system (such as mass, damping, and stiffness) are approximated as discrete, concentrated components (ideal masses, dashpots, and springs) interacting in continuous time.

The state of such a physical system at any instant is represented by a state vector  $x(t)$ , which evolves according to a set of non-linear differential equations:  $\dot{x}(t) = f(x(t), u(t), t) + w(t)$  where  $u(t)$  represents the continuous control input (actuation forces) and  $w(t)$  represents stochastic perturbations or environmental disturbances.

For example, the dynamics of an aerial vehicle require tracking both the translational positions and velocities in the inertial frame, as well as the rotational Euler angles (roll  $\phi$ , pitch  $\theta$ , yaw  $\psi$ ) and angular rates ( $p$ ,  $q$ ,  $r$ ) in the body frame. To simulate this, the continuous-time physics engine numerically integrates these differential equations using advanced solvers (such as Runge-Kutta methods).

The primary objective of the SAL and AAL is to discretize these continuous variables  $x(t)$  for the cyber domain's sensors, and to convert the cyber domain's discrete outputs into the continuous control inputs  $u(t)$  applied to the physical model.

## 11.3 The Sensor Abstraction Layer (SAL)

The Sensor Abstraction Layer (SAL) is responsible for intercepting the exact state variables  $x(t)$  computed by the physics engine and transforming them into the binary formats expected by the firmware's device drivers. This process replicates the entire signal chain of a physical transducer, signal conditioning circuit, and Analog-to-Digital Converter (ADC).

### 11.3.1 Transduction and Analog-to-Digital Conversion (ADC)

In the physical world, sensors rely on electrostatic, electromagnetic, or piezoresistive transducers to convert a physical property (such as pressure, temperature, or acceleration) into a proportional analog electrical signal (voltage or current).

Once transduced to an analog voltage, the signal must be digitized. The SAL models the behavior of a hardware ADC through two primary mathematical operations: **Sampling** (discretization in time) and **Quantization** (discretization in amplitude).

1. **Sampling:** The SAL does not update the sensor registers continuously. Instead, it updates them at a specific frequency, defined by the sensor's sampling rate. If the firmware attempts to read the sensor register faster than the sampling rate, it will simply read the same stale value twice, exactly as it would on real hardware.
2. **Quantization:** A physical ADC maps a continuous voltage range into a finite number of discrete binary bins. An  $N$ -bit ADC provides  $2^N$  possible output values. The resolution of the sensor, or the Least Significant Bit (LSB) weight, is calculated as: 
$$\text{Resolution} = \frac{V_{\max} - V_{\min}}{2^N}$$
 The SAL takes the continuous floating-point variable from the physics engine, scales it according to the sensor's physical sensitivity (e.g.,  $16.4 \text{ LSB per degree/second}$  for a gyroscope), and truncates the floating-point mantissa to force the value into an  $N$ -bit integer. This intentional introduction of **quantization error** is crucial, as it fundamentally limits the precision of the feedback loop in the cyber domain.

### 11.3.2 Simulating Sensor Noise and Stochastic Perturbations

A naive co-simulation environment would simply pass the quantized, ground-truth physics value to the firmware. However, in reality, physical systems exhibit randomness and are perturbed by random perturbations, such as noise, interference, and other stationary and nonstationary stochastic processes.

If a control algorithm is developed and tested against perfect, noise-free sensor data, it will almost certainly fail when deployed to physical hardware. Therefore, the SAL must actively inject noise into the signal before passing it to the MMIO registers.

The SAL models sensor imperfections using several stochastic noise profiles:

- **White Noise (Gaussian):** High-frequency thermal and electrical noise is modeled by generating random variables from a Gaussian distribution with a zero mean and a specific variance ( $\sigma^2$ ). This noise is added to the continuous signal prior to quantization.
- **Bias (Offset):** Manufacturing imperfections mean that a sensor will rarely read exactly zero when the physical state is zero. The SAL adds a constant (or slowly varying) bias offset. For example, a simulated accelerometer resting flat on a table might report  $0.02g$  on the X-axis instead of  $0.00g$ .
- **Random Walk (Drift):** MEMS gyroscopes famously suffer from bias instability, where the sensor bias wanders randomly over time. The SAL simulates this by accumulating small Gaussian random variables at each macro-step, creating a Brownian motion drift profile that forces the student's firmware to implement sensor fusion algorithms (like a Kalman filter) to estimate and subtract the drift.

By mathematically simulating these electrostatic and electromagnetic transducer characteristics, the SAL forces the software engineer to design robust cyber algorithms that can handle the harsh realities of the physical world.

## 11.4 The Actuator Abstraction Layer (AAL)

While the SAL handles data flowing from the physical world to the cyber world, the Actuator Abstraction Layer (AAL) handles the reverse: translating the discrete commands generated by the firmware into continuous physical forces, torques, or voltages  $u(t)$  required by the physics engine.

### 11.4.1 Pulse Width Modulation (PWM) and DAC

In modern embedded systems, it is rare for a microcontroller to output a true continuous analog voltage using a Digital-to-Analog Converter (DAC) to drive heavy loads. Instead, processors control physical actuators—such as heaters, hydraulic valves, and electric motors—using **Pulse Width Modulation (PWM)**.

A PWM controller is a hardware peripheral that generates a high-frequency digital square wave. The firmware configures the PWM peripheral by writing to two MMIO registers: the Period Register (which defines the frequency of the wave) and the Compare Register (which defines how long the wave remains HIGH during each period). The ratio of the HIGH time to the total period is the **duty cycle**.

To an electrical load like a DC motor, the high-frequency switching of the PWM signal is naturally low-pass filtered by the motor's internal inductance and mechanical inertia. Thus, the motor responds to the *average* voltage of the PWM signal, which is directly proportional to the duty cycle.

The AAL in VirtMCU intercepts the firmware's writes to the MMIO PWM Compare Registers. It mathematically computes the duty cycle percentage (e.g., a Compare value of 250 out of a Period of 1000 yields a 25% duty cycle). If the simulated motor is connected to a 12V power supply, the AAL translates this 25% duty cycle into an effective continuous voltage of 3.0V.

### 11.4.2 Actuator Dynamics and Saturation

A common mistake in software engineering is assuming that an actuator responds instantly to a command. In the physical domain, actuators have their



own internal governing differential equations.

Consider an electric drive actuated by a geared permanent magnet DC motor. The motor is not an instantaneous torque generator; it is a complex electro-mechanical system. Its electrical dynamics are governed by its armature resistance ( $r_a$ ) and inductance ( $L_a$ ). When the AAL applies the 3.0V effective voltage, the current  $i_a(t)$  ramps up according to the electrical time constant ( $L_a / r_a$ ). The generated mechanical torque is proportional to this current ( $T = k_t i_a$ ).

Simultaneously, the motor's mechanical dynamics are opposed by the back-electromotive force (back-EMF), which grows as the motor's angular velocity  $\omega_r$  increases, as well as mechanical viscous friction ( $B_m$ ) and the moment of inertia ( $J$ ).

The AAL must account for these dynamics. While the primary physics engine (MuJoCo) handles the rigid body dynamics of the drone or robot, the AAL often computes a localized, high-speed sub-simulation of the actuator's internal electro-magnetic state to produce the final, accurate torque applied to the mechanical joint.

Furthermore, the AAL enforces **control limits (saturation)**. Control limits  $u_{\min} \leq u(t) \leq u_{\max}$  are described by real-valued bounded functions. If a PID controller mathematically demands 50 Amps of current to stabilize a drone, but the physical motor driver can only source 10 Amps, the AAL physically clips the output at 10 Amps. This saturation can cause "integral windup" in poorly designed firmware, providing an excellent pedagogical lesson in non-linear control limits.

## 11.5 Co-Simulation via Zero-Copy Shared Memory

The SAL and AAL perform intense mathematical transformations, but they must exchange this data with the external physics engine at every macro-step of the simulation. Because the physics engine (e.g., MuJoCo) and the cyber emulator (VirtMCU) run as completely separate operating system processes, they must utilize Inter-Process Communication (IPC).

Standard IPC mechanisms, such as TCP/IP sockets or Unix pipes, involve copying data from the user space of one process, into the host kernel space, and then back into the user space of the receiving process. When running a hard real-time co-simulation executing a 1-millisecond macro-step, the latency and context-switching overhead of socket-based IPC destroys simulation throughput.

To resolve this, VirtMCU and the physics engine are coupled through **zero-copy shared memory**. Using the POSIX `shm_open` and `mmap` system calls, the host operating system allocates a single contiguous block of physical RAM that is mapped simultaneously into the virtual address space of both the VirtMCU process and the MuJoCo process.

The shared memory region is structured into three distinct zones:

1. **Synchronization Primitives:** Atomic lock variables and condition variables used by the External Clock Authority to orchestrate the macro-step lock-stepping discussed in Chapter 10.
2. **Physics  $\rightarrow$  Cyber Zone:** Written exclusively by MuJoCo. Contains arrays of continuous variables (forces, angles, positions). The SAL reads directly from these memory addresses, applies its noise and quantization models, and writes the results to the emulated MMIO registers.
3. **Cyber  $\rightarrow$  Physics Zone:** Written exclusively by VirtMCU. Contains the continuous control inputs  $u(t)$  (e.g., applied torques and voltages) computed by the AAL. MuJoCo reads directly from these addresses during its integration step.

Because the data never passes through the host kernel, the exchange occurs in tens of nanoseconds, allowing the co-simulation to execute millions of macro-steps per second without IPC bottlenecking.

## 11.6 Physical Modeling: The Inverted Pendulum

To solidify the concepts of SAL, AAL, and lumped-parameter modeling, we turn to a classic benchmark in control theory: the inverted pendulum on a cart. This

system perfectly encapsulates the multi-axis coupling and instability inherent to cyber-physical systems.

The physical system consists of a motorized cart of mass  $M$  that moves horizontally along a linear track (position  $x$ ). A pendulum of mass  $m$  and length  $l$  is attached to the cart via a frictionless hinge. The angle of the pendulum relative to the vertical upright position is denoted by  $\theta$ .

The goal of the cyber system is to maintain the pendulum in the unstable upright position ( $\theta = 0$ ) by applying a horizontal force  $u(t)$  to the cart. If the pendulum begins to fall forward, the cart must accelerate forward to slide the base “under” the falling mass to catch it.

Using Newtonian mechanics or the Lagrange equations of motion, we derive the non-linear governing differential equations for this system. Summing the forces on the cart and the moments about the pivot, and accounting for viscous friction  $B_m$ , yields a coupled system of equations. For example, the rotational dynamics of the inverted arm on the moving platform can be expressed as: 
$$ml^2 \ddot{\theta} + ml \ddot{x} \cos \theta - mgl \sin \theta + B_m \dot{\theta} = 0$$
 (Note: Gravity  $g$  creates a destabilizing moment when the pendulum deviates from the vertical, hence the  $-mgl \sin \theta$  term when linearized around the upright position).

To design a linear control law (like a PID controller), engineers typically linearize these non-linear differential equations around the unstable equilibrium point ( $\theta = 0, \dot{\theta} = 0, x = 0, \dot{x} = 0$ ). For small angles, we apply the small-angle approximations:  $\sin \theta \approx \theta$  and  $\cos \theta \approx 1$ .

In the co-simulation environment:

1. **The Physics Engine (MuJoCo)** numerically integrates the full, non-linear differential equations in continuous time.
2. **The SAL** reads the continuous cart position  $x(t)$  and pendulum angle  $\theta(t)$  from shared memory. It adds Gaussian noise to simulate optical encoder jitter, quantizes the angles into 12-bit integers, and writes them to the virtual Quadrature Encoder Interface (QEI) MMIO registers in VirtMCU.

3. **The Firmware** executes its control loop, reads the noisy integer angles, computes the proportional, integral, and derivative errors, and writes a duty cycle value to the PWM MMIO register.
4. **The AAL** reads the PWM register, converts the duty cycle to a continuous horizontal force  $u(t)$  (accounting for motor saturation limits), and writes this force back to the shared memory.
5. **The Physics Engine** applies  $u(t)$  to the cart mass  $M$  for the next integration step.

## 11.7 VirtMCU Homework: Inverted Pendulum Simulation

In standard software engineering curricula, writing a sorting algorithm allows you to verify your logic with an array of integers. In Cyber-Physical Systems, you verify your logic by preventing a physical structure from collapsing.

For this chapter's homework, you will utilize the VirtMCU FirmwareStudio linked with the MuJoCo physics engine via the SAL/AAL shared memory interface. You will write the firmware required to balance the inverted pendulum described in Section 11.6.

### Your Assignment:

1. **Driver Initialization:** Write C code to initialize the memory-mapped virtual peripherals. Configure the simulated ARM Cortex-M hardware timer to generate an interrupt exactly every 5.0 milliseconds. This sets your discrete control loop period ( $T_s$ ).
2. **Sensor Acquisition (SAL):** Inside the timer Interrupt Service Routine (ISR), read the raw 16-bit integer values from the virtual Quadrature Encoder registers mapping to the cart position and pendulum angle. Convert these raw integers back into standard floating-point engineering units (meters and radians), accounting for the quantization scaling documented in the hardware specification.
3. **Digital Control Law:** Implement a cascaded Proportional-Integral-Derivative (PID) control algorithm. Your outer loop should compute a

target angle to move the cart to a desired setpoint, while the inner loop computes the stabilizing force required to maintain the pendulum at that target angle. You must manually tune the  $k_p$ ,  $k_i$ , and  $k_d$  feedback gains to achieve stability.

4. **Actuation (AAL):** Convert your computed force request into a PWM duty cycle integer. Write this integer to the virtual motor controller's MMIO Compare Register.
5. **Co-Simulation Execution:** Launch the VirtMCU environment in `slaved-icount` mode alongside MuJoCo. Observe the 3D rendering of the cart. If your mathematics are correct, the cart will catch the falling pendulum and balance it perfectly.
6. **Robustness Testing:** The simulation framework includes a Python script to inject random physical disturbances (simulated wind gusts) into the MuJoCo environment. You must document how your PID control law rejects these unmodeled stochastic perturbations and prevents the pendulum from falling.

# Chapter 12: The RISC-V Architecture and Instruction Set

## 12.1 Introduction: The Need for Open Architectures in CPS

In the preceding chapters, we explored how standard microprocessors interact with the physical world. Through the Sensor/Actuator Abstraction Layer (SAL/AAL) discussed in Chapter 11, we demonstrated how continuous-time physical phenomena are translated into discrete-time memory-mapped I/O (MMIO) registers. However, when developing advanced Cyber-Physical Systems (CPS)—such as high-frequency drone flight controllers or ultra-low-latency robotic actuators—the processing overhead of reading MMIO registers, shifting bits, and executing standard arithmetic can become a bottleneck.

To achieve maximum performance, computer architects often need to modify the processor itself, adding custom hardware accelerators and new machine instructions tailored specifically for the physical application. Historically, this has been nearly impossible for academic researchers and independent engineers because architectures like ARM and x86 are strictly proprietary.

This paradigm shifted dramatically with the introduction of **RISC-V**. RISC-V is an entirely open-source specification for a reduced instruction set computer architecture. Because it is open and royalty-free, hardware designers are explicitly encouraged to modify the silicon and extend the instruction set to meet the unique requirements of their specific domain. In this chapter, we will examine the RISC-V architecture, explore its elegant instruction set, and discuss how its open nature enables the creation of custom silicon for Cyber-Physical Systems.

## 12.2 The Origins and Philosophy of RISC-V

The RISC-V architecture was publicly announced in 2011 after being developed at the University of California, Berkeley, by a team including Yunsup Lee, Krste Asanović, David A. Patterson, and Andrew Waterman. The “V” represents the Roman numeral five, signifying that this was the fifth major RISC architectural design project undertaken at UC Berkeley.

Unlike many commercial architectures that are burdened by decades of legacy compatibility requirements, the RISC-V project began as a clean sheet design. The architects established several principal goals to ensure the ISA could scale across the entire spectrum of computing:

- **Broad Scalability:** The architecture must efficiently support everything from ultra-low-power microcontrollers operating at the IoT edge to extreme-scale computing nodes in warehouse-scale cloud server farms.
- **Modularity and Extensions:** The base instruction set is deliberately minimal. Advanced capabilities are provided via optional ISA extensions, which support features like floating-point mathematics, atomic memory operations, and hardware multiplication and division.
- **Privilege Levels:** Additional extensions support privileged execution modes (similar to the x86 and ARM models discussed in earlier chapters) necessary to run secure operating systems.
- **Code Density:** To accommodate memory-constrained embedded systems, RISC-V supports a compressed instruction set extension. This allows 16-bit versions of common 32-bit instructions to be freely interspersed with standard 32-bit instructions in memory, significantly reducing the overall code footprint.
- **Address Space Expansion:** The architecture provides optional extensions to support 64-bit, and even 128-bit, processor word sizes and paged virtual memory for massive multiprocessing configurations.

Additionally, like modern ARM processors, the RISC-V architecture supports **bi-endianness**. This allows the operating system to select between big-endian or little-endian mode under software control, though most modern operating systems choose little-endian mode by default.

## 12.3 The RISC-V Programmer's Model

To program a RISC-V processor, a software engineer must understand its architectural state. At the core of the RISC-V base ISA is a remarkably orthogonal and simple register set.

The base architecture provides 32 general-purpose registers, named `x0` through `x31`. In the standard 32-bit architecture (RV32), each of these registers is exactly 32 bits wide. Registers `x1` through `x31` are completely unrestricted and have no special functions forced upon them by the processor hardware.

However, the `x0` register is unique: it is hardwired to the constant value zero. Any attempt to read `x0` will always return zero, and any value written to `x0` is immediately discarded by the hardware. As we will see in the next section, this hardwired zero register allows the instruction set to be drastically simplified.

### 12.3.1 The Application Binary Interface (ABI)

Although the processor hardware treats registers `x1` through `x31` as functionally interchangeable, building complex software requires strict agreements on how registers are used across different functions and libraries. Software compatibility requires that we specify which register serves as the stack pointer, which registers contain function arguments, and which contain return values.

These standard rules form the **Application Binary Interface (ABI)**. By standardizing register usage and calling conventions, the ABI ensures that compiled C code can successfully link with handwritten assembly routines on any compatible implementation. To aid programmers, RISC-V assemblers recognize alternate ABI names for the registers. The primary ABI designations are:

- **ra (x1)**: The Return Address register, used to hold the address to jump back to after a function call completes.
- **sp (x2)**: The Stack Pointer, pointing to the top of the current activation record in memory.



- **a0–a7 ( x10–x17 )**: The Argument registers, used to pass parameters into functions and to return results back to the caller.
- **t0–t6** : Temporary registers, which are volatile (caller-saved).
- **s0–s11** : Saved registers, which are non-volatile (callee-saved).

## 12.4 The Instruction Set and Pseudo-Instructions

The philosophy of RISC-V is to keep the hardware as simple as possible. Consequently, the base ISA lacks many distinct instructions found in other architectures (such as explicit `MOV` or `NOT` operations). Instead of forcing the hardware to decode dozens of unique opcodes, the architecture relies on the assembler program to synthesize missing operations using combinations of the base instructions and the hardwired `x0` register.

These synthesized aliases are known as **pseudo-instructions**. When the programmer types a pseudo-instruction, the assembler automatically translates it into the corresponding base RISC-V machine instruction. Table 12.1 outlines some of the most common pseudo-instructions:

- **nop (No Operation)**: Translated to `addi x0, x0, 0`. It adds zero to zero and discards the result in `x0`, consuming a clock cycle without altering the machine state.
- **mv rd, rs (Move)**: Translated to `addi rd, rs, 0`. It copies the value of the source register `rs` into the destination `rd` by adding zero to it.
- **not rd, rs (Bitwise NOT)**: Translated to `xori rd, rs, -1`. Exclusive-ORing a value with all 1s (-1) perfectly inverts all the bits.
- **neg rd, rs (Negate)**: Translated to `sub rd, x0, rs`. It mathematically negates `rs` by subtracting it from the `x0` register (zero).
- **j offset (Unconditional Jump)**: Translated to `jal x0, offset`. It jumps to the offset but discards the return address into `x0`.
- **beqz rs, offset (Branch if Equal to Zero)**: Translated to `beq rs, x0, offset`.

By pushing these translations into the software assembler, the silicon required to decode instructions in the CPU is significantly reduced, lowering power

consumption and increasing maximum clock frequencies.

### 12.4.1 Immediate Values and the Sign-Extension Quirk

Loading 32-bit constants or absolute 32-bit memory addresses into a register poses a challenge because RISC-V instructions are only 32 bits wide themselves. There is no physical room to pack a 32-bit opcode alongside a 32-bit constant.

To solve this, RISC-V breaks the operation into two steps using a 20-bit upper immediate and a 12-bit lower immediate. The `lui` (Load Upper Immediate) instruction loads a 20-bit value into the highest 20 bits of a register, clearing the lower 12 bits. The `auipc` (Add Upper Immediate to PC) performs a similar task but adds the 20-bit value to the Program Counter, enabling position-independent code. The programmer then uses a subsequent instruction—like `addi` or a load/store offset—to supply the remaining lower 12 bits.

However, this two-step process introduces a subtle but critical mathematical quirk. In RISC-V, the 12-bit immediate value in the `addi` instruction is always treated as a signed two's-complement number. Before the processor adds this 12-bit value to the upper 20 bits, the hardware automatically sign-extends the most significant bit of the 12-bit value (bit 11) all the way through bit 31.

If bit 11 happens to be a `1`, the 12-bit value is treated as a negative number, effectively *subtracting* from the upper 20 bits. For example, suppose a programmer wants to load the 32-bit value `0xFFFFFFFF` into the `x1` register. If they naively write:

```
lui x1, 0xFFFF    # x1 now equals 0xFFFF000
addi x1, x1, 0xFFF
```

The `addi` instruction will take `0xFFF`, interpret it as a signed number, sign-extend it to `0xFFFFFFFF` (which is -1), and add it to `0xFFFF000`. The resulting sum placed in `x1` will be `0xFFFFEFFF`, which is totally incorrect.

To prevent this, the assembler must perform mathematical compensation when evaluating the `la` (Load Address), `lw` (Load Word), and `sw` (Store Word) pseudo-instructions. If bit 11 of the lower 12-bit payload is a `1`, the assembler

must pre-emptively add 1 to the upper 20-bit value before generating the `lui` or `auipc` instruction. This slight increment in the upper bits perfectly cancels out the subsequent negative sign-extension of the lower bits, resulting in the correct 32-bit address.

## 12.5 Hardware Implementations and Custom Extensions

What makes RISC-V uniquely suited for modern Cyber-Physical Systems is its explicit support for user-defined enhancements. A team designing a robotic actuator does not have to settle for the standard instruction set; they can augment the processor with custom opcodes, hardware coprocessors, and other modifications as long as they comply with the RISC-V customization rules.

Because all the source code and processor design intellectual property required to implement a RISC-V core is freely available on the internet, developers can synthesize their own custom processors without paying licensing fees. By using an open-source RISC-V design as a starting point, engineers can implement custom modifications that are guaranteed to remain compatible with future versions of the RISC-V standard.

For example, the PicoRV32 is a popular open-source RISC-V core that can be easily synthesized and programmed into low-cost Field Programmable Gate Arrays (FPGAs). A budget-friendly FPGA board like the Digilent Cmod A7-35T utilizes a Xilinx Artix-7 chip (XC7A35T). This specific FPGA provides 5,200 logic slices, 41,600 flip-flops, 1,800 kbits of block RAM, and 90 specialized Digital Signal Processing (DSP) slices that natively support high-performance Multiply-Accumulate (MAC) operations.

An engineer can map the PicoRV32 open-source model onto this FPGA using the Vivado synthesis toolchain. Once the baseline RV32IMC core is instantiated, the engineer can write Verilog code to bind the FPGA's physical DSP slices directly to custom RISC-V opcodes. This creates a deeply integrated, hardware-accelerated processor capable of executing complex PID control loops or AI

inference models in a fraction of the time required by standard software execution.

## 12.6 VirtMCU Homework: Custom ISA Extensions for Deterministic Control

In previous chapters, you wrote control algorithms assuming a standard, immutable CPU architecture. However, in the CPS 2.0 paradigm, hardware and software must be co-designed. If a software control loop cannot meet its hard real-time deadlines, you now have the power to alter the hardware to solve the problem.

For this chapter's homework, you will step beyond standard software engineering and venture into computer architecture by modifying the ISA within the **VirtMCU FirmwareStudio**.

1. **Baseline Profiling:** You will be provided with a standard C program that reads three orthogonal sensor registers (X, Y, and Z axes) via MMIO, computes their vector magnitude ( $\sqrt{X^2 + Y^2 + Z^2}$ ), and writes the result to an actuator register. Compile this for the standard RISC-V RV32I base architecture. Run it in VirtMCU's deterministic `slaved-icount` mode and record the exact number of clock cycles required to execute the math loop.
2. **Architectural Modification:** You will dive into the VirtMCU QEMU backend source code. You will define a custom RISC-V instruction, `VMAG rd, rs1, rs2, rs3`, that utilizes an unused opcode space. You will write the emulation logic in C to calculate the vector magnitude in a single simulated clock cycle.
3. **Inline Assembly Implementation:** Modify your original C firmware to bypass the standard math libraries. Use GCC's inline assembly feature to invoke your new `VMAG` machine instruction directly.
4. **Co-Simulation Validation:** Re-run the modified firmware in VirtMCU `slaved-icount` mode. Compare the new cycle count against your baseline profile. You must document the performance speedup achieved by pushing a specific computational bottleneck directly into the silicon

architecture, empirically demonstrating the power of an open ISA in CPS design.

# Chapter 13: Security and Resilience in Cyber-Physical Systems

## 13.1 Introduction: Securing Critical Infrastructure

As established in Chapter 9, the transition to Cyber-Physical Systems (CPS 2.0) deeply integrates Operational Technology (OT) with Information Technology (IT), significantly expanding the attack surface of physical infrastructure. While traditional IT security focuses on data confidentiality and privacy, security in a CPS context is fundamentally about maintaining physical control and ensuring the safety and availability of critical services.

Nowhere is this more evident than in **Critical Energy Infrastructure (CEI)**. Modern energy systems and power grids are constantly under stress from both natural disruptions—such as wildfires and extreme weather events—and sophisticated, targeted cyberattacks. To deal with the ever-increasing challenge posed by these cyber-physical events, we must shift our focus from mere prevention to **resilience**.

Resilience is defined as the ability to withstand and reduce the magnitude and duration of disruptive events. A resilient CPS possesses the capability to anticipate, absorb, adapt to, and rapidly recover from an attack or failure. This chapter explores the methodologies used to model cyber-physical threats, assess risks, and deploy advanced computational techniques—such as robust optimization, decentralized autonomous controls, and distributed topology reconfiguration—to ensure the survival of critical infrastructure under adversarial conditions.

## 13.2 The Cyber-Physical Threat Landscape

To properly secure a CPS, engineers must understand the underlying assets, their vulnerabilities, and the associated threats. In energy and industrial networks, threats are broad, varying from simple localized faults to complex cascading failures.

### 13.2.1 Classes of Cyber-Physical Attacks

Adversaries target the cyber-physical feedback loop through several distinct vectors:

- **Measurement Error Injection:** Attackers inject constant offsets or time-varying errors into sensor readings. If left untreated, these can manipulate control laws and lead to ramp-induced physical attacks.
- **Replay Attacks:** The measurement output is maliciously intercepted and changed to reflect the value at a previous timestamp, blinding the controller to current physical realities and impacting future decisions.
- **State Estimation Corruption:** Attackers maliciously alter the estimated quantities required by the system's centralized dispatcher, causing the system to compute incorrect control actions.
- **Coordinated Attacks:** Attackers gain access to multiple critical information nodes and launch simultaneous threats from multiple sources in a coordinated manner, designed to overwhelm standard fault-tolerance mechanisms.

### 13.2.2 Threat Modeling and Attack Trees

To analyze these complex vulnerabilities, security architects utilize **Attack Trees**. An attack tree is a model that formalizes the structure of well-known cyber-physical attacks, detailing the exact steps an attacker must take.

For example, consider the cascading threat of cutting a city's water supply through the electrical grid. The attack tree demonstrates that the water supply can be cut if the water pumps are stopped and the reservoir is drained. The

pumps stop if the primary power supply is cut *and* the backup power supply is disabled. The power supply itself can be interrupted maliciously by compromising a substation either physically or through cyber means. By mapping these steps using AND/OR logic gates, engineers can identify the critical nodes where design-time countermeasures (e.g., adding redundancy) or run-time mitigations (e.g., dynamic isolation) must be applied.

Furthermore, engineers must continuously evaluate **unknown threats**. These are addressed by exploring various conceptual spaces: the *threat space* (newly discovered vulnerabilities), the *attack space* (new attacker tools and methods), and the *technology space* (novel threats introduced by emerging technologies like autonomous drones or IoT automation).

## 13.3 Risk Management and Evaluation

Once threats are modeled, they must be assessed for their potential impact. The standard risk management process for critical infrastructure is based on **ISO/IEC 27005**. This process includes:

1. **Context Establishment:** Defining the scope, identifying assets, and mapping threats to those assets.
2. **Risk Assessment:** Estimating the risk based on the likelihood of the threat exploiting a vulnerability and the severity of the harmful result.
3. **Risk Treatment & Acceptance:** Preparing a mitigation plan and acknowledging any residual risks.

Risk estimation combines the human, economic, and social impact with the risk's likelihood. A qualitative approach is often used, categorizing impact from "minimal" to "severe" (levels 1 to 5) and likelihood from "rare" to "almost certain" (e.g., 1% to 80% probability).



## 13.4 Advanced Computational Techniques for Resilience

Traditional risk mitigation relies heavily on fortifying networks. However, in CPS 2.0, adversarial events are notoriously hard to predict and reliable historical data is often sparse. Attempting to make a system invulnerable to every possible disruption makes the operation overly conservative and economically impractical. Instead, modern CPS architecture relies on advanced computational techniques to guarantee resilience dynamically.

### 13.4.1 Distributionally Robust Optimization (DRO)

To handle uncertainty in system parameters and adversarial forecasts, systems utilize stochastic optimization. **Distributionally Robust Optimization (DRO)** is a probabilistic method that determines the exact amount of reserves required to deal with cyber-physical events in a resilient and cost-effective manner.

Rather than planning for a single worst-case scenario, DRO considers a bounded set of uncertainty  $\Omega$  and optimizes the system's response across a probability distribution of potential disruptions. By integrating these probabilistic models, the CPS dispatcher ensures sufficient dispatchable resources are available to absorb the impact of an attack without collapsing the grid.

### 13.4.2 Decentralized Autonomous Controls

During severe cyber-physical attacks, centralized communication networks may be congested, compromised, or entirely unavailable. To survive, the decision-making agents within the network must be able to act autonomously and collaboratively. This is achieved through minimally invasive, decentralized control layers:

- **Proactive Autonomy:** These controls continuously guarantee that the system operates within strict safety limits during likely, bounded disturbances. By utilizing mathematical proofs (such as Nagumo's

theorem), proactive autonomy ensures the system's operational viability without requiring communication with a central coordinator.

- **Reactive Autonomy:** When a high-impact, low-probability adversarial event thrusts the system state outside the safe operating regime, reactive autonomy takes over. Reactive control involves computing control inputs in real-time—often by efficiently solving for energy functions—to steer the highly nonlinear dynamics of the system back into a safe region as quickly as possible.

Together, these autonomous controls act as a bridge between myopic local sensors and the centralized dispatcher, ensuring that a communication failure does not result in physical destruction.

## 13.5 Network Reconfiguration and Resilient Learning

In distributed critical infrastructure, such as interconnected microgrids, resilience can be further enhanced by dynamically altering the physical topology of the system to isolate attacked segments.

### 13.5.1 Topology Reconfiguration via Distributed OPF

If a portion of the grid is compromised by a cyberattack, local controllers must determine the optimal physical interconnection topology to island the threat while minimizing generation costs and line losses. This requires solving a **Distributed Optimal Power Flow (OPF)** problem.

Because attackers may simultaneously target the controllers themselves, resilient consensus algorithms are utilized. These algorithms allow each microgrid to exchange optimal objective functions with neighboring grids via Peer-to-Peer (P2P) communication. To solve these complex distributed equations rapidly, techniques such as the **Alternating Direction Method of Multipliers (ADMM)** are heavily utilized, ensuring convergence even when parts of the network are untrusted.

### 13.5.2 Resilient Reinforcement Learning

Modern energy infrastructure heavily integrates inverter-based resources, such as solar panels and battery storage. These utilize **Grid-Forming (GFM)** inverters, modeled as AC voltage sources with local frequency and voltage control signals governed by droop controls.

To optimize these distributed resources under attack conditions where data sharing is limited or unsafe, researchers are deploying Reinforcement Learning (RL) and Federated Learning (FL). These data-driven, machine learning techniques allow the localized controllers to collaboratively learn optimal resilient policies without transmitting sensitive, interceptable state data across the network.

## 13.6 VirtMCU Homework: Attack Tree Modeling and Resilient Recovery

*(Note: This exercise builds on our established course framework using the VirtMCU firmware studio, applying the theoretical concepts of this chapter to a simulated environment).*

In Chapter 11, you designed a PID control loop to stabilize an inverted pendulum. In this chapter's homework, you will subject that continuous-time control loop to a simulated cyber-physical attack to evaluate its resilience.

#### Your Assignment:

1. **Attack Tree Construction:** Draft a formal Attack Tree outlining the steps an attacker would take to destabilize your inverted pendulum. Include nodes for both cyber vectors (e.g., compromising the PWM duty cycle communication via a buffer overflow) and physical vectors (e.g., injecting false quadrature encoder data).
2. **Implementation of a Replay Attack:** Modify your firmware to simulate a state estimation corruption. At exactly 5.0 seconds into the simulation, your code must freeze the angular sensor reading  $\theta(t)$  to its last

known value, blinding the PID controller to the true physical state. Observe how rapidly the system fails.

3. **Reactive Autonomy (Fallback Control):** Implement a secondary, decentralized fallback control law. You will monitor the raw derivative of the cart's wheel velocity. If the discrepancy between the expected cart position and the stale  $\theta(t)$  reading exceeds a threshold, your firmware must automatically trigger a "Reactive Autonomy" state, overriding the primary PID loop with an aggressive braking maneuver to bring the system to a safe halt.
4. **Deterministic Evaluation:** Run the simulation in VirtMCU's `slaved-icount` mode. Because virtual time is strictly deterministic, you must provide a cycle-accurate report of your Reactive Autonomy detection latency (the exact number of instructions between the attack injection and the override engagement), proving your system can reliably recover from a sensor-blinding attack within the physical constraints of the hardware.

# Chapter 14: Domain-Specific Architectures and Cloud/Edge Integration

## 14.1 Introduction: The End of Moore's Law and the Rise of DSAs

In the preceding chapters, we focused extensively on general-purpose microprocessors and their application in Cyber-Physical Systems (CPS). We explored pipelining, out-of-order execution, memory hierarchies, and the real-time operating systems required to enforce deterministic physical control. However, the computing landscape is undergoing a profound structural shift.

For decades, software engineers relied on Moore's Law and Dennard scaling to automatically provide faster, more energy-efficient general-purpose CPUs every two years. Today, this free ride has effectively ended. As Gordon Moore himself noted, "Moore's Law can't continue forever. We have another 10–20 years before we reach a fundamental limit". The cessation of Dennard scaling (the power wall) means that architects can no longer simply increase clock frequencies or pack more active general-purpose cores onto a die without melting the silicon.

Given the end of Moore's Law and Dennard scaling, the only path forward for radically improved performance and energy efficiency is the **Domain-Specific Architecture (DSA)**. Unlike general-purpose CPUs, which are designed to run everything from text editors to web browsers reasonably well, DSAs are custom-tailored to accelerate a specific class of applications. By discarding the immense hardware complexity required for general-purpose out-of-order execution, DSAs devote their silicon area entirely to highly efficient, specialized arithmetic and data routing.

In the context of modern CPS 2.0, the architecture is split across a vast continuum: from ultra-low-power DSAs operating at the physical edge, to massive Warehouse-Scale Computers (WSCs) operating in the cloud. This chapter bridges that divide, examining the hardware architectures that define the modern edge-to-cloud ecosystem.

## 14.2 Architecting for the Edge: Mobile and SoC Architectures

At the very edge of the CPS network are the personal mobile devices (PMDs) and embedded sensors that directly interact with the physical world. These devices are strictly constrained by battery capacity and thermal dissipation limits.

To meet these constraints, edge devices utilize **System-on-Chip (SoC)** architectures. Rather than isolating the CPU, graphics, and memory controllers on separate physical chips, an SoC integrates them onto a single die. This minimizes the parasitic capacitance of external wiring, drastically reducing the electrical power required to move data.

Furthermore, these SoCs are deeply heterogeneous, containing numerous DSAs to offload specific tasks from the main CPU:

- **Image Signal Processors (ISPs):** Hardware accelerators designed specifically to process raw pixel data from dual cameras or IR projectors, performing noise reduction, color space conversion, and edge enhancement in real-time.
- **Digital Signal Processors (DSPs):** Highly optimized cores for processing continuous analog data, such as audio from microphones or RF signals from cellular modems.
- **Touch Controllers:** Dedicated hardware to process capacitive touch sensing. These sensors monitor changes in capacitance caused by the proximity of a conductive object (such as a human finger), allowing the hardware to filter raw measurements and pinpoint multiple simultaneous touchpoints.

By routing data directly from the physical sensor into a dedicated DSA, the SoC can process real-time events while the main general-purpose CPU remains in a deep, low-power sleep state, maximizing battery life.

## 14.3 Warehouse-Scale Computing (WSC)

While edge devices collect telemetry and execute local control loops, the heavy computational lifting—such as training global AI models or processing massive distributed databases—is offloaded to the cloud. The infrastructure that powers the cloud is the **Warehouse-Scale Computer (WSC)**.

It is a misconception to view a modern data center merely as a building full of individual servers. Given the immense processing, networking, and storage capabilities required, it is far more accurate to conceptualize the entire data center as a single, massively parallel computing system. The price-performance and energy efficiency concerns of a single microprocessor apply equally to WSCs, but at a scale where the equipment and the building cost hundreds of millions of dollars.

### 14.3.1 WSC Network Topology

A WSC is organized hierarchically. At the base level, individual servers are mounted in racks. Each server is a reasonably complete computer system containing a processor, RAM, and local storage. These servers operate in a “headless mode,” lacking any direct physical connection to a display or keyboard; all interaction occurs over the network.

Each server connects to a Top-of-Rack (ToR) network switch, typically using a 10 Gbit/s Ethernet cable. Servers within the same rack can communicate with one another at full speed. However, the racks are then grouped into clusters, sharing a second-level cluster switch. Because it is economically unfeasible to provide full bandwidth between every rack in the facility, the network is often **oversubscribed**. For example, an oversubscription factor of 4 means that the external network capacity leaving the rack is only one-quarter of the peak communication speed available inside the rack. Software architects must

carefully place communicating processes within the same rack to avoid saturating the cluster switches.

### 14.3.2 Hardware Fault Management

Because a WSC contains tens of thousands of individual servers, hardware failures are not anomalies; they are continuous, daily occurrences. Even if highly reliable (and expensive) components are used, the sheer scale guarantees component death.

Therefore, WSC architecture relies on software to mask hardware failures. When a higher-level system dispatches a request (such as an internet search query) to a lower-level index server, it actively monitors the response latency. If the lower-level server becomes persistently unresponsive or returns corrupted data, the requesting system assumes a hardware fault has occurred, abandons the failed node, and immediately resends the request to a redundant backup server. The failed server is taken offline automatically so technicians can repair or replace it without interrupting the global service.

## 14.4 Accelerating AI: GPUs and Neural Processing Units

The most significant driver of Domain-Specific Architectures today is Artificial Intelligence (AI). Historically, computer scientists built AI as a massive set of logical rules; today, it is driven by deep neural networks that learn from data. As it turns out, it is much harder to program a computer to *be* intelligent than to program it to *learn* to be intelligent.

Training and executing Deep Neural Networks (DNNs) requires astonishing amounts of computational power. Because neural networks are fundamentally based on linear algebra (specifically, massive matrix multiplications), they exhibit immense **Data-Level Parallelism (DLP)**. General-purpose CPUs, which are optimized for sequential logic and unpredictable branching, are terrible at this.



### 14.4.1 Graphics Processing Units (GPUs)

The first hardware to successfully exploit the DLP of neural networks was the Graphics Processing Unit (GPU), such as the NVIDIA A100 and T4. Originally designed to compute millions of independent pixels for synthetic video scenes, GPUs employ a heavily threaded, Single Instruction, Multiple Data (SIMD) architecture. A modern GPU can execute thousands of floating-point multiply-accumulate (MAC) operations in parallel, making them exceptional AI accelerators.

### 14.4.2 Tensor Processing Units (TPUs) and Rack-Scale Accelerators

While GPUs are fast, they still retain legacy hardware for graphics rendering (such as texture mapping and rasterization) that is useless for pure AI workloads. To achieve maximum efficiency, companies have developed pure AI DSAs.

Google's **Tensor Processing Unit (TPU)** is a custom Application-Specific Integrated Circuit (ASIC) designed exclusively to accelerate DNNs in the data center. The heart of a TPU is a massive systolic array multiplier—a 2D grid of ALUs that passes data directly from one ALU to the next without writing intermediate results back to registers, drastically reducing power consumption while executing tens of thousands of matrix operations per clock cycle.

Intel has taken a similar architectural leap with **Jaguar Shores**, a next-generation rack-scale AI accelerator platform. Rather than functioning as a standalone chip or PCIe card, Jaguar Shores is designed as an entire rack-based solution, providing a complete AI processing system scaled specifically to meet the extreme memory bandwidth and training demands of modern machine learning.

## 14.5 VirtMCU Homework: Designing an Edge AI Accelerator

In previous chapters, you wrote strict PID control loops for an inverted pendulum and a drone using standard ARM AArch64 assembly. However, in modern CPS, we frequently augment classical control with predictive machine learning models at the edge.

If you attempt to execute a Deep Neural Network (DNN) inference pass using standard sequential software in your ISR, the matrix multiplications will consume too many clock cycles. You will exceed your  $T_s$  (sample time) control deadline, causing the simulated drone to crash.

To solve this, you will step into the role of a computer architect and design a Domain-Specific Architecture (DSA) within the **VirtMCU FirmwareStudio**.

1. **Memory-Mapped Accelerator Design:** You will be provided with the C++ source code for the VirtMCU hardware backend. You must design a mock “Tensor Coprocessor” peripheral that attaches to the virtual APB interconnect.
2. **MMIO Interface:** Design the Memory-Mapped I/O (MMIO) registers for your DSA. You will need a `MATRIX_A` port, a `MATRIX_B` port, a `START_INFERENCE` command register, and a `RESULT` memory buffer.
3. **Firmware Integration:** Rewrite your drone flight controller firmware. Instead of performing the neural network matrix multiplication in software, your firmware will use DMA to blast the sensor data into your custom DSA’s MMIO registers, trigger the `START` command, and yield the CPU.
4. **Deterministic Benchmarking:** Run the co-simulation in `slaved-icount` mode. By offloading the DLP workload to your specialized hardware, you will empirically measure the massive reduction in required CPU clock cycles. You must submit a report proving that the introduction of the DSA reduces your Worst-Case Execution Time (WCET) enough to comfortably meet the hard real-time deadlines of the physical control loop.

# Epilogue: The Future of Cyber-Physical Systems

We have traversed the entire vertical stack of modern computing—from the physics of semiconductor transistors, up through the microarchitecture and instruction sets of ARM and RISC-V processors, into the memory hierarchies and interconnect fabrics that shuffle our data, and finally into the software layers of real-time operating systems and distributed Cyber-Physical Systems (CPS).

The computing landscape you are entering as a graduate engineer is fundamentally different from the one that existed just a decade ago. For over fifty years, the industry relied on Moore's Law and Dennard scaling to reliably double transistor density and performance while maintaining constant power densities. Today, those physical scaling laws have effectively collapsed. We have hit the "power wall," meaning we can no longer simply increase clock frequencies without melting the silicon.

To continue achieving performance gains, the future of CPS architecture lies in **Domain-Specific Architectures (DSAs)** and hardware-software co-design. General-purpose processors will increasingly serve as traffic cops, delegating heavy mathematics to specialized Neural Processing Units (NPUs), Tensor Processing Units (TPUs), and custom Field Programmable Gate Arrays (FPGAs) operating at the edge.

Furthermore, as our physical infrastructure—power grids, medical devices, and autonomous transportation—becomes inextricably linked to cloud networks, **security and resilience** must be treated as first-class architectural constraints, not software afterthoughts. Designing systems that can autonomously withstand, isolate, and recover from cyber-physical attacks using distributed machine learning and robust optimization is the grand challenge of your generation.

The tools you have used in this text, from bare-metal debuggers to the deterministic VirtMCU co-simulation framework, are designed to give you

absolute control over both the logical correctness and the continuous-time physical realities of your designs.

# Appendix A: Instruction Set Architecture (ISA) Reference

This appendix provides a quick reference for the primary Instruction Set Architectures utilized in this text: ARM AArch64 and RISC-V (RV32I).

## A.1 ARM AArch64 (A64) Instruction Subset

The ARMv8-A architecture introduced the 64-bit AArch64 execution state.

### Data Transfer (Load/Store)

- `ldr Xd, [Xn, offset]` : Load a 64-bit doubleword from the memory address  $(Xn + offset)$  into register `Xd`.
- `str Wd, [Xn, offset]` : Store a 32-bit word from register `Wd` into memory.
- `ldp / stp` : Load/Store Pair. Pushes or pops two registers simultaneously, heavily used in stack frame prologues and epilogues (e.g., `stp fp, lr, [sp, #-16]!`).

### Arithmetic and Logic

- `add Xd, Xn, Xm` :  $Xd = Xn + Xm$ . (Use `adds` to update PSTATE condition flags).
- `sub Xd, Xn, Xm` :  $Xd = Xn - Xm$ . (Use `subs` to update flags).
- `mul Xd, Xn, Xm` :  $Xd = Xn * Xm$ .
- `and / orr / eor` : Bitwise AND, OR, and Exclusive-OR.

### Control Flow

- `cmp Xn, Xm` : Compare (implicitly performs `subs` and updates flags without storing the result).
- `b label` : Unconditional branch to `label`.

- `beq / bne / blt / bgt`: Conditional branches (Branch if Equal, Not Equal, Less Than, Greater Than).
- `bl label`: Branch with Link. Stores the return address in the Link Register ( `LR / X30` ) and jumps to `label` (used for function calls).
- `ret`: Return from subroutine. Copies `LR` to the Program Counter ( `PC` ).

## A.2 RISC-V (RV32I) Base Instruction Subset

RISC-V is an open-source ISA that heavily utilizes pseudo-instructions to minimize silicon decoder complexity.

### Core Instructions

- `addi rd, rs1, imm`: Add immediate.  $rd = rs1 + imm$ .
- `lui rd, imm`: Load Upper Immediate. Places a 20-bit value into the highest 20 bits of `rd`.
- `auipc rd, imm`: Add Upper Immediate to PC.
- `jal rd, offset`: Jump and Link. Jumps to  $PC + offset$  and saves the return address in `rd` (usually `x1 / ra`).
- `beq / bne rs1, rs2, offset`: Branch if  $rs1 ==$  (or  $!=$ )  $rs2$ .

### Common Pseudo-Instructions

- `nop`: Translates to `addi x0, x0, 0`.
- `mv rd, rs`: Translates to `addi rd, rs, 0`.
- `not rd, rs`: Translates to `xori rd, rs, -1`.
- `j offset`: Translates to `jal x0, offset`.

# Appendix B: Selected Answers to VirtMCU Exercises

**Chapter 3: Stack Tracing and Deterministic Execution** *Solution Note:* When executing the recursive factorial function in VirtMCU `slaved-icount` mode, students should observe that the Link Register ( `x30` ) is pushed to the stack exactly 16 bytes below the previous Frame Pointer ( `x29` ). Because `slaved-icount` ties the virtual clock strictly to retired instructions, inspecting the cycle counter before and after the recursive calls will yield an identical delta on every simulation run, free from host OS jitter.

**Chapter 7: Deterministic Interrupt Handler** *Solution Note:* Students must successfully construct the Vector Table to point to their ISR. Inside the ISR, the code must execute `stp x0, x1, [sp, #-16]!` (and subsequent saves) to preserve the volatile registers. Failing to acknowledge and clear the simulated push-button peripheral's IRQ flag in the MMIO register will result in an infinite interrupt loop upon executing the return instruction.

**Chapter 11: Inverted Pendulum Co-Simulation** *Solution Note:* The continuous-time MuJoCo physics engine relies on the VirtMCU Sensor Abstraction Layer (SAL) to discretize angles into 12-bit simulated quadrature encoder pulses. Students must apply a Proportional-Integral-Derivative (PID) algorithm in C. Saturation in the Actuator Abstraction Layer (AAL) means that if the proportional gain (`$k_p$`) is set too high, the requested PWM duty cycle will clip, causing integral windup and leading to the pendulum crashing.

# Glossary of Terms and Abbreviations

- **ABI (Application Binary Interface):** The low-level standard that dictates how software interacts with the hardware, including register usage conventions (volatile vs. nonvolatile) and parameter passing during function calls.
- **ADC (Analog-to-Digital Converter):** A hardware component that converts continuous physical voltages into discrete digital numbers, limited by sampling rate and quantization resolution.
- **AMBA (Advanced Microcontroller Bus Architecture):** A family of open-standard, on-chip interconnect specifications (including APB, AHB, and AXI) developed by ARM for routing data blocks within a System-on-Chip.
- **AXI (Advanced eXtensible Interface):** A high-performance AMBA interconnect protocol that utilizes five independent, unidirectional channels to allow full-duplex, out-of-order memory transactions.
- **CPI (Cycles Per Instruction):** A fundamental performance metric representing the average number of clock cycles required to execute a machine instruction.
- **CPS (Cyber-Physical System):** An engineered system built from, and dependent upon, the seamless integration of computational algorithms and physical components.
- **DMA (Direct Memory Access):** A hardware controller that autonomously transfers blocks of data between peripherals and main memory, bypassing the CPU execution pipeline to save clock cycles.
- **DSA (Domain-Specific Architecture):** Custom silicon tailored to accelerate a specific class of applications (e.g., neural networks), trading general-purpose flexibility for massive efficiency gains.
- **ISA (Instruction Set Architecture):** The boundary between hardware and software; the complete set of instructions, registers, and memory models understood by a specific microprocessor.
- **MMIO (Memory-Mapped I/O):** A system design technique where peripheral control and data registers are mapped into the CPU's standard physical memory address space, allowing hardware to be controlled via standard load/store instructions.



- **NVIC (Nested Vectored Interrupt Controller):** A hardware component in ARM architectures that prioritizes and manages incoming physical interrupts, automatically forcing the PC to the correct Interrupt Service Routine.
- **PWM (Pulse Width Modulation):** A technique for simulating an analog voltage output by rapidly switching a digital signal HIGH and LOW. The effective output is proportional to the duty cycle.
- **RTOS (Real-Time Operating System):** An operating system designed to guarantee strict deterministic timing and meet hard deadlines, rather than optimizing for average throughput.
- **TLB (Translation Lookaside Buffer):** An ultra-fast hardware cache located inside the Memory Management Unit (MMU) that stores recent virtual-to-physical page table translations.
- **TLM (Transaction Level Modeling):** A high-level simulation technique (often used in SystemC) that abstracts away individual clock cycles and pins, modeling bus communication via object-oriented method calls to dramatically speed up simulation.

# Bibliography / Further Reading

Feng, C., et al. *Cyber-Physical Virtual Power Plants: Modeling, Operation, and Control*. IEEE / Springer, 2025.

Arm Limited. *Fundamentals of System-on-Chip Design on Arm Cortex-M Microcontrollers*. Arm Education Media, 2021.

*VirtMCU FirmwareStudio*. RefractSystems / GitHub Repository.  
<https://github.com/RefractSystems/virtmcu>

Große, D., and Drechsler, R. *Quality-Driven SystemC Design*. Springer Science+Business Media, 2010.

Harris, D. M., and Harris, S. L. *Digital Design and Computer Architecture*, 2nd Edition. Morgan Kaufmann / Elsevier, 2013.

Hennessey, J. L., and Patterson, D. A. *Computer Architecture: A Quantitative Approach*, 7th Edition. Morgan Kaufmann / Elsevier, 2026.

Hyde, R. *The Art of ARM Assembly Vol 1: 64-Bit ARM Machine Organization and Programming*. No Starch Press, 2025.

Jha, A., et al. *Cyber Physical System 2.0: Communication and Computational Technologies*. 2025.

Ledin, J. *Modern Computer Architecture and Organization*, 3rd Edition. Packt Publishing, 2026.

Lyshevski, S. E. *Design and Control of Physical and Cyber-Physical Systems*. 2026.

Greaves, D. J. *Modern System-on-Chip Design on Arm*. Arm Education Media, 2021.

Patel, H. D., and Shukla, S. K. *SystemC Kernel Extensions for Heterogeneous System Modeling: A Framework for Multi-MoC Modeling & Simulation*. Kluwer Academic Publishers, 2005.

Plantz, R. G. *Introduction to Computer Organization: ARM Edition*. No Starch Press, 2025.

Poonia, R. C., et al. *Smart Cyber-Physical Systems: Innovations and Real-World Implications*. 2025.

Soldatos, J., et al. *Cyber-Physical Threat Intelligence for Critical Infrastructures Security*. Now Publishers, 2020.

Tanenbaum, A. S., and Austin, T. *Structured Computer Organization*, 6th Edition. Pearson, 2016.

Yan, J., et al. *Secure Detection and Control in Cyber-Physical Systems*. Taylor & Francis, 2026.

# Index

- **A**

- Addressing Modes, Ch. 3, Ch. 5, App. A
- Advanced eXtensible Interface (AXI), Ch. 6, Glossary
- Application Binary Interface (ABI), Ch. 3, Ch. 12, Glossary
- Arithmetic Logic Unit (ALU), Ch. 4
- Artificial Intelligence (Edge/Cloud), Ch. 9, Ch. 14

- **B**

- Backpressure, Ch. 6
- Branch Prediction, Ch. 4
- Bubble Sort, Ch. 4 (Homework), Ch. 10

- **C**

- Cache Coherence, Ch. 5
- Cache Memory (L1, L2, L3), Ch. 5
- Compare and Branch, Ch. 3
- Context Switching, Ch. 7, Ch. 8
- Cyber-Physical System (CPS 1.0 & 2.0), Ch. 9, Glossary

- **D**

- Deadlocks, Ch. 8
- Direct Memory Access (DMA), Ch. 6, Glossary
- Discrete-Event (DE) Simulation, Ch. 10
- Domain-Specific Architectures (DSAs), Ch. 14, Epilogue
- Dynamic RAM (DRAM), Ch. 5

- **E**

- Edge Computing, Ch. 9, Ch. 14
- Exceptions and Faults, Ch. 7

- **G**

- General-Purpose Registers, Ch. 3, Ch. 12
- **I**
  - Instruction-Level Parallelism (ILP), Ch. 4
  - Interrupt Service Routine (ISR), Ch. 7
  - Inverted Pendulum, Ch. 11, Ch. 13
- **M**
  - Memory Management Unit (MMU), Ch. 5
  - Memory-Mapped I/O (MMIO), Ch. 6, Glossary
  - Moore's Law, Ch. 14, Epilogue
- **P**
  - Paging and Virtual Memory, Ch. 5
  - Pipelining, Ch. 4
  - Priority Inversion, Ch. 8
  - Pulse Width Modulation (PWM), Ch. 11, Glossary
- **Q**
  - Quantization Error, Ch. 11
- **R**
  - Rate Monotonic Scheduling (RMS), Ch. 8
  - Reorder Buffer (ROB), Ch. 4
  - RISC-V Architecture, Ch. 12, App. A
- **S**
  - Semaphores and Mutexes, Ch. 8
  - Sensor Abstraction Layer (SAL) / Actuator Abstraction Layer (AAL), Ch. 11
  - Superscalar Execution, Ch. 4
  - SystemC, Ch. 10
- **T**

- Transaction Level Modeling (TLM), Ch. 10, Glossary
- Translation Lookaside Buffer (TLB), Ch. 5, Glossary

- **V**

- Vector Table, Ch. 7
- Virtual Power Plants (VPP), Ch. 9, Ch. 10
- VirtMCU FirmwareStudio, Ch. 3, Ch. 4, Ch. 7, Ch. 10, Ch. 11, App. B