

Welcome to the blueprint for our new graduate-level textbook and course structure. Designing this curriculum requires us to bridge a significant gap: taking graduate students who are highly proficient in software engineering and grounding them in the realities of hardware, computer architecture, and the continuous-time physical world.

To achieve this, we will use **VirtMCU** as our core pedagogical tool, allowing students to build deterministic, multi-node digital twins without the scheduling jitter or physical hardware limitations that traditionally plague CPS education.

Below is the comprehensive structural plan for the textbook, formatted natively in `.md` (Markdown). It is designed to cover 28 lectures across 14 chapters.

Book Title: Foundations of Computer Architecture and Cyber-Physical Systems 2.0: A Deterministic Approach Using VirtMCU

Version: 1.0.0 (Commit: `b4eba5f`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Version: 1.0.0 (Commit: `f14d70c`)

Book Metadata and Author Directives

- **Target Audience:** Graduate students with a software background transitioning into Computer Architecture and Cyber-Physical Systems (CPS).
 - **Format:** Markdown (`.md`).
 - **Length Constraints:** Total size $\geq$ 180 pages. Each of the 14 chapters will be strictly $\geq$ 8 pages and contain a minimum of 2,500 words.
 - **Course Mapping:** 14 Chapters, designed to be taught over 28 lectures (2 lectures per chapter).
 - **Pedagogical Tool:** **VirtMCU** (FirmwareStudio). The book leverages VirtMCU's deterministic QEMU-based simulation, Sensor/Actuator Abstraction Layer (SAL/AAL), and MuJoCo physics engine lock-stepping to provide hands-on homework.
-

Part I: Hardware Foundations for Software Engineers (Lectures 1–6)

Objective: Strip away the software abstractions and teach the students how electrons become 1s and 0s, and how 1s and 0s become instructions.

Chapter 1: The Digital Abstraction and Logic Design

- **Lecture 1:** From Transistors to Logic Gates. We begin at the lowest level of abstraction: the physics of semiconductors, CMOS transistors, and logic gates (AND, OR, NOT).
- **Lecture 2:** Combinational and Sequential Logic. Students learn how gates are combined to create multiplexers, adders, and ALUs. We introduce

sequential logic—latches, flip-flops, and clocks—explaining how circuits maintain state and memory.

- **VirtMCU Homework:** Introduction to the VirtMCU environment. Students compile a basic C program into a bare-metal binary and observe simulated clock cycles.

Chapter 2: Data Representation and Machine Organization

- **Lecture 3:** Data Storage and Number Systems. Binary, hexadecimal, signed/unsigned integers, two's complement, and IEEE 754 floating-point formats.
- **Lecture 4:** Memory and Machine Organization. Big-endian vs. little-endian byte ordering, memory addresses, and the von Neumann vs. Harvard architectures.
- **VirtMCU Homework:** Memory inspection. Students use `gdb` connected to the VirtMCU QEMU instance to step through memory and observe byte ordering and data representation live.

Chapter 3: Instruction Set Architecture (ISA) and ARM AArch64

- **Lecture 5:** The AArch64 Programmer's Model. Introduction to the 64-bit ARM architecture. We cover the 32 general-purpose registers (X0-X30, W0-W30), the program counter, and the stack pointer.
 - **Lecture 6:** Assembly Language and the Call Stack. Translating high-level C flow control (if/else, loops) into ARM assembly. We dive deep into function calls, the ARM Application Binary Interface (ABI), parameter passing, and stack frame discipline.
 - **VirtMCU Homework:** Students rewrite a simple C function in AArch64 assembly, load it into VirtMCU, and trace the register modifications and stack allocations.
-

Part II: Modern System-on-Chip (SoC) Architecture (Lectures 7–12)

Objective: Build a modern processor and SoC from the ground up, focusing on timing, memory bottlenecks, and interconnects.

Chapter 4: Processor Microarchitecture

- **Lecture 7:** The Datapath and Control Unit. Designing a single-cycle processor. Instruction fetching, decoding, execution, and write-back phases.
- **Lecture 8:** Pipelining and Superscalar Execution. We explore the classic five-stage RISC pipeline, data/control hazards, and branch prediction. We briefly discuss modern out-of-order execution and multi-core simultaneous multithreading.
- **VirtMCU Homework:** Analyzing pipeline stalls. Using VirtMCU's `slaved-icount` mode (which provides exact nanosecond virtual time), students measure the execution time differences between optimized and unoptimized assembly loops.

Chapter 5: The Memory Hierarchy

- **Lecture 9:** Caches and Spatial/Temporal Locality. The growing gap between CPU and DRAM speeds. Direct-mapped, set-associative, and fully associative caches. Cache coherency protocols (MOESI).
- **Lecture 10:** Virtual Memory and the MMU. Paging, page tables, the Translation Lookaside Buffer (TLB), and how the Memory Management Unit maps virtual addresses to physical addresses.
- **VirtMCU Homework:** Memory access profiling. Students write firmware that intentionally causes cache thrashing and TLB misses, observing the latency footprint.

Chapter 6: SoC Interconnects and Peripherals

- **Lecture 11:** Buses and Interconnect Protocols. Advanced Microcontroller Bus Architecture (AMBA). We cover AHB-Lite, APB for low-power peripherals, and AXI for high-speed packet-switched network-on-chip routing.
 - **Lecture 12:** I/O, UART, and Direct Memory Access (DMA). Memory-mapped I/O (MMIO) versus port-mapped I/O. We cover how peripherals interact with the CPU and how DMA controllers move data without CPU intervention.
 - **VirtMCU Homework:** Students write a bare-metal UART driver. They utilize VirtMCU's deterministic multi-node UART (`hw/rust/comms/chardev`), ensuring their driver respects virtual-timestamped delivery.
-

Part III: The Hardware-Software Interface (Lectures 13–18)

Objective: Transition from hardware design to systems programming, teaching students how software manages hardware events.

Chapter 7: Interrupts, Exceptions, and Bare-Metal Execution

- **Lecture 13:** Booting and the Exception Model. What happens at power-on. The vector table, Nested Vector Interrupt Controller (NVIC), and transitioning between privilege levels (User mode vs. Kernel/Supervisor mode).
- **Lecture 14:** Interrupt-Driven Programming. Context switching, saving machine state, and the difference between polling and interrupt-driven I/O.
- **VirtMCU Homework:** Building an interrupt handler. Students configure VirtMCU to simulate a hardware button press, writing an interrupt service routine (ISR) to handle the event.

Chapter 8: Real-Time Operating Systems (RTOS)

- **Lecture 15:** RTOS Fundamentals. Soft vs. firm vs. hard real-time systems. The importance of deterministic response times over high average throughput.
- **Lecture 16:** Task Scheduling and Synchronization. Priority-based preemption, round-robin scheduling, semaphores, and avoiding priority inversion and deadlocks.
- **VirtMCU Homework:** Students port a lightweight RTOS kernel onto their VirtMCU simulated board and schedule two conflicting tasks, measuring exact context-switch jitter.

Chapter 9: Introduction to Cyber-Physical Systems (CPS 2.0)

- **Lecture 17:** The Evolution to CPS 2.0. Transitioning from isolated embedded systems (CPS 1.0) to CPS 2.0. We cover the 3Cs: Computation, Communication, and Control, and how AI, Edge Computing, and SDN are reshaping infrastructure.
 - **Lecture 18:** Anatomy of a Cyber-Physical System. The feedback loop: sensing the physical world, processing data, and actuating changes.
 - **VirtMCU Homework:** System architecture design. Students map a physical scenario (e.g., a drone flight controller) to a CPS 2.0 architecture layout.
-

Part IV: VirtMCU and Continuous-Time Physics (Lectures 19–24)

Objective: Integrate software execution with physical mechanics, utilizing VirtMCU's unique capabilities to bridge the discrete-time cyber and continuous-time physical worlds.

Chapter 10: Co-Simulation and Virtual Time

- **Lecture 19:** The Problem with Standard Emulators. Why standard QEMU fails for CPS. The necessity of deterministic virtual time, external clock authority, and eliminating wall-clock scheduling jitter.
- **Lecture 20:** VirtMCU Clock Slaving. Understanding the `hw/rust/backbone/clock` module. We explore `slaved-suspend` mode for high throughput and `slaved-icount` mode for exact microsecond-precision hardware timers.
- **VirtMCU Homework:** Students run a firmware simulation that requires strict 10ms control loops. They will compare standard QEMU (where the loop drifts due to host OS jitter) against VirtMCU's deterministic clock.

Chapter 11: The Sensor/Actuator Abstraction Layer (SAL/AAL)

- **Lecture 21:** Bridging MMIO and Physics. How to translate raw firmware register reads/writes into continuous physical properties (e.g., force, acceleration, angle).
- **Lecture 22:** Physics Engine Lock-Stepping. Integrating VirtMCU with the MuJoCo physics engine via zero-copy shared memory. Simulating transfer functions and sensor noise.
- **VirtMCU Homework:** Inverted Pendulum Simulation. Students write a firmware controller that reads simulated gyroscope registers (SAL) and writes to motor PWM registers (AAL) to balance a pendulum simulated in MuJoCo.

Chapter 12: Distributed Control and Multi-Node Networks

- **Lecture 23:** Deterministic Networking. The challenge of distributed control in large-scale CPS (e.g., Virtual Power Plants). Designing for latency, packet loss, and semantic communication.
- **Lecture 24:** VirtMCU Federation Bus. Multi-node simulation using `hw/rust/comms/netdev`. Students learn how Ethernet frames are buffered and injected into the guest NIC only when virtual time reaches the stamped arrival time, avoiding UDP multicast jitter.

- **VirtMCU Homework:** Multi-node Virtual Power Plant. Students spin up two VirtMCU instances representing Distributed Energy Resources (DERs). They communicate over the deterministic bus to balance a simulated power load.
-

Part V: Advanced Topics in CPS Architecture (Lectures 25–28)

Objective: Expand the students' horizons into the modern challenges of security, data processing, and hardware acceleration.

Chapter 13: Security and Resilience in CPS

- **Lecture 25:** Vulnerabilities in Critical Infrastructure. The impact of cyberattacks on physical safety. Protection strategies for industrial serial networks and transitioning to software-defined networking (SDN) for resilience.
- **Lecture 26:** Hardware Security and Firmware Updates. Implementing TrustZone, memory protection units, and secure boot. Strategies for securing Firmware Over-The-Air (FOTA) updates with cryptographic keys and safe rollback mechanisms.
- **VirtMCU Homework:** Students implement an MPU policy in their firmware to restrict a malicious payload from accessing the actuator memory-mapped registers.

Chapter 14: Hardware Acceleration and the Future of CPS

- **Lecture 27:** SIMD, NEON, and DSP. Applying Single Instruction, Multiple Data (SIMD) for parallel processing of sensor arrays. Introduction to ARM NEON registers for vector mathematics and DSP filtering algorithms.
- **Lecture 28:** GPUs, Autonomous Vehicles, and the Edge. The architecture of modern self-driving vehicle computers (e.g., Tesla HW3/HW4). Fusing

LiDAR, RADAR, and vision via Convolutional Neural Networks (CNNs) processing at the edge.

- **VirtMCU Homework:** Vectorized Sensor Processing. Students rewrite a sensor data fusion loop using ARM NEON intrinsics, verifying the execution speedup within the VirtMCU environment.
-

Summary of Pedagogical Alignment

This book structure directly addresses your requirement to write a textbook for software-oriented graduate students. By starting with **Digital Logic (Chapter 1)**, we bridge their software knowledge to the hardware level. We subsequently move through **Architecture (Chapters 4-6)** and the **OS Machine Level (Chapters 7-8)**. Finally, the second half of the book is heavily dedicated to **CPS 2.0 (Chapters 9-14)**, relying completely on the **VirtMCU** framework to provide reproducible, zero-jitter, multi-node homework assignments that interact with continuous-time physics engines.

Preface: How to Use This Book

The exponential increase in computing power and the pervasive integration of digital systems into the physical world have fundamentally transformed engineering. We are no longer simply designing isolated embedded systems; we are engineering **Cyber-Physical Systems (CPS 2.0)** that require a deep, integrated understanding of both hardware microarchitecture and continuous-time physical dynamics.

This textbook was written to bridge the gap between traditional computer architecture and modern cyber-physical system design. It is designed to take students on a journey from the fundamental digital logic that makes up a processor, through the operating systems that schedule its tasks, and finally to the distributed networks, deterministic co-simulations, and domain-specific architectures that define the modern computing landscape.

Target Audience and Prerequisites

This book is intended for senior undergraduate and graduate students in computer engineering, electrical engineering, and computer science. Because modern CPS design sits at the intersection of hardware and software, the text assumes readers are familiar with the basics of computer science and have completed a first programming course in C or C++. Furthermore, since we explore how hardware controls physical actuators, a basic understanding of continuous-time mathematics and introductory electronics will be highly beneficial.

Organization of the Book

The material is systematically divided into distinct parts to manage the complexity of the subject matter:

- **Part I: Machine Organization and Microarchitecture (Chapters 1–6)** builds the foundation. It covers digital logic, the ARM AArch64 Instruction

Set Architecture (ISA), pipelined microarchitecture, the memory hierarchy, and System-on-Chip (SoC) interconnects.

- **Part II: The Hardware-Software Interface (Chapters 7–8)** crosses the boundary into software, focusing on bare-metal execution, exceptions, interrupts, and the critical scheduling algorithms of Real-Time Operating Systems (RTOS).
- **Part III: Cyber-Physical Systems and Co-Simulation (Chapters 9–11)** introduces the CPS 2.0 paradigm. It explores the 3Cs (Computation, Communication, and Control), the necessity of deterministic virtual time, and the Sensor/Actuator Abstraction Layers (SAL/AAL) required to couple discrete-event cyber simulators with continuous-time physics engines.
- **Part IV: Advanced Architectures and Security (Chapters 12–14)** looks to the future, examining the open-source RISC-V architecture, cyber-physical threat intelligence and resilience, and Domain-Specific Architectures (DSAs) that accelerate AI at the edge and in the cloud.

Using This Book in a Course

This book contains enough material for a rigorous, rapid-paced, single-semester introduction to computer architecture and CPS, or it can be expanded into a two-semester sequence to allow students more time to experiment in the lab.

For a Single-Semester Course (14-15 Weeks):

- **Weeks 1–5:** Focus heavily on **Chapters 1 through 5** to establish a strong hardware and microarchitecture foundation.
- **Weeks 6–8:** Transition to software with **Chapters 6 through 8**, emphasizing interrupts, memory-mapped I/O, and RTOS determinism.
- **Weeks 9–12:** Introduce the CPS paradigm using **Chapters 9 through 11**. This is a great time to introduce mid-term laboratory projects involving co-simulation.
- **Weeks 13–15:** Conclude with advanced topics based on instructor interest. An instructor focused on hardware design might prioritize **Chapter 12 (RISC-V)** and **Chapter 14 (DSAs)**, while an instructor focused

on networks and infrastructure might focus on **Chapter 13 (Security and Resilience)**.

For a Two-Semester Sequence:

- **Semester 1 (Computer Architecture):** Dedicate the entire first term to **Chapters 1 through 8**, allowing students to deeply digest the instruction set principles, pipelining, and memory hierarchies. The labs can focus on translating C code to ARM assembly, observing cache hits/misses, and writing bare-metal device drivers.
- **Semester 2 (Cyber-Physical Systems):** Dedicate the second term to **Chapters 9 through 14**. This allows for a much deeper exploration of control algorithms, distributed optimal power flow, and modifying the RISC-V instruction set.

Laboratory and Software Tools

The best way to learn digital design and system architecture is to do it. However, verifying hard real-time CPS constraints on standard operating systems is often frustrated by wall-clock scheduling jitter. To resolve this, this book includes companion homework assignments designed around the **VirtMCU FirmwareStudio**.

Detailed instructions on how to use the software tools in a course are provided alongside the text. The VirtMCU framework allows students to execute ARM and RISC-V binaries in a deterministic, `slaved-icount` mode. By pairing VirtMCU with external physics engines (such as MuJoCo) via zero-copy shared memory, students can gain hands-on experience balancing an inverted pendulum or stabilizing a drone using the precise cycle-accurate timings required in the real world.

Chapter 1: Fundamentals of Computer Architecture and Digital Logic

1.1 Introduction: The Art of Managing Complexity

The design of modern computer systems is a remarkable feat of engineering. A contemporary microprocessor contains billions of microscopic transistors, each switching billions of times per second. No human being could possibly understand or design such a system by writing continuous mathematical equations describing the movement of electrons through each individual transistor. To build these vastly complex systems, computer scientists and engineers rely on a systematic approach to managing complexity: the use of strict **hierarchical abstraction**.

Abstraction involves hiding lower-level details when they are not important to the current problem, allowing engineers to construct complex systems layer by layer. At the very bottom of the computing hierarchy lies solid-state physics, which dictates how electrons move through semiconductor materials. Above physics is the device level, consisting of individual transistors. Transistors are wired together to form analog and digital circuits, which are in turn organized into logic gates (such as AND, OR, and NOT). These logic gates are combined to build microarchitectural components like adders, multiplexers, and memories, which ultimately implement the Instruction Set Architecture (ISA)—the hardware-software interface. Operating systems and high-level application software sit at the very top of this abstraction stack.

1.2 The Digital Abstraction

In the physical world, variables such as voltage, frequency, and mass are continuous. However, designing a computer to perform precise mathematical calculations using continuous analog voltages is highly prone to noise and signal degradation. To overcome this, architects use the **digital abstraction**. Digital systems intentionally restrict their signals to a discrete set of values.

In a binary digital system, information is represented using just two states: 0 and 1, which correspond to the Boolean logic values FALSE and TRUE, or electrical states LOW and HIGH. By disciplining ourselves to use only these two discrete states, we can combine simple components into highly reliable and sophisticated systems.

Beneath this digital abstraction, the physical work is performed by Complementary Metal-Oxide-Semiconductor (CMOS) transistors. A transistor behaves fundamentally as an electrically controlled switch.

- **nMOS Transistors:** Turn ON (conduct current) when a high voltage (logic 1) is applied to the gate, and turn OFF when a low voltage (logic 0) is applied.
- **pMOS Transistors:** Turn ON when a low voltage (logic 0) is applied to the gate, and turn OFF when a high voltage (logic 1) is applied.

By arranging these nMOS and pMOS switches into complementary pull-up and pull-down networks, engineers create robust logic gates. Because digital logic circuits interpret continuous voltages falling within specific ranges as either a strict 0 or 1, they inherently reject minor analog noise, ensuring perfect signal reconstruction across billions of operations.

1.3 Historical Milestones in Computing

To appreciate modern architectures, it is helpful to look back at the technological milestones that brought us here.

The conceptual foundation of the digital computer predates the discovery of the electron. In the 1830s, Charles Babbage designed the **Analytical Engine**, widely considered the first attempt to build a general-purpose digital computer. The machine was designed to be entirely mechanical, constructed of brass and powered by a steam engine. It featured a mechanism for executing algorithms and outputting results, foreshadowing modern programmable architecture, though it was ultimately defeated by the inadequate manufacturing technology of its time.

Modern computer history truly began in 1946 with the completion of the **ENIAC** (Electronic Numerical Integrator and Computer) by J. Presper Eckert and John Mauchly. Composed of 18,000 vacuum tubes, ENIAC was the first general-purpose electronic computer. However, programming ENIAC required physically rewiring huge patch panels, a tedious and inflexible process.

The solution to this inflexibility came from John von Neumann, who helped develop the **Stored-Program Concept** with the EDVAC and IAS machines. Von Neumann realized that a program's instructions could be represented in digital form and stored in the computer's memory alongside the data it manipulated. This architecture—where memory holds both instructions and data, and a central processing unit fetches and executes those instructions sequentially—remains the foundation of almost every modern processor.

Over subsequent decades, the industry evolved rapidly from room-sized mainframes to minicomputers, and eventually to personal computers (PCs) and mobile devices. Today, computing spans an enormous continuum: from miniature embedded microcontrollers operating in the Internet of Things (IoT), to Personal Mobile Devices (PMDs) like smartphones, all the way up to massive Warehouse-Scale Computers (WSCs) powering global cloud infrastructure.

1.4 Moore's Law and the Power Wall

The astonishing proliferation of computing power over the past half-century was driven by two empirical observations: Moore's Law and Dennard Scaling.

In 1965, Intel co-founder Gordon Moore predicted that the number of transistors that could be economically integrated onto a single silicon chip would double approximately every 18 to 24 months. For decades, **Moore's Law** held true, granting architects a seemingly infinite budget of transistors to implement advanced features like superscalar execution and deep caches.

Simultaneously, the industry benefited from **Dennard Scaling**. In 1974, Robert Dennard observed that as the linear dimensions of a transistor shrunk, the operating voltage and current could scale proportionally. Since the dynamic power consumption of a CMOS circuit is proportional to $C \times V^2 \times f$ (where C is capacitance, V is voltage, and f is frequency), Dennard scaling implied that computer architects could continuously shrink transistors, pack more of them onto a die, and aggressively increase the clock frequency *without* increasing overall power density.

However, exponential laws eventually hit physical limits. Dennard scaling officially ended in the mid-2000s. As transistors became microscopic, threshold voltages could not be lowered further without causing massive leakage currents, establishing a non-scaling baseline for power consumption. The result was the **Power Wall**: chips became thermally limited, meaning architects could no longer simply increase clock frequencies to gain performance.

Moore's Law has also slowed significantly; the cadence between new semiconductor process nodes has stretched from 2 years up to 3 to 5 years. To continue improving performance within strict power limits, modern architectures rely heavily on multicore designs (placing multiple processors on a single chip) and Domain-Specific Architectures (DSAs) that aggressively optimize efficiency for specific workloads.

1.5 Quantitative Principles of Computer Design

Because transistor budgets are constrained by power limits and cost, modern computer architecture is an empirical science. Design decisions are made by analyzing the quantitative trade-offs between performance, price, and energy consumption.

When analyzing and optimizing architectures, computer scientists rely on several fundamental, quantitative principles:

1.5.1 The Principle of Locality

Programs do not access their code and data uniformly. Empirical observations reveal a widely held rule of thumb: a program typically spends 90% of its execution time in only 10% of its code. This is known as the **Principle of Locality**, and it manifests in two distinct ways:

1. **Temporal Locality:** If a memory location is accessed, it is highly likely to be accessed again in the near future (e.g., instructions inside a loop, or variables used as counters).
2. **Spatial Locality:** If a memory location is accessed, locations with adjacent or nearby addresses are likely to be accessed soon (e.g., sequentially executing code, or traversing an array of data).

Computer architects aggressively exploit locality by implementing memory hierarchies. Fast, small, and expensive SRAM caches are placed close to the processor to hold the active “working set” of data, while slower, cheaper DRAM holds the remainder.

1.5.2 Amdahl's Law

When considering an architectural enhancement (such as adding a specialized hardware accelerator), engineers must calculate the actual speedup it will provide to the overall system. This is governed by **Amdahl's Law**, which states that the overall performance improvement gained by optimizing a specific part of a system is strictly limited by the fraction of the execution time that the optimized part is actually used.

If an enhancement makes a specific operation N times faster, but that operation only accounts for 20% of the program's execution time, the remaining 80% of the execution time remains entirely unaffected. Ignoring Amdahl's Law is a common pitfall; engineers must always measure where a program spends its time *before* expending the resources to optimize a specific

hardware feature, lest they fall prey to what is often called “Amdahl’s heartbreaking law”.

1.5.3 Dependability via Redundancy

Computers must not only be fast and efficient; they must also be dependable. Any physical device will eventually fail. Therefore, architecture dictates that dependability is achieved through **redundancy**. If a single component (such as a power supply or a memory chip) is prone to failure, designers must include redundant components that can automatically take over. In modern systems, this ranges from Error-Correcting Codes (ECC) in memory arrays to fully redundant servers in a warehouse-scale datacenter.

With these foundational rules established, we are ready to venture into the details of digital logic. In Chapter 2, we will explore the mathematical language of computers—Boolean algebra and binary numbering—and demonstrate how physical transistors are combined into the logic gates that perform computation.

Chapter 2: Data Representation and Machine Organization

2.1 Introduction: The Meaning of Bits

In the previous chapter, we stripped away the software abstractions to reveal the physical reality of computation: metal-oxide-semiconductor field-effect transistors (MOSFETs) acting as voltage-controlled switches,. We established that modern digital systems operate on the **digital abstraction**, restricting continuous analog voltages into two discrete, reliable states representing logic 0 and logic 1.

However, a single bit—a solitary 0 or 1—carries very little information. Furthermore, the hardware itself has no inherent understanding of what these bits represent. Inside the computer, a number is simply a physical pattern of high and low voltages; it is the human programmer and the software architecture that assign meaning to these bit strings. Depending on the context of the program, a specific sequence of 32 bits might represent a signed integer, a floating-point fractional number, a string of text characters, an RGBA color pixel, or even an instruction dictating the next operation the processor must perform.

For software engineers transitioning into computer architecture and cyber-physical systems (CPS), understanding precisely how data is represented, organized, and accessed in hardware is paramount. When writing high-level software, the compiler hides the physical layout of memory. In CPS and embedded systems, however, you will frequently interact directly with hardware registers, communication protocols, and sensor data. A misinterpretation of data formats—such as confusing a signed integer for an unsigned one, or reading bytes in the wrong order—can lead to catastrophic failures in the physical world. This chapter bridges the gap between raw binary logic and structured data, exploring number systems, memory architectures, and machine-level organization.

2.2 Number Systems and Positional Notation

You are already intimately familiar with the decimal numbering system, which is a base-10 positional notation system utilizing the digits 0 through 9. The base, or radix, defines the number of unique symbols used. In any positional notation system, the value of a digit is determined by its position relative to the radix point. Moving a digit one place to the left increases its value by a factor of the base, while moving it one place to the right decreases its value by the same factor.

For example, the decimal number \$123.456\$ mathematically represents: $1 \times 10^2 + 2 \times 10^1 + 3 \times 10^0 + 4 \times 10^{-1} + 5 \times 10^{-2} + 6 \times 10^{-3}$.

2.2.1 Binary and Hexadecimal

Because computer hardware relies on the two states of electronic switches, computers natively use the **binary numbering system** (base-2), consisting only of the digits 0 and 1. In binary, each column's weight is a power of 2. For instance, the binary number \$11001010_2\$ translates to decimal as: $1 \times 128 + 1 \times 64 + 0 \times 32 + 0 \times 16 + 1 \times 8 + 0 \times 4 + 1 \times 2 + 0 \times 1 = 202_{10}$.

While binary perfectly matches the physical hardware, it is exceedingly verbose and difficult for humans to read. Representing the decimal number 202 requires eight binary digits but only three decimal digits. To mitigate this verbosity without losing the direct mapping to underlying bits, computer architects rely heavily on the **hexadecimal numbering system** (base-16).

Hexadecimal uses 16 distinct symbols: the digits 0 through 9, followed by the letters A through F to represent the decimal values 10 through 15. The mathematical beauty of hexadecimal lies in its relationship with binary: because $16 = 2^4$, exactly four binary digits (bits) map perfectly to one hexadecimal digit. By convention, hexadecimal numbers are often prefixed with `0x` in C-style languages or assembly (e.g., `0xDEADBEEF`) to distinguish them from decimal values.

Converting between binary and hexadecimal is a trivial process of grouping bits. To convert a binary number to hexadecimal, simply pad the binary number with leading zeros until its length is a multiple of four, separate it into groups of four bits, and substitute the corresponding hexadecimal digit. For example, the binary string `1010101111001101` groups into `1010_1011_1100_1101`, which directly translates to `0xABCD`.

2.2.2 Data Groupings: Bits, Nibbles, Bytes, and Words

Processors rarely manipulate individual bits in isolation. Instead, bits are grouped into standardized widths to form addressable data structures.

- **Nibble:** A grouping of 4 bits. A nibble can represent 16 distinct values (2^4) and perfectly corresponds to a single hexadecimal digit. It is also the size required to represent a single Binary-Coded Decimal (BCD) digit.
- **Byte:** A grouping of 8 bits. The byte is arguably the most fundamental data structure in modern computing, as main memory and I/O addresses are almost universally byte-addressable. An 8-bit byte can represent 256 distinct values (2^8), typically ranging from 0 to 255 for unsigned integers.
- **Half-Word, Word, and Double-Word:** While a byte is definitively 8 bits, the definition of a “word” depends strictly on the specific architecture of the processor—it represents the fundamental unit of data the CPU routinely handles. In the 64-bit ARM architecture (AArch64), a **half-word** is 16 bits, a **word** is 32 bits, and a **double-word** is 64 bits,. In contrast, in the x86 architecture, a word is 16 bits and a double-word is 32 bits, a holdover from the 16-bit 8086 processor era.

When working with groups of bits, we must establish a numbering convention. The rightmost bit is known as bit 0, or the **Least Significant Bit (LSB)**, because it carries the smallest mathematical weight (2^0). The leftmost bit is the **Most Significant Bit (MSB)**, carrying the greatest weight.

2.3 Integer Representation

Software engineers must acutely understand how the hardware interprets combinations of bits to represent mathematical integers, both positive and negative.

2.3.1 Unsigned Integers

If a data structure uses all available bits to represent a positive magnitude, it is an **unsigned integer**. An N -bit unsigned number represents one of 2^N possibilities, spanning the strictly positive range from 0 to $2^N - 1$. For a 32-bit word, this provides a range from 0 to $4,294,967,295$. While straightforward, unsigned integers cannot represent negative values, which are essential for many physical-world calculations, such as temperatures dropping below zero or reverse motor actuation.

2.3.2 Signed Integers: The Quest for Negativity

To represent negative numbers, the most intuitive human approach is to use a minus sign followed by a magnitude. Early computer architects attempted to mimic this with **Sign and Magnitude** representation, reserving the MSB as a sign bit (0 for positive, 1 for negative) and using the remaining $N-1$ bits for the absolute value,. While easy for humans to read, this system suffers from severe hardware drawbacks. First, it requires different, highly complex ALU logic to perform addition depending on the signs of the operands. Second, it creates two distinct representations for zero (a positive zero and a negative zero), which forces software to make double checks whenever testing for equality with zero.

An alternative historical approach was **Ones' Complement**, where a negative number is formed by simply inverting every bit of its positive counterpart (changing 0s to 1s, and 1s to 0s). This made negation extremely fast but still suffered from the "two zeros" problem (00000000 and 11111111 both representing zero in an 8-bit system),.

2.3.3 Two's Complement Notation

Modern computer systems universally solve these problems using **Two's Complement** representation,. Two's complement overcomes the shortcomings of earlier systems: zero has exactly one binary representation, and ordinary binary addition hardware works seamlessly for both positive and negative numbers without requiring separate subtraction circuits.

In a two's complement system, the MSB still acts as a sign bit (0 for positive, 1 for negative). However, the mathematical weight of the MSB is negative. For an N -bit number, the MSB has a weight of -2^{N-1} instead of 2^{N-1} .

To convert a positive binary number to its negative two's complement equivalent, you apply a two-step algorithm:

1. **Invert all the bits** (apply a logical NOT operation).
2. **Add 1** to the result, ignoring any carry out of the MSB.

For example, to find the 8-bit two's complement representation of -5_{10} :

- Start with positive 5: `0000_0101`
- Invert all bits: `1111_1010`
- Add 1: `1111_1011` .

If you add $+5$ (`0000_0101`) and -5 (`1111_1011`) together in hardware, the binary sum is `1_0000_0000` . Because the system only holds 8 bits, the ninth carry bit is discarded, leaving exactly `0000_0000` , the single, correct representation of zero.

Because there is no negative zero, an N -bit two's complement system represents an asymmetric range: $[-2^{N-1}$ to $2^{N-1} - 1]$,. For an 8-bit byte, this range is -128 to $+127$. The most negative number, -128 (`1000_0000`), is an anomaly often called the "weird number" because attempting to negate it via the two's complement algorithm (invert to `0111_1111` , add 1 to get `1000_0000`) simply yields -128 again; it has no positive counterpart in the 8-bit space,.

2.3.4 Sign Extension vs. Zero Extension

When moving data between variables of different widths—such as loading an 8-bit sensor reading into a 32-bit CPU register—the hardware must pad the extra 24 bits. If the 8-bit value is an *unsigned* integer, the hardware performs **zero extension**, simply filling the new upper bits with 0s. For instance, `0x82` becomes `0x00000082`. However, if the 8-bit value is a *signed* two's complement integer, zero-extending a negative number would accidentally convert it into a large positive number. Instead, the hardware must perform **sign extension**, copying the original sign bit (the MSB) into all the new upper bit positions. An 8-bit `-5` (`0xFB`) sign-extended to 32 bits becomes `0xFFFFF8`. Both represent the mathematical value -5 , ensuring that subsequent ALU operations remain accurate.

2.4 Fractional Numbers: Fixed and Floating Point

While integers are sufficient for counting, cyber-physical systems deal with continuous physical phenomena—velocity, voltage, mass—that require fractional representation.

2.4.1 Fixed-Point Representation

The simplest method to represent real numbers in binary is **fixed-point notation**, where an implicit, immovable binary point separates the integer bits from the fractional bits. For example, in an 8-bit system using 4 integer bits and 4 fractional bits (often denoted as Q4.4 format), the binary value `0110.1100` translates to: $0 \times 2^3 + 1 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 + 1 \times 2^{-1} + 1 \times 2^{-2} + 0 \times 2^{-3} + 0 \times 2^{-4} = 4 + 2 + 0.5 + 0.25 = 6.75_{10}$.

Fixed-point arithmetic is highly efficient because it utilizes standard integer ALU adders and multipliers (requiring only a subsequent bit-shift to realign the implicit binary point). This makes it ideal for power-constrained microcontrollers without dedicated math hardware. However, fixed-point

numbers suffer from a severely constrained dynamic range; you cannot represent a very large number and a very small precise fraction simultaneously.

2.4.2 Floating-Point Representation (IEEE 754)

To accommodate the vast dynamic ranges required in scientific computing and complex CPS control algorithms, architectures use **floating-point** representation, which operates similarly to scientific notation. A floating-point number represents a value as: $(-1)^{\text{Sign}} \times \text{Mantissa} \times 2^{\text{Exponent}}$.

Modern processors almost universally adhere to the **IEEE 754** floating-point standard, established in 1985,. The standard defines several precision formats, most notably:

- **Single-Precision (32-bit):** 1 sign bit, an 8-bit exponent, and a 23-bit mantissa (fraction).
- **Double-Precision (64-bit):** 1 sign bit, an 11-bit exponent, and a 52-bit mantissa.

To maximize precision, IEEE 754 assumes the mantissa is **normalized**. In binary, a normalized scientific number always has a single 1 to the left of the binary point (e.g., 1.011×2^3). Because this leading 1 is a guarantee for all normalized numbers, the IEEE 754 standard drops it from the physical bit storage, gaining an extra bit of precision “for free.” This implied bit, combined with the explicitly stored fraction, is known as the **significand**,.

The exponent is stored using an **excess (or biased) notation** rather than two’s complement. For single-precision, the bias is 127. This means an actual exponent of 0 is stored as 127 (0111_1111), an exponent of +1 is stored as 128, and an exponent of -1 is stored as 126. This bias guarantees that the stored exponent is always a positive, unsigned binary number, which allows the CPU hardware to compare two floating-point numbers using fast, standard integer comparison logic,.

The IEEE 754 standard also defines bit patterns for vital edge cases that commonly arise in sensor data processing:

- **Zero:** Represented by an exponent of 0 and a fractional field of 0. (The standard uniquely defines both a $+0.0$ and a -0.0),.
- **Infinity:** Represented by an exponent of all 1s and a fraction of 0. This handles overflow conditions gracefully, allowing calculations like $X / 0.0$ to yield $+\infty$ rather than crashing the system,,.
- **NaN (Not a Number):** Represented by an exponent of all 1s and a non-zero fraction. NaNs propagate through calculations to indicate mathematical impossibilities, such as $\sqrt{-1}$ or ∞ / ∞ ,.
- **Denormalized Numbers:** If a calculation underflows (becomes too small to represent with a normalized exponent), the hardware transitions to denormalized numbers, where the exponent is forced to 0 and the implied leading bit becomes 0. This provides a “graceful underflow” that trades precision to represent incredibly tiny values near zero.

2.5 Character Data Representation

Cyber-physical systems must occasionally communicate with human operators via logging, network interfaces, or diagnostic terminals. This requires representing alphanumeric characters as binary numbers.

The most ubiquitous legacy standard is **ASCII** (American Standard Code for Information Interchange), developed in 1963,. ASCII is fundamentally a 7-bit code mapped into an 8-bit byte, defining 128 unique characters,. The first 32 codes (0x00 to 0x1F) are non-printable control characters originally designed for teletype machines (e.g., Carriage Return 0x0D , Line Feed 0x0A),. The numeric digits ‘0’ through ‘9’ are mapped to 0x30 through 0x39 . Conveniently, the lowercase letters differ from their uppercase equivalents by exactly one bit (bit 5); for instance, ‘A’ is 0x41 (0100_0001) and ‘a’ is 0x61 (0110_0001), allowing software to toggle case with a simple bitwise XOR operation,.

Because 128 characters are entirely inadequate for international languages and technical symbols, modern systems primarily use **Unicode**, which assigns a unique 16-bit or 32-bit code point to virtually every character in human existence,,,,. To maintain backward compatibility with legacy C code and serial terminal protocols, Unicode is most frequently transmitted using **UTF-8**

encoding. UTF-8 is a variable-length encoding: standard ASCII characters consume only 1 byte, while complex symbols span 2, 3, or 4 bytes. UTF-8 is self-synchronizing, meaning that if a byte is corrupted over a noisy network line, the receiving software can easily resynchronize at the start of the next valid character.

2.6 Machine Organization: Architectures and Memory Maps

With our data formatted into bytes and words, we must examine how the processor architecture organizes and accesses this data. At a high level, general-purpose computer architectures fall into two primary structural families based on how they route data and instructions.

2.6.1 The Von Neumann Architecture

Proposed by John von Neumann in 1945, the **von Neumann architecture** utilizes a single, unified memory space that holds both the executable program instructions and the variable data the program operates upon,. The processor interfaces with this unified memory over a single set of shared buses (address, data, and control),,,.

This design is incredibly efficient from a hardware perspective, minimizing the physical pin count on the processor and simplifying memory allocation for the operating system. However, it suffers from the infamous “**von Neumann bottleneck.**” Because the CPU must fetch an instruction and then subsequently read or write the data for that instruction over the same bus, it requires multiple sequential clock cycles, severely limiting peak throughput,. Furthermore, placing executable code and writable data in the same unified memory space exposes the system to severe security vulnerabilities, such as buffer overflow attacks where malicious data is injected and subsequently executed as code,.

2.6.2 The Harvard and Modified Harvard Architectures

To bypass the bottleneck, the **Harvard architecture** employs two completely physically isolated memory spaces: one strictly for program instructions and one strictly for data,. Each memory has its own dedicated address and data buses. This allows the CPU to fetch the next instruction simultaneously while reading or writing data for the current instruction, enabling single-cycle execution for many operations,. Strict Harvard architectures are prevalent in dedicated Digital Signal Processors (DSPs), where deterministic, high-speed mathematical throughput is prioritized over programming flexibility,.

Modern System-on-Chip (SoC) microprocessors—such as the ARM Cortex-A series—utilize a **Modified Harvard Architecture**,. At the highest level, the main DRAM is unified (von Neumann), allowing the OS flexibility in loading programs. However, inside the processor core, the memory immediately adjacent to the execution unit is split into a separate **Level 1 Instruction Cache (L1 I-cache)** and **Level 1 Data Cache (L1 D-cache)**,. This grants the CPU the immense parallel bandwidth of the Harvard architecture during typical execution, falling back to the von Neumann bus only when a cache miss forces a read from main DRAM,.

2.6.3 Memory Maps and Direct Addressing

From the software engineer's perspective, whether the underlying architecture is von Neumann or Harvard, the CPU presents memory as a linear array of byte-sized "drawers," each numbered sequentially from 0 to $2^N - 1$, where N is the width of the address bus,. This logical arrangement is known as the **memory map**.

In embedded SoC development, the memory map does not exclusively point to RAM. Using **Memory-Mapped I/O (MMIO)**, the physical addresses decode to various distinct hardware components on the silicon die,. For example, addresses `0x0000_0000` to `0x000F_FFFF` might map to non-volatile Flash memory (for boot code), addresses `0x2000_0000` to `0x2003_FFFF` might map to high-speed volatile SRAM, and addresses starting at `0x4000_0000` might map

directly to the control registers of physical peripherals, such as UARTs, timers, and pulse-width modulation (PWM) motor controllers,.

2.7 Endianness: The Byte Ordering Battle

We now arrive at a subtle machine organization detail that routinely haunts software engineers writing network stacks and low-level drivers: **Endianness**.

As established, a 32-bit CPU operates on 32-bit words, but memory is heavily byte-addressable. This means a single 32-bit (4-byte) word spans four distinct memory addresses,. If a software engineer writes the 32-bit hexadecimal value `0x12345678` into memory starting at address `0x1000`, the hardware must decide *which* byte goes into address `0x1000`, `0x1001`, `0x1002`, and `0x1003` ,.

- **Little-Endian:** The *Least Significant Byte* (LSB, `0x78` in our example) is stored at the lowest memory address (`0x1000`). The Most Significant Byte (MSB, `0x12`) is stored at the highest address (`0x1003`),,. The x86 architecture is famously little-endian,.
- **Big-Endian:** The *Most Significant Byte* (MSB, `0x12`) is stored at the lowest memory address (`0x1000`). The LSB (`0x78`) is stored at the highest address (`0x1003`),,. Many classic network protocols (TCP/IP) and historical IBM mainframes mandate big-endian byte ordering,.

The terms derive from Jonathan Swift's 1726 satire *Gulliver's Travels*, in which two warring nations violently disagree over which end of a soft-boiled egg should be cracked first,,. While internally consistent as long as a processor only talks to itself, endianness becomes a critical point of failure when transmitting multi-byte data across a network or shared memory bus to a processor of the opposite endianness. If an x86 processor (little-endian) sends a 32-bit integer over a network to a classic SPARC processor (big-endian), the SPARC will read the bytes backward, interpreting `0x12345678` as `0x78563412` ,,,.

Many modern architectures, including ARMv8 (AArch64) and RISC-V, support **bi-endianness**, allowing the operating system to configure the processor to access memory in either big- or little-endian mode dynamically, though little-endian remains the default for most consumer operating systems,.

Furthermore, endianness directly impacts **unaligned memory transfers**. An aligned transfer means a 32-bit word is read from an address evenly divisible by 4 (e.g., `0x1000` , `0x1004`). If a programmer attempts to read a 32-bit word from an unaligned address (e.g., `0x1001`), some strict RISC processors will immediately crash with a hardware fault,,. More forgiving processors (like x86) will handle the unaligned read by silently issuing two separate, sequential hardware memory reads and stitching the bytes together in the CPU, incurring a massive unseen performance penalty that varies wildly based on the endianness configuration,,.

2.8 VirtMCU Homework: Inspecting Memory and Data Types

In the physical world, bugs related to data types and endianness do not always crash a program immediately; they often manifest as corrupted actuator outputs or bizarre sensor readings. To bridge the gap between C code and raw hardware organization, this chapter's homework relies on the **VirtMCU FirmwareStudio**.

You will be provided with a firmware binary compiled for the VirtMCU ARM Cortex-M target. This binary declares several global variables in C: an array of 8-bit characters, a signed 32-bit integer, and a 32-bit IEEE 754 floating-point number.

1. **Connecting the Debugger:** You will launch the VirtMCU QEMU instance and attach the GNU Debugger (`gdb`) to the live simulation port.
2. **Memory Inspection:** Instead of evaluating the variables by their C names, you will use `gdb` 's raw memory examination commands (e.g., `x/4xb &my_integer`) to dump the exact bytes residing at the physical memory addresses mapped by VirtMCU.
3. **Endianness Discovery:** By comparing the known 32-bit value in your C code to the physical byte sequence dumped by `gdb` , you will empirically verify the endianness of the simulated ARM core.
4. **Floating-Point Analysis:** You will dump the raw hexadecimal bytes of the floating-point variable, manually decode the IEEE 754 Sign, Biased

Exponent, and Significand bits, and calculate the decimal equivalent to prove the hardware's internal representation matches the mathematical theory discussed in this chapter.

Through VirtMCU, you will bypass the compiler's abstractions and witness firsthand how raw binary data is organized, stored, and retrieved in an active machine architecture.

Here is Chapter 3 of the textbook. I have meticulously expanded the explanations, provided concrete assembly language examples, and thoroughly explored the ARM64 Application Binary Interface (ABI) and call stack to ensure the text provides the necessary academic depth for your graduate students while strictly meeting your length constraints.

Chapter 3: Instruction Set Architecture (ISA) and ARM AArch64

3.1 Introduction to the Instruction Set Architecture

In the preceding chapters, we explored how continuous voltages are abstracted into discrete digital logic and how those digital states are grouped into bytes and words to represent meaningful data. However, data without manipulation is inert. To perform computation, a computer must be commanded by a sequence of instructions. The vocabulary of commands understood by a specific microprocessor is known as its **Instruction Set Architecture (ISA)**.

The ISA serves as the critical contract between the hardware engineers who design the silicon and the software engineers who write the compilers and applications. From the programmer's perspective, the ISA defines the architectural state of the processor: the memory model, the available registers, the data types supported, and the specific binary encodings required to force the hardware to perform operations such as addition, branching, and data movement. Because the ISA abstracts away the underlying microarchitecture—the physical arrangement of logic gates, pipelines, and caches—software written for a specific ISA can run unmodified on entirely different hardware implementations, provided both adhere to the same ISA standard.

Historically, ISAs fall into two primary design philosophies: **Complex Instruction Set Computer (CISC)** and **Reduced Instruction Set Computer**

(RISC). CISC architectures, most notably the x86 and x64 families that dominate personal computers, provide a vast, complex vocabulary of variable-length instructions. A single CISC instruction might fetch data from memory, perform a mathematical operation, and write the result back to memory. While this approach produces compact program code, the variable-length nature of the instructions requires highly complex silicon to decode and execute.

In contrast, **RISC** architectures, such as ARM and RISC-V, were designed from a clean slate to execute instructions as rapidly and simply as possible. RISC architectures follow a strict **load/store model**. Computational instructions operate strictly on data already held within the processor's internal registers. Accessing system memory is isolated entirely to explicit load and store instructions. Although a RISC processor requires more total instructions to perform a memory-to-memory calculation, the instructions themselves are of a fixed length (typically 32 bits), allowing the hardware decoder to be significantly simpler, the clock frequencies to be higher, and the execution pipelines to be deeper and more efficient.

3.2 The Evolution of ARM and AArch64

The ARM architecture, originally standing for Acorn RISC Machine, was first developed in 1983. Over the subsequent decades, ARM processors became the dominant architecture in embedded systems, mobile phones, and tablets due to their exceptional balance of high performance, low power consumption, and small physical footprint.

For most of its history, ARM was a 32-bit architecture. However, as mobile and computing workloads evolved to require processing massive datasets and accessing more than 4 Gigabytes of memory, the 32-bit address space became a critical bottleneck. In 2011, ARM announced the ARMv8-A architecture, introducing a full 64-bit architecture known as **AArch64**.

The AArch64 execution state provides a 64-bit virtual address space and introduces a completely new instruction set named **A64**. Unlike previous transitions in other architectures (such as x86 to x64) where the new architecture was merely an extension of the old, the A64 instruction set was

designed as a distinct, modern 64-bit ISA. While modern ARMv8-A and later processors maintain backward compatibility by supporting an AArch32 execution state for legacy 32-bit code, the two instruction sets are fundamentally incompatible; you cannot mix A64 and 32-bit ARM instructions in the same execution stream. Because modern high-performance computing, deep learning, and advanced Cyber-Physical Systems require the vast memory and throughput capabilities of 64-bit systems, this book focuses exclusively on the AArch64 state and the A64 instruction set.

3.3 The AArch64 Programmer's Model

Software engineers accustomed to high-level languages like C or Python are used to declaring an arbitrary number of variables. At the ISA level, however, the processor provides only a small, fixed number of ultra-fast memory locations directly on the silicon core. These are known as **registers**. Because AArch64 relies on a load/store architecture, registers act as the middleman for almost every calculation. Understanding the layout and purpose of these registers is the first step in mastering the programmer's model.

3.3.1 General-Purpose Registers (X0-X30, W0-W30)

The AArch64 architecture provides thirty-one general-purpose 64-bit registers, accessible in assembly language as **X0 through X30**. These registers can be used for arithmetic, logical operations, and memory address calculation.

To efficiently handle 32-bit data types (such as standard `int` variables in C), the architecture allows programmers to access the lower 32 bits of any X register by using the **W0 through W30** aliases. The X and W registers physically map to the exact same silicon. If an assembly instruction modifies a 32-bit W register, the processor automatically zero-extends the result, clearing the upper 32 bits of the corresponding 64-bit X register to zero. For example, writing the 32-bit value `0xFFFFFFFF` into W0 automatically forces X0 to become

`0x00000000FFFFFFFF`.

3.3.2 Special-Purpose Registers (SP, LR, FP, PC, PSTATE)

In addition to the general-purpose registers, the ARM CPU tracks its execution state using a set of dedicated special-purpose registers:

- **Stack Pointer (SP / X31):** The processor maintains a hardware stack in memory to track function calls and local variables. The 64-bit SP register points to the top of this active stack. Interestingly, the AArch64 architecture multiplexes the register index 31. Depending on the specific instruction context, referring to register 31 either accesses the Stack Pointer (SP) or acts as a dedicated Zero Register (**XZR** for 64-bit, **WZR** for 32-bit) which always reads as zero and discards any writes.
- **Program Counter (PC):** The 64-bit PC register contains the absolute memory address of the instruction currently being executed. Because A64 instructions are uniformly 32 bits (4 bytes) wide, the hardware automatically increments the PC by 4 upon the completion of each sequential instruction. Unlike in the 32-bit ARM architecture, user code cannot write directly to the PC register to cause a jump; it must use dedicated branch instructions.
- **Link Register (LR / X30):** When the processor executes a function call, it must remember the address to return to once the function finishes. The AArch64 architecture dedicates X30 as the Link Register (LR) for this purpose.
- **Frame Pointer (FP / X29):** Complex functions establish “stack frames” or activation records in memory. X29 is conventionally reserved as the Frame Pointer (FP) to serve as a reliable base address for accessing local variables and parameters allocated on the stack.
- **Process State Register (PSTATE):** The PSTATE register holds various system configurations and the highly critical **Condition Code Flags**. The flags indicate the results of the most recently executed arithmetic or logical instruction. The four primary flags are:
 - **N (Negative):** Set to 1 if the result of an operation is mathematically negative (i.e., the most significant bit is 1).
 - **Z (Zero):** Set to 1 if the result of the operation is exactly zero.
 - **C (Carry):** Set to 1 if an unsigned addition results in a carry-out, or an unsigned subtraction does not require a borrow.

- **V (Overflow):** Set to 1 if a signed arithmetic operation produces a result too large (or too negative) to fit in the destination register.

3.4 Assembly Language and Load/Store Mechanics

High-level variables must be transferred between the processor's registers and main memory. AArch64 utilizes the `ldr` (Load Register) and `str` (Store Register) instructions for this task.

3.4.1 Addressing Modes

To access memory, the hardware calculates an **Effective Address** using various addressing modes.

- **Register-Indirect:** The simplest addressing mode fetches data using an absolute memory address held in a 64-bit register. `ldr x0, [x1]` loads the 64-bit value located at the address held in X1 into the destination register X0.
- **Indirect-Plus-Offset:** It is incredibly common to access data sequentially, such as in an array. This mode adds a fixed offset to the base register. `ldr w0, [x1, #4]` calculates the effective address as $X1 + 4$ bytes, loading the resulting 32-bit word into W0.
- **Pre-indexed:** This mode adds an offset to the base register, accesses the memory, and then *permanently updates* the base register with the new address. This is denoted by an exclamation mark `!. str x0, [sp, #-16]!` drops the stack pointer by 16 bytes and then stores X0 at that new memory location. This is the fundamental mechanism for pushing data onto the stack.
- **Post-indexed:** This mode accesses the memory at the base register's *current* address, and only updates the base register with the offset *after* the access completes. `ldr x0, [sp], #16` reads 64 bits from the top of the stack into X0, and subsequently raises the stack pointer by 16 bytes. This is the fundamental mechanism for popping data off the stack.

3.4.2 Arithmetic and Data Movement

To copy data between registers, or to load a small constant into a register, AArch64 provides the `mov` instruction: `mov x1, x0` (Copies X0 to X1) `mov x2, #10` (Loads the immediate constant 10 into X2).

For mathematics, the `add` and `sub` instructions perform addition and subtraction. By default, these operations *do not* affect the PSTATE condition flags. This is a deliberate RISC design choice, allowing the compiler to reorder instructions without accidentally destroying flag states needed by subsequent branches. To force the hardware to update the N, Z, C, and V flags, the software must append an `s` suffix to the instruction mnemonic, utilizing `adds` or `subs`.

3.5 Translating High-Level Flow Control

Computers derive their power from the ability to make decisions. In high-level languages like C, this is achieved via `if/else` statements and `while` or `for` loops. In AArch64 assembly, all control flow boils down to comparing values and conditionally altering the Program Counter (PC) to branch to new memory addresses.

3.5.1 The Compare and Branch Instructions

Decision-making requires two steps: evaluating a condition, and branching based on the result. The **cmp (Compare)** instruction subtracts its right operand from its left operand, discards the mathematical result, and updates the PSTATE flags (N, Z, C, V). For example, `cmp x0, x1` performs $X0 - X1$. If $X0$ exactly equals $X1$, the subtraction yields 0, and the Z (Zero) flag is set to 1.

Following the `cmp`, the program uses a **Conditional Branch** instruction to test the flags and potentially jump to a new statement label. Common conditional branches include:

- `beq` (Branch if Equal): Jumps if the Z flag is 1.
- `bne` (Branch if Not Equal): Jumps if the Z flag is 0.

- `bgt` (Branch if Greater Than): Jumps if signed greater than (evaluates N and V flags).
- `blt` (Branch if Less Than): Jumps if signed less than.
- `bhi` / `blo` (Branch if Higher / Lower): Jumps based on unsigned comparisons.
- `b` or `bal` (Branch Unconditional): Forces a jump regardless of flags, equivalent to a `goto` statement.

3.5.2 Simulating If-Then-Else

Consider a standard C block:

```
if (x == y) {
    c = d;
} else {
    c = 0;
}
```

At the machine level, the compiler generates assembly language that evaluates the opposite condition to skip blocks of code. Assuming `x` is in X0, `y` is in X1, `c` is in X2, and `d` is in X3, the equivalent AArch64 code is:

```
    cmp x0, x1          // Compare x and y
    bne else_block      // If x != y, jump directly to the else block
    mov x2, x3          // THEN block: c = d
    b end_if            // Unconditionally jump over the else block
else_block:
    mov x2, xzr         // ELSE block: c = 0 (using the Zero
Register)
end_if:
```

Software engineers translating complex compound booleans (e.g., `if (x == y && z < t)`) rely on **short-circuit evaluation**. Once the first condition `x == y` is proven false, the program immediately jumps to the `else` block without wasting CPU cycles evaluating the `z < t` condition. The AArch64 architecture also provides a highly advanced `ccmp` (Conditional Compare) instruction designed specifically to chain Boolean evaluations natively in hardware, minimizing expensive branch operations.

3.5.3 Loops and Optimization

Loops execute a block of code repeatedly. A `while` loop checks a termination condition at the beginning of the block. If the condition is false, it branches past the loop body. If true, it executes the body and uses an unconditional branch `b` to jump back to the test.

However, computer architects and compiler writers employ strict optimization techniques to increase pipeline efficiency. The most common optimization is **moving the termination condition to the end of the loop**. A `do-while` (or `repeat-until`) loop evaluates its condition at the bottom. By compiling standard `for` and `while` loops to behave like `do-while` loops (preceded by a single initial check to ensure the loop should execute at least once), the compiler eliminates one unconditional branch instruction per iteration.

Furthermore, loops counting downward to zero are inherently faster than loops counting upward. Executing `subs x0, x0, #1` inherently sets the PSTATE flags, allowing a subsequent `bne` (Branch if Not Equal) to evaluate the loop termination without requiring an explicit `cmp` instruction against a target boundary. Understanding these hardware-level behaviors allows software engineers to write highly performant C code that maps gracefully to ISA optimizations.

3.6 The Application Binary Interface (ABI)

As software engineering modularized, the ability to write a function in C and reliably call a function written in C++, Assembly, or Rust became mandatory. For this interoperability to exist, there must be strict, universally agreed-upon rules governing how functions interact at the hardware level. This contract is the **Application Binary Interface (ABI)**.

While the ISA dictates how the silicon executes instructions, the ABI dictates how software uses the registers and memory. If a caller routine and a callee routine do not agree on where parameters are stored, the program will instantly crash or corrupt memory.

3.6.1 Parameter Passing and Return Values

The ARM64 ABI mandates that the first eight arguments of a function call are passed in the general-purpose registers **X0 through X7**. Passing parameters in registers is tremendously fast because it bypasses the slower main memory completely. If a function requires more than eight arguments, or the arguments are too large to fit in 64-bit registers (such as large C `struct`s), the additional parameters must be pushed onto the hardware stack in memory.

When the callee function finishes its work, the ABI dictates that integer and pointer return values must be placed in the **X0** register.

3.6.2 Volatile vs. Nonvolatile Registers

When Function A calls Function B, Function A might have critical calculations currently held in its registers. If Function B overwrites those registers, Function A's logic is destroyed upon return. The ABI solves this by classifying registers as either **volatile** (caller-saved) or **nonvolatile** (callee-saved).

- **Volatile Registers (X0-X15):** The ABI states that these registers are completely disposable across function calls. Function B is free to overwrite X0-X15 at any time. If Function A needs the data in X0-X15 to survive the function call, Function A must push those registers to the stack *before* calling Function B, and pop them back off afterward.
- **Nonvolatile Registers (X19-X28):** The ABI dictates that Function B *must not* destroy the values in these registers. If Function B wishes to use X19 for its own mathematics, Function B must push the original value of X19 to the stack upon entry, use the register, and pop the original value back into X19 before returning to Function A.

3.7 The Call Stack and Stack Frames

With the ABI rules established, we can examine the mechanics of invoking a function. A procedure call alters the flow of control, but unlike a standard

branch, it must return to the exact instruction immediately following the call once it completes.

3.7.1 Branch with Link (bl) and Return (ret)

In AArch64, a function call is initiated via the **bl (Branch with Link)** instruction. When the CPU executes `bl my_function`, it does two things simultaneously:

1. It jumps the Program Counter (PC) to the first instruction of `my_function`.
2. It copies the address of the *next* sequential instruction (the return address) directly into the **Link Register (LR / X30)**.

Once `my_function` completes its task, it executes the **ret (Return)** instruction. The `ret` instruction seamlessly copies the 64-bit address held in the Link Register back into the Program Counter, instantly returning execution to the caller.

However, this creates a major architectural challenge. Because the processor only has one Link Register, what happens if Function A calls Function B, and Function B needs to call Function C? When Function B executes its `bl` instruction, the hardware will overwrite the Link Register with the new return address for C, permanently obliterating the return address for A. If Function B then attempts to `ret`, it will enter an infinite loop or crash the system.

3.7.2 Building the Activation Record

To solve the destroyed Link Register problem, and to allocate memory for localized variables, functions must build an **Activation Record**, more commonly known as a **Stack Frame**, in main memory. The stack frame isolates the execution state of the current function from all other functions in the system.

Upon entry, an ABI-compliant function executes a “prologue” sequence. The most critical step is preserving both the Frame Pointer (FP) and the Link Register (LR) by pushing them to the stack. Because the ARM architecture

strictly demands that the Stack Pointer (SP) remains aligned on a **16-byte boundary** at all times, pushing the 8-byte FP and the 8-byte LR simultaneously is highly efficient.

The standard AArch64 stack frame prologue looks like this:

```
my_function:
    // 1. Push FP and LR to the stack, decrementing SP by 16 bytes
    stp fp, lr, [sp, #-16]!

    // 2. Set the Frame Pointer to the bottom of the new frame
    mov fp, sp

    // 3. Allocate N bytes of space for local variables (N must be a
    multiple of 16)
    sub sp, sp, #32
```

With this stack frame established, the function can freely use the `bl` instruction to call other subroutines, knowing its original return address is safely preserved in RAM. It can also use the Frame Pointer (`fp`) to reliably access its local variables (e.g., `ldr w0, [fp, #16]`).

When the function finishes, it executes an “epilogue” sequence to deallocate the stack frame, pop the original Frame Pointer and Link Register back into the CPU registers, and return:

```
// 1. Deallocate local variables
mov sp, fp

// 2. Pop FP and LR from the stack, incrementing SP by 16 bytes
ldp fp, lr, [sp], #16

// 3. Return to caller using the restored LR
ret
```

The stack frame discipline enables the most critical concept in computer science: **recursion**. Because each function invocation generates an entirely new, isolated stack frame in memory containing its own local variables and specific return address, a function can call itself infinitely (until physical memory is exhausted) without overwriting its previous state.

3.8 VirtMCU Homework: Stack Tracing and Deterministic Execution

In traditional operating systems, inspecting stack frames and observing execution state at the instruction level is obfuscated by OS schedulers and virtual memory mapping. To bridge this gap, you will utilize the **VirtMCU** framework for your Chapter 3 homework.

VirtMCU instantiates dynamic ARM machines relying on QEMU, utilizing dynamic QOM plugins to bypass standard emulation limits. Because VirtMCU implements *Cooperative Time Slaving*, the emulated processor executes deterministically.

1. **Translating C to AArch64:** You will be provided with a recursive C function that computes the factorial of an integer. Your first task is to manually translate this C code into strict AArch64 assembly, adhering completely to the ARM64 ABI. You must properly classify registers as volatile or nonvolatile, handle parameter passing via X0, and manually construct and destruct the stack frame (`stp fp, lr, [sp, #-16]!`) to preserve the Link Register during recursion.
2. **Deterministic Step Execution:** You will compile your bare-metal binary and load it into a VirtMCU instance running in `slaved-icount` mode. Because this mode provides exact nanosecond virtual time and locks the execution loop, you will be able to attach `gdb` and step through the binary cycle-by-cycle without any background host OS jitter interfering with the execution timing.
3. **Stack Frame Inspection:** As your recursive function calls itself, you will use the `gdb` memory inspection tools (introduced in Chapter 2) to manually dump the stack memory. You will trace the 16-byte alignments, identify where the `stp` instruction has written the sequential Link Register return addresses, and empirically verify the isolated local variables residing at fixed offsets from the Frame Pointer (`x29`).

Welcome back to class. In our previous lectures, we established the digital foundations of our hardware and explored the Instruction Set Architecture (ISA)—specifically ARM AArch64—which serves as the contract between software and hardware. Today, we cross the boundary from what the hardware *promises* to do, to how the hardware *actually does it*.

We are now entering Part II of our course: Modern System-on-Chip (SoC) Architecture. In Chapter 4, we will look inside the black box of the processor. We will design a processor microarchitecture from the ground up, starting with a simple single-cycle datapath, evolving it into a pipelined architecture, and finally examining the advanced superscalar and out-of-order execution engines that power modern high-performance computational nodes.

Here is the complete text for Chapter 4, prepared in Markdown format.

Chapter 4: Processor Microarchitecture

4.1 Introduction: Architecture vs. Microarchitecture

To a software engineer, the Instruction Set Architecture (ISA) defines the entirety of the processor. The ISA specifies the registers, the memory model, and the instructions available to the programmer. However, the ISA does not dictate how those instructions are physically executed.

The specific physical arrangement of registers, memories, Arithmetic Logic Units (ALUs), and multiplexers used to implement an ISA is known as the **microarchitecture**. A single architecture, such as ARMv8, can be implemented by many different microarchitectures. One microarchitecture might be optimized for ultra-low power consumption in a battery-operated cyber-physical system, while another might be optimized for massive throughput in a

cloud server. Though they share the same ISA and run the exact same compiled C code, their internal hardware designs differ vastly.

At the microarchitecture level, a processor is fundamentally a highly complex synchronous sequential digital circuit—essentially a giant finite state machine. The processor’s job is to read instructions and data, compute results, and update the architectural state on the rising edge of the system clock.

4.2 The Physics of Performance

Before designing our processor, we must define how we measure its performance. The execution time of a program is determined by the following fundamental performance equation:

$$\text{Execution Time} = \text{Instructions} \times \text{CPI} \times T_c$$

1. **Instructions:** The total number of instructions executed by the program. This is determined by the compiler and the ISA.
2. **CPI (Cycles Per Instruction):** The average number of clock cycles required to execute each instruction.
3. **T_c (Clock Period):** The duration of a single clock cycle (in seconds), which is the inverse of the clock frequency.

The challenge of the microarchitect is to choose a design that minimizes the overall execution time. As we will see, this is a delicate balancing act. Making the hardware simpler might decrease the Clock Period (T_c) but increase the CPI. Conversely, adding complex hardware to execute multiple instructions simultaneously will decrease the CPI, but the added logic delay might require a longer Clock Period.

4.3 The Single-Cycle Datapath and Control Unit

The simplest way to build a processor is to design a microarchitecture that executes one entire instruction in a single clock cycle. This is known as a single-

cycle processor. To build it, we separate the design into two interacting parts: the **datapath**, which contains the structural components that hold and manipulate data, and the **control unit**, which decodes the instruction and commands the datapath.

4.3.1 The Fetch-Decode-Execute Cycle

The execution of any instruction on a von Neumann architecture universally follows a sequence of discrete steps known as the fetch-decode-execute cycle:

1. **Fetch:** The processor uses the Program Counter (PC) to fetch the 32-bit instruction from the instruction memory,. Simultaneously, the processor calculates the address of the next instruction (typically $PC + 4$).
2. **Decode:** The processor parses the instruction bits to determine the operation type. Simultaneously, it reads the required source operands from the register file.
3. **Execute:** The Arithmetic Logic Unit (ALU) performs the requested operation, such as addition, subtraction, or logical AND, on the operands. For load/store instructions, the ALU adds the base register value to a sign-extended immediate offset to calculate the effective memory address.
4. **Memory:** If the instruction is a load or a store, the processor reads from or writes to the data memory at the address computed by the ALU.
5. **Write-back:** The result of the ALU computation or the data read from memory is written back to the destination register in the register file.

4.3.2 The Single-Cycle Control Unit

The datapath contains multiple interconnected hardware blocks, and multiplexers are used to route data between them. The control unit acts as the “brain,” driving the select lines of these multiplexers.

The control unit is purely combinational logic. It receives the opcode and function fields of the fetched instruction and instantly generates the necessary control signals. For example, if the instruction is a memory store (*sw*), the control unit will assert the `MemWrite` signal to enable writing to the RAM, but it

will de-assert `RegWrite` because a store instruction does not update the processor's registers. If the instruction is a conditional branch (`beq`), the control unit commands the ALU to subtract the two source registers. If the ALU outputs a `Zero` flag (indicating the registers are equal), the control unit switches a multiplexer to load the PC with the computed branch target address instead of the standard $PC + 4$.

4.3.3 The Single-Cycle Timing Bottleneck

In a single-cycle microarchitecture, the architectural state (the PC, the register file, and memory) is updated strictly on the rising edge of the clock. This means that the clock period (T_{clk}) must be long enough to allow electrical signals to propagate through the entire datapath.

The critical path—the longest and slowest sequence of logic—typically occurs during a memory load instruction. The signal must travel through the PC, into the instruction memory, through the register file, through the ALU to calculate the address, into the data memory to fetch the value, through a multiplexer, and finally back to the setup pins of the register file.

Because the clock cycle must be stretched to accommodate the slowest possible instruction, all other, faster instructions (like a simple addition) are artificially delayed. In modern deep sub-micron silicon, this makes the single-cycle processor incredibly inefficient, limiting the maximum clock speed to a few tens of megahertz.

4.4 Pipelining: Increasing Throughput

To overcome the bottleneck of the single-cycle design, computer architects employ **pipelining**. Pipelining divides the instruction execution process into a sequence of smaller, faster stages.

Consider the analogy of an automobile assembly line. If a team of mechanics builds one entire car from start to finish before beginning the next, the factory produces one car per day. However, if the factory is divided into stations—

chassis, engine, doors, wheels, paint—multiple cars can be assembled simultaneously. A finished car rolls off the line every hour, even though any individual car still takes a full day to build.

4.4.1 The Classic Five-Stage RISC Pipeline

Modern processors slice the datapath into multiple stages separated by hardware registers known as pipeline registers. The classic RISC pipeline consists of five stages,:

1. **Fetch (IF):** Read the instruction from the L1 Instruction Cache.
2. **Decode (ID):** Decode the instruction, resolve register addresses, and read operands from the Register File.
3. **Execute (EX):** The ALU performs the computation or address calculation.
4. **Memory (MEM):** Access the L1 Data Cache for load/store operations.
5. **Write-back (WB):** Write the final result back to the Register File,.

Because the pipeline registers isolate the stages, five different instructions can be in flight simultaneously.

4.4.2 Latency vs. Throughput

It is critical to understand that pipelining *does not* reduce the execution time of an individual instruction; in fact, the addition of pipeline registers introduces a slight propagation delay that actually increases the latency of a single instruction.

However, microprocessors process billions of instructions per second, so we prioritize **throughput** over single-instruction latency. In a perfectly balanced 5-stage pipeline, the clock frequency can be increased to nearly five times that of a single-cycle processor, because the clock period is now determined only by the delay of the single slowest pipeline stage, rather than the sum of all stages. Ideally, one instruction completes and retires every single clock cycle, achieving an ideal CPI of 1.

4.5 Pipeline Hazards and Mitigation

In reality, the CPI of a pipelined processor is always greater than 1. This is because instructions in a program are rarely completely independent. When one instruction depends on the result of a previous instruction that is still in the pipeline, a **hazard** occurs.

4.5.1 Data Hazards and Forwarding

A **data hazard** occurs when an instruction attempts to read a register that has not yet been written back to the register file by a preceding instruction. The most common data hazard is the Read-After-Write (RAW) dependency,.

Consider this sequence:

```
add x1, x2, x3    // x1 = x2 + x3
and x4, x1, x5     // x4 = x1 AND x5
```

The `add` instruction calculates the result in the EX stage, but it does not write the value into `x1` until the WB stage. The subsequent `and` instruction attempts to read `x1` during its ID stage. If we allow the pipeline to proceed normally, the `and` instruction will read the stale, outdated value of `x1` from the register file, causing catastrophic computational failure.

To resolve this, modern processors use **forwarding** (or bypassing),. The processor adds a Hazard Detection Unit that monitors the pipeline. If it detects that an instruction in the EX stage needs a value currently held in the MEM or WB pipeline registers, it activates multiplexers that intercept the data and route it directly into the ALU input, bypassing the register file entirely,.

4.5.2 Load-Use Stalls (Bubbles)

Forwarding solves most data hazards, but it cannot solve all of them. Consider a **Load-Use hazard**:

```
ldr x1, [x2]      // Load memory at x2 into x1
add x4, x1, x5     // x4 = x1 + x5
```

The `ldr` instruction does not actually retrieve the data from memory until the end of the MEM stage. However, the subsequent `add` instruction requires that data at the beginning of the EX stage. Because the data physically does not exist inside the processor yet, it is impossible to forward it back in time.

In this scenario, the hardware must physically **stall** the pipeline. The Hazard Detection Unit forces the `add` instruction to wait in the Decode stage for an extra clock cycle. To prevent the Execute stage from executing garbage, the control unit zeroes out the control signals, inserting a “bubble” (a hardware No-Operation) into the pipeline. Stalls degrade performance, increasing the CPI above 1,.

4.5.3 Control Hazards and Branch Prediction

The most severe threat to pipeline performance is the **control hazard**. When the processor fetches a conditional branch instruction (like `beq`), it does not actually know if the branch will be taken until the ALU evaluates the condition in the Execute stage. But by that time, the processor has already fetched two subsequent instructions into the pipeline.

If the pipeline simply pauses to wait for the branch to resolve, it would suffer massive delays. If it guesses the branch is not taken and fetches sequentially, it will execute the wrong code if the branch is actually taken. When this happens, the processor must **flush** the pipeline, throwing away the mistakenly fetched instructions and converting them into bubbles,. The cycles wasted fetching the wrong instructions are called the branch misprediction penalty.

To mitigate this, processors utilize complex **Branch Predictors**. A dynamic branch predictor utilizes a Branch Target Buffer (BTB) to cache the destination addresses of recent branches, alongside a state machine that tracks the history of the branch. A common design uses a 2-bit finite state machine for each branch. The FSM transitions between “Strongly Not Taken”, “Weakly Not Taken”, “Weakly Taken”, and “Strongly Taken”. This hysteresis ensures that a loop that

executes 100 times will only mispredict on the very last iteration when the loop finally exits. Modern branch predictors achieve accuracy rates exceeding 90%.

4.6 Advanced Microarchitecture: Superscalar and Out-of-Order Execution

The 5-stage pipeline represents a scalar processor, meaning it issues exactly one instruction per clock cycle. The absolute maximum throughput of a scalar processor is an IPC (Instructions Per Cycle) of 1. To achieve greater performance, engineers exploit **Instruction-Level Parallelism (ILP)**, executing entirely independent instructions simultaneously,.

4.6.1 Superscalar Issue

A **superscalar** microarchitecture duplicates the internal datapath hardware to fetch, decode, and issue multiple instructions in a single clock cycle,. A modern high-performance processor (such as the ARM Cortex-A series or Intel Core architectures) might fetch up to four or six instructions per cycle.

To support this, the execution unit is populated with multiple ALUs, Floating-Point Units (FPUs), and load/store units,. Superscalar processors boast an ideal IPC greater than 1 (or a CPI of a fraction of a cycle),.

4.6.2 Out-of-Order (OoO) Execution

In a strictly in-order pipeline, if an instruction stalls due to a cache miss, every single instruction behind it in the program is blocked, even if those subsequent instructions have their operands ready and are completely independent,.

Out-of-Order (OoO) execution allows the processor to look ahead across a vast “instruction window”. When a micro-operation is decoded, it is dispatched to a queue known as a Reservation Station. The scheduler monitors the reservation stations and dispatches instructions to the ALUs the exact moment

their data operands are ready, regardless of their original sequence in the compiled program code,.

Because instructions finish in a chaotic order, the processor utilizes a **Reorder Buffer (ROB)**. Instructions are retired from the ROB strictly in their original program order. This ensures that if a hardware exception or interrupt occurs, the architectural state of the machine remains perfectly precise and easily recoverable,,.

4.6.3 Register Renaming

Out-of-order execution introduces new hazards known as Write-After-Write (WAW) and Write-After-Read (WAR) dependencies. These are “false” dependencies caused by the compiler reusing a limited number of architectural registers (like `x0` through `x30` in ARM).

To eliminate these false dependencies, OoO processors implement **Register Renaming**. The processor contains a massive pool of hidden, internal physical registers. When an instruction decodes, a hardware mapping table dynamically assigns the architectural destination register (e.g., `x1`) to an unused physical register. This allows multiple instructions to write to `x1` simultaneously in the pipeline without corrupting each other’s data, allowing true ILP to flourish,.

4.7 Multithreading and Multi-core Architectures

There is a hard physical limit to how much ILP exists in a single program thread. Even with massively wide superscalar execution and deep OoO windows, processors eventually run out of independent instructions to schedule, leaving expensive functional units completely idle,.

4.7.1 Hardware Multithreading

To keep the execution units saturated, microarchitects implemented **hardware multithreading**, allowing a single physical processor core to present the

illusion of multiple logical processors to the operating system. The processor duplicates the Program Counter (PC) and the Register File for each thread, but shares the heavy execution hardware (the ALUs and memory buses),.

- **Coarse-grained Multithreading:** The processor runs one thread until it suffers a major stall (like an L2 cache miss), at which point it instantly switches context to a second thread to hide the memory latency,.
- **Fine-grained Multithreading:** The processor switches between threads every single clock cycle in a round-robin fashion, inherently masking data and control hazards.
- **Simultaneous Multithreading (SMT):** Found in almost all modern high-performance cores, SMT allows the superscalar scheduler to issue instructions from multiple independent threads into the execution units during the *exact same clock cycle*,.

As thermal dissipation limits (the “power wall”) prevented further increases in single-core clock frequencies, the industry transitioned to **multi-core** architectures, placing multiple independent CPU cores onto a single silicon die. Modern SoCs utilize a combination of multicore chips, superscalar pipelines, and SMT to achieve vast computational throughput.

4.8 VirtMCU Homework: Analyzing Pipeline Stalls and Deterministic Execution

The advanced microarchitectural optimizations we just discussed—deep pipelining, out-of-order execution, branch prediction, and multi-level caching—are brilliant for increasing average throughput. However, for Cyber-Physical Systems (CPS), these features create a nightmare scenario: **timing unpredictability**.

When controlling an aircraft actuator or a robotic joint, missing a control loop deadline by a few microseconds can cause catastrophic physical failure. In a standard OS running on a modern OoO core, it is virtually impossible to guarantee exact microsecond execution bounds due to the dynamic nature of branch mispredictions and cache misses.

In this homework, you will use the **VirtMCU FirmwareStudio** to analyze instruction latency at the bare-metal microarchitectural level.

1. **Code Profiling:** You will be provided with two versions of a C array-processing loop. The first is an unoptimized sequence containing severe Read-After-Write (RAW) Load-Use data hazards. The second version utilizes loop unrolling and instruction scheduling to place independent math instructions between the `ldr` and data consumption, eliminating the pipeline stall.
2. **Simulation under Jitter:** First, compile the unoptimized binary and run it in a standard, free-running QEMU emulator. Observe how the host operating system's scheduler creates wall-clock jitter, making it difficult to measure the exact nanosecond penalty of the data hazards.
3. **Deterministic `slaved-icount` Execution:** You will then load the binaries into VirtMCU configured in `slaved-icount` mode. Because VirtMCU strictly lock-steps the virtual clock to the exact number of retired instructions (bypassing host OS jitter), you will use the emulated hardware timer registers to measure the exact cycle count of the execution loops.
4. **Pipeline Analysis:** Calculate the empirical CPI of both the unoptimized and optimized assembly loops. You must submit a report identifying the exact assembly instructions where the processor was forced to insert a pipeline bubble due to a Load-Use data hazard, verifying your analysis against the cycle counts retrieved from the VirtMCU deterministic timer.

Chapter 5: The Memory Hierarchy

5.1 Introduction: The Memory Wall

In the preceding chapter, we examined the internal microarchitecture of the processor, demonstrating how pipelining, superscalar execution, and out-of-order scheduling can dramatically increase the instruction execution rate. However, a processor can only execute instructions as fast as it can be supplied with code and data. Historically, microprocessor speeds have doubled every two years, whereas the access latency of main memory has decreased at only a fraction of that rate. This ever-expanding divergence between CPU execution speed and memory latency is known as the “memory wall”.

If a modern processor running at 3 GHz (where one clock cycle takes 0.33 nanoseconds) were forced to access main memory directly for every instruction, it would have to wait hundreds of clock cycles for the data to arrive. The execution pipeline would starve, and the massive throughput gains of the microarchitecture would be entirely wasted.

To overcome this bottleneck, computer architects do not rely on a single, monolithic pool of memory. Instead, they construct a **memory hierarchy**. The hierarchy is organized as a pyramid: at the very top, residing inside the CPU core, are the processor registers, offering instant access. Below the registers is the **cache memory**, which is small but extremely fast, operating at speeds comparable to the processor. Below the cache is the main memory, which is much larger but significantly slower. At the base of the pyramid lies secondary storage (magnetic hard disk drives or solid-state drives), offering vast capacity but agonizingly slow access times measured in milliseconds. By intelligently migrating data up and down this hierarchy, the memory system provides the software engineer with an illusion: a memory pool that possesses the speed of the fastest cache and the immense capacity of the largest disk.

5.2 Primary Storage Technologies: SRAM and DRAM

To understand the trade-offs in the memory hierarchy, we must first look at the physical semiconductor technologies used to implement it. There are two primary types of volatile memory used in computer architectures: Static Random Access Memory (SRAM) and Dynamic Random Access Memory (DRAM).

5.2.1 Static RAM (SRAM)

SRAM is the fastest memory technology available and is utilized primarily for processor caches. A standard SRAM bit cell is constructed from six MOSFET transistors arranged as two cross-coupled inverters (forming a flip-flop) plus two access transistors,. Because the data is actively driven by a continuous feedback loop between the inverters, SRAM retains its state as long as power is applied without any need for external intervention,.

The flip-flop design provides incredibly low latency, allowing read and write operations to keep pace with high-frequency CPU clocks. Furthermore, reading an SRAM cell is non-destructive. The major downside to SRAM is its size and cost. Because it requires six transistors per bit, an SRAM array consumes a massive amount of silicon die area,. Additionally, the transistors in the cross-coupled inverters continuously leak a small amount of current, resulting in static power dissipation that grows severely as transistor sizes shrink.

5.2.2 Dynamic RAM (DRAM)

Main memory is almost universally implemented using DRAM. A DRAM bit cell is astonishingly simple: it consists of exactly one transistor and one microscopic capacitor,. A logical 1 or 0 is represented by the presence or absence of an electrical charge on the capacitor. Because it only uses one transistor per bit, DRAM achieves enormous data density, making it vastly cheaper per megabyte than SRAM.

However, this density comes with severe performance penalties. The capacitor naturally leaks charge over time, meaning the memory will “forget” its data within milliseconds,. To prevent data loss, the memory controller must continually read and rewrite the contents of the entire memory, a process known as **refreshing**,. Furthermore, reading a DRAM cell requires the bit-line to share charge with the microscopic capacitor, a delicate analog process that takes significant time and actually destroys the charge in the capacitor,. As a result, every DRAM read must be immediately followed by a hardware rewrite to restore the data. To manage this, DRAM chips use a multiplexed addressing scheme where a row (page) is activated first, followed by column accesses. This multi-step process makes DRAM latency hundreds of times slower than SRAM,.

5.3 Cache Memory Fundamentals

Because SRAM is too large and expensive to serve as main memory, architects insert small SRAM caches between the CPU and the DRAM,. The effectiveness of a cache relies entirely on the mathematical probability that a required piece of data is already present in the SRAM. This probability is driven by two empirical principles of software behavior known as **locality of reference**:

1. **Temporal Locality**: If a memory location is accessed, it is highly likely to be accessed again in the near future,,. A software engineer iterating a counter variable inside a `for` loop is demonstrating temporal locality. The cache exploits this by keeping a copy of recently accessed data.
2. **Spatial Locality**: If a memory location is accessed, memory locations with nearby addresses are likely to be accessed soon,,. A software engineer traversing an array of integers demonstrates spatial locality. The cache exploits this by never fetching just a single byte from main memory; instead, it fetches an entire contiguous **cache line** (or cache block), typically 32 or 64 bytes wide,.

When the processor requests data, it searches the cache. If the data is found, a **cache hit** occurs, and the data is served at SRAM speeds,. If the data is absent, a **cache miss** occurs, and the processor pipeline must stall while the hardware fetches the entire cache line from the slower main memory,.

The overall efficiency of the memory hierarchy is quantified by the **Average Memory Access Time (AMAT)**: $AMAT = Hit_Time + (Miss_Rate \times Miss_Penalty)$. To reduce AMAT, modern SoCs utilize multiple levels of cache. The **Level 1 (L1)** cache is tightly coupled to the processor pipeline, incredibly fast (often 1-3 cycles), but very small (e.g., 32 KB). It is usually split into an L1 Instruction Cache (I-cache) and L1 Data Cache (D-cache) to prevent structural hazards and enable simultaneous instruction fetches and data loads—a hallmark of the Harvard architecture. Below the L1 is the **Level 2 (L2)** cache, which is larger (e.g., 256 KB to 1 MB), somewhat slower, and usually unified (holding both instructions and data). Many high-performance systems add a massive **Level 3 (L3)** cache shared among multiple processor cores.

5.4 Cache Organization and Placement Policies

When a cache fetches a 64-byte block from main memory, it must decide exactly where in its internal SRAM array to place it. Because the main memory is gigabytes in size and the cache is only kilobytes, countless memory blocks must map to the same cache slots. The hardware rules governing this mapping dictate the cache's placement policy.

5.4.1 Direct-Mapped Caches

A direct-mapped cache is the simplest organization. The cache is divided into a one-dimensional array of "sets," where each set can hold exactly one cache line. A physical memory address is mapped to a specific set using a hardwired modulo operation.

To look up data, the cache controller splits the memory address into three fields:

- **Byte Offset:** The lowest bits specify which exact byte within the 64-byte cache line is requested.
- **Index (Set):** The middle bits determine which specific set in the cache array to check.

- **Tag:** The uppermost bits identify the exact main memory block currently residing in that set,.

Because each memory address maps to exactly one set, finding data is fast; the cache checks the set's Tag against the address's Tag. If they match, and the "Valid" bit is set, it is a hit,.

However, direct-mapped caches suffer from severe **conflict misses**. If a program actively bounces between two different memory addresses that happen to map to the exact same cache set (e.g., array elements spaced exactly by a multiple of the cache size), the cache will continually evict and reload the same set,. This pathological condition, known as thrashing, destroys performance.

5.4.2 Fully Associative Caches

To eliminate conflict misses entirely, a cache can be made **fully associative**. In this design, a memory block can be placed in *any* available slot in the cache. The address is only split into a Tag and a Byte Offset; there is no Index.

While this prevents arbitrary collisions, looking up data becomes a hardware nightmare. The cache controller must compare the requested Tag against the Tags of *every single line in the entire cache simultaneously*. This requires massive amounts of comparator logic (AND-gates and OR-gates), which consumes excessive die area and power, making it entirely impractical for large L1 or L2 caches,,.

5.4.3 N-Way Set-Associative Caches

Modern processor caches strike a compromise using **set-associative** mapping. The cache is divided into sets, but each set contains N separate slots (or "ways"),. For example, in an 8-way set-associative cache, a memory address still hashes to one specific set, but the cache line can be placed into any of the 8 available ways within that set.

During a lookup, the hardware reads all N tags in the indexed set in parallel and compares them to the requested address,. This effectively mitigates the “birthday paradox” of cache collisions without requiring the massive hardware overhead of a fully associative search. Most L1 caches today are 4-way or 8-way set-associative, offering an optimal balance between hit rate, hardware complexity, and access latency.

5.5 Cache Replacement and Write Policies

When a set-associative cache experiences a miss and all N ways in the target set are currently full, the cache controller must choose one existing line to evict.

The gold standard for eviction is the **Least Recently Used (LRU)** policy,. LRU relies on the principle of temporal locality, assuming that the cache line which has gone the longest without being accessed is the safest to discard,. True LRU is easy to implement for 2-way caches but becomes exponentially complex for 8-way or 16-way caches. Therefore, modern processors often use **Pseudo-LRU** (using a tree of binary flags to approximate age) or a simple Not-Recently-Used (NRU) clock algorithm,.

Handling memory stores (writes) introduces further complexity. If the CPU updates a variable, it modifies the copy in the cache. The cache and main memory are now out of sync. There are two primary write policies to resolve this:

- **Write-Through:** Every write to the cache is simultaneously propagated out to the main memory,,. This ensures memory is always perfectly up-to-date, making recovery easy. However, it severely wastes memory bus bandwidth. If a program increments a counter 1,000 times, write-through initiates 1,000 slow DRAM transactions.
- **Write-Back:** Writes only update the cache,,. A special metadata flag called the **dirty bit** is flipped to indicate the cache line contains modified data,. The line is only written back to DRAM when it is eventually evicted. This technique, called write coalescing, drastically improves performance by absorbing thousands of intermediate writes,.

Furthermore, if a CPU attempts to write to an address that is *not* currently in the cache, it must decide whether to fetch the block first. Modern processors usually employ a **write-allocate** policy, dragging the missing cache line from DRAM into the cache before modifying it, assuming the software will likely write to adjacent variables in the same block shortly.

5.6 Cache Coherence in Multicore Systems

As discussed in Chapter 4, to increase throughput, modern architectures place multiple processor cores on a single silicon die (Symmetric Multiprocessing or SMP). Each core possesses its own private L1 cache. If we utilize a high-performance write-back policy, we inevitably trigger the **cache coherence problem**.

Suppose Core A and Core B both load a shared memory variable X (currently 0) into their respective L1 caches. Core A executes $x = 5$. Because Core A uses a write-back cache, the new value 5 is stored only in Core A's L1 cache; main memory remains 0, and Core B's cache remains 0. If Core B subsequently reads X , it will see stale, invalid data. This violates the sequential consistency expected by software engineers writing multithreaded applications.

To prevent this, the hardware implements a **cache coherence protocol**. In smaller SMP systems, this is achieved via **snooping**. All cache controllers constantly monitor (snoop) the shared memory bus. When a core wishes to write to a shared variable, it broadcasts a signal.

The baseline snooping protocol is **MSI**, which assigns one of three states to every cache line:

1. **Modified (M)**: The line is dirty. This core owns the only valid copy in the system.
2. **Shared (S)**: The line is clean. Multiple cores may have copies. The core can read it, but cannot write to it without asking permission.
3. **Invalid (I)**: The line contains garbage.

If a line is Shared, and Core A wants to write to it, Core A must broadcast an “Invalidate” message on the bus (BusRdX). Core B hears this snoop message and downgrades its copy of XX to Invalid. Core A then transitions to Modified and performs the write. If Core B later wants to read XX , it broadcasts a read request. Core A snoops the request, pauses, flushes its Modified line back to DRAM, and then both cores transition back to Shared.

Modern processors (like the Core i7) use the extended **MESI** protocol, which adds an **Exclusive (E)** state. Exclusive indicates the line is clean, but this core is the *only* one holding it. If the core decides to write, it can transition from E to M silently, without broadcasting an expensive invalidate message on the bus.

Advanced ARM processors use the **MOESI** protocol, which adds an **Owned (O)** state,. The Owned state solves a major inefficiency: if Core A has a Modified line and Core B wants to read it, standard MESI forces Core A to flush the line to slow DRAM before Core B can read it. Under MOESI, Core A simply transmits the dirty data directly to Core B over the fast internal bus. Core A becomes the “Owner” (responsible for eventually writing it back to DRAM), and Core B enters the Shared state. This peer-to-peer cache transfer entirely bypasses the main memory bottleneck.

For enormous systems with hundreds of cores where broadcast snooping would saturate the bus, architects use **directory-based coherence**,. A centralized directory tracks exactly which cores hold which cache lines, allowing point-to-point invalidation messages rather than global broadcasts,.

5.7 Virtual Memory

In the early days of computing, software ran directly on physical memory. If an application needed 64 MB of RAM but the machine only had 32 MB, the programmer had to manually chop the program into “overlays,” shuttling segments of code back and forth from the disk—a tremendously tedious process.

To solve this, architects introduced **Virtual Memory**. The operating system and hardware conspire to give each running process the illusion of an immense,

private, contiguous address space (e.g., 4 GB on a 32-bit system or 256 TB on a 64-bit system), regardless of how much physical RAM is actually installed,.

This illusion is sustained via **paging**. The vast virtual address space is sliced into fixed-size blocks called **virtual pages** (typically 4 KB),. The physical RAM is similarly sliced into 4 KB chunks called **page frames**. At any given time, only the actively used virtual pages (the “working set”) reside in physical page frames. The rest are parked on the incredibly slow secondary storage (the hard disk or SSD), in an area known as the swap file.

If the CPU attempts to read a virtual address whose corresponding page is currently on disk, the hardware triggers an exception called a **page fault**,. The operating system traps the fault, pauses the program, evicts an old physical page frame (often using the Clock or LRU algorithm), fetches the requested 4 KB page from the hard drive, and then restarts the instruction,. From the software engineer’s perspective, the memory read simply took a very long time; the underlying mechanics are entirely transparent.

Beyond giving the illusion of infinite memory, paging provides strict **memory protection**. Because each process has its own isolated mapping, Process A cannot physically formulate an address that maps to Process B’s page frames. The MMU simply forbids it, throwing an access violation exception if attempted,.

5.8 Address Translation and the MMU

The translation of a programmer’s virtual address into a hardware physical address is performed by the **Memory Management Unit (MMU)**.

For a 32-bit system using 4 KB (2^{12} bytes) pages, a virtual address is mathematically divided into two parts:

- **Page Offset:** The lowest 12 bits. This defines the exact byte within the 4 KB page. Because the page size and frame size are identical, the offset requires absolutely no translation. It passes directly to the physical address,.

- **Virtual Page Number (VPN):** The upper 20 bits. The MMU uses the VPN as an index into a **Page Table** to find the corresponding Physical Page Number (PPN),.

The Page Table is a data structure maintained by the Operating System residing in physical memory. Each Page Table Entry (PTE) contains the mapped Physical Page Number, alongside metadata bits: a Valid bit (is it in RAM or on disk?), a Dirty bit (has it been modified?), and protection bits (Read/Write/Execute permissions),.

A flat page table mapping 4 GB of space requires 2^{20} (over 1 million) entries per process. To avoid wasting megabytes of contiguous RAM just for mapping tables, modern architectures use **multilevel page tables**,. The 20-bit VPN is split again (e.g., into a 10-bit Page Directory Index and a 10-bit Page Table Index),. The CPU first looks up the Page Directory, which points to a secondary Page Table, which finally points to the physical frame. This allows sparse allocation, where page tables for unused memory regions do not need to be instantiated,.

5.9 The Translation Lookaside Buffer (TLB)

Hierarchical page tables solve the memory footprint issue but introduce a catastrophic performance problem. Because the page tables are stored in main memory, translating a single virtual address requires two separate physical memory reads (one for the page directory, one for the page table) *before* the CPU can even attempt to read the actual data,. This would triple the execution time of every instruction.

To bypass this penalty, the MMU relies on a specialized cache called the **Translation Lookaside Buffer (TLB)**,,. The TLB is an ultra-fast, fully associative cache placed directly inside the MMU that stores the most recently used Virtual-to-Physical translations,.

When the CPU generates a virtual address, the MMU first blasts the Virtual Page Number into the TLB,. If a TLB Hit occurs (which happens >99% of the time due to spatial and temporal locality), the MMU instantly retrieves the Physical Page

Number and synthesizes the physical address in a single clock cycle, completely skipping the page table walk,.

If a **TLB Miss** occurs, the hardware (or a microcode software trap, depending on the architecture) must perform the expensive “page walk” through the hierarchical tables in DRAM, locate the translation, and load it into the TLB, evicting an older entry,. Because a TLB is fully associative to maximize hit rates and minimize conflict misses, it is hardware-intensive and usually quite small (holding only 64 to 512 entries),.

For a software engineer, understanding the TLB is critical for high-performance code. If a program randomly accesses memory across thousands of distinct 4 KB pages (e.g., traversing a poorly constructed linked list spread across a large heap), it will cause “TLB thrashing.” The TLB will constantly miss, forcing the hardware to perform page table walks, reducing the CPU’s throughput to a crawl even if the data itself resides comfortably in the L2 cache.

5.10 VirtMCU Homework: Memory Access Profiling

In standard software development on a host operating system, observing the exact microarchitectural penalties of cache misses and TLB thrashing is incredibly difficult. The host OS scheduler interrupts your process randomly, and background tasks continually pollute the L1/L2 caches and TLB, injecting massive amounts of wall-clock jitter into any performance measurement.

For this chapter’s homework, you will bypass these abstractions using **VirtMCU FirmwareStudio**. You will analyze memory hierarchy latency utilizing VirtMCU’s deterministic `slaved-icount` mode, which guarantees exact, jitter-free cycle counts.

1. **Cache Thrashing:** You will be provided with a C program that performs matrix multiplication. You will write two versions of the inner loop: one that iterates through the matrices in row-major order (respecting spatial locality) and one that iterates in column-major order (violating spatial locality).

2. **TLB Miss Measurement:** You will write a second program that initializes a large array of structs. One version of your code will traverse the structs sequentially. The second version will traverse the structs using a randomized stride that deliberately jumps across 4 KB page boundaries on every access, intentionally exceeding the TLB's 64-entry capacity.
3. **Deterministic Cycle Profiling:** You will compile these binaries for bare-metal ARM execution and load them into VirtMCU. By reading the virtual hardware timers in `slaved-icount` mode before and after your execution loops, you will measure the exact nanosecond penalty introduced by cache line fetches and page-table walks. You will document the performance divergence, proving empirically how software data structures interact with physical cache associativity and MMU translation buffers.

Chapter 6: SoC Interconnects and Peripherals

6.1 Introduction: Breaking Out of the Core

Welcome back. In our previous chapters, we constructed the processor pipeline and designed a deep memory hierarchy to keep the execution units fed with instructions and data. However, a processor that only talks to its own memory is effectively a brain in a jar; it can compute, but it cannot interact with the universe. For a Cyber-Physical System (CPS), interaction is the entire point. The system must sense physical properties, process them, and actuate mechanical responses.

To bridge this gap, modern microprocessors are integrated with numerous communication interfaces and hardware controllers onto a single silicon die, creating a System-on-Chip (SoC). A SoC is not merely a collection of isolated parts; it is a complex, distributed network of functional blocks. These blocks, ranging from low-speed sensors to high-speed graphics accelerators, must all communicate flawlessly.

In this chapter, we explore how data physically moves across the silicon. We will abandon the simplified abstraction of a monolithic “system bus” and dive into the mechanics of modern SoC interconnects. We will examine the Advanced Microcontroller Bus Architecture (AMBA), detailing the APB, AHB-Lite, and AXI protocols. Next, we will transition to the hardware-software interface, explaining how your C code interacts with these hardware peripherals through Memory-Mapped I/O (MMIO), interrupts, and Direct Memory Access (DMA). Finally, we will look at the Universal Asynchronous Receiver-Transmitter (UART) as a fundamental case study for external serial communication.

6.2 The Principles of SoC Interconnects

Historically, computer components were distributed across multiple chips on a printed circuit board and connected via shared, parallel copper wires—a traditional “bus”. In a classic shared bus, multiple devices physically share the same wires. To prevent electrical short circuits (a “bus fight”), devices use tri-state logic gates to disconnect themselves from the wires when not transmitting. However, shared buses suffer from severe physical and performance limitations. Only one device can transmit at a time, meaning bandwidth is inherently limited. Furthermore, as clock frequencies increase, the capacitance and physical length of shared wires cause signal degradation and propagation delays that cap the maximum speed.

In a modern SoC, the “bus” is no longer a single set of shared wires; it is an active **interconnect fabric**. The interconnect fabric multiplexes and routes transactions between components using dedicated, unidirectional point-to-point wiring.

6.2.1 Initiators and Targets

Communication across an interconnect always involves two roles:

- **Initiator (or Manager / Requester):** The device that starts a transaction by supplying an address and command. The CPU core and the DMA controller are the most common initiators.
- **Target (or Subordinate / Completer):** The device that responds to the transaction. Memory controllers, UART peripherals, and timers are common targets.

The interconnect fabric sits between the initiators and the targets. It intercepts all transactions issued by the initiators and routes them to the appropriate target based on the address.

6.2.2 Address Decoding and Routing

Every target connected to the fabric is assigned a specific range of addresses, defining the system's **memory map**. When an initiator issues a transaction, it drives the target address onto the interconnect.

The fabric contains an internal **address decoder**. The decoder compares the high-order bits of the requested address against a routing table to determine which target the initiator is trying to reach. For example, if a SoC places a UART controller at address `0x4000_0000`, any read or write starting with `0x40` will be routed specifically to the UART's hardware port.

Once the target is identified, the interconnect fabric uses multiplexers to physically connect the initiator's data and control lines to the target's receiving lines. To avoid complex hardware subtractors in the decoding logic, SoC memory maps generally place peripherals at boundaries that are multiples of powers of two. This allows the interconnect to route traffic simply by looking at a specific prefix of the address bits, functioning much like an IP subnet mask in computer networking.

6.2.3 Backpressure and Flow Control

In a heterogeneous SoC, an initiator running at 1 GHz may attempt to send data to a peripheral running at 10 MHz. If the initiator sends data faster than the target can process it, data will be lost. To prevent this, interconnect protocols implement **backpressure**.

Backpressure allows a slower component to stall a faster component. This is universally handled using a "ready-valid" handshake mechanism. The initiator asserts a `VALID` signal when its data and address are stable. The target asserts a `READY` signal when it is capable of accepting the data. The transaction only completes on the exact clock edge where both `VALID` and `READY` are asserted simultaneously. If the target is busy, it deasserts `READY`, forcing the initiator to hold its data on the lines and wait, effectively stalling the initiator's pipeline.

6.3 The AMBA Interconnect Protocols

To allow intellectual property (IP) blocks from different vendors (e.g., an ARM CPU, a Synopsys memory controller, and a custom accelerator) to communicate, the semiconductor industry relies on standardized interconnect protocols. The most ubiquitous standard is ARM's **Advanced Microcontroller Bus Architecture (AMBA)**. AMBA is not a single bus; it is a family of protocols designed for different performance, power, and area trade-offs. We will examine three critical AMBA protocols: APB, AHB-Lite, and AXI.

6.3.1 APB: The Advanced Peripheral Bus

Not all peripherals require high-bandwidth communication. Devices like timers, GPIO controllers, and UARTs have very low data throughput requirements. Providing these slow devices with complex, high-speed interconnect logic would waste precious silicon area and increase static power consumption.

The **Advanced Peripheral Bus (APB)** is designed specifically for these low-bandwidth control interfaces. APB is a non-pipelined, state-machine-based protocol. Every transaction takes a minimum of two clock cycles to complete.

The APB protocol utilizes a strict three-state machine: **IDLE**, **SETUP**, and **ACCESS**:

1. **IDLE**: The bus is inactive.
2. **SETUP**: When a transaction begins, the requester asserts a select signal (**PSEL**) for the target peripheral, along with the address (**PADDR**) and a read/write control signal (**PWRITE**). The bus remains in this state for exactly one clock cycle, giving the target peripheral time to wake up and prepare its internal logic.
3. **ACCESS**: On the next clock edge, the requester asserts the enable signal (**PENABLE**). If it is a write, the data is driven onto **PWDATA** . The target evaluates the request. If the target needs more time, it holds its **PREADY** signal low, extending the **ACCESS** state and applying backpressure. Once the target is finished, it drives **PREADY** high (and supplies read data on

`PRDATA` if applicable). The transaction completes, and the bus returns to `IDLE` or begins a new `SETUP` phase.

By guaranteeing that signals remain completely stable during the `SETUP` to `ACCESS` transition, APB eliminates unwanted circuit switching (glitches), significantly reducing dynamic power consumption.

6.3.2 AHB-Lite: The Advanced High-performance Bus

While APB is excellent for slow peripherals, it is far too slow for main memory or high-speed DMA transfers. For high-bandwidth components, designers use the **Advanced High-performance Bus (AHB-Lite)**.

AHB-Lite achieves high throughput by abandoning the slow state-machine approach of APB in favor of a strictly **pipelined** architecture. In AHB-Lite, every transaction is split into two distinct phases: an **Address Phase** and a **Data Phase**.

- **Address Phase:** Lasts exactly one clock cycle. The manager drives the address (`HADDR`), the transfer direction (`HWRITE`), and the size of the transfer (`HSIZE`).
- **Data Phase:** Begins immediately after the Address Phase. The manager drives write data (`HWDATA`), or the subordinate provides read data (`HRDATA`).

Because of the pipelining, the *Data Phase of the current transaction overlaps exactly with the Address Phase of the next transaction*. This means that while the memory is fetching data for address `$A$`, the manager is already transmitting address `$A+4$` on the bus. This overlapping allows AHB-Lite to achieve 100% bus utilization during sequential reads or writes, effectively transferring one word per clock cycle.

To apply backpressure, an AHB-Lite subordinate pulls the `HREADY` signal low during the Data Phase. Because the pipeline is locked together, stalling the Data Phase of the current transaction inherently forces the Address Phase of the subsequent transaction to stall as well.

AHB-Lite also supports **Bursts** (**HBURST**). A burst allows a manager to inform the subordinate that it plans to perform a sequence of 4, 8, or 16 consecutive transfers. This allows advanced memory controllers, like SDRAM, to optimize their internal row-buffer accesses before the subsequent addresses are even placed on the bus.

6.3.3 AXI: The Advanced eXtensible Interface

As processor speeds pushed into the gigahertz range and multi-core SoCs emerged, even the pipelined AHB-Lite protocol became a bottleneck. In AHB-Lite, the Address and Data phases are rigidly locked together; if a memory read stalls, the entire bus pipeline is blocked, preventing any other transactions from occurring. Furthermore, AHB-Lite uses a single set of shared control signals for both reading and writing, meaning a system cannot read and write simultaneously (half-duplex).

To solve this, ARM introduced the **Advanced eXtensible Interface (AXI)**. AXI is a radical departure from traditional bus architectures. Instead of a single monolithic bus, AXI defines **five completely independent, unidirectional channels** between a manager and a subordinate:

1. **Read Address Channel (AR):** Carries the address and control info for a read.
2. **Read Data Channel (R):** Carries the requested data from the subordinate back to the manager.
3. **Write Address Channel (AW):** Carries the address and control info for a write.
4. **Write Data Channel (W):** Carries the payload data from the manager to the subordinate.
5. **Write Response Channel (B):** Carries a success/failure acknowledgment from the subordinate back to the manager.

Each of these five channels operates completely independently, using its own **VALID** and **READY** handshake signals. This physical decoupling yields immense performance benefits.

First, it enables **Full-Duplex** operation; a DMA controller can write data on the W channel at the exact same time it is receiving completely unrelated read data on the R channel, doubling peak throughput.

Second, AXI supports **Outstanding Transactions** and **Out-of-Order Completion**. A CPU can blast five different read addresses down the AR channel without waiting for the data to return. The memory controller receives all five addresses and can reorder the reads based on which physical DRAM banks are currently open. When the memory controller finds the data, it sends it back on the R channel. Because the data might return out of order, every transaction is tagged with a unique Transaction ID (`ARID` / `RID`). The CPU matches the incoming data to the original request using this ID, much like a restaurant patron matching their receipt number to an order.

Because AXI is complex, ARM also provides **AXI-Lite**, a stripped-down version that removes burst capabilities and IDs. AXI-Lite provides AXI's channel decoupling but is simple enough to interface with standard registers.

6.4 The Hardware-Software Interface

Having established how data traverses the silicon via interconnects like AXI and APB, we must examine how the software engineer actually commands this hardware. In high-level languages like C, we manipulate variables in memory. But how do we tell a physical UART peripheral to transmit a byte?

6.4.1 Memory-Mapped I/O (MMIO) vs. Port-Mapped I/O

There are two architectural paradigms for interacting with peripherals: Port-Mapped I/O and Memory-Mapped I/O.

In older architectures like the Intel x86, peripheral devices exist in a completely separate, isolated address space. To access them, the CPU provides special, dedicated assembly instructions like `IN` and `OUT`. When the CPU executes an `OUT` instruction, it asserts a specific physical pin on the processor (the `M/IO`

pin) to tell the hardware, “I am talking to a peripheral, not main memory”. This is known as **Port-Mapped I/O**.

The ARM architecture (and almost all modern RISC designs) abandons this approach in favor of **Memory-Mapped I/O (MMIO)**. In MMIO, there are no special I/O instructions. The peripheral’s control and data registers are mapped directly into the standard, unified physical memory address space. From the CPU’s perspective, a peripheral register is indistinguishable from a standard RAM location. To send data to a peripheral, the software simply executes a standard assembly `STR` (Store Register) instruction targeting the specific memory address assigned to that device.

This allows the C programmer to interact with hardware simply by casting an integer address to a volatile pointer:

```
#define UART_TX_REG (*(volatile uint32_t*) 0x40001000)
UART_TX_REG = 'A'; // Hardware automatically routes this to the UART
```

The `volatile` keyword is critical here. It explicitly commands the C compiler *never* to optimize away the memory access or cache the value in a CPU register, guaranteeing that every read or write genuinely reaches the hardware. Furthermore, MMIO regions are explicitly marked as “Device Memory” in the MMU page tables, bypassing the L1/L2 caches to ensure immediate, strongly-ordered electrical delivery to the peripheral.

6.4.2 Programmed I/O (Polling)

The simplest method to manage a peripheral is **Programmed I/O** (often called polling). In this approach, the CPU executes a continuous loop, reading the peripheral’s status register until a specific bit indicates the device is ready.

For example, to receive a byte from a UART, the CPU repeatedly reads the UART’s line status register checking the `Data Ready` flag. Once the flag becomes `1`, the CPU reads the data register to retrieve the byte. While trivial to program, polling is highly inefficient. The CPU wastes thousands or millions of clock cycles spinning in a tight loop, unable to perform useful computation or

enter low-power sleep states. In modern, multitasking operating systems, polling is strictly avoided for routine operations.

6.4.3 Interrupt-Driven I/O

To eliminate the massive waste of CPU cycles caused by polling, systems rely on **Interrupt-Driven I/O**. Instead of the CPU constantly asking the peripheral if it is ready, the peripheral notifies the CPU when an event occurs.

Peripherals are wired directly to the CPU's **Nested Vector Interrupt Controller (NVIC)**. When a UART receives a byte, its hardware asserts an interrupt request (IRQ) line. The NVIC detects this, immediately pauses the CPU's current execution thread, pushes the CPU state to the stack, and forces the Program Counter to jump to an **Interrupt Service Routine (ISR)**. The ISR reads the byte from the UART's data register, places it into an operating system queue, clears the hardware interrupt flag, and returns execution to the original thread.

Interrupt-driven I/O allows the CPU to execute completely unrelated applications, or drop into deep sleep modes, waking only for a few microseconds when a physical event requires attention.

6.4.4 Direct Memory Access (DMA)

While interrupts solve the problem of polling, they are still too slow for massive data transfers. If a gigabit Ethernet controller receives a 1,500-byte packet, generating 1,500 individual CPU interrupts to copy the data byte-by-byte into RAM would overwhelm the processor, a problem known as "interrupt storming".

To handle bulk data, SoCs utilize a **Direct Memory Access (DMA) Controller**. A DMA controller is an independent hardware engine that sits on the interconnect (typically as an AXI Manager) capable of generating its own read and write transactions.

When the CPU needs to transfer data, it configures the DMA controller's MMIO registers, providing a source address, a destination address, and a block size.

The CPU then goes back to its normal tasks. The DMA controller autonomously arbitrates for interconnect bandwidth and transfers the data directly between the peripheral and main memory, completely bypassing the CPU execution pipeline. Because the DMA controller shares the memory interconnect, it performs **cycle stealing**, slipping its transactions onto the bus during cycles when the CPU is idle or accessing L1 cache. Once the entire multi-kilobyte block is successfully transferred, the DMA controller fires a single interrupt to notify the CPU that the memory buffer is ready for software processing.

6.5 Core SoC Peripherals

With an understanding of how they connect to the system, we can briefly examine the internal architecture of the most ubiquitous CPS peripherals.

6.5.1 General-Purpose Input/Output (GPIO)

The most fundamental interface is the **GPIO** controller. GPIOs allow software to directly manipulate the physical voltage levels on the external pins of the SoC chip. A GPIO block contains several MMIO registers:

- **Direction Register:** Sets each pin as an input (read sensor) or output (drive actuator).
- **Data Register:** If configured as an output, writing a 1 or 0 physically drives the pin to V_{DD} or GND. If configured as an input, reading this register samples the current digital voltage on the pin.
- **Pull-up/Pull-down Configuration:** Activates internal resistors to prevent “floating” inputs that could cause CMOS shoot-through current.

6.5.2 The Universal Asynchronous Receiver-Transmitter (UART)

While GPIO is useful for simple logic, exchanging complex data with external microcontrollers or diagnostic terminals requires serial communication. The

UART is a ubiquitous, legacy protocol that transmits data sequentially over a single wire without a shared clock signal.

Because there is no clock wire to synchronize the sender and receiver, a UART relies on strict timing agreements. Both ends must be configured to the exact same **baud rate** (symbols per second).

A standard UART transmission frame operates as follows:

1. **Idle State:** The line is held at logic HIGH (1).
2. **Start Bit:** The transmitter pulls the line LOW (0) for exactly one baud period. This alerts the receiver that a frame is beginning, allowing it to synchronize its internal clock timers.
3. **Data Bits:** The transmitter sends 5 to 8 bits of payload data, usually Least Significant Bit (LSB) first.
4. **Parity Bit (Optional):** A hardware-generated bit used for basic error detection (e.g., ensuring the total number of 1s is even).
5. **Stop Bit:** The line is driven HIGH (1) to return to the idle state and guarantee a clear transition for the next start bit.

Because modern CPUs run billions of times faster than standard UART baud rates (e.g., 115,200 bps), UART IP blocks incorporate hardware **FIFO (First-In, First-Out) buffers**. When the CPU wishes to print a string, it blasts the entire string into the UART's TX FIFO using MMIO in just a few nanoseconds. The hardware state machine then sequentially shifts the bits out over the physical wire at the slow baud rate, eventually raising an interrupt when the FIFO is empty and ready for more data.

6.5.3 Synchronous Serial: SPI and I2C

While UART is asynchronous, faster peripheral communication is achieved using synchronous protocols that include a dedicated clock wire.

- **SPI (Serial Peripheral Interface):** A full-duplex protocol utilizing four wires: Master Out Slave In (MOSI), Master In Slave Out (MISO), Serial Clock (SCK), and Chip Select (CS). Because the master explicitly drives the clock,

SPI can easily run at tens of megahertz, making it ideal for interfacing with external flash memory or LCD screens.

- **I2C (Inter-Integrated Circuit):** A half-duplex, multi-master protocol utilizing only two wires: Serial Data (SDA) and Serial Clock (SCL). I2C uses open-drain signaling with pull-up resistors and requires devices to broadcast target addresses over the bus before data transfer. It is slower than SPI but drastically reduces pin count, making it standard for reading low-speed sensor arrays (e.g., accelerometers, temperature probes).

6.6 VirtMCU Homework: Building a Deterministic UART Driver

In traditional embedded systems education, writing a bare-metal UART driver introduces students to the harsh realities of physical timing. If your C code relies on polling, and the host operating system scheduling your emulator preempts your thread for a few milliseconds, the simulated UART peripheral will drop characters, leading to corrupted strings and immense frustration. Standard virtual machines fail to replicate the strict, continuous-time relationships required for asynchronous serial communication.

To resolve this, this chapter's homework relies on the **VirtMCU** framework. Because VirtMCU utilizes a dynamic QOM plugin architecture to enforce deterministic virtual time via the `hw/rust/backbone/clock` module, time *stops* for the simulated SoC when the simulator pauses.

1. **Memory Map Investigation:** You will inspect the VirtMCU hardware definition files to locate the base address of the APB-attached UART peripheral.
2. **MMIO Driver Implementation:** You will write a C driver utilizing volatile pointers to interact with the UART's memory-mapped control, status, and data registers. First, you will implement a polled transmission function, writing a string to the TX register and spinning until the `TX_EMPTY` flag is asserted.
3. **Deterministic Evaluation:** You will launch the VirtMCU instance and connect a simulated serial terminal to the `hw/rust/comms/chardev`

deterministic backend. Because VirtMCU strictly lock-steps the UART baud-rate generation to the exact number of retired instructions in the `slaved-icount` mode, you will observe flawless character delivery. You will then intentionally configure a baud rate mismatch between your driver and the terminal to observe framing errors, empirically proving how asynchronous protocols fail without precise timing synchronization.

Welcome to Part III of our course: The Hardware-Software Interface. We have spent the first half of this textbook exploring the digital logic that forms a processor, the microarchitecture that executes instructions, and the complex memory and interconnect hierarchies that move data across the silicon. We now reach a pivotal transition. Hardware without software is merely a space heater; to create a functional Cyber-Physical System (CPS), we must write software that commands the silicon.

Because our target audience consists of graduate students transitioning from high-level software engineering, you are likely accustomed to writing applications that run on top of a sophisticated operating system (OS) like Linux or Windows. In those environments, the OS provides a comfortable, abstracted virtual machine. If you want to read from a file or send data over a network, you invoke a system call, and the OS handles the complex, messy interactions with the physical hardware.

However, in the world of embedded systems and CPS, you are often programming “bare-metal.” There is no operating system hiding the hardware from you; your code executes directly on the processor, meaning you are responsible for everything from the moment power is applied to the chip. This chapter bridges the gap between hardware and software by exploring how a processor boots, how it handles unexpected events through the exception model, how it interacts with peripherals via interrupts, and how it isolates trusted code using privilege levels.

Here is the complete text for Chapter 7, formatted in Markdown.

Chapter 7: Interrupts, Exceptions, and Bare-Metal Execution

7.1 Introduction: The Bare-Metal Environment

In a high-level programming environment, a software engineer rarely concerns themselves with the physical state of the machine at power-on. When a user runs a C program, the OS allocates virtual memory, sets up the execution environment, and smoothly transfers control to the `main()` function.

In bare-metal programming, however, `main()` does not execute by magic. The hardware strictly defines what happens the instant power is applied to the processor. A system must undergo a highly choreographed sequence of hardware and software initializations—collectively known as the boot process—to prepare the environment for high-level language execution.

Furthermore, once the application is running, it must interact with the continuous-time physical world. In Chapter 6, we introduced the concept of programmed I/O, or polling, where the CPU continuously loops to check the status of a peripheral. Polling is incredibly inefficient because it wastes thousands or millions of clock cycles spinning in a tight loop, preventing the CPU from performing useful computation or dropping into low-power sleep states. To create responsive, real-time Cyber-Physical Systems, we must abandon polling in favor of asynchronous, interrupt-driven event handling, where the hardware forcibly alerts the software that an event has occurred.

7.2 Booting: What Happens at Power-On

When an embedded processor is first powered up, or when a hardware reset button is pressed, the physical reset signal (often an active-low signal denoted as `nRST` or `RST#`) is asserted. This forces the processor's internal state machine into a known, rigid initial state.

But where does the processor find its first instruction? It cannot rely on the main memory (DRAM) because volatile memory contains garbage data at power-on. Instead, the processor is hardwired to look at a highly specific, predefined address in a non-volatile memory region, such as an on-chip Flash memory or a Boot ROM.

7.2.1 The Vector Table

Rather than placing the raw executable code directly at the boot address, modern architectures typically place a data structure known as the **Vector Table** at this location. The vector table is an array of addresses (pointers) that dictate where the processor should jump to handle specific hardware events.

For example, in the ARM Cortex-M architecture, the vector table is located at the very beginning of memory. The hardware expects the first two entries of this table to contain the most critical information needed to boot the system:

1. **Initial Stack Pointer (Offset 0x00):** The first 32-bit entry is not an instruction; it is the physical memory address of the top of the stack. When the CPU powers on, it immediately reads this value and loads it directly into the hardware Stack Pointer (SP) register. This is a brilliant architectural optimization, as it guarantees that a valid stack exists before the processor executes a single line of code, allowing the boot code to safely make function calls.
2. **Reset Vector (Offset 0x04):** The second entry contains the memory address of the **Reset Handler**. The hardware automatically loads this address into the Program Counter (PC) and begins fetching instructions.

Because the vector table is typically stored in read-only, non-volatile memory, its contents are established at compile-time. In some advanced implementations, the base address of the vector table can be dynamically relocated by software later in the execution cycle by writing to a specific control register, allowing the operating system to install its own set of handlers.

7.2.2 The C Runtime Initialization (CRT0)

The Reset Handler is a small, specialized block of assembly code—often called the startup code or `crt0` (C runtime zero)—responsible for preparing the environment so that C/C++ code can run safely.

Software engineers are used to writing C code containing global variables with initial values (e.g., `int x = 5;`). However, non-volatile Flash memory cannot be modified at runtime, and volatile RAM contains random garbage at power-on. To solve this, the linker organizes the compiled program into distinct segments:

- **.text:** Contains the executable machine instructions.
- **.data:** Contains initialized global and static variables.
- **.bss:** Contains uninitialized global and static variables (which the C standard dictates must default to zero).

During the boot process, the startup code must physically copy the initial values for the `.data` segment from the non-volatile ROM into the volatile RAM. Next, it must loop through the entire `.bss` segment in RAM and write zeros to every memory location. Finally, the startup code configures the start and limit addresses for the heap memory (used for dynamic allocations like `malloc`), and only then does it execute a branch instruction to transfer control to the user's `main()` function.

7.3 The Exception Model

In the normal course of execution, the processor fetches instructions sequentially, occasionally altering the Program Counter (PC) due to software branch instructions. However, the processor must possess mechanisms to handle unplanned, asynchronous events that demand immediate attention. This mechanism is known as the **Exception Model**.

An exception is conceptually similar to an unscheduled, forced function call. Exceptions can be broadly divided into two categories based on their origin: hardware and software.

7.3.1 Hardware Interrupts

A hardware exception, commonly referred to as an **interrupt**, is triggered by physical events external to the CPU execution pipeline. Examples include a user pressing a key on a keyboard, a timer reaching zero, or a network interface controller receiving a packet.

When a peripheral requires attention, it asserts a physical interrupt request (IRQ) line wired to the processor. Because these events happen independently of the software's execution flow, they are asynchronous. The processor must finish executing the current instruction, halt its normal sequence, and jump to an **Interrupt Service Routine (ISR)** to handle the device.

7.3.2 Software Exceptions (Traps and Faults)

Software exceptions are synchronous events generated directly by the execution of a specific instruction within the pipeline.

- **Faults:** Occur when the processor encounters an error condition it cannot resolve natively. Examples include attempting to execute an undefined or illegal instruction opcode, attempting to divide a number by zero, arithmetic overflow, or an MMU page fault (attempting to access virtual memory that is not currently mapped to physical RAM).
- **Traps:** Are intentional, scheduled exceptions triggered by specific assembly instructions (such as `SVC` or `SWI` in ARM, or `SYS / TRAP` in other architectures). Traps are the fundamental mechanism used by user-space applications to request services from the operating system, a concept we will explore further when discussing privilege levels.

7.4 Handling Exceptions: The Hardware Perspective

When an exception occurs, the processor hardware must perform several immediate, automated steps before any software handler can run.

7.4.1 Recording the Exception State

To handle the exception and eventually return to the interrupted program, the hardware must save the state of the machine. At a minimum, this involves saving the current Program Counter (so the CPU knows where to return) and a code indicating the reason for the exception.

Architectures handle this differently. In the MIPS architecture, for example, the processor utilizes a dedicated set of system management registers known as **Coprocessor 0**. When an exception fires, the hardware automatically copies the current PC into the **Exception Program Counter (EPC)** register and writes an identification code into the **Cause** register. The processor then forces the PC to jump to a hardcoded global exception handler address (e.g., `0x80000180`), where the OS must read the Cause register to determine what actually happened.

In contrast, the ARM architecture utilizes a vectored approach via the Nested Vectored Interrupt Controller (NVIC). Each specific type of exception or peripheral interrupt is assigned a unique index number. The hardware uses this index to directly look up the specific handler address in the Vector Table we discussed earlier, branching immediately to the correct software routine without requiring a massive, centralized `switch/case` statement to decode the cause.

7.4.2 Precise Interrupts and Pipeline Flushing

Handling exceptions becomes incredibly complex in modern microarchitectures featuring deep pipelines, superscalar issue, and out-of-order execution.

If a processor is executing five instructions simultaneously in a pipeline and an external interrupt fires, the processor cannot simply stop instantly. A fundamental requirement of modern systems is the **precise interrupt**. For an interrupt to be precise, the hardware must guarantee that:

1. All instructions appearing in the program code *before* the interrupted instruction have fully completed and retired their results to the

architectural state.

2. All instructions appearing *after* the interrupted instruction have completely discarded any partial computations, leaving no trace in the architectural state.

To achieve this in an out-of-order processor like the Intel Core i7, instructions write their intermediate results to hidden physical registers, but they only commit their results to the visible architectural registers (a process called retirement) in strict, sequential program order. If a fault or interrupt occurs, the Retirement Unit halts, preserving the precise state of the machine up to that exact instruction, and the pipeline is flushed of all subsequent speculative instructions. This ensures that when the ISR finishes, the program resumes exactly where it left off, oblivious to the fact that it was ever interrupted.

7.5 The Interrupt Service Routine (ISR)

Once the hardware has verified the precise state and branched to the correct address in the vector table, the software **Interrupt Service Routine (ISR)** takes over. Writing an ISR requires extreme discipline, as the routine is hijacking the CPU from an active thread.

As outlined in fundamental operating systems theory, a robust ISR must perform the following distinct software actions:

1. **Context Save:** The very first action the ISR must take is to save the contents of all general-purpose registers it intends to use. If the ISR modifies a register (like `x0` or `R1`) without saving it, it will corrupt the mathematical calculations of the interrupted background program. Registers are typically pushed onto the current stack. *(Note to students: While the specific ARM AArch64 assembly instructions for context switching are standard industry practice and draw on our knowledge from outside the provided foundational texts, you are encouraged to independently verify the exact syntax in the ARM Architecture Reference Manual. Typically, this involves using the `STP` (Store Pair) instruction to push registers to the stack).*
2. **Acknowledge and Clear:** The ISR must communicate with the hardware peripheral (e.g., the UART or Timer) that triggered the interrupt, reading

its status registers to identify the specific sub-event, and explicitly writing a command to clear the hardware interrupt flag. If the flag is not cleared, the CPU will immediately re-trigger the interrupt the moment the ISR returns, resulting in an infinite loop.

3. **Process the Data:** The ISR executes its specific logic—such as copying a received byte from a UART data register into an in-memory ring buffer, or decrementing a counter.
4. **Context Restore:** Once the work is complete, the ISR pops the saved values off the stack, restoring the general-purpose registers to their exact state prior to the interrupt.
5. **Return:** The ISR executes a special machine instruction (often called `RTI`, `IRET`, or a specific branch return in ARM) that atomically restores the processor's status flags from the hardware exception registers and copies the Exception Program Counter back into the active Program Counter, resuming the background task.

7.5.1 Interrupt Priorities and Nesting

In a complex system, multiple peripherals may attempt to interrupt the CPU simultaneously. Furthermore, a high-priority event (like an emergency motor shutoff) might occur while the CPU is already in the middle of executing a low-priority ISR (like a keyboard press).

If an ISR disables all other interrupts during its execution to prevent race conditions, the system becomes completely sequential, risking the loss of time-critical data if the first ISR takes too long. To resolve this, modern architectures assign strict mathematical priorities to different interrupts.

When an interrupt fires, the CPU elevates its current execution priority. If a new interrupt arrives with a *lower* priority, it is held in a pending state until the current ISR finishes. However, if an interrupt arrives with a *higher* priority, the CPU immediately preempts the active ISR, pushes its state to the stack, and services the more critical event—a concept known as **Nested Interrupts**. The ARM Nested Vectored Interrupt Controller (NVIC) handles these complex priority evaluations entirely in hardware, significantly reducing software latency.

7.6 Privilege Levels and Protection

Interrupts and exceptions are not only used for handling hardware events; they form the absolute foundation of system security and operating system architecture.

Modern processors implement multiple hardware privilege levels, or execution modes. The Intel Core i7, for instance, supports four privilege rings (Ring 0 to Ring 3), though typical operating systems like Linux and Windows only utilize two: Ring 0 (Kernel Mode) and Ring 3 (User Mode).

When the processor is in User Mode, the hardware Memory Management Unit (MMU) strictly restricts what the software can do. User Mode code is absolutely prohibited from executing privileged machine instructions (like disabling hardware interrupts) or accessing memory-mapped I/O peripheral registers. If a User Mode application attempts to directly modify a hardware timer, the MMU intervenes, throwing a synchronous fault exception that terminates the offending program.

But if a user program cannot access hardware, how does it ever print to the screen or save a file? It must ask the highly trusted Operating System to do the work on its behalf.

7.6.1 The System Call

Because User Mode code cannot simply execute a `CALL` instruction to branch directly into the Kernel's memory space (which is protected by the MMU), it must use a controlled hardware gateway.

To request a service, the user program loads a specific service number into a general-purpose register and executes a software trap instruction (e.g., `sys` or `svc`). This instruction intentionally causes a hardware exception. The processor instantly elevates its privilege level to Kernel Mode and forces the PC to jump to a highly secured, predetermined vector address containing the OS System Call Handler.

The OS examines the requested service number, securely validates any passed pointers, performs the privileged hardware I/O operation, and then executes a return-from-exception instruction to demote the processor back to User Mode and resume the application. This strict separation ensures that buggy or malicious user software cannot crash or compromise the underlying Cyber-Physical System.

7.7 Context Switching and Parallel Processing

By combining hardware timer interrupts with the privilege and exception model, operating systems can create the illusion of true parallel processing on a single CPU core.

If multiple programs (processes or threads) reside in memory, an RTOS (Real-Time Operating System) configures a hardware timer peripheral to generate a high-priority interrupt at a fixed interval (e.g., every 1 millisecond).

1. The CPU is running Task A.
2. The timer fires, preempting Task A and forcing a jump to the OS Scheduler ISR.
3. The OS saves the exact register state of Task A onto Task A's private stack.
4. The OS changes the CPU's hardware Stack Pointer to point to Task B's private stack.
5. The OS pops Task B's previously saved registers off Task B's stack.
6. The OS executes the return-from-exception instruction, which restores Task B's Program Counter.
7. The CPU resumes executing Task B.

This incredibly rapid process is known as a **Context Switch**. Because it relies entirely on the precise interrupt mechanism preserving and restoring architectural state, neither Task A nor Task B is aware that it was paused and swapped out of the execution pipeline.

7.8 VirtMCU Homework: Building a Deterministic Interrupt Handler

In a standard embedded systems curriculum, writing bare-metal interrupt handlers introduces massive frustration. If a student configures a timer interrupt in a standard software emulator (like a vanilla QEMU instance), the emulator relies on the host operating system's wall-clock to trigger the event. Because the host OS (Linux/Windows) introduces unpredictable scheduling delays, context switches, and network latencies, the simulated timer interrupt suffers from immense **wall-clock jitter**,.

This jitter makes it impossible to reliably debug precise hardware timing dependencies, a fatal flaw when designing Cyber-Physical Systems where missing a control deadline by microseconds can crash a drone.

To solve this, your homework relies on the **VirtMCU FirmwareStudio**. VirtMCU completely abandons the host OS wall-clock. Instead, it utilizes a dynamic QOM plugin architecture to enforce the **Big QEMU Lock (BQL)** constraints and implements strict cooperative time slaving.

1. **Architecture Review:** You will begin by reviewing the VirtMCU design rationale in `docs/architecture/01-system-overview.md` to understand how the five implementation pillars guarantee cycle-accurate deterministic execution regardless of host computer speed.
2. **Writing the Handler:** You will be provided with a bare-metal C framework containing an uninitialized vector table. You must write an assembly language Interrupt Service Routine (ISR) that adheres to the strict context save/restore requirements discussed in Section 7.5.
3. **Peripheral Interfacing:** Using the `hw/rust/common/rust-dummy/` module as a template, you will instantiate a simulated push-button peripheral that triggers an IRQ line mapped to the NVIC.
4. **Execution:** You will compile your binary and run it within VirtMCU in `slaved-icount` mode. Because VirtMCU strictly lock-steps the virtual clock to the exact number of retired instructions (bypassing host OS jitter), you will use the `gdb` debugger to step through your ISR cycle-by-cycle. You will empirically verify that your interrupt fires at the exact same

deterministic virtual nanosecond across multiple simulation runs, proving the stability necessary for rigorous CPS 2.0 design.

Chapter 8: Real-Time Operating Systems (RTOS)

8.1 Introduction: Beyond the Super-Loop

In Chapter 7, we breached the hardware-software interface, examining how a bare-metal processor boots and how it utilizes the exception model to asynchronously respond to physical events via hardware interrupts. For highly simplistic embedded devices—such as a microwave oven controller—software engineers often rely on a “super-loop” architecture. In a super-loop, the `main()` function consists of an infinite `while(1)` loop that continuously polls sensors, computes logic, and writes to actuators, occasionally pausing to service a hardware interrupt.

However, as we transition into modern Cyber-Physical Systems (CPS)—such as autonomous drone flight controllers, automotive anti-lock braking systems, and interconnected smart manufacturing robotics—the super-loop rapidly collapses under the weight of its own complexity. When a system must simultaneously sample fifty different sensors at varying frequencies, manage an encrypted TCP/IP network stack, and execute complex motor control algorithms, a single execution thread riddled with interrupt service routines (ISRs) becomes impossible to maintain, debug, or verify.

To manage this immense software complexity, we must introduce an Operating System (OS). For software engineers accustomed to writing applications for Windows or Linux, an OS is viewed as a massive resource manager that abstracts the hardware and provides virtual memory, file systems, and fair multi-tasking. However, general-purpose operating systems (GPOS) are fundamentally designed to maximize **average throughput** and maintain fairness among users. If a desktop PC delays the rendering of a web browser by 200 milliseconds to flush a disk cache, the human user barely notices.

In a Cyber-Physical System, time is a logical correctness criterion. If an automotive airbag controller computes the mathematically perfect deployment force but delivers that signal 50 milliseconds too late, the result is catastrophic. Therefore, CPS architectures rely on a **Real-Time Operating System (RTOS)**. An RTOS strips away the heavy abstractions of a GPOS and focuses obsessively on **determinism**, guaranteeing that critical tasks will execute within strict, mathematically provable time constraints.

8.2 The Real-Time Paradigm: Soft, Firm, and Hard

Before examining the internal architecture of an RTOS, we must establish the vocabulary used to classify real-time constraints. A real-time system is one in which the correctness of a computation depends not only upon the logical result but also upon the time at which the result is produced. Real-time systems are categorized by the severity of the consequences when a deadline is missed:

- **Soft Real-Time:** Missing a deadline degrades the performance of the system, but the late result still retains some value, and the system can recover. An example is a user interface update on a smart thermostat or a streaming video decoder; a dropped frame is annoying but not fatal.
- **Firm Real-Time:** Missing a deadline means the computed result is now entirely useless and must be discarded. However, missing the deadline does not cause catastrophic system failure. An example is a high-frequency financial trading algorithm; if the execution engine misses the optimal microsecond window to execute a trade, the opportunity is gone, but the server itself remains operational.
- **Hard Real-Time:** Missing a deadline constitutes an absolute, catastrophic system failure, often resulting in physical destruction or loss of life. In these systems, a late answer is identically equivalent to a wrong answer. Examples include aerospace fly-by-wire controls, pacemaker timing circuits, and industrial robotic arm actuation.

An RTOS is specifically engineered to support hard real-time requirements. It achieves this not necessarily by being “fast” (a 3 GHz desktop processor running Linux will execute raw math much faster than a 100 MHz RTOS-driven

microcontroller), but by being **predictable**. An RTOS minimizes scheduling jitter, bounded interrupt latency, and context-switch overhead to ensure that worst-case execution times (WCET) can be strictly calculated and guaranteed.

8.3 Anatomy of an RTOS: Tasks, Threads, and States

In a modern GPOS, the fundamental unit of execution is the process, which possesses its own isolated virtual memory address space mapped by the MMU. To switch between processes, the OS must flush the Translation Lookaside Buffer (TLB) and swap page tables, an incredibly expensive operation.

In a lightweight RTOS, memory protection is often relaxed to achieve deterministic speed. The RTOS typically operates in a single, unified physical address space. Therefore, the fundamental unit of execution is the **thread** or **task**. Multiple threads run in the same address space, meaning they can trivially share global variables, but they maintain their own independent execution flow.

8.3.1 The Task Control Block (TCB)

To sustain the illusion that multiple tasks are executing simultaneously on a single CPU core, the RTOS kernel must track the exact machine state of every task. It does this using a fundamental data structure known as the **Task Control Block (TCB)**. The TCB contains the task's unique ID, its scheduling priority, a pointer to its private stack memory, and its current execution state.

At any given moment in time, an RTOS task exists in one of several states:

1. **Running:** The task is currently executing instructions in the CPU pipeline. On a single-core processor, only one task can be in the Running state at any given nanosecond.
2. **Ready:** The task is fully prepared to execute and is simply waiting for the scheduler to grant it access to the CPU.

3. **Blocked (Waiting):** The task cannot execute because it is waiting for a specific event to occur. This could be waiting for a hardware timer to expire, waiting for a packet to arrive on a network interface, or waiting for a shared resource to become available.
4. **Suspended:** The task has been explicitly halted by the programmer or the system and will not be scheduled until it is explicitly resumed.

8.3.2 The Context Switch

When the RTOS decides to stop running Task A and start running Task B, it performs a **context switch**. This is an intricately choreographed sequence of assembly language instructions that relies heavily on the processor's exception model.

First, a hardware timer—often called the SysTick timer in ARM architectures—fires a periodic interrupt (e.g., every 1 millisecond). This interrupt preempts the Running task (Task A) and forces the CPU to jump to the RTOS Scheduler ISR. The scheduler must preserve the exact architectural state of Task A so it can be resumed later. It does this by pushing Task A's general-purpose registers (X0-X30 in AArch64), its Program Counter (PC), and its Process State Register (PSTATE) onto Task A's private stack.

The scheduler then updates Task A's TCB with the new stack pointer address. Next, the scheduler selects Task B from the Ready queue, reads Task B's TCB to find its stack pointer, and points the hardware Stack Pointer (SP) register to Task B's stack. Finally, the scheduler pops Task B's saved registers off its stack and executes a return-from-exception instruction. The CPU pipeline seamlessly resumes executing Task B, entirely unaware that a switch occurred.

Because context switching consumes CPU cycles without performing any application-level work, RTOS developers optimize this assembly code ruthlessly, ensuring the context switch overhead is highly deterministic and consumes only a few dozen clock cycles.

8.4 RTOS Scheduling Algorithms

The brain of the RTOS is the scheduler, the algorithm that decides which Ready task should transition to the Running state.

8.4.1 Round-Robin and Time-Slicing

In a pure **Round-Robin** scheduling algorithm, every task is assigned an equal, fixed time quantum (e.g., 5 milliseconds). The scheduler cycles through a circular queue of Ready tasks, giving each task the CPU for its allotted quantum before preempting it and moving to the next. While this guarantees fairness and prevents any single task from starving, it is disastrous for hard real-time systems. If a critical motor-control task becomes Ready, it cannot execute until all other tasks in the queue have exhausted their time slices, destroying determinism.

8.4.2 Priority-Based Preemptive Scheduling

Virtually all modern RTOS kernels utilize **Priority-Based Preemptive Scheduling**. The software engineer explicitly assigns a strict mathematical priority to every task. The scheduling rule is absolute: *The CPU will always execute the highest-priority task that is in the Ready state.*

If a low-priority task is currently Running, and a hardware interrupt wakes up a high-priority task (moving it from Blocked to Ready), the scheduler will immediately preempt the low-priority task and context-switch to the high-priority task. The low-priority task will only resume once the high-priority task voluntarily yields the CPU or enters the Blocked state.

8.4.3 Rate Monotonic Scheduling (RMS)

Assigning priorities manually can lead to missed deadlines if the software engineer guesses incorrectly. To mathematically guarantee that a set of

periodic tasks will always meet their hard real-time deadlines, systems architects use analytic algorithms like **Rate Monotonic Scheduling (RMS)**.

RMS assigns static priorities to tasks based on their frequency of execution: the shorter the task's period (the more frequently it must run), the higher its priority. To determine if a set of n independent, preemptable tasks is schedulable under RMS without missing any deadlines, we calculate the total CPU utilization U . Let C_k be the maximum worst-case execution time of task k , and T_k be its execution period. The system is provably schedulable if the following condition is met:

$$U = \sum_{k=1}^n \frac{C_k}{T_k} \leq n(2^{1/n} - 1)$$

Consider an example of three critical CPS threads:

- Thread 1: Execution Time $C_1 = 50$, Period $T_1 = 100$
- Thread 2: Execution Time $C_2 = 100$, Period $T_2 = 500$
- Thread 3: Execution Time $C_3 = 120$, Period $T_3 = 1000$

We evaluate the left side of the RMS formula: $U = \frac{50}{100} + \frac{100}{500} + \frac{120}{1000} = 0.50 + 0.20 + 0.12 = 0.82$ Next, we evaluate the theoretical upper bound for $n=3$ threads on the right side of the equation: $3(2^{1/3} - 1) \approx 3(1.2599 - 1) = 0.7798$ Because $0.82 \not\leq 0.7798$, this specific set of threads is mathematically *not* guaranteed to be schedulable under pure RMS, meaning under worst-case phasing, a deadline will be missed. By performing this analysis at compile-time, engineers can redesign the system before deploying it to physical hardware.

8.5 Inter-Task Communication and Synchronization

Because threads in an RTOS share the same memory space, they can easily pass data back and forth by reading and writing global variables. However, uncontrolled interaction between concurrent processes leads to catastrophic **race conditions**.

Suppose Task A and Task B both increment a global counter tracking physical motor rotations. The C code `counter++` translates to three machine instructions: Load the variable from RAM to a register, increment the register, and Store the register back to RAM. If Task A is preempted by Task B precisely between the Load and Store instructions, both tasks will read the exact same original value, increment it, and write back identical values. One physical rotation is permanently lost in the software state.

The segment of code where a task accesses a shared variable is known as a **critical section** or critical region. To prevent race conditions, the RTOS must provide synchronization primitives that enforce **mutual exclusion**, ensuring that if one task is inside the critical section, no other task can enter it.

8.5.1 Semaphores and Mutexes

The most fundamental synchronization primitive is the **semaphore**, invented by E.W. Dijkstra in 1965. A semaphore is an integer variable managed securely by the RTOS kernel. It supports two atomic operations: `down` (historically *P*) and `up` (historically *V*).

When a task wishes to enter a critical section, it executes a `down` operation on the semaphore. If the semaphore's value is greater than zero, the RTOS decrements the value and allows the task to proceed. If the value is 0, the RTOS immediately removes the task from the Running state, places it in the Blocked queue, and schedules a different task. When the task finishes its critical section, it executes an `up` operation, which increments the semaphore. If one or more tasks are sleeping in the Blocked queue waiting for this semaphore, the RTOS wakes up one of them, moving it back to the Ready queue.

A **Mutex** (Mutual Exclusion object) is a specialized binary semaphore initialized to 1. It operates as an electronic "key." A task must acquire the mutex before accessing a shared memory block and must release it when finished. Mutexes often include concept of "ownership," meaning the exact task that acquired the mutex is the only one legally permitted to release it, preventing rogue tasks from unlocking resources they do not own.

Semaphores are also heavily used for the **producer-consumer problem**, where one task generates data (e.g., an ADC sampling voltages) and places it into a shared ring buffer, while another task consumes the data. Semaphores inherently track the number of full slots and empty slots in the buffer, elegantly putting the producer to sleep if the buffer fills up, and waking it the instant the consumer removes an item.

8.6 The Hazards of Concurrency: Deadlocks and Priority Inversion

While semaphores solve the race condition problem, they introduce entirely new classes of systemic failure.

8.6.1 Deadlock

A **deadlock** occurs when a set of tasks are permanently blocked, each waiting for a resource that is held by another task in the set. This is famously illustrated by the Dining Philosophers problem, where multiple philosophers sit at a circular table requiring two forks to eat, but there is only one fork between each plate. If every philosopher picks up their left fork simultaneously, they will all block forever waiting for the right fork, creating a circular wait dependency.

In CPS architecture, deadlocks are identified by the presence of the four Coffman conditions: Mutual exclusion, hold and wait, no preemption, and circular wait. RTOS designers must break these conditions structurally—for instance, by mathematically mandating a strict global hierarchy for acquiring multiple mutexes, preventing a circular dependency chain from ever forming.

8.6.2 Priority Inversion

Priority Inversion is perhaps the most notorious software hazard in Cyber-Physical Systems, famously causing the near-loss of NASA's Mars Pathfinder rover in 1997.

Priority inversion occurs under a specific combination of priority-based preemptive scheduling and mutexes. Consider three tasks: High (H), Medium (M), and Low (L).

1. Task L begins executing and acquires a Mutex protecting a shared communications bus.
2. Task H wakes up, preempting Task L. Task H attempts to acquire the same Mutex. Because Task L holds it, Task H is put into the Blocked state.
3. This is normal so far; however, Task M suddenly wakes up. Because Task M has a higher priority than Task L, and does *not* need the Mutex, the scheduler grants the CPU to Task M.

Task M begins executing a long computation. Task L cannot run because it is preempted by Task M. Task H cannot run because it is blocked waiting for Task L to release the Mutex. Thus, the High-priority task is effectively suspended indefinitely while the Medium-priority task dominates the CPU. The priorities have been inverted!

To solve this, modern RTOS kernels implement the **Priority Inheritance Protocol (PIP)**. When Task H blocks on a Mutex held by Task L, the RTOS temporarily elevates the priority of Task L to match the priority of Task H. Because Task L now has “High” priority, it cannot be preempted by Task M. Task L finishes its critical section, releases the Mutex, and its priority instantly drops back to “Low.” Task H acquires the Mutex and resumes execution, preventing the inversion.

8.7 Industry Standards: CMSIS-RTOS

In the embedded ecosystem, learning a new proprietary RTOS API for every microchip vendor creates severe friction. To standardize software portability across the vast array of ARM Cortex-M microcontrollers, ARM developed the Cortex Microcontroller Software Interface Standard (CMSIS).

CMSIS provides a hardware abstraction layer that defines standardized functions for core CPU access (CMSIS-Core), DSP math functions (CMSIS-DSP), and importantly, a standardized RTOS API called **CMSIS-RTOS**.

CMSIS-RTOS is not an RTOS itself, but rather a wrapper API that sits on top of commercial or open-source kernels like FreeRTOS or Keil RTX. It defines a universal set of C macros and functions to manage RTOS task parameters, inter-process communication, synchronization, and scheduling configurations. By adhering strictly to the CMSIS-RTOS API, a software engineer can write a complex multithreaded drone controller on a small Cortex-M3 processor, and effortlessly recompile and port that exact same codebase to a massive multi-core Cortex-M7 SoC without rewriting the thread management or semaphore logic.

8.8 VirtMCU Homework: Deterministic Context-Switch Analysis

Understanding the theoretical concepts of priority inversion and RMS scheduling is trivial; identifying and debugging them in an active execution pipeline is enormously difficult.

If you attempt to trace a context switch or a priority inversion bug using a standard QEMU emulator running on Linux or Windows, you will fail. Standard emulators execute instructions using a free-running host clock. The host OS scheduler interrupts the emulator unpredictably to service its own background tasks, injecting massive amounts of wall-clock scheduling jitter into the emulation. Measuring the exact microsecond penalty of an RTOS context-switch becomes impossible because the measurement is polluted by the host machine's background noise.

To overcome this, you will utilize the **VirtMCU FirmwareStudio**. Because VirtMCU uses a dynamic QOM plugin architecture to lock-step the virtual clock precisely to the number of retired instructions (`slaved-icount` mode), it creates a perfectly deterministic simulation environment devoid of host OS jitter.

1. **RTOS Porting:** You will be provided with a bare-metal C implementation of a lightweight RTOS kernel. Your first task is to examine the context

switch assembly routine, validating how it pushes X0-X30, the FP, and the LR onto the task stack.

2. **Jitter Measurement:** You will configure two threads to toggle a GPIO pin. You will run this in standard QEMU and observe the massive timing variance using a virtual logic analyzer. Then, you will run the identical binary in VirtMCU `slaved-icount` mode, using the deterministic virtual hardware timers to measure the *exact* number of clock cycles consumed by the RTOS scheduler during the context switch.
3. **Provoking Priority Inversion:** You will write an application containing three tasks (High, Medium, Low) and a single shared Mutex. You will intentionally orchestrate a priority inversion scenario. By pausing the VirtMCU simulation at exact cycle counts, you will use GDB to dump the RTOS Ready queues and empirically verify that the High-priority task is incorrectly starving.
4. **Protocol Verification:** Finally, you will enable the Priority Inheritance Protocol feature in your RTOS configuration. Rerunning the deterministic simulation, you will observe the kernel actively promoting the Low-priority task's TCB priority field, proving how software mechanisms resolve the hazards of asynchronous CPS concurrency.

Chapter 9: Introduction to Cyber-Physical Systems (CPS 2.0)

9.1 Introduction: Beyond the Embedded System

Welcome to Part III of our course. Up to this point, we have focused extensively on the internal architecture of the computing node. We progressed from the physics of transistors to the pipeline of the ARM AArch64 processor, navigated the complexities of the memory hierarchy, and finally, in Chapter 8, deployed a Real-Time Operating System (RTOS) to guarantee deterministic thread scheduling.

If our curriculum ended here, you would be well-equipped to design a traditional embedded system. However, the modern engineering landscape demands more. We are no longer simply building isolated microcontrollers that blink LEDs or manage singular engines. We are engineering **Cyber-Physical Systems (CPS)**.

Cyber-Physical System theory emerged as a direct response to the increasingly complex interactions between computer algorithms and physical infrastructure. The goal of a CPS is to seamlessly bundle sensory data and advanced mathematical models to improve the control, automation, and decision-making of physical assets. In the late 2000s, as these systems took the form of next-generation embedded networks, the pioneer Wayne Wolf coined the term “control-computing codesign” to describe this paradigm. Wolf argued that distributed computing and control elements needed a collective, integrated form of examination rather than being isolated into separate hardware and software layers.

In a cyber-physical system, computation is inherently tied to the continuous-time physics of the real world. A desktop computer computes data to produce more data (e.g., calculating a spreadsheet). A CPS computes data to produce *physical action* (e.g., actuating an aircraft’s rudder, deploying a medical

pacemaker, or braking an autonomous vehicle). Consequently, a CPS is evaluated not just on its logical correctness, but on its physical timeliness and safety.

This chapter introduces the evolution of CPS from its foundational state to what is now known as **CPS 2.0**. We will explore the interplay of communication and computational technologies that underpin the fabric of CPS 2.0. We will dissect the anatomy of these systems, examining the feedback loops that bind the cyber and physical domains. Finally, we will outline the “3Cs” that define all modern CPS architectures: Computation, Communication, and Control.

9.2 The Evolution to CPS 2.0

To understand where we are, we must understand how we got here. The evolution of cyber-physical systems can be broadly categorized into distinct generational phases.

9.2.1 CPS 1.0: Isolated Automation and the Industrial IoT

First-generation industrial CPS—often retrospectively labeled CPS 1.0—focused heavily on the basic automation and control of manufacturing machines and isolated lines. These systems were typically characterized by localized Programmable Logic Controllers (PLCs) governing physical processes via closed-loop feedback.

As networking technologies advanced, these systems transitioned into the Industrial Internet of Things (IIoT). The primary objective during this phase was data telemetry: capturing vast amounts of industrial sensor data and transmitting it to remote cloud services to be visualized, logged, and analyzed by human operators. While this provided unprecedented visibility into physical operations, the systems remained largely reactive. The cloud was used for storage and offline analysis, not for real-time, closed-loop physical control, largely due to the prohibitive latency and jitter of internet-based communication.

9.2.2 CPS 2.0: Intelligence, Convergence, and Metasystems

As related technologies matured—specifically in edge computing, deterministic networking, and artificial intelligence—new opportunities opened, leading to the second generation: **CPS 2.0**.

CPS 2.0 requires entirely new software platforms and architectures to be deployed. The defining characteristic of CPS 2.0 is the profound integration of three traditionally isolated domains:

1. **Operational Technology (OT):** The hardware and software that detects or causes a change through the direct monitoring and/or control of physical devices, processes, and events.
2. **Information Technology (IT):** The networking, storage, and processing infrastructure used to manage enterprise data.
3. **Artificial Intelligence (AI):** The application of machine learning, neural networks, and high-frequency data models to enable predictive and autonomous decision-making.

In CPS 2.0, the system transcends traditional boundaries, giving rise to “metasystems” that permanently blur the lines between the physical and virtual worlds. Rather than merely sending data to the cloud for offline analysis, a CPS 2.0 metasystem utilizes high-frequency (HF) data streams to generate AI models on the fly, pushing these models down to the edge (directly onto the local microcontrollers) to influence real-time physical control.

These advanced architectures span multiple domains. We see CPS 2.0 manifesting in smart health-care systems deployed across smart cities, in cloud-enabled smart manufacturing processes (often referred to as Industry 4.0), and in distributed infrastructure like Virtual Power Plants (VPPs). In autonomous vehicles (AVs), CPS 2.0 allows the vehicle to integrate tangible, physical kinematics with advanced cyber decision-making, navigating complex environments autonomously.

9.3 Anatomy of a Cyber-Physical System

At the core of every cyber-physical system, regardless of its generational label, is the **cyber-physical feedback loop**. This loop is the fundamental mechanism by which a system observes the environment, reasons about it, and takes action to alter it.

9.3.1 Sensing the Physical World

The loop begins in the physical domain. The physical world is governed by continuous-time dynamics—classical mechanics, Newtonian motion, thermodynamics, and electromagnetics. To bring these continuous physical phenomena into the discrete cyber domain, the system relies on sensors (transducers) that convert physical energy into electrical signals.

In a modern CPS, sensing goes far beyond reading a simple analog voltage. Consider the field of Structural Health Monitoring (SHM) for civil infrastructure, such as bridges. In an SHM-oriented CPS, a network of accelerometers continuously monitors the vibration response of the bridge. This raw sensory data is processed using signal feature extraction techniques, such as Fast Fourier Transforms (FFT) or Discrete Wavelet Transforms (DWT). By analyzing high-frequency data, the cyber system can identify decaying modal frequencies—for instance, observing a frequency drop from 2.93 Hz to 1.56 Hz—which physically translates to a reduction in the bridge's structural stiffness due to earthquake damage or aging.

9.3.2 Model-Driven Processing and Decision Making

Once the physical data is ingested, the computing core must process it. In CPS 2.0, this relies heavily on **model-driven** processing. The raw data is fed into a mathematical representation (a “digital twin”) of the physical asset.

Returning to our bridge example, the cyber system maintains a baseline finite element model of the bridge. The sensor data is used to continually calibrate this model, updating uncertain parameters such as mass-stiffness features and

boundary conditions to reflect the real-world ground truth. By linking the physical asset to the digital model via this calibration protocol, the CPS can perform structural reliability assessments and risk analyses. If the model predicts an impending failure under a simulated seismic load, the system can autonomously trigger mitigation decisions—such as dynamically optimizing transportation networks to route heavy traffic away from the compromised structure,.

This demonstrates the power of a CPS: it does not just report data; it extrapolates existing sensory knowledge through physics-informed strategies to make intelligent decisions.

9.3.3 Actuation and Physical Control

The final stage of the feedback loop is actuation. Based on the decisions made by the cyber algorithms, the system must interact with the physical world to change its state. This is achieved through actuators—such as electrostatic or electromagnetic transducers, variable reluctance motors, and hydraulic valves.

Actuation is fundamentally tied to control laws. As the cyber system commands an actuator, it alters the physical domain, which in turn alters the readings of the sensors, closing the feedback loop. System functionality and capability are accomplished by means of rigorous optimization and control engineering to ensure physical controllability.

9.4 The 3Cs of CPS: Computation, Communication, and Control

To formalize the architecture of a CPS 2.0 metasystem, engineers analyze the design through the lens of the “3Cs”: Computation, Communication, and Control. Let us examine how each pillar underpins the fabric of modern cyber-physical infrastructure.

9.4.1 Computation

Computation in a CPS is the processing engine that executes the control laws and AI models. However, computation in the physical world is vastly different from computation in standard IT.

In a CPS, the state of the system is often defined by non-linear differential equations representing the physical kinematics (e.g., the rigid body dynamics and equations of motion for an aircraft). The processor must solve these equations and execute tracking control laws in real-time. For example, controlling the multi-axis coupling, resonance, and flexible modes of a robotic arm requires matrix mathematics and Proportional-Integral-Derivative (PID) control algorithms to be executed with strict microsecond deadlines,.

In CPS 2.0, computation is heavily decentralized. The architecture often employs **Edge Computing**, where the heavy computational lifting (such as high-frequency signal feature extraction and dimensionality reduction) is performed directly on the local microcontroller rather than being offloaded to the cloud. This guarantees that the control loop maintains its real-time determinism. The cloud is reserved for global orchestration, training heavy AI models, and performing long-term reliability assessments,.

9.4.2 Communication

Communication is the lifeblood of any distributed Cyber-Physical System. In a massive system like a Virtual Power Plant (VPP)—which coordinates hundreds of distributed energy resources, servers, and loads—the architecture relies entirely on the timely exchange of telemetry and command data.

Traditional TCP/IP networks are optimized for average throughput, making them highly susceptible to packet loss and latency spikes (jitter). In a hard real-time CPS, a delayed network packet containing a braking command is functionally equivalent to a lost packet, potentially resulting in catastrophic physical failure.

To resolve this, CPS 2.0 embraces **Software-Defined Networking (SDN)**. SDN enables a new era of flexibility by separating the network's control plane from

the data plane. It provides dynamic network management, enhancing system resilience while addressing the imperative requirement to maintain high-performance, deterministic communication in real-time operations. By utilizing coordinated communication control strategies, engineers can systematically analyze and mitigate the adverse effects of packet loss and latency, optimizing performance under real-world physical constraints.

9.4.3 Control

Control is the application of cybernetics principles to govern the behavior of the physical asset. The objective of the control system is to force the physical plant to follow a desired reference trajectory, despite external physical disturbances or internal system noise.

In a closed-loop system, the control law minimizes the tracking error—the difference between the desired physical state and the actual measured state. For instance, in the torque control of an electromagnetic induction motor, engineers design optimal tracking control laws that rely on system state matrices and feedback gains. By minimizing quadratic performance functionals and tuning the proportional, integral, and derivative (k_p , k_i , k_d) gains of a PID controller, the cyber system guarantees that the motor's settling time remains within strict physical bounds.

In the CPS 2.0 paradigm, control systems are evolving into **self-adapting systems**. Using holonic reference architectures, the control layers are becoming decentralized and autonomous, capable of dynamically adjusting their own control laws in response to changing physical environments or hardware degradation.

9.5 Security and Resilience in CPS 2.0

As we transition to CPS 2.0 and deeply integrate Operational Technology (OT) with internet-facing Information Technology (IT), the attack surface of our physical infrastructure expands exponentially.

In isolated embedded systems (CPS 1.0), security was largely physical; you protected the PLC by putting a lock on the factory door. In CPS 2.0, systems are connected to cloud environments, exposing critical infrastructure—such as power grids, medical cyber-physical systems, and transportation networks—to sophisticated cyber threats,. A cyber-attack on a traditional IT database results in a loss of privacy or data. A cyber-physical attack on a CPS results in a loss of physical control, which can lead to equipment destruction or human casualties.

Therefore, CPS 2.0 architecture mandates “security by design.” This requires the implementation of Cyber-Physical Threat Intelligence, anomaly detection models, and robust access controls. It also requires the system to possess physical resilience—the ability of the control system to detect a cyber-induced hardware fault (such as an attacker feeding spoofed sensor data to the PID controller) and safely degrade operations to prevent physical catastrophe,. We will explore these advanced security topics deeply in Chapter 13.

9.6 VirtMCU Homework: CPS 2.0 Architecture Design

In a standard university software environment, designing and testing a true Cyber-Physical System is exceptionally difficult. If you attempt to simulate a drone flight controller relying on a host operating system’s networking stack, the wall-clock scheduling jitter will destroy your PID control loops. Your drone will crash in simulation not because your control mathematics were wrong, but because the host OS delayed a sensor packet by 50 milliseconds to process a background task.

To bridge this gap, this chapter’s homework relies on the **VirtMCU FirmwareStudio**. Because VirtMCU enforces deterministic virtual time and physics engine lock-stepping via the `hw/rust/backbone/clock` module, it completely eliminates wall-clock jitter. Time inside the VirtMCU environment moves strictly in accordance with retired instructions and precise hardware timers.

Your Assignment: You are tasked with designing the high-level CPS 2.0 architecture for a multi-rotor autonomous drone.

1. **Map the 3Cs:** You must provide a formal block diagram mapping the Computation, Communication, and Control layers.
2. **Sensing and Actuation (Control):** Identify the physical continuous-time variables you must measure (e.g., gyroscopic pitch, roll, yaw) and the actuators you will drive (e.g., PWM motor controllers). Map these physical properties to the VirtMCU Sensor/Actuator Abstraction Layer (SAL/AAL).
3. **Network Design (Communication):** The drone will utilize a distributed architecture where each motor has its own microcontroller, communicating with a central flight computer. You will design the message format required to transmit control data over the VirtMCU deterministic bus, ensuring you account for the timing strictness required by the physics engine.
4. **Edge AI (Computation):** Describe how a predictive AI model (simulated within the VirtMCU environment) could be deployed at the edge to dynamically adjust the PID feedback gains in response to a simulated physical disturbance, such as a sudden loss of lift on one rotor.

By mapping this architecture into the deterministic confines of VirtMCU, you will transition your understanding from isolated software algorithms to holistic, continuous-time control-computing codesign.

Chapter 10: Co-Simulation and Virtual Time

10.1 Introduction: The Illusion of Time in Computing

In the first nine chapters of this book, we constructed a comprehensive understanding of computing nodes from the digital logic level up to the Real-Time Operating System (RTOS) and the overarching architecture of Cyber-Physical Systems (CPS 2.0). We established that in a true CPS, time is not merely a performance metric—it is a strict criterion for logical correctness. If an autonomous vehicle's braking controller computes the mathematically perfect stopping force but delivers it 50 milliseconds too late, the result is catastrophic physical failure.

However, a fundamental impedance mismatch exists between the cyber and physical domains regarding the concept of time. In the physical world, time is a continuous, monotonically advancing variable governed by differential equations and the laws of physics. Physical components—such as masses, springs, induction motors, and fluids—react instantly and continuously to applied forces.

Conversely, the cyber domain is entirely discrete. Software engineers generally view time as a sequence of logically ordered events. A program executes line one, then line two, then line three. The exact wall-clock time it takes to transition between these lines is abstracted away by the hardware pipeline, the memory hierarchy, and the operating system scheduler. When software is tested in standard development environments, it runs as fast as the host processor allows, completely decoupled from the real-world timeline.

When designing a CPS 2.0 architecture, we cannot test the cyber algorithms and the physical mechanics in isolation. We must evaluate them together to observe their interactions. This requirement introduces the discipline of **co-**

simulation, which fuses a discrete-event software simulator with a continuous-time physics engine. To achieve this without introducing massive timing errors, we must strip the concept of “time” away from the host computer’s operating system and replace it with a mathematically rigorous, deterministic **virtual time** clock. This chapter explores the theory of discrete-event simulation, the implementation of Transaction Level Modeling (TLM), and the mechanisms VirtMCU uses to achieve deterministic co-simulation for CPS.

10.2 The Problem with Standard Emulation

To test embedded firmware before physical hardware is manufactured, engineers have traditionally relied on software emulators, the most famous being QEMU (Quick Emulator). QEMU utilizes Dynamic Binary Translation (DBT) to convert blocks of guest instruction set architecture (e.g., ARM AArch64) into the host instruction set architecture (e.g., x86_64) on the fly. This provides incredibly fast execution, allowing developers to boot entire guest operating systems in seconds.

However, standard emulation environments are built for functional software testing, not for Cyber-Physical System validation. They suffer from a critical flaw: **wall-clock scheduling jitter**.

In a standard QEMU setup, the simulated processor’s timer peripherals and interrupt controllers are tied to the host operating system’s real-time clock. When the guest firmware configures a hardware timer to fire an interrupt every 10 milliseconds, QEMU asks the Linux or Windows host OS to set a software timer. The host OS is a general-purpose, multitasking system. It is simultaneously managing disk I/O, network traffic, background services, and graphical user interfaces. If the host OS scheduler decides to preempt the QEMU process to handle an incoming gigabit network packet, the emulator is physically paused.

When QEMU is eventually rescheduled and resumes execution, it checks the host clock, realizes that 15 milliseconds have passed, and immediately fires the guest’s 10ms timer interrupt. From the perspective of the guest firmware, the hardware timer took 15 milliseconds to count to 10 milliseconds. This timing

violation propagates through the guest's RTOS scheduler, delaying the execution of critical control loops.

If this emulator is loosely coupled to a physics engine simulating a drone, the physics engine will continue to advance its continuous-time dynamics while the emulator is paused. When the emulator finally sends its delayed actuator command, the drone's physical state has already degraded, causing the simulated drone to crash. The drone crashed not because the student's Proportional-Integral-Derivative (PID) control math was incorrect, but because the host OS injected scheduling jitter into the simulation. This lack of reproducibility renders standard emulators useless for hard real-time CPS verification.

10.3 Principles of Discrete-Event (DE) Simulation

To build a simulation environment devoid of wall-clock jitter, we must formalize how digital hardware simulates time. At the system level, hardware is typically modeled using **Discrete-Event (DE) Simulation** semantics. SystemC, the industry standard C++ library for hardware modeling, utilizes a specialized DE simulation kernel.

Unlike continuous-time physics, where variables change constantly, digital hardware state changes only at discrete points in time, typically triggered by clock edges or signal transitions. A Discrete-Event kernel models this by maintaining a centralized **Event Queue** ordered strictly by virtual timestamps.

The simulation does not execute continuously; it jumps from one event timestamp to the next. The DE simulation semantics rely heavily on a two-phase execution model known as the **Evaluate-Update paradigm**. This paradigm is designed to accurately simulate the parallel nature of hardware on a sequential host CPU.

1. **Evaluate Phase:** The simulation kernel selects all processes (threads or methods) that are sensitive to the current timestamp or triggered by events at this exact virtual time. The kernel executes these processes one by one. Crucially, during this phase, any assignments made to hardware

signals *do not immediately update the signal's value*. Instead, the new values are scheduled to be updated in the future.

2. **Update Phase:** Once all runnable processes have finished evaluating, the kernel applies all the scheduled signal updates simultaneously.

This two-phase approach guarantees deterministic behavior. If Process A and Process B both read Signal X and write to Signal Y at the exact same virtual nanosecond, the Evaluate-Update paradigm ensures that both processes read the old value of X before either can update Y. The order in which the host CPU executes Process A and Process B during the Evaluate phase becomes completely irrelevant, ensuring reproducibility.

10.3.1 Delta Cycles

A unique artifact of the DE simulation kernel is the **delta cycle**. A delta cycle represents an infinitesimal step in virtual time. If the Update phase alters a signal, that alteration might trigger a *new* event that is scheduled to occur at the *exact same virtual timestamp* (zero-delay).

When this happens, the simulation kernel does not advance the virtual clock. Instead, it initiates a new Evaluate phase followed by a new Update phase, all within the same virtual nanosecond. This loop of Evaluate-Update stages at a constant virtual time is called a delta cycle. Delta cycles allow the simulator to resolve zero-delay combinational logic paths (such as a cascade of logic gates) in a causally correct sequence without advancing physical time. Only when the event queue contains no more zero-delay events does the kernel advance the virtual clock to the timestamp of the next scheduled event.

10.4 Transaction Level Modeling (TLM) and Time

While cycle-accurate DE simulation is highly precise, simulating every individual clock edge, pin transition, and delta cycle for a massive System-on-Chip (SoC) interconnect is incredibly slow. To achieve the simulation speeds necessary to boot an RTOS and run complex CPS software, architects rely on **Transaction Level Modeling (TLM)**.

TLM abstracts away the low-level physical pins and wires of the SoC interconnect (such as AMBA AXI or APB). Instead of simulating the clock-by-clock toggling of `VALID` and `READY` signals, TLM uses object-oriented method calls (like C++ functions) to transfer data between initiators (e.g., CPUs, DMA controllers) and targets (e.g., memories, UARTs).

In the widely adopted OSCI TLM 2.0 standard, these interactions are modeled as transactions containing a payload (the data, address, and command). TLM provides multiple levels of timing abstraction:

- **Untimed (UT):** Used for pure functional software development. Transactions execute instantly in zero virtual time. There is no concept of bus contention or propagation delay.
- **Loosely Timed (LT):** Used for RTOS and driver development. Transactions are annotated with estimated delays, but the simulation relies on **temporal decoupling**. In LT, the CPU is allowed to run ahead of the global virtual clock for a specific “quantum” of time. It accumulates its own local time and only synchronizes with the rest of the system when its quantum expires. This provides immense simulation speedups by avoiding constant context switching with the DE kernel.
- **Approximately Timed (AT):** Used for architectural exploration. This model breaks transactions into multiple phases to accurately simulate bus contention, pipelining, and backpressure. For example, a write transaction is modeled using a four-phase protocol: *Begin Request*, *End Request*, *Begin Response*, and *End Response*. The AT model interacts heavily with the DE kernel's event queue to enforce these timing boundaries, resulting in highly accurate but slower simulations.

For VirtMCU and CPS co-simulation, we must strike a careful balance. We need the speed of Loosely Timed TLM to execute millions of instructions, but we require strict timing boundaries to interact with continuous-time physics.

10.5 Co-Simulation: Bridging the Cyber and Physical

Co-simulation in CPS 2.0 is the simultaneous execution of two fundamentally different simulation domains: the discrete-event cyber simulator (e.g., VirtMCU/QEMU executing firmware) and the continuous-time physical simulator (e.g., a MuJoCo physics engine simulating kinematics).

Because these two engines possess entirely different mathematical foundations and concepts of time, they cannot run asynchronously. They must be coupled through an **External Clock Authority**.

In a co-simulation environment, time is quantized into strictly bounded **macro-steps** (e.g., 1 millisecond).

1. The External Clock Authority grants permission for the physics engine to integrate its differential equations forward by 1ms.
2. Simultaneously, it grants the cyber simulator permission to execute exactly 1ms worth of virtual CPU instructions and timer ticks.
3. At the end of the 1ms macro-step, both simulators pause.
4. The physical domain samples its state (e.g., joint angles, velocities) and maps these continuous values to discrete sensor registers (the Sensor Abstraction Layer, SAL).
5. The cyber domain samples its memory-mapped output registers (the Actuator Abstraction Layer, AAL) and maps these discrete commands (e.g., PWM duty cycles) into physical forces.
6. The state variables are exchanged via zero-copy shared memory, and the External Clock Authority issues the command to begin the next 1ms macro-step.

This lock-step synchronization guarantees strict causality. The software cannot run ahead of the physics, and the physics cannot outpace the software. However, enforcing this macro-step boundary on a fast binary-translating emulator like QEMU requires aggressive structural modifications to how the emulator manages its internal execution loop.

10.6 VirtMCU Clock Slaving

VirtMCU solves the emulator timing problem by entirely overriding QEMU's internal clock infrastructure. Instead of querying the host OS wall-clock to manage guest timers, VirtMCU utilizes a dynamic QOM (QEMU Object Model) plugin architecture to enforce deterministic virtual time via the `hw/rust/backbone/clock` module.

VirtMCU implements a concept known as **Cooperative Time Slaving**. Under this architecture, the virtual CPU execution loop is subordinated to an external synchronizer. VirtMCU provides two primary modes of operation to achieve this determinism: `slaved-suspend` and `slaved-icount`.

10.6.1 Slaved-Suspend Mode

In `slaved-suspend` mode, VirtMCU is designed for high-throughput temporal decoupling. The CPU executes instructions as fast as the host hardware allows. However, the `hw/rust/backbone/clock` module intercepts the QEMU main loop and enforces a strict barrier at the macro-step boundary. If the external physics engine dictates a 1ms step, the CPU is allowed to execute instructions until its accumulated internal virtual time reaches 1ms. Once it hits this exact boundary, the virtual CPU thread is physically suspended (blocked) by the VirtMCU backbone. It will remain suspended indefinitely until the external physics engine completes its own 1ms integration and sends an IPC (Inter-Process Communication) signal to resume for the next step.

10.6.2 Slaved-Icount Mode

While `slaved-suspend` forces alignment at the macro-step, the virtual time within that step might still be somewhat abstracted. For cycle-accurate, microsecond-precision validation, VirtMCU utilizes `slaved-icount` mode.

In `slaved-icount` (Instruction Count) mode, virtual time is completely decoupled from any concept of physical time or host execution speed. Instead, time advances purely as a strict mathematical function of the number of retired

instructions. If the simulated embedded processor is defined to run at 100 MHz, the execution of exactly one machine instruction advances the global virtual clock by precisely 10 nanoseconds.

Because the `icount` virtual clock is the absolute source of truth for the entire SoC, all internal hardware timers, watchdogs, and peripheral baud-rate generators are driven directly by this instruction-derived clock. If the firmware configures a 10ms hardware timer, the VirtMCU backbone calculates that exactly 1,000,000 instructions must retire before the interrupt can fire. The execution loop will execute exactly 1,000,000 instructions, stop, inject the hardware IRQ into the NVIC, and then execute the first instruction of the ISR.

This creates a perfectly deterministic, 100% reproducible execution trace. If you run the simulation on a massive server, it might finish in two seconds. If you run it on a slow, 10-year-old laptop, it might take five minutes. However, in both scenarios, the simulated control loop will execute the exact same number of instructions, and the hardware timer will fire on the exact same instruction boundary.

10.7 Determinism and Reproducibility in CPS Validation

Why go to such extreme lengths to eliminate jitter? In the context of Cyber-Physical Systems, non-deterministic timing behaviors can easily mask fatal architectural flaws.

Consider the operational dynamics of a Cyber-Physical Virtual Power Plant (VPP), which coordinates hundreds of Distributed Energy Resources (DERs) via complex communication networks. To maintain grid stability during an emergency frequency support event, a centralized VPP must issue precision demand response commands to isolated loads and energy storage systems. The timing of these command deliveries is critical. If a command packet to a storage inverter is delayed by communication jitter or processing stalls, the response power changes unexpectedly.

As highlighted in distributed system literature, due to the heterogeneity of the information transmitted, changes in the information arrival state result in unequal covariate coefficients in the system's economic and stability models. If a response command to DER A is delayed, it might cause a system cost increment of \$30, whereas delaying a command to DER B might only cost \$10. If the validation environment suffers from wall-clock scheduling jitter, the simulations will produce different economic and physical stability profiles on every single run. It becomes impossible to mathematically verify the control algorithms or optimize the communication scheduling policies because the baseline data is corrupted by simulation noise.

By deploying VirtMCU's deterministic clock slaving, engineers can create a multi-node simulation where the network delays, interrupt latencies, and RTOS context switches are absolutely rigid. If a priority inversion bug exists in the RTOS that causes a 50-microsecond delay in the VPP frequency response task, VirtMCU guarantees that this 50-microsecond delay will appear consistently in every simulation, allowing the engineer to attach a debugger, halt time, and inspect the exact state of the pipeline and memory map.

10.8 VirtMCU Homework: Control Loop Simulation

Understanding the theoretical impact of simulation jitter is different from observing it directly in an active execution pipeline. To fully internalize the difference between traditional software emulation and strict co-simulation, this chapter's homework requires you to deploy a real-time control loop within the VirtMCU environment.

1. **Firmware Implementation:** You will be provided with a bare-metal C framework. You must write an Interrupt Service Routine (ISR) that is triggered by a hardware timer configured to fire exactly every 10.00 milliseconds. Inside the ISR, you will increment a global volatile counter and read the current value of a free-running microsecond hardware counter to measure the elapsed time since the last invocation.

2. **Execution under Jitter:** First, compile the firmware and load it into a standard, unmodified QEMU ARM emulator running on your host OS. To simulate a noisy industrial environment, you will concurrently run a heavy, multi-threaded CPU stress test in the background of your host OS. You will print the delta between hardware timer invocations. You will observe significant variance (jitter) as the host OS scheduler randomly preempts the emulator to service the stress test.
3. **Deterministic `slaved-icount` Execution:** Next, load the exact same compiled binary into the VirtMCU framework configured in `slaved-icount` mode. Launch the same background CPU stress test on your host OS.
4. **Analysis:** You will observe that regardless of the immense load on your physical host machine, the delta between ISR invocations in VirtMCU will remain mathematically perfect. Because the virtual clock is slaved exclusively to the instruction count and not the wall clock, the VirtMCU environment completely isolates the cyber-physical state from external computing noise. You must submit a brief report containing the histogram of your timing measurements, proving empirically the necessity of virtual time in CPS 2.0 design.

Chapter 11: The Sensor/Actuator Abstraction Layer (SAL/AAL)

11.1 Introduction: The Cyber-Physical Boundary

In Chapter 10, we established the temporal foundation required for Cyber-Physical System (CPS) co-simulation. We demonstrated that standard software emulators fail to provide the rigorous timing guarantees required for CPS validation due to host operating system scheduling jitter. By implementing Cooperative Time Slaving and deterministic `slaved-icount` execution in VirtMCU, we successfully lock-stepped the discrete-time cyber domain (the processor) with the continuous-time physical domain (the physics engine).

However, synchronizing time is only half of the co-simulation equation. Even if the virtual processor and the physics engine advance their clocks in perfect harmony, they do not speak the same language. A physics engine like MuJoCo calculates the state of the world using continuous floating-point variables representing Newtonian mechanics—positions, velocities, accelerations, and forces. Conversely, the microcontroller firmware operates purely in the discrete digital domain, reading and writing binary integers to memory-mapped I/O (MMIO) registers.

A simulated drone’s flight controller cannot magically “read” a floating-point quaternion representing its physical orientation directly from the physics engine’s memory. In reality, the drone’s processor communicates with a physical Micro-Electro-Mechanical System (MEMS) gyroscope over an I2C or SPI serial bus. The gyroscope senses the continuous physical rotation, converts it to an analog voltage, digitizes it, and stores it in an internal hardware register, which the CPU then reads.

To bridge this semantic gap in our simulation environment, we must construct a rigorous translation interface. In the VirtMCU framework, this is achieved through the **Sensor Abstraction Layer (SAL)** and the **Actuator Abstraction**

Layer (AAL). These layers sit between the raw MMIO hardware emulation and the external physics engine, translating continuous-time physical variables into discrete-time cyber registers (and vice versa) while injecting real-world artifacts such as quantization error, sensor noise, and actuator saturation.

11.2 Mathematical Modeling of Physical Systems

Before we can translate physical phenomena into digital data, we must formalize how the continuous physical world is modeled. The physical domain is governed by the laws of physics, which yield governing differential equations and constitutive relationships.

For mechanical systems, engineers typically employ lumped-parameter models analyzed via classical mechanics, specifically Newtonian, Lagrangian, or Hamiltonian equations of motion. In a lumped-parameter model, the spatially distributed properties of a physical system (such as mass, damping, and stiffness) are approximated as discrete, concentrated components (ideal masses, dashpots, and springs) interacting in continuous time.

The state of such a physical system at any instant is represented by a state vector $x(t)$, which evolves according to a set of non-linear differential equations: $\dot{x}(t) = f(x(t), u(t), t) + w(t)$ where $u(t)$ represents the continuous control input (actuation forces) and $w(t)$ represents stochastic perturbations or environmental disturbances.

For example, the dynamics of an aerial vehicle require tracking both the translational positions and velocities in the inertial frame, as well as the rotational Euler angles (roll ϕ , pitch θ , yaw ψ) and angular rates (p , q , r) in the body frame. To simulate this, the continuous-time physics engine numerically integrates these differential equations using advanced solvers (such as Runge-Kutta methods).

The primary objective of the SAL and AAL is to discretize these continuous variables $x(t)$ for the cyber domain's sensors, and to convert the cyber domain's discrete outputs into the continuous control inputs $u(t)$ applied to the physical model.

11.3 The Sensor Abstraction Layer (SAL)

The Sensor Abstraction Layer (SAL) is responsible for intercepting the exact state variables $x(t)$ computed by the physics engine and transforming them into the binary formats expected by the firmware's device drivers. This process replicates the entire signal chain of a physical transducer, signal conditioning circuit, and Analog-to-Digital Converter (ADC).

11.3.1 Transduction and Analog-to-Digital Conversion (ADC)

In the physical world, sensors rely on electrostatic, electromagnetic, or piezoresistive transducers to convert a physical property (such as pressure, temperature, or acceleration) into a proportional analog electrical signal (voltage or current).

Once transduced to an analog voltage, the signal must be digitized. The SAL models the behavior of a hardware ADC through two primary mathematical operations: **Sampling** (discretization in time) and **Quantization** (discretization in amplitude).

1. **Sampling:** The SAL does not update the sensor registers continuously. Instead, it updates them at a specific frequency, defined by the sensor's sampling rate. If the firmware attempts to read the sensor register faster than the sampling rate, it will simply read the same stale value twice, exactly as it would on real hardware.
2. **Quantization:** A physical ADC maps a continuous voltage range into a finite number of discrete binary bins. An N -bit ADC provides 2^N possible output values. The resolution of the sensor, or the Least Significant Bit (LSB) weight, is calculated as: $\text{Resolution} = \frac{V_{\max} - V_{\min}}{2^N}$. The SAL takes the continuous floating-point variable from the physics engine, scales it according to the sensor's physical sensitivity (e.g., $16.4 \text{ LSB per degree/second}$ for a gyroscope), and truncates the floating-point mantissa to force the value into an N -bit integer. This intentional introduction of **quantization error** is crucial, as it fundamentally limits the precision of the feedback loop in the cyber domain.

11.3.2 Simulating Sensor Noise and Stochastic Perturbations

A naive co-simulation environment would simply pass the quantized, ground-truth physics value to the firmware. However, in reality, physical systems exhibit randomness and are perturbed by random perturbations, such as noise, interference, and other stationary and nonstationary stochastic processes.

If a control algorithm is developed and tested against perfect, noise-free sensor data, it will almost certainly fail when deployed to physical hardware. Therefore, the SAL must actively inject noise into the signal before passing it to the MMIO registers.

The SAL models sensor imperfections using several stochastic noise profiles:

- **White Noise (Gaussian):** High-frequency thermal and electrical noise is modeled by generating random variables from a Gaussian distribution with a zero mean and a specific variance (σ^2). This noise is added to the continuous signal prior to quantization.
- **Bias (Offset):** Manufacturing imperfections mean that a sensor will rarely read exactly zero when the physical state is zero. The SAL adds a constant (or slowly varying) bias offset. For example, a simulated accelerometer resting flat on a table might report $0.02g$ on the X-axis instead of $0.00g$.
- **Random Walk (Drift):** MEMS gyroscopes famously suffer from bias instability, where the sensor bias wanders randomly over time. The SAL simulates this by accumulating small Gaussian random variables at each macro-step, creating a Brownian motion drift profile that forces the student's firmware to implement sensor fusion algorithms (like a Kalman filter) to estimate and subtract the drift.

By mathematically simulating these electrostatic and electromagnetic transducer characteristics, the SAL forces the software engineer to design robust cyber algorithms that can handle the harsh realities of the physical world.

11.4 The Actuator Abstraction Layer (AAL)

While the SAL handles data flowing from the physical world to the cyber world, the Actuator Abstraction Layer (AAL) handles the reverse: translating the discrete commands generated by the firmware into continuous physical forces, torques, or voltages $u(t)$ required by the physics engine.

11.4.1 Pulse Width Modulation (PWM) and DAC

In modern embedded systems, it is rare for a microcontroller to output a true continuous analog voltage using a Digital-to-Analog Converter (DAC) to drive heavy loads. Instead, processors control physical actuators—such as heaters, hydraulic valves, and electric motors—using **Pulse Width Modulation (PWM)**.

A PWM controller is a hardware peripheral that generates a high-frequency digital square wave. The firmware configures the PWM peripheral by writing to two MMIO registers: the Period Register (which defines the frequency of the wave) and the Compare Register (which defines how long the wave remains HIGH during each period). The ratio of the HIGH time to the total period is the **duty cycle**.

To an electrical load like a DC motor, the high-frequency switching of the PWM signal is naturally low-pass filtered by the motor's internal inductance and mechanical inertia. Thus, the motor responds to the *average* voltage of the PWM signal, which is directly proportional to the duty cycle.

The AAL in VirtMCU intercepts the firmware's writes to the MMIO PWM Compare Registers. It mathematically computes the duty cycle percentage (e.g., a Compare value of 250 out of a Period of 1000 yields a 25% duty cycle). If the simulated motor is connected to a 12V power supply, the AAL translates this 25% duty cycle into an effective continuous voltage of 3.0V.

11.4.2 Actuator Dynamics and Saturation

A common mistake in software engineering is assuming that an actuator responds instantly to a command. In the physical domain, actuators have their

own internal governing differential equations.

Consider an electric drive actuated by a geared permanent magnet DC motor. The motor is not an instantaneous torque generator; it is a complex electro-mechanical system. Its electrical dynamics are governed by its armature resistance (r_a) and inductance (L_a). When the AAL applies the 3.0V effective voltage, the current $i_a(t)$ ramps up according to the electrical time constant (L_a / r_a). The generated mechanical torque is proportional to this current ($T = k_t i_a$).

Simultaneously, the motor's mechanical dynamics are opposed by the back-electromotive force (back-EMF), which grows as the motor's angular velocity ω_r increases, as well as mechanical viscous friction (B_m) and the moment of inertia (J).

The AAL must account for these dynamics. While the primary physics engine (MuJoCo) handles the rigid body dynamics of the drone or robot, the AAL often computes a localized, high-speed sub-simulation of the actuator's internal electro-magnetic state to produce the final, accurate torque applied to the mechanical joint.

Furthermore, the AAL enforces **control limits (saturation)**. Control limits $u_{\min} \leq u(t) \leq u_{\max}$ are described by real-valued bounded functions. If a PID controller mathematically demands 50 Amps of current to stabilize a drone, but the physical motor driver can only source 10 Amps, the AAL physically clips the output at 10 Amps. This saturation can cause "integral windup" in poorly designed firmware, providing an excellent pedagogical lesson in non-linear control limits.

11.5 Co-Simulation via Zero-Copy Shared Memory

The SAL and AAL perform intense mathematical transformations, but they must exchange this data with the external physics engine at every macro-step of the simulation. Because the physics engine (e.g., MuJoCo) and the cyber emulator (VirtMCU) run as completely separate operating system processes, they must utilize Inter-Process Communication (IPC).

Standard IPC mechanisms, such as TCP/IP sockets or Unix pipes, involve copying data from the user space of one process, into the host kernel space, and then back into the user space of the receiving process. When running a hard real-time co-simulation executing a 1-millisecond macro-step, the latency and context-switching overhead of socket-based IPC destroys simulation throughput.

To resolve this, VirtMCU and the physics engine are coupled through **zero-copy shared memory**. Using the POSIX `shm_open` and `mmap` system calls, the host operating system allocates a single contiguous block of physical RAM that is mapped simultaneously into the virtual address space of both the VirtMCU process and the MuJoCo process.

The shared memory region is structured into three distinct zones:

1. **Synchronization Primitives:** Atomic lock variables and condition variables used by the External Clock Authority to orchestrate the macro-step lock-stepping discussed in Chapter 10.
2. **Physics $\rightarrow$ Cyber Zone:** Written exclusively by MuJoCo. Contains arrays of continuous variables (forces, angles, positions). The SAL reads directly from these memory addresses, applies its noise and quantization models, and writes the results to the emulated MMIO registers.
3. **Cyber $\rightarrow$ Physics Zone:** Written exclusively by VirtMCU. Contains the continuous control inputs $u(t)$ (e.g., applied torques and voltages) computed by the AAL. MuJoCo reads directly from these addresses during its integration step.

Because the data never passes through the host kernel, the exchange occurs in tens of nanoseconds, allowing the co-simulation to execute millions of macro-steps per second without IPC bottlenecking.

11.6 Physical Modeling: The Inverted Pendulum

To solidify the concepts of SAL, AAL, and lumped-parameter modeling, we turn to a classic benchmark in control theory: the inverted pendulum on a cart. This

system perfectly encapsulates the multi-axis coupling and instability inherent to cyber-physical systems.

The physical system consists of a motorized cart of mass M that moves horizontally along a linear track (position x). A pendulum of mass m and length l is attached to the cart via a frictionless hinge. The angle of the pendulum relative to the vertical upright position is denoted by θ .

The goal of the cyber system is to maintain the pendulum in the unstable upright position ($\theta = 0$) by applying a horizontal force $u(t)$ to the cart. If the pendulum begins to fall forward, the cart must accelerate forward to slide the base “under” the falling mass to catch it.

Using Newtonian mechanics or the Lagrange equations of motion, we derive the non-linear governing differential equations for this system. Summing the forces on the cart and the moments about the pivot, and accounting for viscous friction B_m , yields a coupled system of equations. For example, the rotational dynamics of the inverted arm on the moving platform can be expressed as:
$$ml^2 \ddot{\theta} + ml \ddot{x} \cos \theta - mgl \sin \theta + B_m \dot{\theta} = 0$$
 (Note: Gravity g creates a destabilizing moment when the pendulum deviates from the vertical, hence the $-mgl \sin \theta$ term when linearized around the upright position).

To design a linear control law (like a PID controller), engineers typically linearize these non-linear differential equations around the unstable equilibrium point ($\theta = 0, \dot{\theta} = 0, x = 0, \dot{x} = 0$). For small angles, we apply the small-angle approximations: $\sin \theta \approx \theta$ and $\cos \theta \approx 1$.

In the co-simulation environment:

1. **The Physics Engine (MuJoCo)** numerically integrates the full, non-linear differential equations in continuous time.
2. **The SAL** reads the continuous cart position $x(t)$ and pendulum angle $\theta(t)$ from shared memory. It adds Gaussian noise to simulate optical encoder jitter, quantizes the angles into 12-bit integers, and writes them to the virtual Quadrature Encoder Interface (QEI) MMIO registers in VirtMCU.

3. **The Firmware** executes its control loop, reads the noisy integer angles, computes the proportional, integral, and derivative errors, and writes a duty cycle value to the PWM MMIO register.
4. **The AAL** reads the PWM register, converts the duty cycle to a continuous horizontal force $u(t)$ (accounting for motor saturation limits), and writes this force back to the shared memory.
5. **The Physics Engine** applies $u(t)$ to the cart mass M for the next integration step.

11.7 VirtMCU Homework: Inverted Pendulum Simulation

In standard software engineering curricula, writing a sorting algorithm allows you to verify your logic with an array of integers. In Cyber-Physical Systems, you verify your logic by preventing a physical structure from collapsing.

For this chapter's homework, you will utilize the VirtMCU FirmwareStudio linked with the MuJoCo physics engine via the SAL/AAL shared memory interface. You will write the firmware required to balance the inverted pendulum described in Section 11.6.

Your Assignment:

1. **Driver Initialization:** Write C code to initialize the memory-mapped virtual peripherals. Configure the simulated ARM Cortex-M hardware timer to generate an interrupt exactly every 5.0 milliseconds. This sets your discrete control loop period (T_s).
2. **Sensor Acquisition (SAL):** Inside the timer Interrupt Service Routine (ISR), read the raw 16-bit integer values from the virtual Quadrature Encoder registers mapping to the cart position and pendulum angle. Convert these raw integers back into standard floating-point engineering units (meters and radians), accounting for the quantization scaling documented in the hardware specification.
3. **Digital Control Law:** Implement a cascaded Proportional-Integral-Derivative (PID) control algorithm. Your outer loop should compute a

target angle to move the cart to a desired setpoint, while the inner loop computes the stabilizing force required to maintain the pendulum at that target angle. You must manually tune the k_p , k_i , and k_d feedback gains to achieve stability.

4. **Actuation (AAL):** Convert your computed force request into a PWM duty cycle integer. Write this integer to the virtual motor controller's MMIO Compare Register.
5. **Co-Simulation Execution:** Launch the VirtMCU environment in `slaved-icount` mode alongside MuJoCo. Observe the 3D rendering of the cart. If your mathematics are correct, the cart will catch the falling pendulum and balance it perfectly.
6. **Robustness Testing:** The simulation framework includes a Python script to inject random physical disturbances (simulated wind gusts) into the MuJoCo environment. You must document how your PID control law rejects these unmodeled stochastic perturbations and prevents the pendulum from falling.

Chapter 12: The RISC-V Architecture and Instruction Set

12.1 Introduction: The Need for Open Architectures in CPS

In the preceding chapters, we explored how standard microprocessors interact with the physical world. Through the Sensor/Actuator Abstraction Layer (SAL/AAL) discussed in Chapter 11, we demonstrated how continuous-time physical phenomena are translated into discrete-time memory-mapped I/O (MMIO) registers. However, when developing advanced Cyber-Physical Systems (CPS)—such as high-frequency drone flight controllers or ultra-low-latency robotic actuators—the processing overhead of reading MMIO registers, shifting bits, and executing standard arithmetic can become a bottleneck.

To achieve maximum performance, computer architects often need to modify the processor itself, adding custom hardware accelerators and new machine instructions tailored specifically for the physical application. Historically, this has been nearly impossible for academic researchers and independent engineers because architectures like ARM and x86 are strictly proprietary.

This paradigm shifted dramatically with the introduction of **RISC-V**. RISC-V is an entirely open-source specification for a reduced instruction set computer architecture. Because it is open and royalty-free, hardware designers are explicitly encouraged to modify the silicon and extend the instruction set to meet the unique requirements of their specific domain. In this chapter, we will examine the RISC-V architecture, explore its elegant instruction set, and discuss how its open nature enables the creation of custom silicon for Cyber-Physical Systems.

12.2 The Origins and Philosophy of RISC-V

The RISC-V architecture was publicly announced in 2011 after being developed at the University of California, Berkeley, by a team including Yunsup Lee, Krste Asanović, David A. Patterson, and Andrew Waterman. The “V” represents the Roman numeral five, signifying that this was the fifth major RISC architectural design project undertaken at UC Berkeley.

Unlike many commercial architectures that are burdened by decades of legacy compatibility requirements, the RISC-V project began as a clean sheet design. The architects established several principal goals to ensure the ISA could scale across the entire spectrum of computing:

- **Broad Scalability:** The architecture must efficiently support everything from ultra-low-power microcontrollers operating at the IoT edge to extreme-scale computing nodes in warehouse-scale cloud server farms.
- **Modularity and Extensions:** The base instruction set is deliberately minimal. Advanced capabilities are provided via optional ISA extensions, which support features like floating-point mathematics, atomic memory operations, and hardware multiplication and division.
- **Privilege Levels:** Additional extensions support privileged execution modes (similar to the x86 and ARM models discussed in earlier chapters) necessary to run secure operating systems.
- **Code Density:** To accommodate memory-constrained embedded systems, RISC-V supports a compressed instruction set extension. This allows 16-bit versions of common 32-bit instructions to be freely interspersed with standard 32-bit instructions in memory, significantly reducing the overall code footprint.
- **Address Space Expansion:** The architecture provides optional extensions to support 64-bit, and even 128-bit, processor word sizes and paged virtual memory for massive multiprocessing configurations.

Additionally, like modern ARM processors, the RISC-V architecture supports **bi-endianness**. This allows the operating system to select between big-endian or little-endian mode under software control, though most modern operating systems choose little-endian mode by default.

12.3 The RISC-V Programmer's Model

To program a RISC-V processor, a software engineer must understand its architectural state. At the core of the RISC-V base ISA is a remarkably orthogonal and simple register set.

The base architecture provides 32 general-purpose registers, named `x0` through `x31`. In the standard 32-bit architecture (RV32), each of these registers is exactly 32 bits wide. Registers `x1` through `x31` are completely unrestricted and have no special functions forced upon them by the processor hardware.

However, the `x0` register is unique: it is hardwired to the constant value zero. Any attempt to read `x0` will always return zero, and any value written to `x0` is immediately discarded by the hardware. As we will see in the next section, this hardwired zero register allows the instruction set to be drastically simplified.

12.3.1 The Application Binary Interface (ABI)

Although the processor hardware treats registers `x1` through `x31` as functionally interchangeable, building complex software requires strict agreements on how registers are used across different functions and libraries. Software compatibility requires that we specify which register serves as the stack pointer, which registers contain function arguments, and which contain return values.

These standard rules form the **Application Binary Interface (ABI)**. By standardizing register usage and calling conventions, the ABI ensures that compiled C code can successfully link with handwritten assembly routines on any compatible implementation. To aid programmers, RISC-V assemblers recognize alternate ABI names for the registers. The primary ABI designations are:

- **ra (x1)**: The Return Address register, used to hold the address to jump back to after a function call completes.
- **sp (x2)**: The Stack Pointer, pointing to the top of the current activation record in memory.

- **a0–a7 (x10–x17)**: The Argument registers, used to pass parameters into functions and to return results back to the caller.
- **t0–t6** : Temporary registers, which are volatile (caller-saved).
- **s0–s11** : Saved registers, which are non-volatile (callee-saved).

12.4 The Instruction Set and Pseudo-Instructions

The philosophy of RISC-V is to keep the hardware as simple as possible. Consequently, the base ISA lacks many distinct instructions found in other architectures (such as explicit `MOV` or `NOT` operations). Instead of forcing the hardware to decode dozens of unique opcodes, the architecture relies on the assembler program to synthesize missing operations using combinations of the base instructions and the hardwired `x0` register.

These synthesized aliases are known as **pseudo-instructions**. When the programmer types a pseudo-instruction, the assembler automatically translates it into the corresponding base RISC-V machine instruction. Table 12.1 outlines some of the most common pseudo-instructions:

- **nop (No Operation)**: Translated to `addi x0, x0, 0`. It adds zero to zero and discards the result in `x0`, consuming a clock cycle without altering the machine state.
- **mv rd, rs (Move)**: Translated to `addi rd, rs, 0`. It copies the value of the source register `rs` into the destination `rd` by adding zero to it.
- **not rd, rs (Bitwise NOT)**: Translated to `xori rd, rs, -1`. Exclusive-ORing a value with all 1s (-1) perfectly inverts all the bits.
- **neg rd, rs (Negate)**: Translated to `sub rd, x0, rs`. It mathematically negates `rs` by subtracting it from the `x0` register (zero).
- **j offset (Unconditional Jump)**: Translated to `jal x0, offset`. It jumps to the offset but discards the return address into `x0`.
- **beqz rs, offset (Branch if Equal to Zero)**: Translated to `beq rs, x0, offset`.

By pushing these translations into the software assembler, the silicon required to decode instructions in the CPU is significantly reduced, lowering power

consumption and increasing maximum clock frequencies.

12.4.1 Immediate Values and the Sign-Extension Quirk

Loading 32-bit constants or absolute 32-bit memory addresses into a register poses a challenge because RISC-V instructions are only 32 bits wide themselves. There is no physical room to pack a 32-bit opcode alongside a 32-bit constant.

To solve this, RISC-V breaks the operation into two steps using a 20-bit upper immediate and a 12-bit lower immediate. The `lui` (Load Upper Immediate) instruction loads a 20-bit value into the highest 20 bits of a register, clearing the lower 12 bits. The `auipc` (Add Upper Immediate to PC) performs a similar task but adds the 20-bit value to the Program Counter, enabling position-independent code. The programmer then uses a subsequent instruction—like `addi` or a load/store offset—to supply the remaining lower 12 bits.

However, this two-step process introduces a subtle but critical mathematical quirk. In RISC-V, the 12-bit immediate value in the `addi` instruction is always treated as a signed two's-complement number. Before the processor adds this 12-bit value to the upper 20 bits, the hardware automatically sign-extends the most significant bit of the 12-bit value (bit 11) all the way through bit 31.

If bit 11 happens to be a `1`, the 12-bit value is treated as a negative number, effectively *subtracting* from the upper 20 bits. For example, suppose a programmer wants to load the 32-bit value `0xFFFFFFFF` into the `x1` register. If they naively write:

```
lui x1, 0xFFFF    # x1 now equals 0xFFFF000
addi x1, x1, 0xFFF
```

The `addi` instruction will take `0xFFF`, interpret it as a signed number, sign-extend it to `0xFFFFFFFF` (which is -1), and add it to `0xFFFF000`. The resulting sum placed in `x1` will be `0xFFFFEFFF`, which is totally incorrect.

To prevent this, the assembler must perform mathematical compensation when evaluating the `la` (Load Address), `lw` (Load Word), and `sw` (Store Word) pseudo-instructions. If bit 11 of the lower 12-bit payload is a `1`, the assembler

must pre-emptively add 1 to the upper 20-bit value before generating the `lui` or `auipc` instruction. This slight increment in the upper bits perfectly cancels out the subsequent negative sign-extension of the lower bits, resulting in the correct 32-bit address.

12.5 Hardware Implementations and Custom Extensions

What makes RISC-V uniquely suited for modern Cyber-Physical Systems is its explicit support for user-defined enhancements. A team designing a robotic actuator does not have to settle for the standard instruction set; they can augment the processor with custom opcodes, hardware coprocessors, and other modifications as long as they comply with the RISC-V customization rules.

Because all the source code and processor design intellectual property required to implement a RISC-V core is freely available on the internet, developers can synthesize their own custom processors without paying licensing fees. By using an open-source RISC-V design as a starting point, engineers can implement custom modifications that are guaranteed to remain compatible with future versions of the RISC-V standard.

For example, the PicoRV32 is a popular open-source RISC-V core that can be easily synthesized and programmed into low-cost Field Programmable Gate Arrays (FPGAs). A budget-friendly FPGA board like the Digilent Cmod A7-35T utilizes a Xilinx Artix-7 chip (XC7A35T). This specific FPGA provides 5,200 logic slices, 41,600 flip-flops, 1,800 kbits of block RAM, and 90 specialized Digital Signal Processing (DSP) slices that natively support high-performance Multiply-Accumulate (MAC) operations.

An engineer can map the PicoRV32 open-source model onto this FPGA using the Vivado synthesis toolchain. Once the baseline RV32IMC core is instantiated, the engineer can write Verilog code to bind the FPGA's physical DSP slices directly to custom RISC-V opcodes. This creates a deeply integrated, hardware-accelerated processor capable of executing complex PID control loops or AI

inference models in a fraction of the time required by standard software execution.

12.6 VirtMCU Homework: Custom ISA Extensions for Deterministic Control

In previous chapters, you wrote control algorithms assuming a standard, immutable CPU architecture. However, in the CPS 2.0 paradigm, hardware and software must be co-designed. If a software control loop cannot meet its hard real-time deadlines, you now have the power to alter the hardware to solve the problem.

For this chapter's homework, you will step beyond standard software engineering and venture into computer architecture by modifying the ISA within the **VirtMCU FirmwareStudio**.

1. **Baseline Profiling:** You will be provided with a standard C program that reads three orthogonal sensor registers (X, Y, and Z axes) via MMIO, computes their vector magnitude ($\sqrt{X^2 + Y^2 + Z^2}$), and writes the result to an actuator register. Compile this for the standard RISC-V RV32I base architecture. Run it in VirtMCU's deterministic `slaved-icount` mode and record the exact number of clock cycles required to execute the math loop.
2. **Architectural Modification:** You will dive into the VirtMCU QEMU backend source code. You will define a custom RISC-V instruction, `VMAG rd, rs1, rs2, rs3`, that utilizes an unused opcode space. You will write the emulation logic in C to calculate the vector magnitude in a single simulated clock cycle.
3. **Inline Assembly Implementation:** Modify your original C firmware to bypass the standard math libraries. Use GCC's inline assembly feature to invoke your new `VMAG` machine instruction directly.
4. **Co-Simulation Validation:** Re-run the modified firmware in VirtMCU `slaved-icount` mode. Compare the new cycle count against your baseline profile. You must document the performance speedup achieved by pushing a specific computational bottleneck directly into the silicon

architecture, empirically demonstrating the power of an open ISA in CPS design.

Chapter 13: Security and Resilience in Cyber-Physical Systems

13.1 Introduction: Securing Critical Infrastructure

As established in Chapter 9, the transition to Cyber-Physical Systems (CPS 2.0) deeply integrates Operational Technology (OT) with Information Technology (IT), significantly expanding the attack surface of physical infrastructure. While traditional IT security focuses on data confidentiality and privacy, security in a CPS context is fundamentally about maintaining physical control and ensuring the safety and availability of critical services.

Nowhere is this more evident than in **Critical Energy Infrastructure (CEI)**. Modern energy systems and power grids are constantly under stress from both natural disruptions—such as wildfires and extreme weather events—and sophisticated, targeted cyberattacks. To deal with the ever-increasing challenge posed by these cyber-physical events, we must shift our focus from mere prevention to **resilience**.

Resilience is defined as the ability to withstand and reduce the magnitude and duration of disruptive events. A resilient CPS possesses the capability to anticipate, absorb, adapt to, and rapidly recover from an attack or failure. This chapter explores the methodologies used to model cyber-physical threats, assess risks, and deploy advanced computational techniques—such as robust optimization, decentralized autonomous controls, and distributed topology reconfiguration—to ensure the survival of critical infrastructure under adversarial conditions.

13.2 The Cyber-Physical Threat Landscape

To properly secure a CPS, engineers must understand the underlying assets, their vulnerabilities, and the associated threats. In energy and industrial networks, threats are broad, varying from simple localized faults to complex cascading failures.

13.2.1 Classes of Cyber-Physical Attacks

Adversaries target the cyber-physical feedback loop through several distinct vectors:

- **Measurement Error Injection:** Attackers inject constant offsets or time-varying errors into sensor readings. If left untreated, these can manipulate control laws and lead to ramp-induced physical attacks.
- **Replay Attacks:** The measurement output is maliciously intercepted and changed to reflect the value at a previous timestamp, blinding the controller to current physical realities and impacting future decisions.
- **State Estimation Corruption:** Attackers maliciously alter the estimated quantities required by the system's centralized dispatcher, causing the system to compute incorrect control actions.
- **Coordinated Attacks:** Attackers gain access to multiple critical information nodes and launch simultaneous threats from multiple sources in a coordinated manner, designed to overwhelm standard fault-tolerance mechanisms.

13.2.2 Threat Modeling and Attack Trees

To analyze these complex vulnerabilities, security architects utilize **Attack Trees**. An attack tree is a model that formalizes the structure of well-known cyber-physical attacks, detailing the exact steps an attacker must take.

For example, consider the cascading threat of cutting a city's water supply through the electrical grid. The attack tree demonstrates that the water supply can be cut if the water pumps are stopped and the reservoir is drained. The

pumps stop if the primary power supply is cut *and* the backup power supply is disabled. The power supply itself can be interrupted maliciously by compromising a substation either physically or through cyber means. By mapping these steps using AND/OR logic gates, engineers can identify the critical nodes where design-time countermeasures (e.g., adding redundancy) or run-time mitigations (e.g., dynamic isolation) must be applied.

Furthermore, engineers must continuously evaluate **unknown threats**. These are addressed by exploring various conceptual spaces: the *threat space* (newly discovered vulnerabilities), the *attack space* (new attacker tools and methods), and the *technology space* (novel threats introduced by emerging technologies like autonomous drones or IoT automation).

13.3 Risk Management and Evaluation

Once threats are modeled, they must be assessed for their potential impact. The standard risk management process for critical infrastructure is based on **ISO/IEC 27005**. This process includes:

1. **Context Establishment:** Defining the scope, identifying assets, and mapping threats to those assets.
2. **Risk Assessment:** Estimating the risk based on the likelihood of the threat exploiting a vulnerability and the severity of the harmful result.
3. **Risk Treatment & Acceptance:** Preparing a mitigation plan and acknowledging any residual risks.

Risk estimation combines the human, economic, and social impact with the risk's likelihood. A qualitative approach is often used, categorizing impact from "minimal" to "severe" (levels 1 to 5) and likelihood from "rare" to "almost certain" (e.g., 1% to 80% probability).

13.4 Advanced Computational Techniques for Resilience

Traditional risk mitigation relies heavily on fortifying networks. However, in CPS 2.0, adversarial events are notoriously hard to predict and reliable historical data is often sparse. Attempting to make a system invulnerable to every possible disruption makes the operation overly conservative and economically impractical. Instead, modern CPS architecture relies on advanced computational techniques to guarantee resilience dynamically.

13.4.1 Distributionally Robust Optimization (DRO)

To handle uncertainty in system parameters and adversarial forecasts, systems utilize stochastic optimization. **Distributionally Robust Optimization (DRO)** is a probabilistic method that determines the exact amount of reserves required to deal with cyber-physical events in a resilient and cost-effective manner.

Rather than planning for a single worst-case scenario, DRO considers a bounded set of uncertainty Ω and optimizes the system's response across a probability distribution of potential disruptions. By integrating these probabilistic models, the CPS dispatcher ensures sufficient dispatchable resources are available to absorb the impact of an attack without collapsing the grid.

13.4.2 Decentralized Autonomous Controls

During severe cyber-physical attacks, centralized communication networks may be congested, compromised, or entirely unavailable. To survive, the decision-making agents within the network must be able to act autonomously and collaboratively. This is achieved through minimally invasive, decentralized control layers:

- **Proactive Autonomy:** These controls continuously guarantee that the system operates within strict safety limits during likely, bounded disturbances. By utilizing mathematical proofs (such as Nagumo's

theorem), proactive autonomy ensures the system's operational viability without requiring communication with a central coordinator.

- **Reactive Autonomy:** When a high-impact, low-probability adversarial event thrusts the system state outside the safe operating regime, reactive autonomy takes over. Reactive control involves computing control inputs in real-time—often by efficiently solving for energy functions—to steer the highly nonlinear dynamics of the system back into a safe region as quickly as possible.

Together, these autonomous controls act as a bridge between myopic local sensors and the centralized dispatcher, ensuring that a communication failure does not result in physical destruction.

13.5 Network Reconfiguration and Resilient Learning

In distributed critical infrastructure, such as interconnected microgrids, resilience can be further enhanced by dynamically altering the physical topology of the system to isolate attacked segments.

13.5.1 Topology Reconfiguration via Distributed OPF

If a portion of the grid is compromised by a cyberattack, local controllers must determine the optimal physical interconnection topology to island the threat while minimizing generation costs and line losses. This requires solving a **Distributed Optimal Power Flow (OPF)** problem.

Because attackers may simultaneously target the controllers themselves, resilient consensus algorithms are utilized. These algorithms allow each microgrid to exchange optimal objective functions with neighboring grids via Peer-to-Peer (P2P) communication. To solve these complex distributed equations rapidly, techniques such as the **Alternating Direction Method of Multipliers (ADMM)** are heavily utilized, ensuring convergence even when parts of the network are untrusted.

13.5.2 Resilient Reinforcement Learning

Modern energy infrastructure heavily integrates inverter-based resources, such as solar panels and battery storage. These utilize **Grid-Forming (GFM)** inverters, modeled as AC voltage sources with local frequency and voltage control signals governed by droop controls.

To optimize these distributed resources under attack conditions where data sharing is limited or unsafe, researchers are deploying Reinforcement Learning (RL) and Federated Learning (FL). These data-driven, machine learning techniques allow the localized controllers to collaboratively learn optimal resilient policies without transmitting sensitive, interceptable state data across the network.

13.6 VirtMCU Homework: Attack Tree Modeling and Resilient Recovery

(Note: This exercise builds on our established course framework using the VirtMCU firmware studio, applying the theoretical concepts of this chapter to a simulated environment).

In Chapter 11, you designed a PID control loop to stabilize an inverted pendulum. In this chapter's homework, you will subject that continuous-time control loop to a simulated cyber-physical attack to evaluate its resilience.

Your Assignment:

1. **Attack Tree Construction:** Draft a formal Attack Tree outlining the steps an attacker would take to destabilize your inverted pendulum. Include nodes for both cyber vectors (e.g., compromising the PWM duty cycle communication via a buffer overflow) and physical vectors (e.g., injecting false quadrature encoder data).
2. **Implementation of a Replay Attack:** Modify your firmware to simulate a state estimation corruption. At exactly 5.0 seconds into the simulation, your code must freeze the angular sensor reading $\theta(t)$ to its last

known value, blinding the PID controller to the true physical state. Observe how rapidly the system fails.

3. **Reactive Autonomy (Fallback Control):** Implement a secondary, decentralized fallback control law. You will monitor the raw derivative of the cart's wheel velocity. If the discrepancy between the expected cart position and the stale $\theta(t)$ reading exceeds a threshold, your firmware must automatically trigger a "Reactive Autonomy" state, overriding the primary PID loop with an aggressive braking maneuver to bring the system to a safe halt.
4. **Deterministic Evaluation:** Run the simulation in VirtMCU's `slaved-icount` mode. Because virtual time is strictly deterministic, you must provide a cycle-accurate report of your Reactive Autonomy detection latency (the exact number of instructions between the attack injection and the override engagement), proving your system can reliably recover from a sensor-blinding attack within the physical constraints of the hardware.

Chapter 14: Domain-Specific Architectures and Cloud/Edge Integration

14.1 Introduction: The End of Moore's Law and the Rise of DSAs

In the preceding chapters, we focused extensively on general-purpose microprocessors and their application in Cyber-Physical Systems (CPS). We explored pipelining, out-of-order execution, memory hierarchies, and the real-time operating systems required to enforce deterministic physical control. However, the computing landscape is undergoing a profound structural shift.

For decades, software engineers relied on Moore's Law and Dennard scaling to automatically provide faster, more energy-efficient general-purpose CPUs every two years. Today, this free ride has effectively ended. As Gordon Moore himself noted, "Moore's Law can't continue forever. We have another 10–20 years before we reach a fundamental limit". The cessation of Dennard scaling (the power wall) means that architects can no longer simply increase clock frequencies or pack more active general-purpose cores onto a die without melting the silicon.

Given the end of Moore's Law and Dennard scaling, the only path forward for radically improved performance and energy efficiency is the **Domain-Specific Architecture (DSA)**. Unlike general-purpose CPUs, which are designed to run everything from text editors to web browsers reasonably well, DSAs are custom-tailored to accelerate a specific class of applications. By discarding the immense hardware complexity required for general-purpose out-of-order execution, DSAs devote their silicon area entirely to highly efficient, specialized arithmetic and data routing.

In the context of modern CPS 2.0, the architecture is split across a vast continuum: from ultra-low-power DSAs operating at the physical edge, to massive Warehouse-Scale Computers (WSCs) operating in the cloud. This chapter bridges that divide, examining the hardware architectures that define the modern edge-to-cloud ecosystem.

14.2 Architecting for the Edge: Mobile and SoC Architectures

At the very edge of the CPS network are the personal mobile devices (PMDs) and embedded sensors that directly interact with the physical world. These devices are strictly constrained by battery capacity and thermal dissipation limits.

To meet these constraints, edge devices utilize **System-on-Chip (SoC)** architectures. Rather than isolating the CPU, graphics, and memory controllers on separate physical chips, an SoC integrates them onto a single die. This minimizes the parasitic capacitance of external wiring, drastically reducing the electrical power required to move data.

Furthermore, these SoCs are deeply heterogeneous, containing numerous DSAs to offload specific tasks from the main CPU:

- **Image Signal Processors (ISPs):** Hardware accelerators designed specifically to process raw pixel data from dual cameras or IR projectors, performing noise reduction, color space conversion, and edge enhancement in real-time.
- **Digital Signal Processors (DSPs):** Highly optimized cores for processing continuous analog data, such as audio from microphones or RF signals from cellular modems.
- **Touch Controllers:** Dedicated hardware to process capacitive touch sensing. These sensors monitor changes in capacitance caused by the proximity of a conductive object (such as a human finger), allowing the hardware to filter raw measurements and pinpoint multiple simultaneous touchpoints.

By routing data directly from the physical sensor into a dedicated DSA, the SoC can process real-time events while the main general-purpose CPU remains in a deep, low-power sleep state, maximizing battery life.

14.3 Warehouse-Scale Computing (WSC)

While edge devices collect telemetry and execute local control loops, the heavy computational lifting—such as training global AI models or processing massive distributed databases—is offloaded to the cloud. The infrastructure that powers the cloud is the **Warehouse-Scale Computer (WSC)**.

It is a misconception to view a modern data center merely as a building full of individual servers. Given the immense processing, networking, and storage capabilities required, it is far more accurate to conceptualize the entire data center as a single, massively parallel computing system. The price-performance and energy efficiency concerns of a single microprocessor apply equally to WSCs, but at a scale where the equipment and the building cost hundreds of millions of dollars.

14.3.1 WSC Network Topology

A WSC is organized hierarchically. At the base level, individual servers are mounted in racks. Each server is a reasonably complete computer system containing a processor, RAM, and local storage. These servers operate in a “headless mode,” lacking any direct physical connection to a display or keyboard; all interaction occurs over the network.

Each server connects to a Top-of-Rack (ToR) network switch, typically using a 10 Gbit/s Ethernet cable. Servers within the same rack can communicate with one another at full speed. However, the racks are then grouped into clusters, sharing a second-level cluster switch. Because it is economically unfeasible to provide full bandwidth between every rack in the facility, the network is often **oversubscribed**. For example, an oversubscription factor of 4 means that the external network capacity leaving the rack is only one-quarter of the peak communication speed available inside the rack. Software architects must

carefully place communicating processes within the same rack to avoid saturating the cluster switches.

14.3.2 Hardware Fault Management

Because a WSC contains tens of thousands of individual servers, hardware failures are not anomalies; they are continuous, daily occurrences. Even if highly reliable (and expensive) components are used, the sheer scale guarantees component death.

Therefore, WSC architecture relies on software to mask hardware failures. When a higher-level system dispatches a request (such as an internet search query) to a lower-level index server, it actively monitors the response latency. If the lower-level server becomes persistently unresponsive or returns corrupted data, the requesting system assumes a hardware fault has occurred, abandons the failed node, and immediately resends the request to a redundant backup server. The failed server is taken offline automatically so technicians can repair or replace it without interrupting the global service.

14.4 Accelerating AI: GPUs and Neural Processing Units

The most significant driver of Domain-Specific Architectures today is Artificial Intelligence (AI). Historically, computer scientists built AI as a massive set of logical rules; today, it is driven by deep neural networks that learn from data. As it turns out, it is much harder to program a computer to *be* intelligent than to program it to *learn* to be intelligent.

Training and executing Deep Neural Networks (DNNs) requires astonishing amounts of computational power. Because neural networks are fundamentally based on linear algebra (specifically, massive matrix multiplications), they exhibit immense **Data-Level Parallelism (DLP)**. General-purpose CPUs, which are optimized for sequential logic and unpredictable branching, are terrible at this.

14.4.1 Graphics Processing Units (GPUs)

The first hardware to successfully exploit the DLP of neural networks was the Graphics Processing Unit (GPU), such as the NVIDIA A100 and T4. Originally designed to compute millions of independent pixels for synthetic video scenes, GPUs employ a heavily threaded, Single Instruction, Multiple Data (SIMD) architecture. A modern GPU can execute thousands of floating-point multiply-accumulate (MAC) operations in parallel, making them exceptional AI accelerators.

14.4.2 Tensor Processing Units (TPUs) and Rack-Scale Accelerators

While GPUs are fast, they still retain legacy hardware for graphics rendering (such as texture mapping and rasterization) that is useless for pure AI workloads. To achieve maximum efficiency, companies have developed pure AI DSAs.

Google's **Tensor Processing Unit (TPU)** is a custom Application-Specific Integrated Circuit (ASIC) designed exclusively to accelerate DNNs in the data center. The heart of a TPU is a massive systolic array multiplier—a 2D grid of ALUs that passes data directly from one ALU to the next without writing intermediate results back to registers, drastically reducing power consumption while executing tens of thousands of matrix operations per clock cycle.

Intel has taken a similar architectural leap with **Jaguar Shores**, a next-generation rack-scale AI accelerator platform. Rather than functioning as a standalone chip or PCIe card, Jaguar Shores is designed as an entire rack-based solution, providing a complete AI processing system scaled specifically to meet the extreme memory bandwidth and training demands of modern machine learning.

14.5 VirtMCU Homework: Designing an Edge AI Accelerator

In previous chapters, you wrote strict PID control loops for an inverted pendulum and a drone using standard ARM AArch64 assembly. However, in modern CPS, we frequently augment classical control with predictive machine learning models at the edge.

If you attempt to execute a Deep Neural Network (DNN) inference pass using standard sequential software in your ISR, the matrix multiplications will consume too many clock cycles. You will exceed your T_s (sample time) control deadline, causing the simulated drone to crash.

To solve this, you will step into the role of a computer architect and design a Domain-Specific Architecture (DSA) within the **VirtMCU FirmwareStudio**.

1. **Memory-Mapped Accelerator Design:** You will be provided with the C++ source code for the VirtMCU hardware backend. You must design a mock “Tensor Coprocessor” peripheral that attaches to the virtual APB interconnect.
2. **MMIO Interface:** Design the Memory-Mapped I/O (MMIO) registers for your DSA. You will need a `MATRIX_A` port, a `MATRIX_B` port, a `START_INFERENCE` command register, and a `RESULT` memory buffer.
3. **Firmware Integration:** Rewrite your drone flight controller firmware. Instead of performing the neural network matrix multiplication in software, your firmware will use DMA to blast the sensor data into your custom DSA’s MMIO registers, trigger the `START` command, and yield the CPU.
4. **Deterministic Benchmarking:** Run the co-simulation in `slaved-icount` mode. By offloading the DLP workload to your specialized hardware, you will empirically measure the massive reduction in required CPU clock cycles. You must submit a report proving that the introduction of the DSA reduces your Worst-Case Execution Time (WCET) enough to comfortably meet the hard real-time deadlines of the physical control loop.

Epilogue: The Future of Cyber-Physical Systems

We have traversed the entire vertical stack of modern computing—from the physics of semiconductor transistors, up through the microarchitecture and instruction sets of ARM and RISC-V processors, into the memory hierarchies and interconnect fabrics that shuffle our data, and finally into the software layers of real-time operating systems and distributed Cyber-Physical Systems (CPS).

The computing landscape you are entering as a graduate engineer is fundamentally different from the one that existed just a decade ago. For over fifty years, the industry relied on Moore's Law and Dennard scaling to reliably double transistor density and performance while maintaining constant power densities. Today, those physical scaling laws have effectively collapsed. We have hit the "power wall," meaning we can no longer simply increase clock frequencies without melting the silicon.

To continue achieving performance gains, the future of CPS architecture lies in **Domain-Specific Architectures (DSAs)** and hardware-software co-design. General-purpose processors will increasingly serve as traffic cops, delegating heavy mathematics to specialized Neural Processing Units (NPUs), Tensor Processing Units (TPUs), and custom Field Programmable Gate Arrays (FPGAs) operating at the edge.

Furthermore, as our physical infrastructure—power grids, medical devices, and autonomous transportation—becomes inextricably linked to cloud networks, **security and resilience** must be treated as first-class architectural constraints, not software afterthoughts. Designing systems that can autonomously withstand, isolate, and recover from cyber-physical attacks using distributed machine learning and robust optimization is the grand challenge of your generation.

The tools you have used in this text, from bare-metal debuggers to the deterministic VirtMCU co-simulation framework, are designed to give you

absolute control over both the logical correctness and the continuous-time physical realities of your designs.

Appendix A: Instruction Set Architecture (ISA) Reference

This appendix provides a quick reference for the primary Instruction Set Architectures utilized in this text: ARM AArch64 and RISC-V (RV32I).

A.1 ARM AArch64 (A64) Instruction Subset

The ARMv8-A architecture introduced the 64-bit AArch64 execution state.

Data Transfer (Load/Store)

- `ldr Xd, [Xn, offset]` : Load a 64-bit doubleword from the memory address $(Xn + offset)$ into register `Xd`.
- `str Wd, [Xn, offset]` : Store a 32-bit word from register `Wd` into memory.
- `ldp / stp` : Load/Store Pair. Pushes or pops two registers simultaneously, heavily used in stack frame prologues and epilogues (e.g., `stp fp, lr, [sp, #-16]!`).

Arithmetic and Logic

- `add Xd, Xn, Xm` : $Xd = Xn + Xm$. (Use `adds` to update PSTATE condition flags).
- `sub Xd, Xn, Xm` : $Xd = Xn - Xm$. (Use `subs` to update flags).
- `mul Xd, Xn, Xm` : $Xd = Xn * Xm$.
- `and / orr / eor` : Bitwise AND, OR, and Exclusive-OR.

Control Flow

- `cmp Xn, Xm` : Compare (implicitly performs `subs` and updates flags without storing the result).
- `b label` : Unconditional branch to `label`.

- `beq / bne / blt / bgt`: Conditional branches (Branch if Equal, Not Equal, Less Than, Greater Than).
- `bl label`: Branch with Link. Stores the return address in the Link Register (`LR / X30`) and jumps to `label` (used for function calls).
- `ret`: Return from subroutine. Copies `LR` to the Program Counter (`PC`).

A.2 RISC-V (RV32I) Base Instruction Subset

RISC-V is an open-source ISA that heavily utilizes pseudo-instructions to minimize silicon decoder complexity.

Core Instructions

- `addi rd, rs1, imm`: Add immediate. $rd = rs1 + imm$.
- `lui rd, imm`: Load Upper Immediate. Places a 20-bit value into the highest 20 bits of `rd`.
- `auipc rd, imm`: Add Upper Immediate to PC.
- `jal rd, offset`: Jump and Link. Jumps to $PC + offset$ and saves the return address in `rd` (usually `x1 / ra`).
- `beq / bne rs1, rs2, offset`: Branch if `rs1 == (or !=) rs2`.

Common Pseudo-Instructions

- `nop`: Translates to `addi x0, x0, 0`.
- `mv rd, rs`: Translates to `addi rd, rs, 0`.
- `not rd, rs`: Translates to `xori rd, rs, -1`.
- `j offset`: Translates to `jal x0, offset`.

Appendix B: Selected Answers to VirtMCU Exercises

Chapter 3: Stack Tracing and Deterministic Execution *Solution Note:* When executing the recursive factorial function in VirtMCU `slaved-icount` mode, students should observe that the Link Register (`x30`) is pushed to the stack exactly 16 bytes below the previous Frame Pointer (`x29`). Because `slaved-icount` ties the virtual clock strictly to retired instructions, inspecting the cycle counter before and after the recursive calls will yield an identical delta on every simulation run, free from host OS jitter.

Chapter 7: Deterministic Interrupt Handler *Solution Note:* Students must successfully construct the Vector Table to point to their ISR. Inside the ISR, the code must execute `stp x0, x1, [sp, #-16]!` (and subsequent saves) to preserve the volatile registers. Failing to acknowledge and clear the simulated push-button peripheral's IRQ flag in the MMIO register will result in an infinite interrupt loop upon executing the return instruction.

Chapter 11: Inverted Pendulum Co-Simulation *Solution Note:* The continuous-time MuJoCo physics engine relies on the VirtMCU Sensor Abstraction Layer (SAL) to discretize angles into 12-bit simulated quadrature encoder pulses. Students must apply a Proportional-Integral-Derivative (PID) algorithm in C. Saturation in the Actuator Abstraction Layer (AAL) means that if the proportional gain (`$k_p$`) is set too high, the requested PWM duty cycle will clip, causing integral windup and leading to the pendulum crashing.

Glossary of Terms and Abbreviations

- **ABI (Application Binary Interface):** The low-level standard that dictates how software interacts with the hardware, including register usage conventions (volatile vs. nonvolatile) and parameter passing during function calls.
- **ADC (Analog-to-Digital Converter):** A hardware component that converts continuous physical voltages into discrete digital numbers, limited by sampling rate and quantization resolution.
- **AMBA (Advanced Microcontroller Bus Architecture):** A family of open-standard, on-chip interconnect specifications (including APB, AHB, and AXI) developed by ARM for routing data blocks within a System-on-Chip.
- **AXI (Advanced eXtensible Interface):** A high-performance AMBA interconnect protocol that utilizes five independent, unidirectional channels to allow full-duplex, out-of-order memory transactions.
- **CPI (Cycles Per Instruction):** A fundamental performance metric representing the average number of clock cycles required to execute a machine instruction.
- **CPS (Cyber-Physical System):** An engineered system built from, and dependent upon, the seamless integration of computational algorithms and physical components.
- **DMA (Direct Memory Access):** A hardware controller that autonomously transfers blocks of data between peripherals and main memory, bypassing the CPU execution pipeline to save clock cycles.
- **DSA (Domain-Specific Architecture):** Custom silicon tailored to accelerate a specific class of applications (e.g., neural networks), trading general-purpose flexibility for massive efficiency gains.
- **ISA (Instruction Set Architecture):** The boundary between hardware and software; the complete set of instructions, registers, and memory models understood by a specific microprocessor.
- **MMIO (Memory-Mapped I/O):** A system design technique where peripheral control and data registers are mapped into the CPU's standard physical memory address space, allowing hardware to be controlled via standard load/store instructions.

- **NVIC (Nested Vectored Interrupt Controller):** A hardware component in ARM architectures that prioritizes and manages incoming physical interrupts, automatically forcing the PC to the correct Interrupt Service Routine.
- **PWM (Pulse Width Modulation):** A technique for simulating an analog voltage output by rapidly switching a digital signal HIGH and LOW. The effective output is proportional to the duty cycle.
- **RTOS (Real-Time Operating System):** An operating system designed to guarantee strict deterministic timing and meet hard deadlines, rather than optimizing for average throughput.
- **TLB (Translation Lookaside Buffer):** An ultra-fast hardware cache located inside the Memory Management Unit (MMU) that stores recent virtual-to-physical page table translations.
- **TLM (Transaction Level Modeling):** A high-level simulation technique (often used in SystemC) that abstracts away individual clock cycles and pins, modeling bus communication via object-oriented method calls to dramatically speed up simulation.

Bibliography / Further Reading

Feng, C., et al. *Cyber-Physical Virtual Power Plants: Modeling, Operation, and Control*. IEEE / Springer, 2025.

Arm Limited. *Fundamentals of System-on-Chip Design on Arm Cortex-M Microcontrollers*. Arm Education Media, 2021.

VirtMCU FirmwareStudio. RefractSystems / GitHub Repository.
<https://github.com/RefractSystems/virtmcu>

Große, D., and Drechsler, R. *Quality-Driven SystemC Design*. Springer Science+Business Media, 2010.

Harris, D. M., and Harris, S. L. *Digital Design and Computer Architecture*, 2nd Edition. Morgan Kaufmann / Elsevier, 2013.

Hennessey, J. L., and Patterson, D. A. *Computer Architecture: A Quantitative Approach*, 7th Edition. Morgan Kaufmann / Elsevier, 2026.

Hyde, R. *The Art of ARM Assembly Vol 1: 64-Bit ARM Machine Organization and Programming*. No Starch Press, 2025.

Jha, A., et al. *Cyber Physical System 2.0: Communication and Computational Technologies*. 2025.

Ledin, J. *Modern Computer Architecture and Organization*, 3rd Edition. Packt Publishing, 2026.

Lyshevski, S. E. *Design and Control of Physical and Cyber-Physical Systems*. 2026.

Greaves, D. J. *Modern System-on-Chip Design on Arm*. Arm Education Media, 2021.

Patel, H. D., and Shukla, S. K. *SystemC Kernel Extensions for Heterogeneous System Modeling: A Framework for Multi-MoC Modeling & Simulation*. Kluwer Academic Publishers, 2005.

Plantz, R. G. *Introduction to Computer Organization: ARM Edition*. No Starch Press, 2025.

Poonia, R. C., et al. *Smart Cyber-Physical Systems: Innovations and Real-World Implications*. 2025.

Soldatos, J., et al. *Cyber-Physical Threat Intelligence for Critical Infrastructures Security*. Now Publishers, 2020.

Tanenbaum, A. S., and Austin, T. *Structured Computer Organization*, 6th Edition. Pearson, 2016.

Yan, J., et al. *Secure Detection and Control in Cyber-Physical Systems*. Taylor & Francis, 2026.

Index

- **A**

- Addressing Modes, Ch. 3, Ch. 5, App. A
- Advanced eXtensible Interface (AXI), Ch. 6, Glossary
- Application Binary Interface (ABI), Ch. 3, Ch. 12, Glossary
- Arithmetic Logic Unit (ALU), Ch. 4
- Artificial Intelligence (Edge/Cloud), Ch. 9, Ch. 14

- **B**

- Backpressure, Ch. 6
- Branch Prediction, Ch. 4
- Bubble Sort, Ch. 4 (Homework), Ch. 10

- **C**

- Cache Coherence, Ch. 5
- Cache Memory (L1, L2, L3), Ch. 5
- Compare and Branch, Ch. 3
- Context Switching, Ch. 7, Ch. 8
- Cyber-Physical System (CPS 1.0 & 2.0), Ch. 9, Glossary

- **D**

- Deadlocks, Ch. 8
- Direct Memory Access (DMA), Ch. 6, Glossary
- Discrete-Event (DE) Simulation, Ch. 10
- Domain-Specific Architectures (DSAs), Ch. 14, Epilogue
- Dynamic RAM (DRAM), Ch. 5

- **E**

- Edge Computing, Ch. 9, Ch. 14
- Exceptions and Faults, Ch. 7

- **G**

- General-Purpose Registers, Ch. 3, Ch. 12
- **I**
 - Instruction-Level Parallelism (ILP), Ch. 4
 - Interrupt Service Routine (ISR), Ch. 7
 - Inverted Pendulum, Ch. 11, Ch. 13
- **M**
 - Memory Management Unit (MMU), Ch. 5
 - Memory-Mapped I/O (MMIO), Ch. 6, Glossary
 - Moore's Law, Ch. 14, Epilogue
- **P**
 - Paging and Virtual Memory, Ch. 5
 - Pipelining, Ch. 4
 - Priority Inversion, Ch. 8
 - Pulse Width Modulation (PWM), Ch. 11, Glossary
- **Q**
 - Quantization Error, Ch. 11
- **R**
 - Rate Monotonic Scheduling (RMS), Ch. 8
 - Reorder Buffer (ROB), Ch. 4
 - RISC-V Architecture, Ch. 12, App. A
- **S**
 - Semaphores and Mutexes, Ch. 8
 - Sensor Abstraction Layer (SAL) / Actuator Abstraction Layer (AAL), Ch. 11
 - Superscalar Execution, Ch. 4
 - SystemC, Ch. 10
- **T**

- Transaction Level Modeling (TLM), Ch. 10, Glossary
- Translation Lookaside Buffer (TLB), Ch. 5, Glossary

- **V**

- Vector Table, Ch. 7
- Virtual Power Plants (VPP), Ch. 9, Ch. 10
- VirtMCU FirmwareStudio, Ch. 3, Ch. 4, Ch. 7, Ch. 10, Ch. 11, App. B